*make it*

*use it*

*run and adapt it*

**Pretrain**

next token on ~15T tokens

**Scale**

loss falls as a power law

**Post-train**

base model to assistant

**Prompt**

steer it, no training

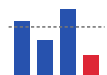

**Evaluate**

benchmarks and judges

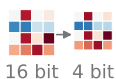


**Quantize**

8B model: 16 → 5 GB

**LoRA**

train < 1% of weights