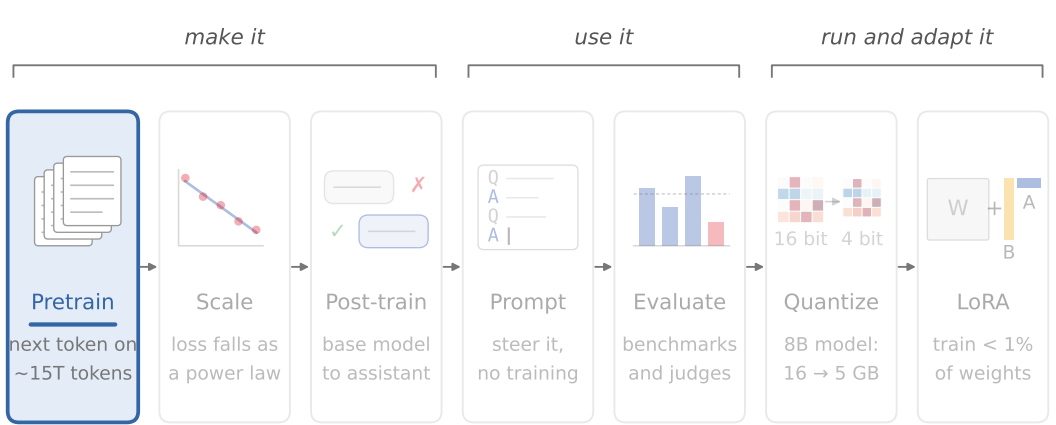
make it
use it
run and adapt it
Pretrain
next token on
~15T tokens
Scale
loss falls as
a power law
Post-train
base model
to assistant
Q
A
Q
A
Prompt
steer it,
no training
Evaluate
benchmarks
and judges
16 bit
4 bit
Quantize
8B model:
16 → 5 GB
W
+
A
B
LoRA
train < 1%
of weights