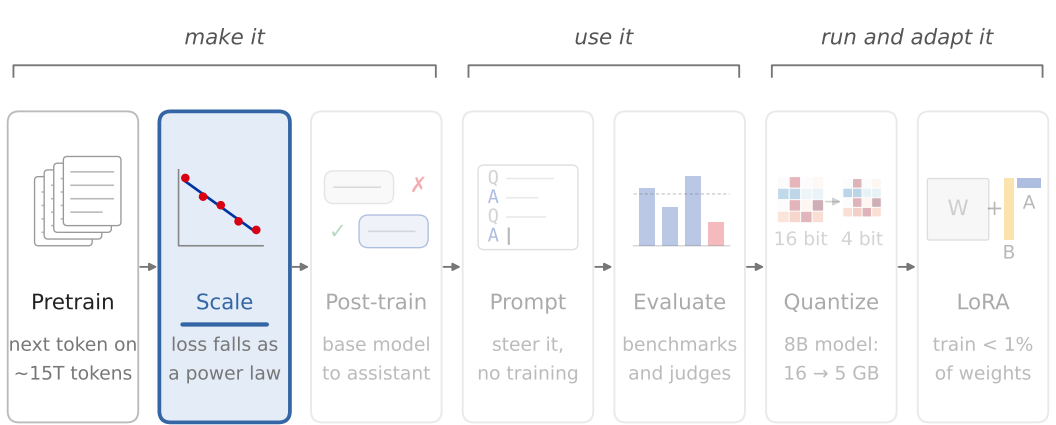
make it
use it
run and adapt it
Pretrain
next token on ~15T tokens
Scale
loss falls as a power law
Post-train
base model to assistant
Q
A
Q
A
Prompt
steer it, no training
Evaluate
benchmarks and judges
16 bit
4 bit
Quantize
8B model: 16 → 5 GB
W
+
A
B
LoRA
train < 1% of weights