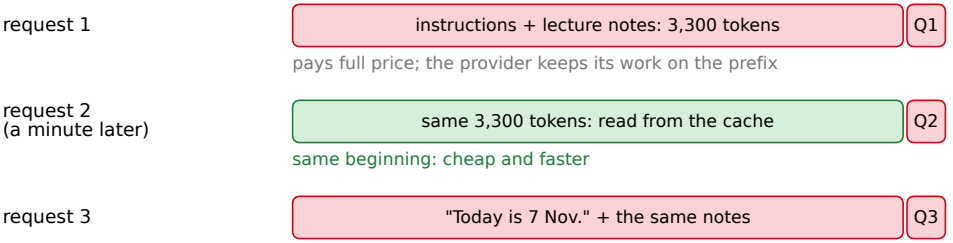
request 1
instructions + lecture notes: 3,300 tokens
Q1
pays full price; the provider keeps its work on the prefix
request 2
(a minute later)
same 3,300 tokens: read from the cache
Q2
same beginning: cheap and faster
request 3
"Today is 7 Nov." + the same notes
Q3
first token differs: no cache hit at all

schematic