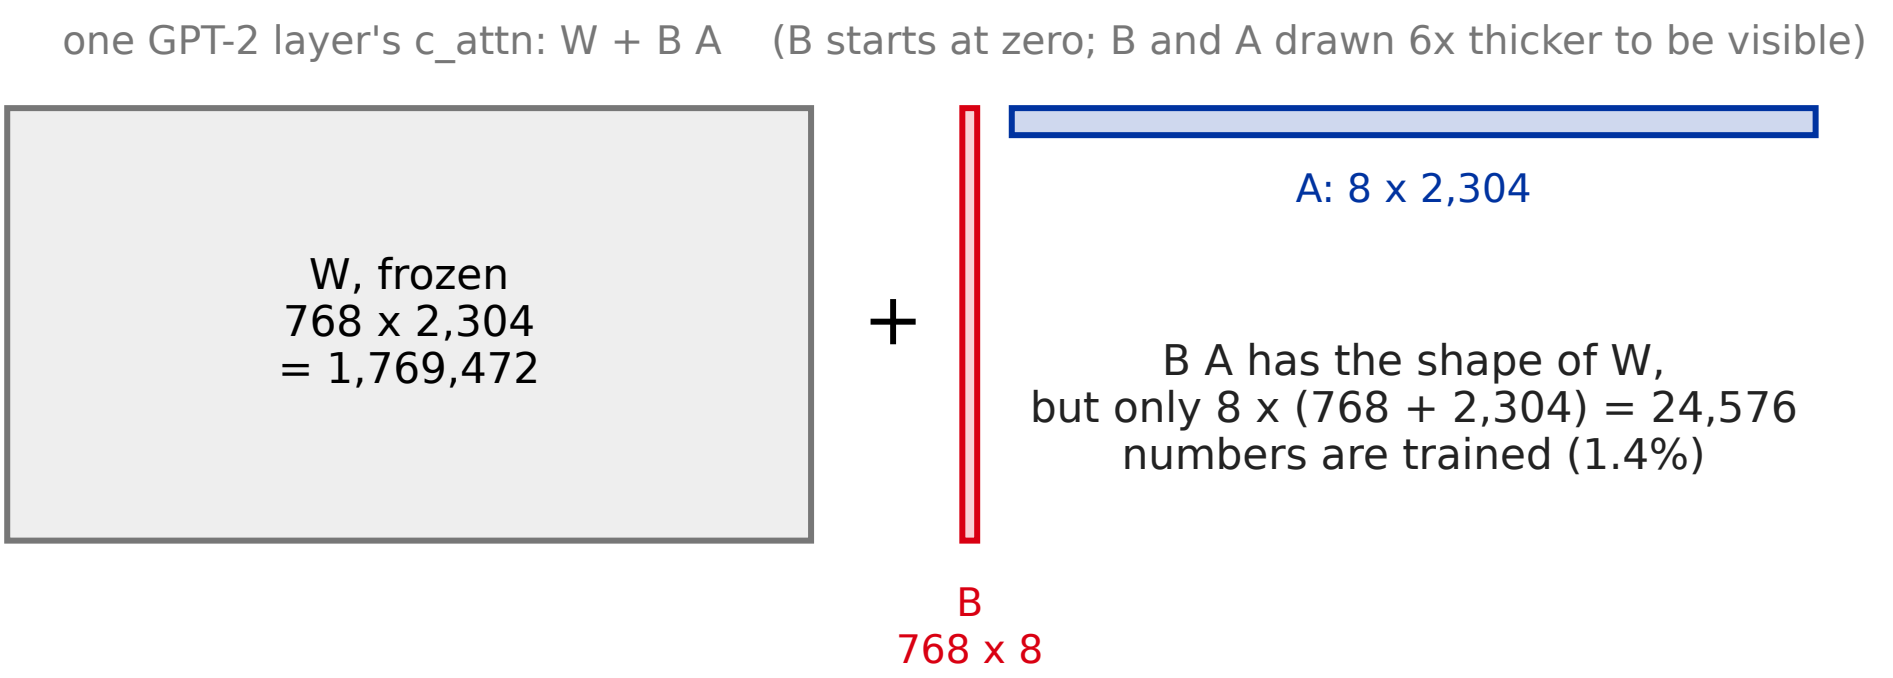
one GPT-2 layer's c_attn: W + B A (B starts at zero; B and A drawn 6x thicker to be visible)
W, frozen
768 x 2,304
= 1,769,472
+
A: 8 x 2,304
B A has the shape of W,
but only 8 x (768 + 2,304) = 24,576
numbers are trained (1.4%)
B
768 x 8