


## Information Theory III: Advanced Topics for ML

Data Processing Inequality ·  $f$ -Divergences · ELBO · Information Bottleneck

# Recap: What We Have So Far



We established: cross-entropy loss = MLE, forward KL for supervised learning, reverse KL for variational inference, MI for feature selection.

**Today:** Four powerful extensions.

1. Data Processing Inequality
2.  $f$ -Divergences & GANs
3. ELBO & VAEs
4. Information Bottleneck


## Data Processing Inequality

Processing data can only **destroy** information, never create it.

# Markov Chains and Information Flow

A **Markov chain**  $X \rightarrow Y \rightarrow Z$  means:  $Z$  depends on  $X$  only through  $Y$ .

$$p(x, y, z) = p(x) p(y | x) p(z | y)$$



## Examples in ML:

- ▶ Raw pixels → convolutional features → class prediction
- ▶ Original data → PCA projection → clustering
- ▶ Text → embedding → classifier output

# The Data Processing Inequality (DPI)

## Data Processing Inequality:

If  $X \rightarrow Y \rightarrow Z$  is a Markov chain, then:

$$I(X; Z) \leq I(X; Y)$$

"No processing of  $Y$  can increase the information that  $Y$  contains about  $X$ ."

# The Data Processing Inequality (DPI)

## Data Processing Inequality:

If  $X \rightarrow Y \rightarrow Z$  is a Markov chain, then:

$$I(X; Z) \leq I(X; Y)$$

"No processing of  $Y$  can increase the information that  $Y$  contains about  $X$ ."

### Proof sketch:


Chain rule:  $I(X; Y, Z) = I(X; Y) + I(X; Z | Y) = I(X; Z) + I(X; Y | Z)$ .

Since  $X \rightarrow Y \rightarrow Z$ :  $I(X; Z | Y) = 0$ , so  $I(X; Y) = I(X; Z) + \underbrace{I(X; Y | Z)}_{\geq 0} \geq I(X; Z)$ .  $\square$

Equality iff  $Z$  is a **sufficient statistic** for  $X$  w.r.t.  $Y$ .

# DPI in Neural Networks

A feedforward network with layers  $\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_L$ :



Each layer forms a Markov chain:  $\mathbf{X} \rightarrow \mathbf{h}_1 \rightarrow \mathbf{h}_2 \rightarrow \dots \rightarrow \mathbf{h}_L \rightarrow \hat{Y}$ .

By DPI applied repeatedly:

$$I(\mathbf{X}; \mathbf{h}_1) \geq I(\mathbf{X}; \mathbf{h}_2) \geq \dots \geq I(\mathbf{X}; \mathbf{h}_L) \geq I(\mathbf{X}; \hat{Y})$$

**Each layer can only lose information about the input.**

The network must learn to keep what matters (about  $Y$ ) and discard what doesn't.

This is exactly the **information bottleneck** idea (coming later).

# DPI: Practical Consequences

**1. Feature engineering matters:** If your features throw away relevant information, no model can recover it. Garbage in, garbage out — **formally**.

**2. Dimensionality reduction has a cost:** PCA, autoencoders, embeddings — all lose information. The question is whether they keep what matters for your task.

**3. Sufficient statistics are special:** A statistic  $T(\mathbf{X})$  is sufficient for  $\theta$  iff  $I(T; \theta) = I(\mathbf{X}; \theta)$ . DPI equality case!

**4. Post-processing can't help:** If model  $A$  has less MI with the target than model  $B$ , no amount of post-processing of  $A$ 's output can beat  $B$ .

## *f*-Divergences

KL divergence is just one member of a large family.  
The right divergence depends on the task.

# The $f$ -Divergence Family

**$f$ -Divergence** (Ali–Silvey, Csiszár, 1963/1967):

$$D_f(p\|q) = \sum_x q(x) f\left(\frac{p(x)}{q(x)}\right)$$

$$\text{where } f \text{ is convex with } f(1) = 0.$$

**Key properties** (inherited from convexity of  $f$ ):

- ▶  $D_f(p\|q) \geq 0$  always, with equality iff  $p = q$
- ▶ **Satisfies DPI:** If  $X \rightarrow Y$ , then  $D_f(p_Y\|q_Y) \leq D_f(p_X\|q_X)$
- ▶ Joint convexity in  $(p, q)$

Different choices of  $f$  give different divergences — each with different sensitivities to where  $p$  and  $q$  disagree.

# The Family Members

Name	$f(t)$	Formula	ML use
KL	$t \log t$	$\sum p \log \frac{p}{q}$	MLE, VI
Reverse KL	$-\log t$	$\sum q \log \frac{q}{p}$	Variational inference
Total Variation	$\frac{1}{2} t - 1 $	$\frac{1}{2} \sum  p - q $	Robustness
Chi-squared	$(t - 1)^2$	$\sum \frac{(p-q)^2}{q}$	Goodness-of-fit
<b>Jensen–Shannon</b>	(see below)	$\frac{1}{2} D_{\text{KL}}(p\ m) + \frac{1}{2} D_{\text{KL}}(q\ m)$	<b>GANs</b>
Hellinger	$(\sqrt{t} - 1)^2$	$\sum (\sqrt{p} - \sqrt{q})^2$	Density estimation

$$\text{JS: } f(t) = -\frac{t+1}{2} \log \frac{t+1}{2} + \frac{t \log t}{2}, \quad m = \frac{p+q}{2}.$$

## Jensen–Shannon Divergence: The Symmetric KL

KL divergence has two problems: it's **asymmetric** and **unbounded**. Jensen–Shannon fixes both:

$$\text{JSD}(p\|q) = \frac{1}{2} D_{\text{KL}}(p \parallel \frac{p+q}{2}) + \frac{1}{2} D_{\text{KL}}(q \parallel \frac{p+q}{2})$$

### Properties:

- ▶ ✓ Symmetric:  $\text{JSD}(p\|q) = \text{JSD}(q\|p)$
- ▶ ✓ Bounded:  $0 \leq \text{JSD} \leq \log 2$
- ▶ ✓  $\sqrt{\text{JSD}}$  is a true metric
- ▶ ✓ Always finite (even when supports differ)

### Compare to KL:

- ▶ ✗ KL: asymmetric
- ▶ ✗ KL: unbounded ( $\rightarrow \infty$ )
- ▶ ✗ KL:  $\infty$  when  $q(x) = 0, p(x) > 0$

**Intuition:** Instead of comparing  $p$  to  $q$  directly, both  $p$  and  $q$  are compared to their average  $m = \frac{p+q}{2}$ .

## *f*-Divergences and GANs

The original GAN (Goodfellow et al., 2014) trains a generator  $G$  and discriminator  $D$ :

$$\min_G \max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

## *f*-Divergences and GANs

The original GAN (Goodfellow et al., 2014) trains a generator  $G$  and discriminator  $D$ :

$$\min_G \max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

**Key result:** For optimal  $D^*$ , the GAN objective becomes:

$$\min_G 2 \text{JSD}(p_{\text{data}} \| p_G) - \log 4$$

**Training a GAN = minimizing JS divergence between real and generated data.**

## *f*-Divergences and GANs

The original GAN (Goodfellow et al., 2014) trains a generator  $G$  and discriminator  $D$ :

$$\min_G \max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

**Key result:** For optimal  $D^*$ , the GAN objective becomes:

$$\min_G 2 \text{JSD}(p_{\text{data}} \| p_G) - \log 4$$

**Training a GAN = minimizing JS divergence between real and generated data.**


***f*-GAN** (Nowozin et al., 2016) generalizes this:

Choose **any** *f*-divergence → get a different GAN variant

KL-GAN, reverse-KL-GAN, Pearson- $\chi^2$ -GAN, Hellinger-GAN, ...

## When Divergences Differ: Support Mismatch

What happens when  $p$  and  $q$  have **different supports**?



Divergence	Disjoint value	Gradient?
$D_{\text{KL}}(p\ q)$	$+\infty$	$\times$ undefined
Total Variation	1 (saturated)	$\times$ zero
$JSD(p\ q)$	$\log 2$ (saturated)	$\times$ zero

This is why early GANs were hard to train! JS gradients vanish when  $p_{\text{data}}$  and  $p_G$  have little overlap. Fix: **Wasserstein distance** — always gives useful gradients.


## ELBO & Variational Autoencoders

The reverse KL from Lecture 2 leads to  
the most important equation in generative modeling.

# The Problem: Intractable Posteriors

In a latent variable model, we want  $p_\theta(\mathbf{x})$  — the marginal likelihood:

$$p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$



## Two problems:

1. The integral over  $\mathbf{z}$  is usually **intractable** (no closed form).
2. The posterior  $p_\theta(\mathbf{z} | \mathbf{x}) = \frac{p_\theta(\mathbf{x}|\mathbf{z}) p(\mathbf{z})}{p_\theta(\mathbf{x})}$  requires  $p_\theta(\mathbf{x})$  — circular!

**Idea:** Approximate the intractable posterior  $p_\theta(\mathbf{z}|\mathbf{x})$  with a simple distribution  $q_\phi(\mathbf{z}|\mathbf{x})$  — using **reverse KL**.

## Deriving the ELBO

Start with the log-marginal likelihood and use any distribution  $q_\phi(\mathbf{z}|\mathbf{x})$ :

$$\begin{aligned}\log p_\theta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \log \int q_\phi(\mathbf{z}|\mathbf{x}) \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \quad (\text{Jensen's inequality})\end{aligned}$$

# Deriving the ELBO

Start with the log-marginal likelihood and use any distribution  $q_\phi(\mathbf{z}|\mathbf{x})$ :

$$\begin{aligned}\log p_\theta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \log \int q_\phi(\mathbf{z}|\mathbf{x}) \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \quad (\text{Jensen's inequality})\end{aligned}$$

**Evidence Lower BOund (ELBO):**

$$\mathcal{L}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right] \leq \log p_\theta(\mathbf{x})$$

The **gap** between ELBO and log-evidence is exactly the reverse KL:

$$\log p_\theta(\mathbf{x}) = \underbrace{\mathcal{L}(\theta, \phi; \mathbf{x})}_{\text{ELBO}} + \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\theta(\mathbf{z}|\mathbf{x}))}_{\geq 0}$$

# ELBO = Reconstruction – KL Penalty

Expanding the ELBO gives a beautiful decomposition:

$$\mathcal{L}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})}[\log p_\theta(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}))$$

**Reconstruction:**  $\mathbb{E}_q[\log p_\theta(\mathbf{x}|\mathbf{z})]$

“How well can we decode  $\mathbf{x}$  from  $\mathbf{z}$ ? ”

**KL penalty:**  $D_{\text{KL}}(q \parallel p(\mathbf{z}))$


“How close is  $q$  to the prior? ”

**Maximize ELBO** = make reconstructions good (high  $\log p_\theta(\mathbf{x}|\mathbf{z})$ ) while keeping  $q_\phi(\mathbf{z}|\mathbf{x})$  close to the prior  $p(\mathbf{z})$ .

The KL penalty acts as a **regularizer** on the latent space.

# The Variational Autoencoder (VAE)

The VAE (Kingma & Welling, 2014) implements the ELBO with neural networks:



**Encoder**  $q_\phi(z|x) = N(\mu_\phi(x), \text{diag}(\sigma_\phi^2(x)))$  — outputs  $\mu, \sigma$

**Reparameterization:**  $z = \mu + \sigma \odot \varepsilon, \varepsilon \sim N(\mathbf{0}, \mathbf{I})$  — enables backprop through sampling

**Decoder**  $p_\theta(x|z)$  — reconstructs input from latent code

# VAE Loss = ELBO in Practice

For Gaussian encoder and prior, the KL term has a **closed form**:

$$D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^d (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1)$$


The reconstruction term depends on the output distribution:

Data type	$p_{\theta}(\mathbf{x} \mathbf{z})$	Reconstruction loss
Continuous	Gaussian	MSE: $\ \mathbf{x} - \hat{\mathbf{x}}\ ^2$
Binary / images	Bernoulli	Binary cross-entropy

$$\text{VAE loss} = \underbrace{\text{Reconstruction error}}_{\text{MSE or BCE}} + \beta \cdot \underbrace{D_{\text{KL}}(q_{\phi}||p)}_{\text{latent regularizer}}$$

$\beta = 1$ : standard VAE.     $\beta > 1$ :  $\beta$ -VAE (disentangled representations).

# The ELBO Landscape



**Maximize ELBO w.r.t.  $\theta$ :** Pushes up the evidence (better generative model).

**Maximize ELBO w.r.t.  $\phi$ :** Shrinks the gap (better approximate posterior).


## The Information Bottleneck

Compress the input. Preserve what matters for the task.

A principled theory of representation learning.

# The Information Bottleneck Principle

Setup: input  $X$ , target  $Y$ , and a representation  $T$  that compresses  $X$ .



**IB Objective** (Tishby et al., 1999):


$$\min_{p(t|x)} I(X; T) - \beta I(T; Y)$$

Minimize info kept about  $X$  (compress), maximize info kept about  $Y$  (predict).

$\beta > 0$  is a Lagrange multiplier controlling the compression–prediction tradeoff.

# The Information Plane

Each representation  $T$  can be plotted as a point  $(I(X; T), I(T; Y))$ :




The IB curve traces optimal tradeoffs. Points below it are suboptimal.

$\beta$  moves you along the curve: small  $\beta \rightarrow$  more compression, large  $\beta \rightarrow$  more prediction.

# IB and Deep Learning (Tishby & Zaslavsky, 2015)

**Claim:** Deep networks implicitly optimize the IB tradeoff.

Each hidden layer  $\mathbf{h}_\ell$  defines a point in the information plane:



**Phase 1 (fitting):**  $I(\mathbf{h}; \mathbf{Y})$  increases (learning).    **Phase 2 (compression):**  $I(\mathbf{X}; \mathbf{h})$  decreases (forgetting irrelevant details).

# IB: The Debate and the Takeaway

## Evidence for IB:

- ▶ Two-phase behavior observed empirically with saturating activations ( $\tanh$ )
- ▶ Compression correlates with generalization
- ▶ IB provides a principled objective for representation learning

## Caveats:


- ▶ Compression phase not always observed (e.g., ReLU networks — Saxe et al., 2018)
- ▶ MI estimation in high dimensions is hard and noisy
- ▶ Deterministic networks:  $I(X; \mathbf{h})$  can be infinite

**Regardless of the debate, IB gives a valuable conceptual framework:**

Good representations compress the input (low  $I(X; T)$ ) while retaining what's relevant for the task (high  $I(T; Y)$ ).

This intuition underlies autoencoders, bottleneck layers, and regularization.

# IB Connects to Everything



IB is a **unifying lens**: many ML techniques can be seen as approximate solutions to the information bottleneck tradeoff.

# Today's Toolbox

Concept	Key Idea	ML Application
DPI	$X \rightarrow Y \rightarrow Z \Rightarrow I(X; Z) \leq I(X; Y)$	Layers lose info, sufficient statistics
$f$ -divergence	$D_f(p\ q) = \sum q f(p/q)$	Unifies KL, TV, $\chi^2$ , Hellinger
Jensen–Shannon	$\text{JSD} = \frac{1}{2}D_{\text{KL}}(p\ m) + \frac{1}{2}D_{\text{KL}}(q\ m)$	Original GAN objective
ELBO	$\log p(\mathbf{x}) = \text{ELBO} + D_{\text{KL}}(q\ p)$	VAEs, variational inference
VAE loss	Reconstruction $-\beta \cdot D_{\text{KL}}(q\ p(\mathbf{z}))$	Generative modeling
Info Bottleneck	$\min I(X; T) - \beta I(T; Y)$	Representation learning

**The common thread:** Information theory gives us the language to reason about

what neural networks learn, what they forget, and how to train them.

DPI says info is lost; IB says lose the right info; ELBO says how to do it in practice.

# Homework

1. **DPI application.** Suppose  $X \rightarrow Y \rightarrow Z$  with  $I(X; Y) = 2$  bits.
  - (a) What is the maximum possible value of  $I(X; Z)$ ?
  - (b) Give an example where  $I(X; Z) = I(X; Y)$  (hint: sufficient statistic).
  - (c) Give an example where  $I(X; Z) = 0$ .
2.  **$f$ -divergences.** Show that the total variation distance  $\text{TV}(p, q) = \frac{1}{2} \sum |p(x) - q(x)|$  is an  $f$ -divergence with  $f(t) = \frac{1}{2}|t - 1|$ . Verify  $f$  is convex and  $f(1) = 0$ .
3. **ELBO derivation.** Starting from  $\log p_\theta(\mathbf{x})$ , derive the ELBO by writing  $\log p_\theta(\mathbf{x}) = \log p_\theta(\mathbf{x}) \cdot \int q_\phi(\mathbf{z}|\mathbf{x}) d\mathbf{z}$  and applying Jensen's inequality. Show that the gap is  $D_{\text{KL}}(q_\phi \| p_\theta(\mathbf{z}|\mathbf{x}))$ .
4. **IB tradeoff.** A 10-class classifier uses 128-dim features from a bottleneck layer.
  - (a) What is the maximum  $I(T; Y)$ ? (Hint:  $H(Y) \leq \log_2 10$ .)
  - (b) If we reduce to 2-dim features, how does the IB tradeoff change?

# Questions?