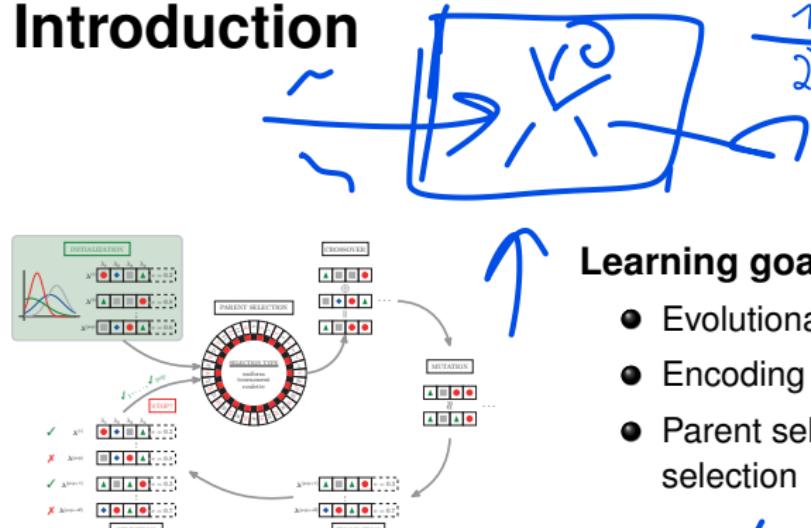


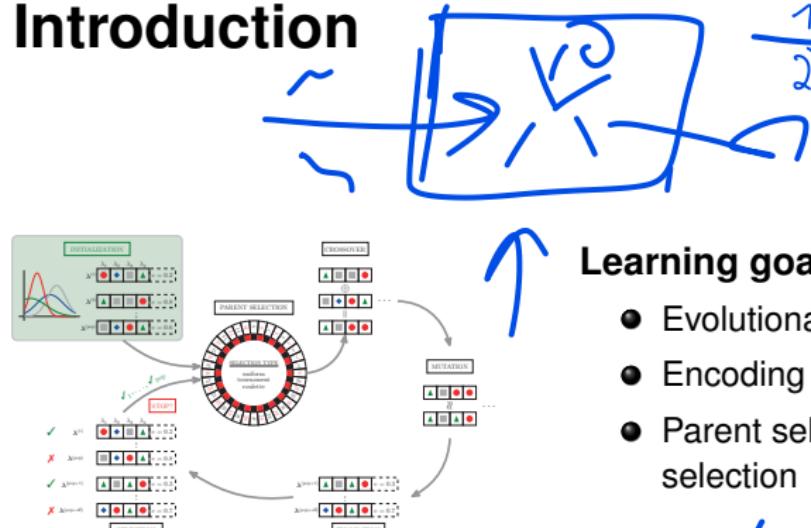
Optimization in Machine Learning

Evolutionary Algorithms Introduction



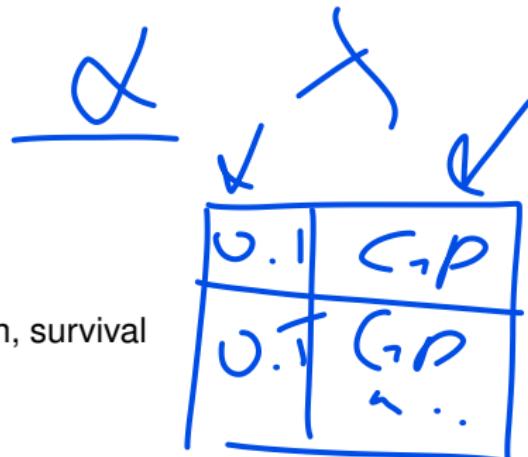
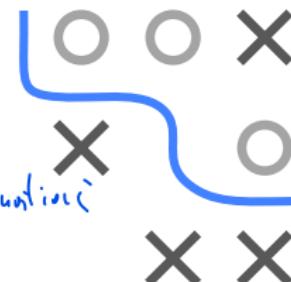
Examples of deriv. free situations

- 1) Pig-nose optimization
- 2) Simulations



Learning goals

- Evolutionary algorithms
- Encoding
- Parent selection, variation, survival selection



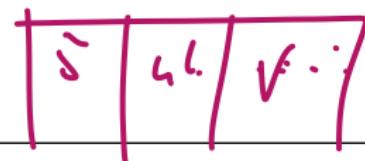
EVOLUTIONARY ALGORITHMS

Evolutionary algorithms (EA) are a class of stochastic, metaheuristic optimization techniques whose mode of operation is inspired by the evolution of natural organisms.

History of evolutionary algorithms:

- **Genetic algorithms**: Use binary problem representation, therefore closest to the biological model of evolution.
- **Evolution strategies**: Use direct problem representation, e.g., vector of real numbers.
- **Genetic programming**: Create structures that convert an input into a fixed output (e.g. computer programs); solution candidates are represented as trees.
- **Evolutionary programming**: Similar to genetic programming, but solution candidates are not represented by trees, but by finite state machines.

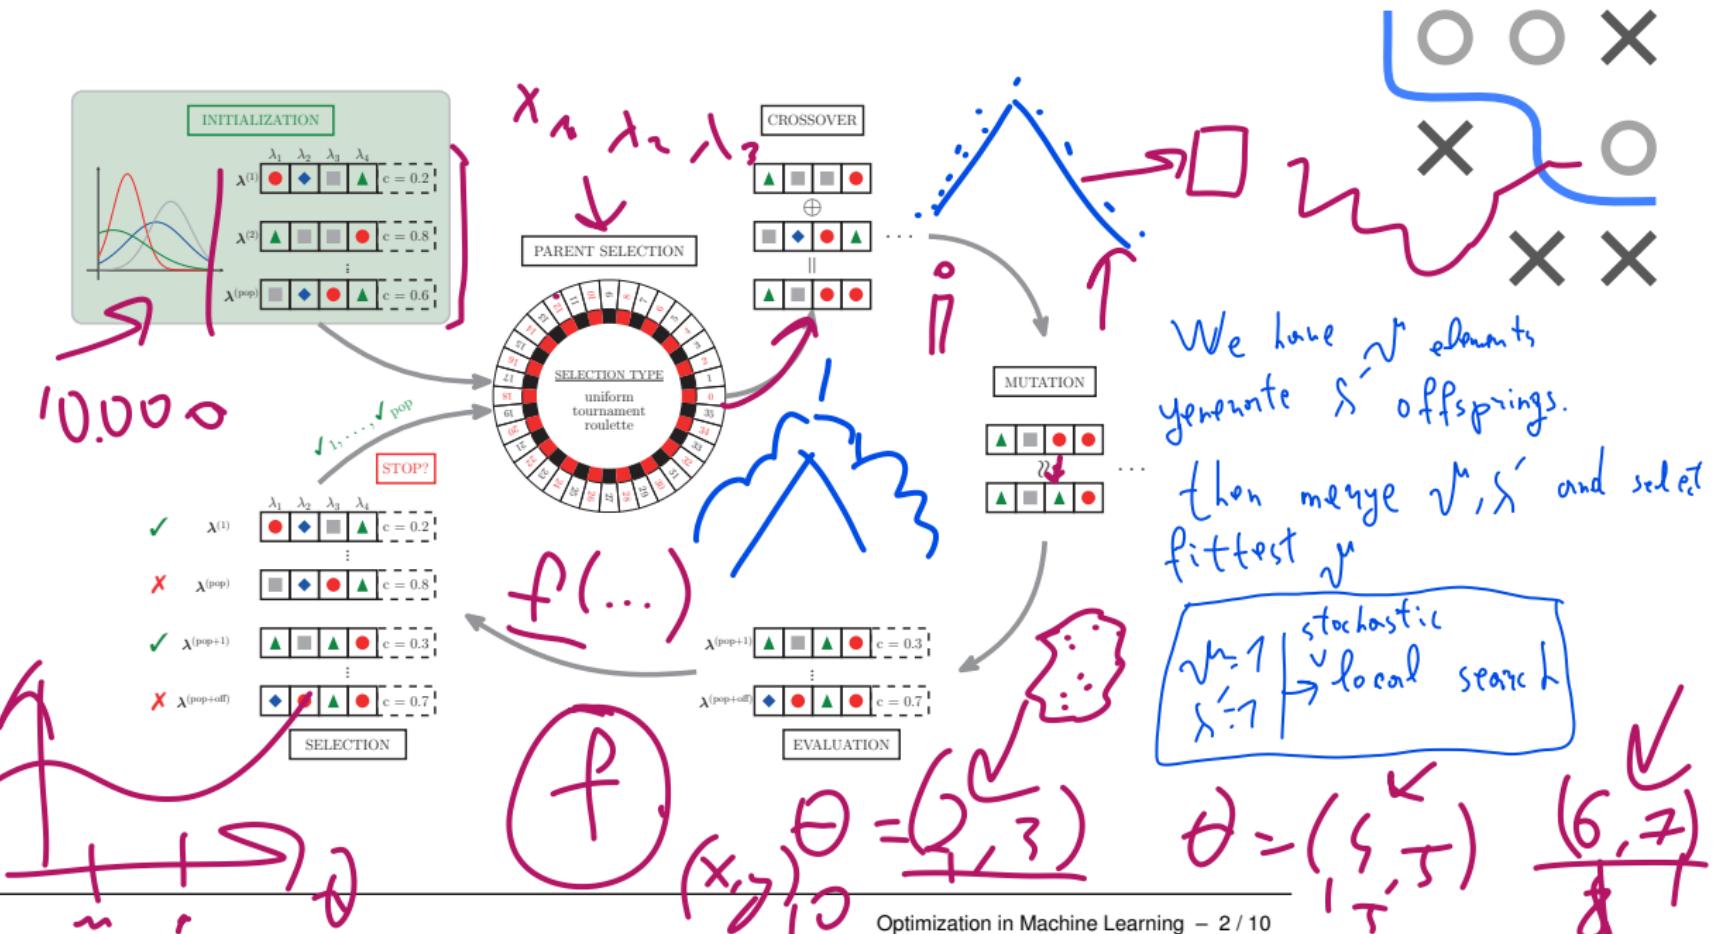
The boundaries between the terms become increasingly blurred and are often used synonymously.



NN < b,

0.1
0.5
0.3
0.3

STRUCTURE OF AN EVOLUTIONARY ALGORITHM



NOTATION AND TERMINOLOGY

(encoding)

- A chromosome is a set of parameters which encodes a proposed solution to the problem that the genetic algorithm is trying to solve. The chromosome is often represented as a binary string, although a wide variety of other data structures are also used.
- The set of all solutions is known as the population.

(2,7)(4,5) (. .)

Symbols	EA Terminology
$x \in S$	solution candidate $x \in S$
x_j	j -th gene of chromosome
P with $\mu = P $	chromosome of an individual population and size
λ	number of generated offsprings
$f : S \rightarrow \mathbb{R}$	fitness function

Note: Unintuitively, we are minimizing fitness because we always minimize f by convention.

ENCODING

Encoding of chromosomes is the first step of solving a problem with EAs. Technically: Mapping from **genotype** to **phenotype**. Encoding depends on the problem, and eventually decides performance of problem solving.



Encoding methods:

- Binary encoding: Strings of 0s and 1s
- Real value encoding: Real values

Genotype:

x_1	x_2
-------	-------

Phenotype:

↓ ↓ ↓ ↓

01101	11001
-------	-------

Binary encoding

13	25
----	----

Real value encoding

When we choose encoding approach we should think how mutation and crossover will look like.

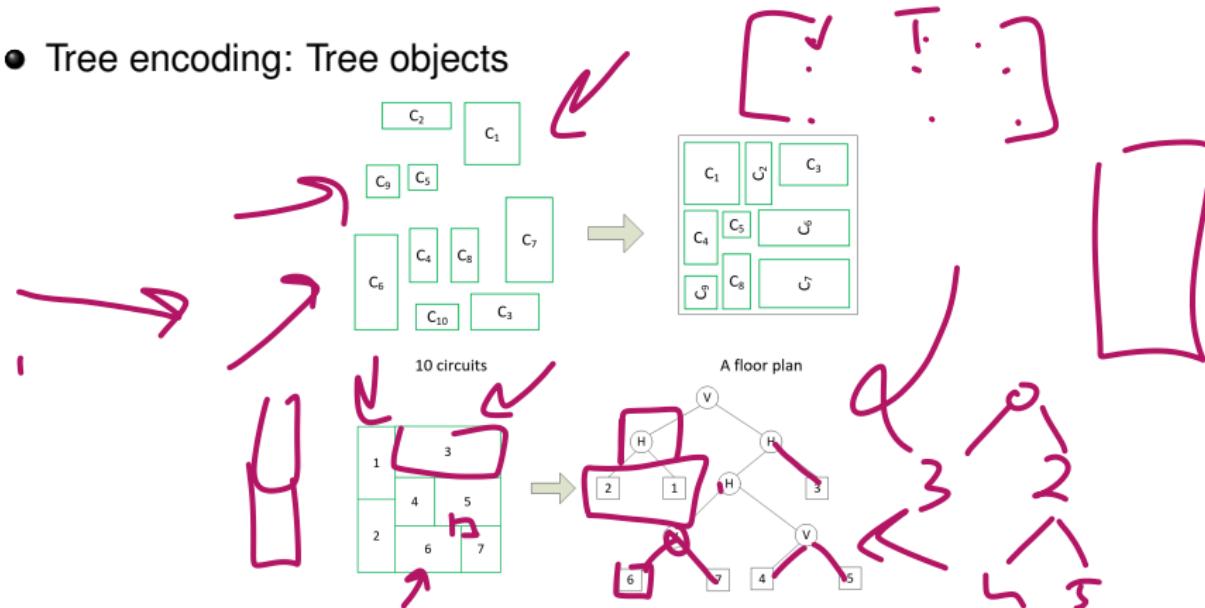
possible

? . .

↗

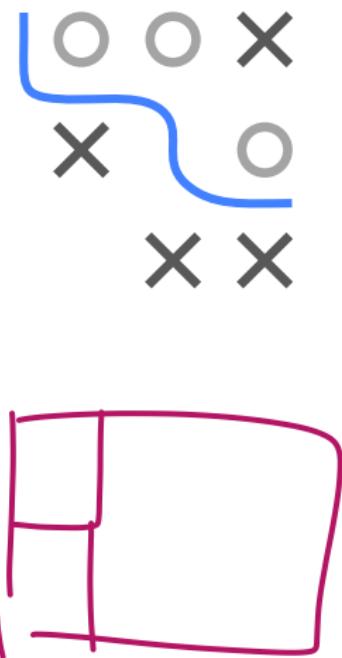
ENCODING

- Tree encoding: Tree objects



Floor planning problem. Given are n circuits of different area requirements. Goal: arrange them into a floor layout so that all circuits are placed in a minimum layout. Each solution candidate can be represented by a tree.

Source: Encoding Techniques in Genetic Algorithms, Debasis Samanta, 2018.



STEP 1: INITIALIZE POPULATION

- Evolutionary algorithms start with generating initial population

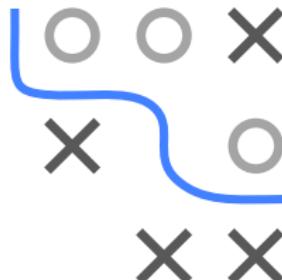
$$P = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\mu)}\}.$$

- Usually: Initialize uniformly at random.
(or use heuristics if available
(domain knowledge))

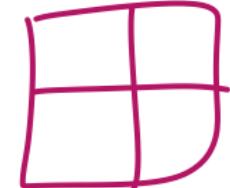
- Introducing prior knowledge possible.

- Population is evaluated: objective function is computed for each initial individual.
(the most expensive step, luckily can be parallelized)

- Initialization influences quality of solution, so many EAs employ *restarts* with new randomly generated initial populations.



$f(x^{(1)})$, $f(x^{(2)})$, ..., $f(x^{(n)})$

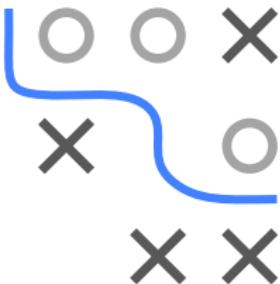


STEP 2: PARENT SELECTION

Choose a number of λ parents pairs creating λ offsprings.



- **Neutral selection:** Draw parents uniformly at random.
- **Fitness-proportional / Roulette wheel selection:** Draw individuals with probability proportional to their fitness.
- **Tournament selection:** Randomly select k individuals for a "tournament group" and pick the best one (according to fitness value).



Left: Fitness-proportional selection. Fitness values of $\mu = 10$ individuals are converted into probabilities. **Right:** Tournament selection.

STEP 3: VARIATION

New individuals (offsprings) are generated from parents.

- Recombination/Crossover: Combine two parents into offspring.
- Mutation: Modify the offspring locally.

Sometimes only one of both operations is performed.

Mutation

x_1	x_2	\dots	x_{d-1}	x_d
1	0	\dots	0	0
1	0	\dots	1	0

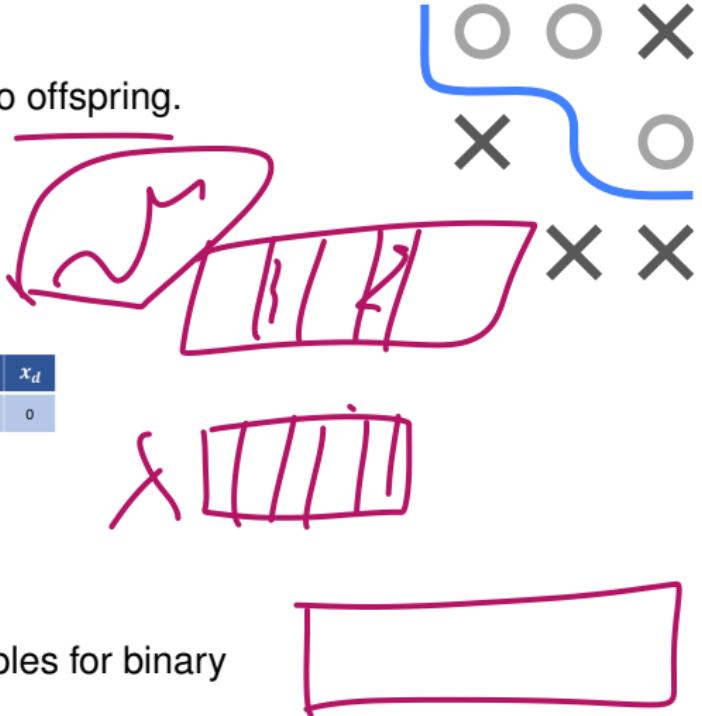
Recombination

x_1	x_2	\dots	x_{d-1}	x_d
0	0	\dots	1	1
1	0	\dots	0	0

x_1	x_2	\dots	x_{d-1}	x_d
1	0	\dots	0	1

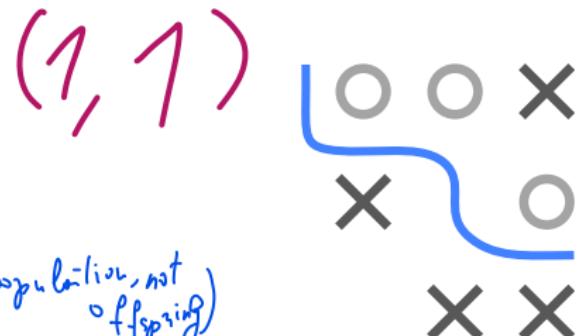


Note: Particular operation depends on encoding. Examples for binary and numeric encodings follow later.



STEP 4: SURVIVAL SELECTION

Choosing surviving individuals. Two common strategies are:



- **(μ, λ) -selection:** Select μ best individuals only from set of offsprings ($\lambda \geq \mu$ necessary).

But: Best individual can get lost! (Because can be from population, not offspring)

- **$(\mu + \lambda)$ -selection:** Select μ best individuals from set of μ parents and λ offsprings

Now: Best individual certainly survives.

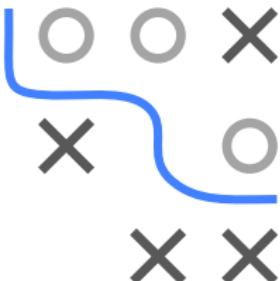
Mostly (maybe always) this is used



EVOLUTIONARY ALGORITHMS

Advantages

- Simple but enough to solve complex problems
- All parameter types possible in general
- Highly parallelizable
- Flexible through different variation operations



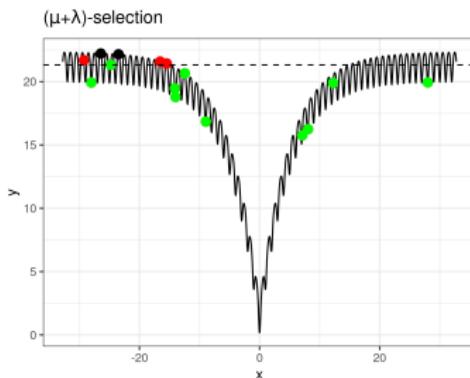
Disadvantages

- Little mathematical rigor (for realistic, complex EAs) *unrigorous*
- Hard to find balance between exploration and exploitation *i.e., no crossover (big steps)* *mutation (short steps)*
- Quite some parameters, hard to determine them
- Customization necessary for complex problems
- Not suitable for expensive problems like HPO as large number of function evaluations necessary

H_y

Optimization in Machine Learning

Evolutionary Algorithms ES / Numerical Encodings



Learning goals

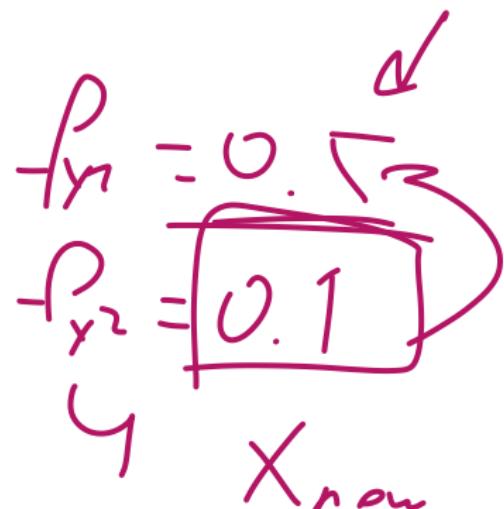
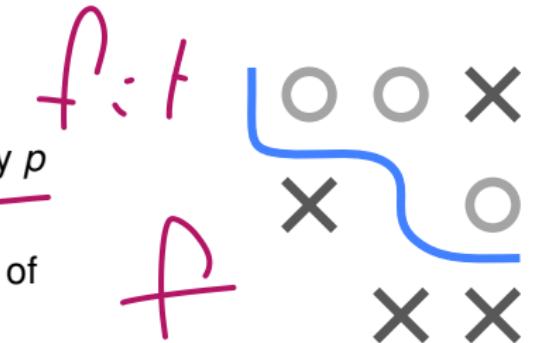
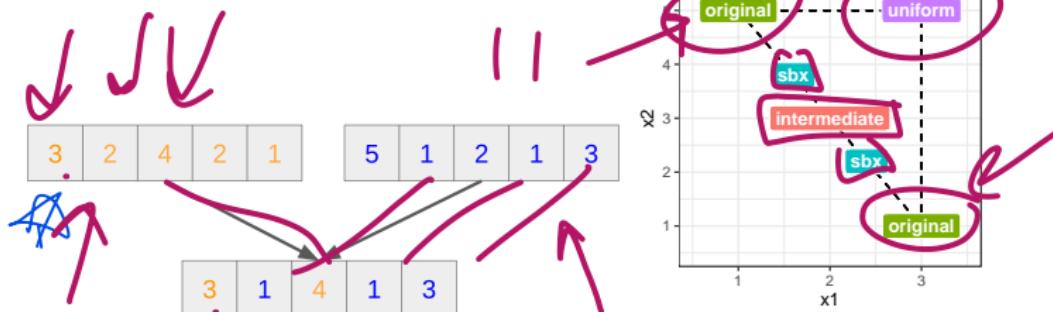
- Recombination
- Mutation
- A few simple examples

RECOMBINATION FOR NUMERIC

Options for recombination of two individuals $\mathbf{x}, \tilde{\mathbf{x}} \in \mathbb{R}^d$:

- **Uniform crossover:** Choose gene j of parent 1 with probability p and of parent 2 with probability $1 - p$
- **Intermediate recombination:** Offspring is created from mean of two parents: $\frac{1}{2}(\mathbf{x} + \tilde{\mathbf{x}})$
- **Simulated Binary Crossover (SBX):** generate **two offspring**
convex combination of 2 parents

$$\bar{\mathbf{x}} \pm \frac{1}{2}\beta(\tilde{\mathbf{x}} - \mathbf{x}), \quad \bar{\mathbf{x}} = \frac{1}{2}(\mathbf{x} + \tilde{\mathbf{x}}), \quad \beta \in [0, 1] \text{ uniformly at random}$$

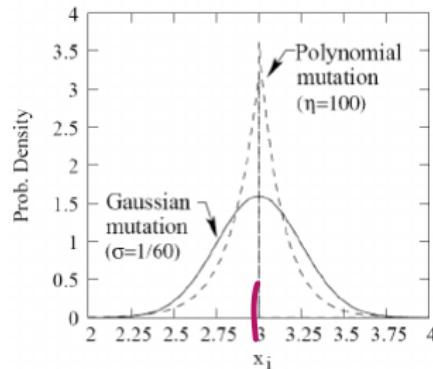


MUTATION FOR NUMERIC

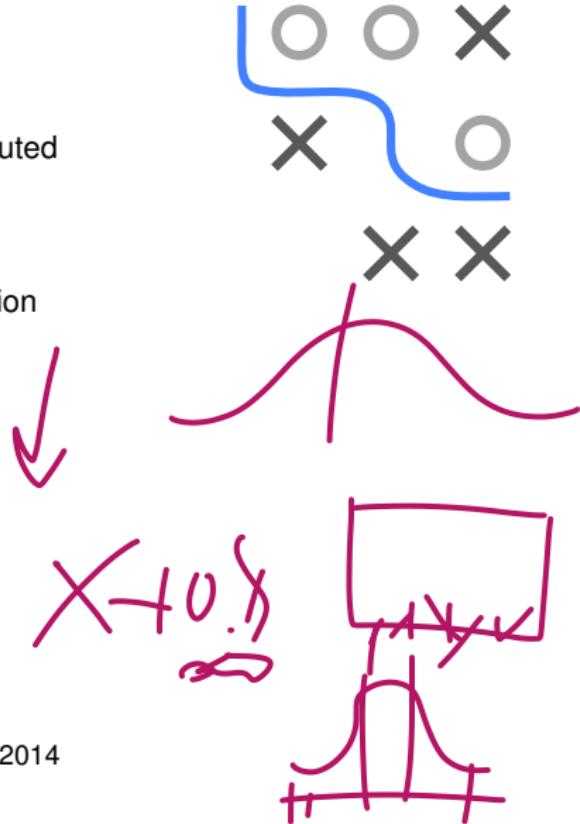
Mutation: Individuals get modified

Example for $\mathbf{x} \in \mathbb{R}^d$:

- **Uniform mutation:** Select random gene x_j and replace it by uniformly distributed value (within feasible range). → may cause a large step
- **Gauss mutation:** $\mathbf{x} \pm \mathcal{N}(0, \sigma \mathbf{I})$
- **Polynomial mutation:** Use a different distribution instead of normal distribution

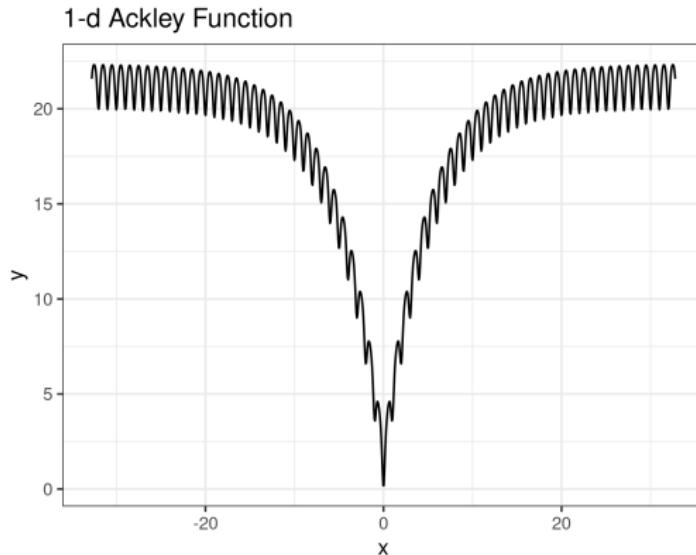


Source: K. Deb, D. Deb. Analysing mutation schemes for real-parameter genetic algorithms, 2014



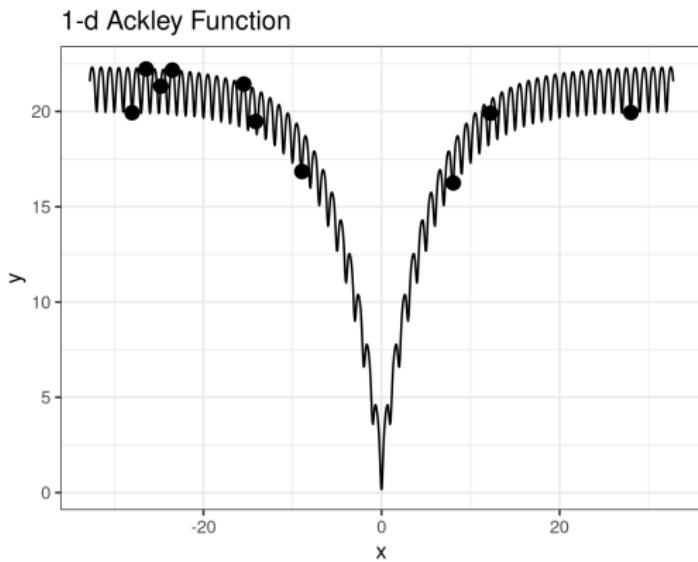
EXAMPLE OF AN EVOLUTIONARY ALGORITHM

(Simple) EA on 1-dim Ackley function on $[-30, 30]$. Usually, for optimizing a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$, individuals are encoded as real vectors $\mathbf{x} \in \mathbb{R}^d$.



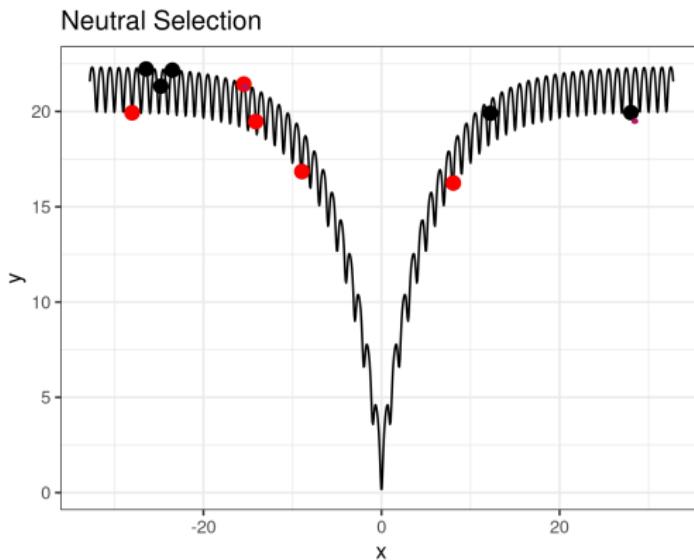
EXAMPLE OF AN EVOLUTIONARY ALGORITHM

Random initial population with size $\mu = 10$



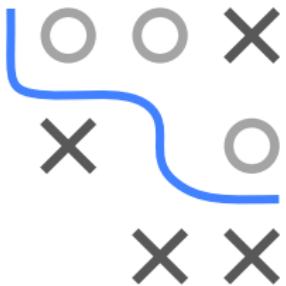
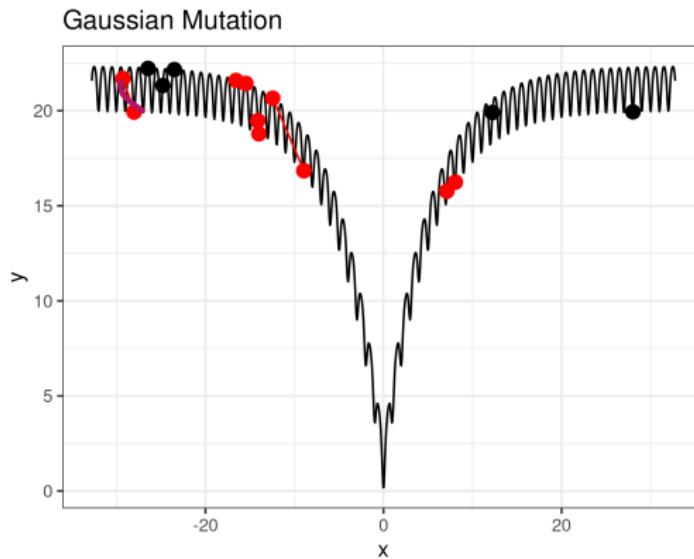
EXAMPLE 1: ACKLEY FUNCTION

We choose $\lambda = 5$ offsprings by neutral selection (red individuals).



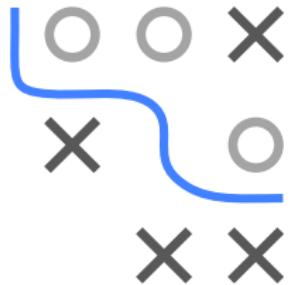
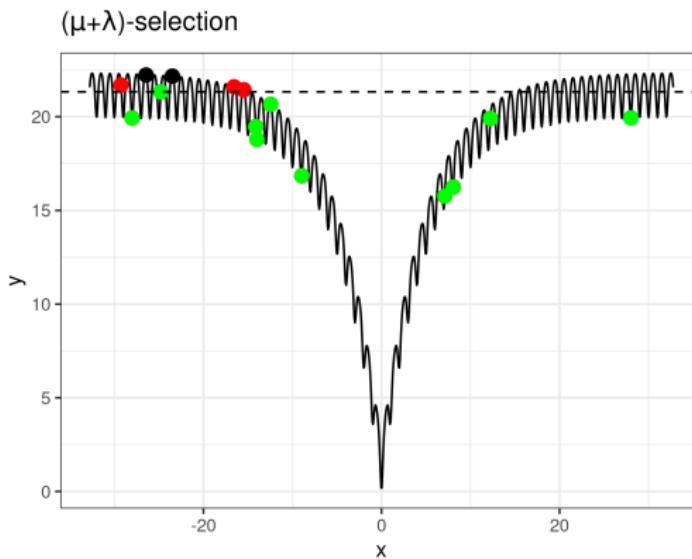
EXAMPLE 1: ACKLEY FUNCTION

Use Gaussian mutation with $\sigma = 2$, but without recombination.



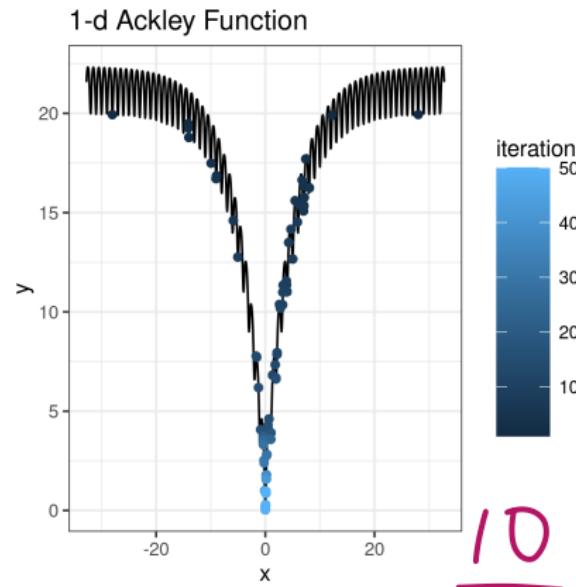
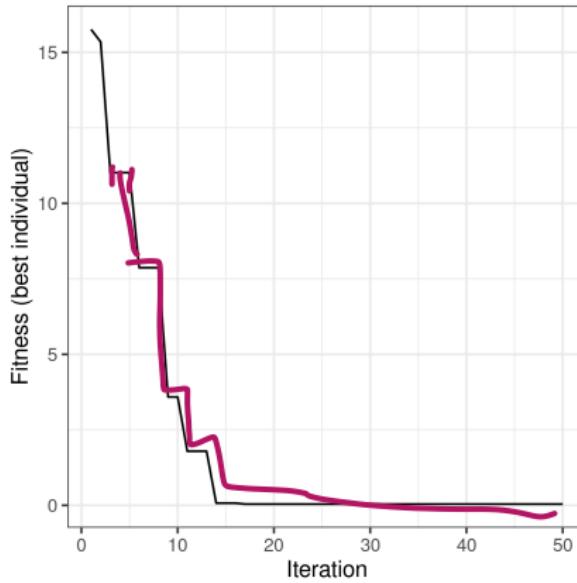
EXAMPLE 1: ACKLEY FUNCTION

Use $(\mu + \lambda)$ selection. Selected individuals are marked in green.



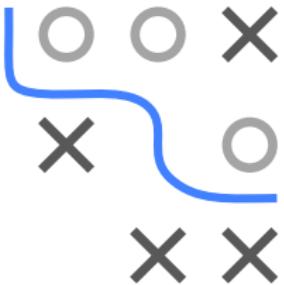
EXAMPLE 1: ACKLEY FUNCTION

After 50 iterations:



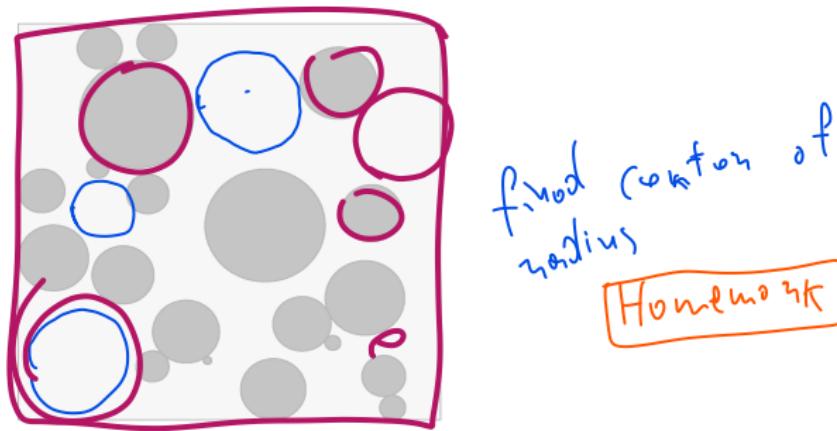
iteration
(generation
producing new
population)

10



EXAMPLE 2: GRID OF BALLS

Consider a grid in which n balls with random radius are placed.



Aim: Find the circle with the largest possible radius in the grid that does **not** intersect with the other existing circles.

- What is the fitness function?
- How is the population defined?

Implementation: <https://juliamb.github.io/balls/>



EXAMPLE 2: GRID OF BALLS

In our example, the chromosome of an individual is the center of a circle, so the chromosomes are encoded as 2-dimensional real vectors $\mathbf{x} = (x_1, x_2) \in \mathbb{R}^2$.

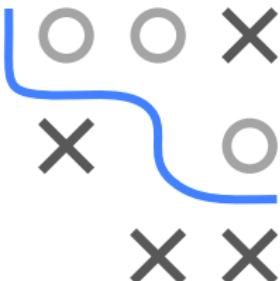
The population $P \subset \mathbb{R}^2$ is given as a set of circle centers.

The fitness function evaluates an individual $\mathbf{x} \in P$ based on the distance to the nearest neighboring gray circle k .

$$f(\mathbf{x}) = \min_{k \in \text{Grid}} \text{distance}(k, \mathbf{x}),$$

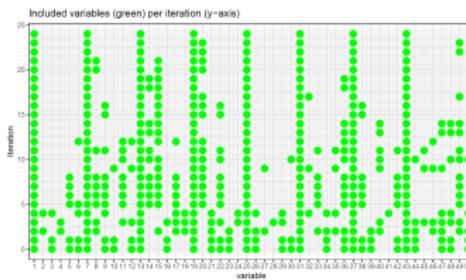
where the distance is defined as 0 if a circle center is within the radius of a circle of the grid.

This function is to be maximized: we are looking for the largest circle that does not touch any of the gray circles.



Optimization in Machine Learning

Evolutionary Algorithms GA / Bit Strings



Learning goals

- Recombination
- Mutation
- Simple examples



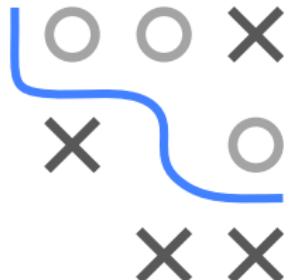
BINARY ENCODING

- In theory: Each problem can be encoded binary
- In practice: Binary not always best representation (e.g., if values are numeric, trees or programs)

We typically encode problems with **binary decision variables** in binary representation.

Examples:

- Scheduling problems
- Integer / binary linear programming
- Feature selection
- ...

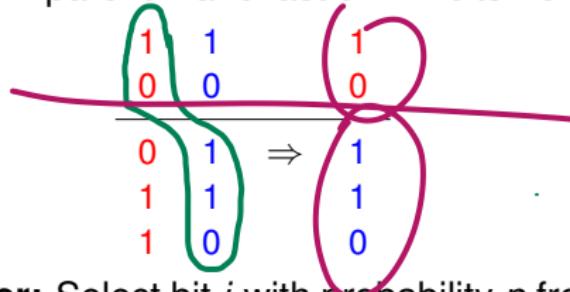


RECOMBINATION FOR BIT STRINGS

Two individuals $\mathbf{x}, \tilde{\mathbf{x}} \in \{0, 1\}^d$ encoded as bit strings can be recombined as follows:

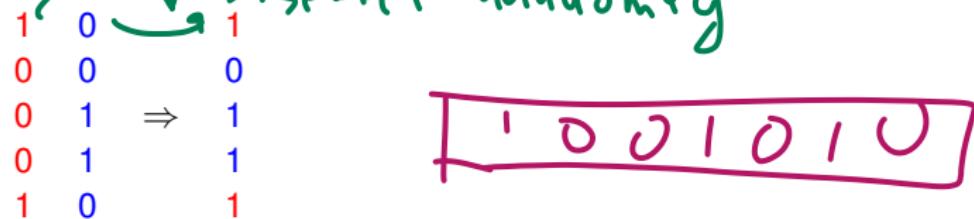
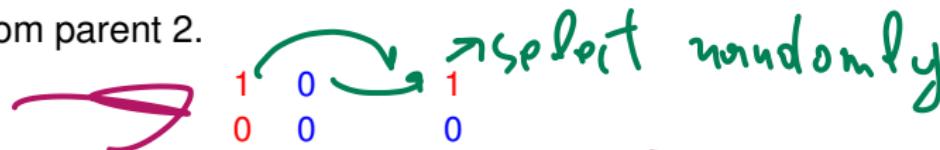
- **1-point crossover:** Select crossover $k \in \{1, \dots, d - 1\}$ randomly.

Take first k bits from parent 1 and last $d - k$ bits from parent 2.



for example
For feature selection
we don't care about
order so 1-point
crossover makes
sense

- **Uniform crossover:** Select bit j with probability p from parent 1 and $1 - p$ from parent 2.



MUTATION FOR BIT STRINGS

Offspring $\mathbf{x} \in \{0, 1\}^d$ encoded as a bit string can be mutated as follows:

- **Bitflip:** Each bit j is flipped with probability $p \in (0, 1)$.



1	0
0	0
0	⇒ 0
0	1
1	1

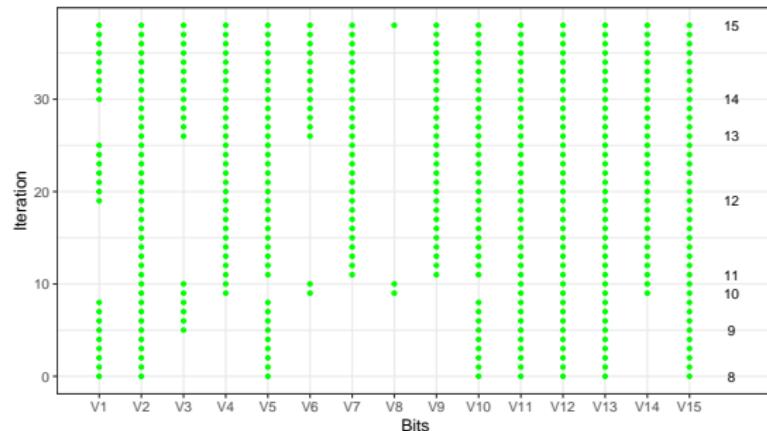
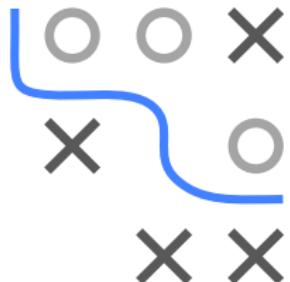
EXAMPLE 1: ONE-MAX EXAMPLE

$\mathbf{x} \in \{0, 1\}^d$, $d = 15$ bit vector representation.

Used as sanity check

Goal: Find the vector with the maximum number of 1's.

- Fitness: $f(\mathbf{x}) = \sum_{i=1}^d x_i$
- $\mu = 15$, $\lambda = 5$, $(\mu + \lambda)$ -strategy, bitflip mutation, no recombination



Green: Representation of best individual per iteration. Right scale shows fitness.

EXAMPLE 2: FEATURE SELECTION

We consider the following toy setting:

- Generate design matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ by drawing $n = 1000$ samples of $p = 50$ independent normally distributed features with $\mu_j = 0$ and $\sigma_j^2 > 0$ varying between 1 and 5 for $j = 1, \dots, p$.
- Linear regression problem with dependent variable \mathbf{y} :

$$\mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \epsilon$$

with $\epsilon \sim \mathcal{N}(0, 1)$.

Parameter $\boldsymbol{\theta}$:

$$\theta_0 = -1.2$$

$$\theta_j = \begin{cases} 1 & \text{for } j \in \{1, 7, 13, 19, 25, 31, 37, 43\} \\ 0 & \text{otherwise} \end{cases}$$

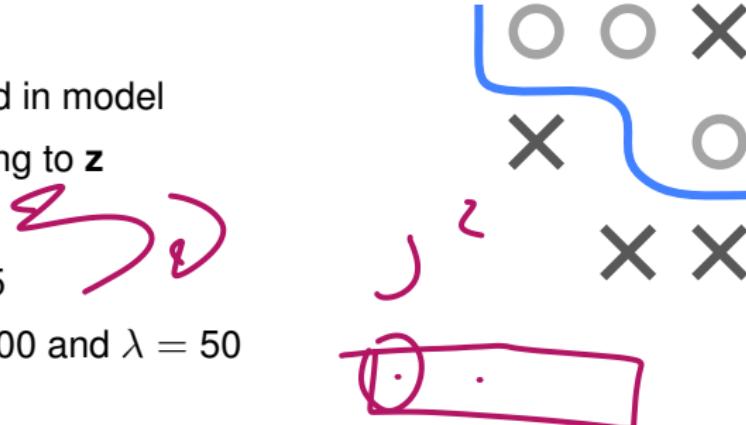


⇒ Only 8 out of 50 equally influential features



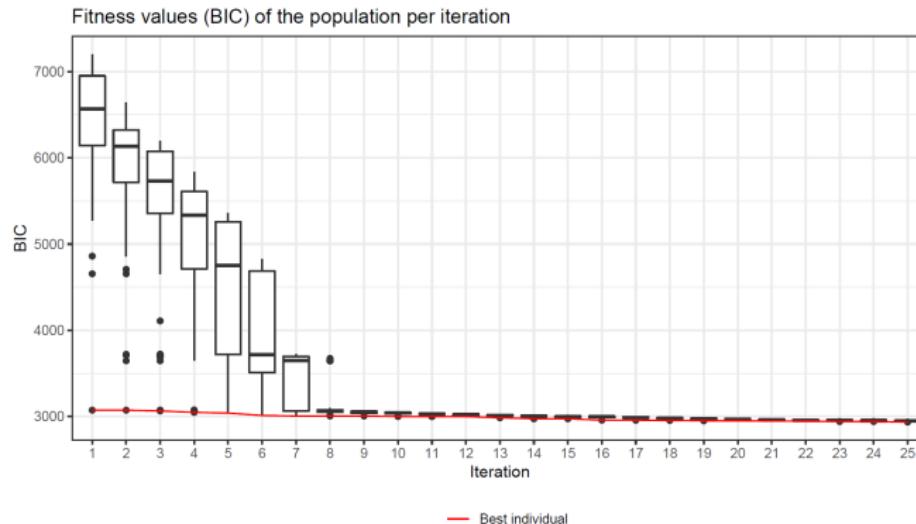
EXAMPLE 2: FEATURE SELECTION

- **Aim:** Find influential features
- **Encoding:** $\mathbf{z} \in \{0, 1\}^P$, $z_j = 1$ means θ_j included in model
- **Fitness function $f(\mathbf{z})$:** BIC of the model belonging to \mathbf{z}
- **Mutation:** Bit flip with $p = 0.3$
- **Recombination:** Uniform crossover with $p = 0.5$
- **Survival selection:** $(\mu + \lambda)$ strategy with $\mu = 100$ and $\lambda = 50$

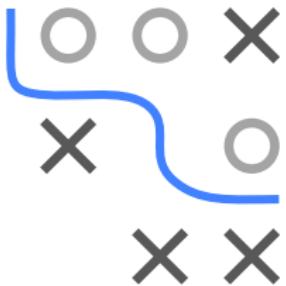
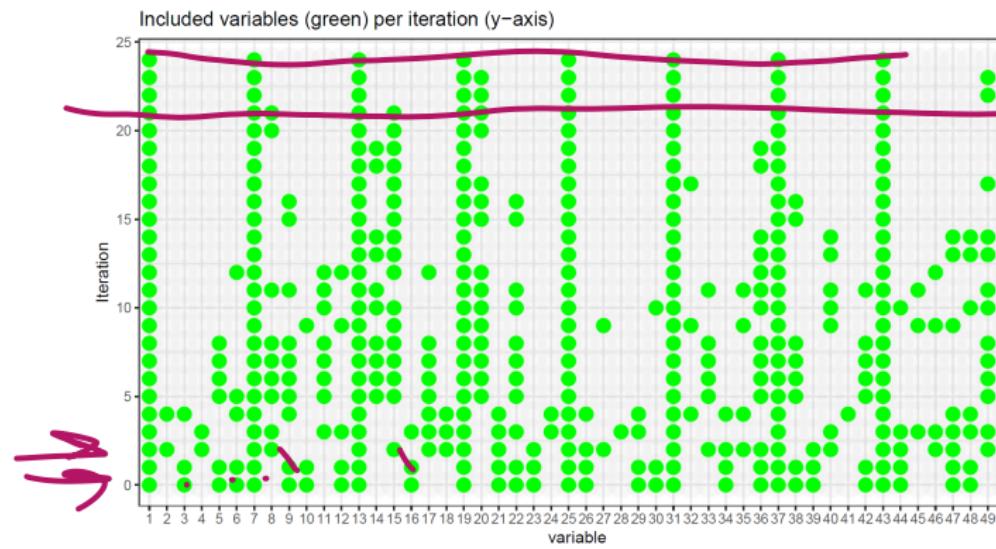


```
## [1] "After 10 iterations:"  
## [1] 1 7 11 13 14 15 19 20 22 25 30 31 36 37 40 43 44 48  
## [19] 49 50  
## [1] "After 20 iterations:"  
## [1] 1 7 8 13 15 19 20 25 31 37 43  
## [1] "Included variables after 24 iterations:"  
## [1] 1 7 13 19 25 31 37 43
```

EXAMPLE 2: FEATURE SELECTION



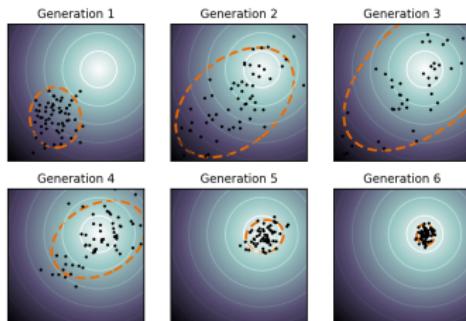
EXAMPLE 2: FEATURE SELECTION



Optimization in Machine Learning

Evolutionary Algorithms CMA-ES Algorithm

~~14.01.29~~

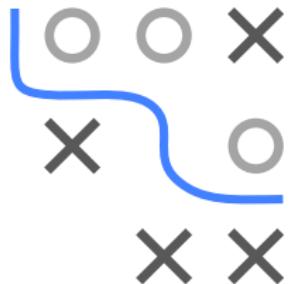


Learning goals

- CMA-ES strategy
- Estimation of distribution
- Step size control

from 2001, still SOTA

Author is Phd in Medicine

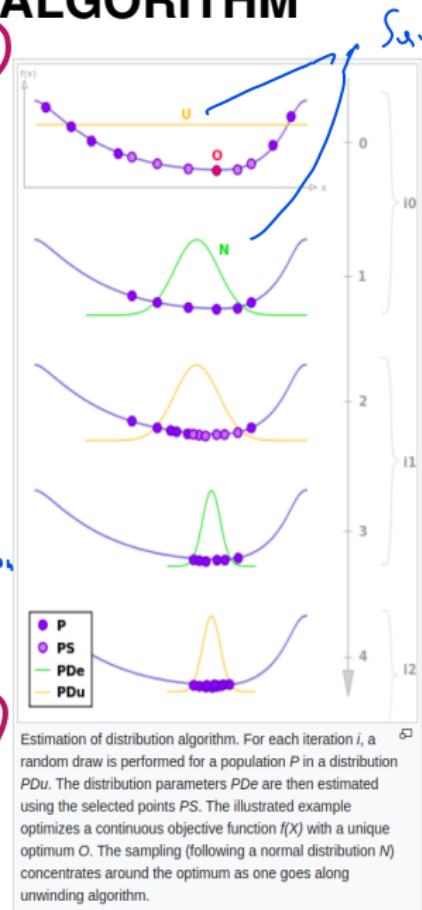


ESTIMATION OF DISTRIBUTION ALGORITHM

- Instead of population, maintain distribution to sample offspring from

- Draw λ offsprings $\mathbf{x}^{(i)}$ from $p(\cdot | \theta^{[t]})$
- Evaluate fitness $f(\mathbf{x}^{(i)})$ (we're doing maximization here)
- Update $\theta^{[t+1]}$ with μ best offsprings

$X_1 = (3, 5)$
 $X_2 = (5, 6)$



Sampling type (uniform...)

Do we switch between Sampling distributions? $U-N \rightarrow U$

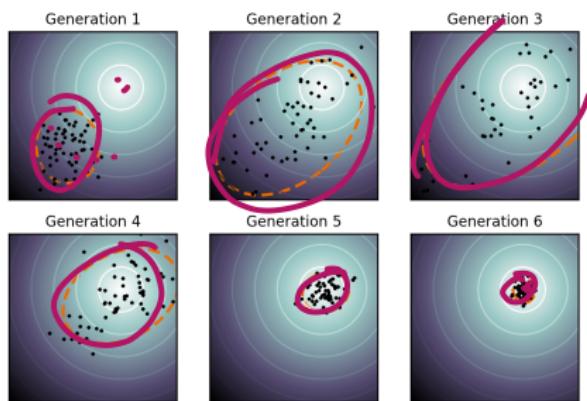
$X_1 = (3, 5)$
 $X_2 = (5, 6)$

COVARIANCE MATRIX ADAPTATION

Sample distribution is multivariate Gaussian

$$\mathbf{x}^{[t+1](i)} \sim \mathbf{m}^{[t]} + \sigma^{[t]} \mathcal{N}(\mathbf{0}, \mathbf{C}^{[t]}) \quad \text{for } i = 1, \dots, \lambda$$

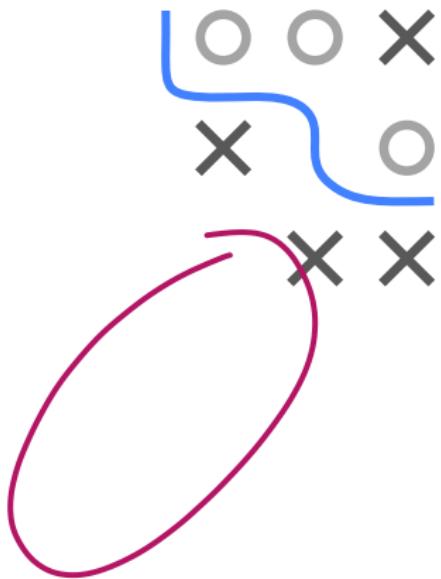
- $\mathbf{x}^{[t+1](i)} \in \mathbb{R}^d$ i -th offspring; $\lambda \geq 2$ number of offspring
- $\mathbf{m}^{[t]}$ $\in \mathbb{R}^d$ mean value and $\mathbf{C}^{[t]} \in \mathbb{R}^{d \times d}$ covariance matrix
- $\sigma^{[t]} \in \mathbb{R}_+$ “overall” standard deviation/step size $\xrightarrow{\text{“mean + stepsize direction”}}$



Question: How to adapt $\mathbf{m}^{[t+1]}, \mathbf{C}^{[t+1]}, \sigma^{[t+1]}$ for next generation $t + 1$?

CMA-ES: BASIC METHOD - ITERATION 1

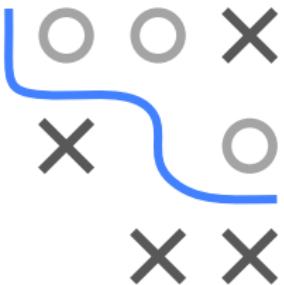
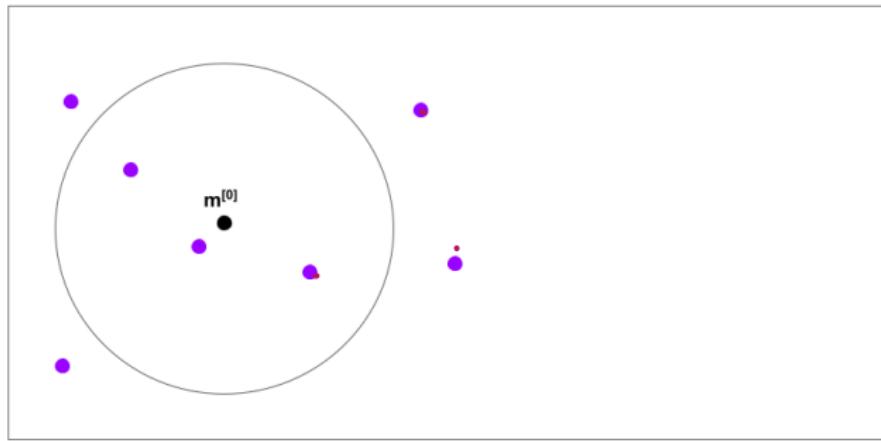
- ➊ Initialize $\mathbf{m}^{[0]}, \sigma^{[0]}$ problem-dependent and $\mathbf{C}^{[0]} = \mathbf{I}_d$



CMA-ES: BASIC METHOD - ITERATION 1

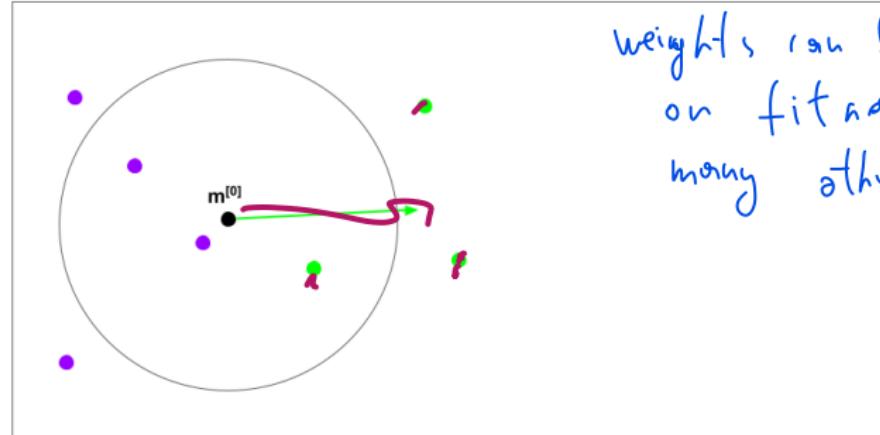
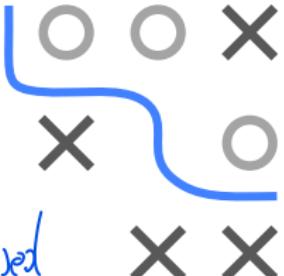
- ➊ Sample λ offsprings from distribution

$$\mathbf{x}^{[1](i)} = \mathbf{m}^{[0]} + \sigma^{[0]} \mathcal{N}(\mathbf{0}, \mathbf{C}^{[0]})$$



CMA-ES: BASIC METHOD - ITERATION 1

- ② Selection and recombination of $\mu < \lambda$ best-performing offspring using fixed weights $w_1 \geq \dots \geq w_\mu > 0$, $\sum_{i=1}^\mu w_i = 1$.
 $\mathbf{x}_{i:\lambda}$ is i -th ranked solution, ranked by $f(\mathbf{x}_{i:\lambda})$.



weights can be adjusted based
on fitness values, but there are
many other approaches

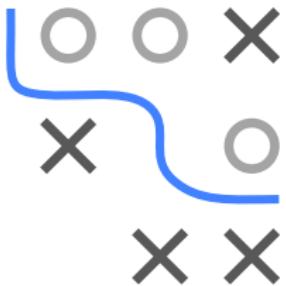
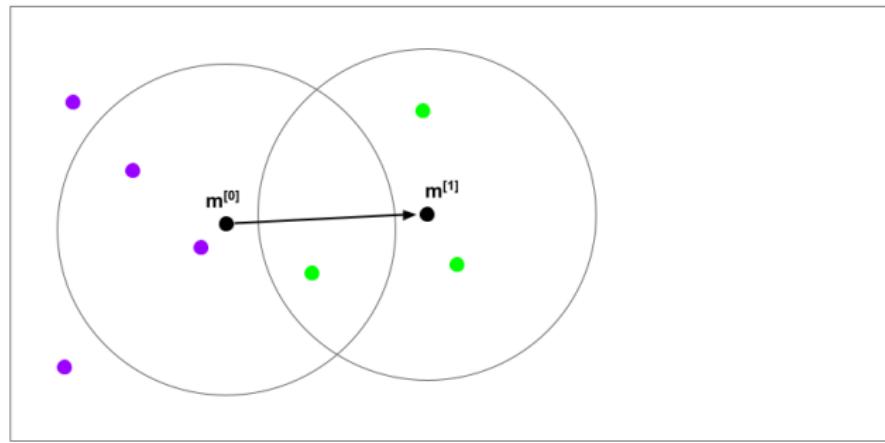
Calculation of auxiliary variables ($\mu = 3$ points)

$$\mathbf{y}_w^{[1]} := \sum_{i=1}^\mu w_i (\mathbf{x}_{i:\lambda}^{[1]} - \mathbf{m}^{[0]}) / \sigma^{[0]} := \sum_{i=1}^\mu w_i \mathbf{y}_{i:\lambda}^{[1]}$$

normalizing

CMA-ES: BASIC METHOD - ITERATION 1

③ Update mean



Movement towards the new distribution with mean

$$\mathbf{m}^{[1]} = \mathbf{m}^{[0]} + \sigma^{[0]} \mathbf{y}_w^{[1]}$$

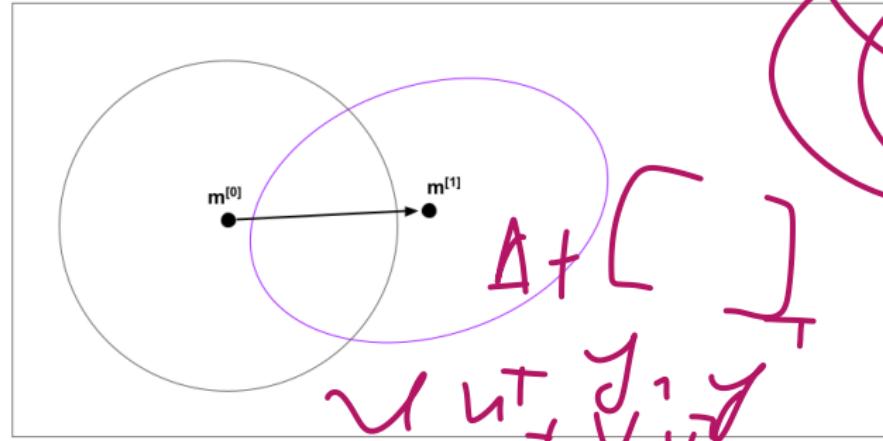
→ 6 controls out

CMA-ES: BASIC METHOD - ITERATION 1

④ Update covariance matrix

Roughly: elongate density ellipsoid in direction of successful steps.

$\mathbf{C}^{[1]}$ reproduces successful points with higher probability than $\mathbf{C}^{[0]}$.



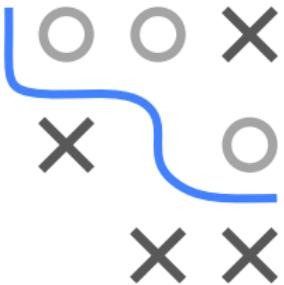
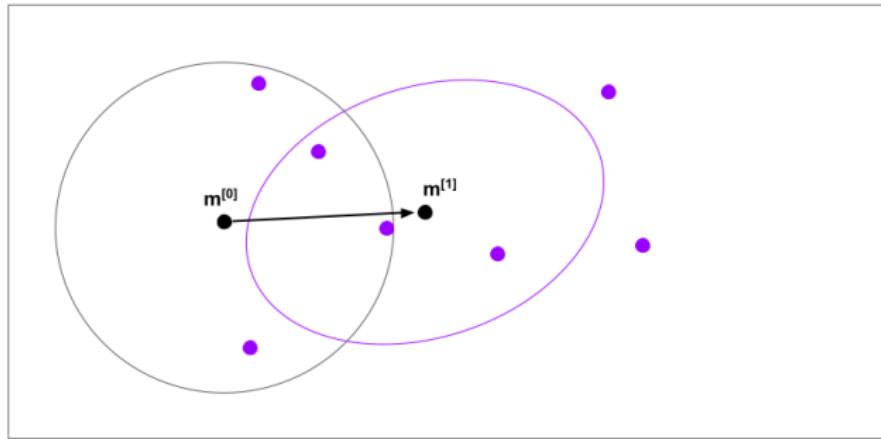
Update $\mathbf{C}^{[0]}$ using sum of outer products and parameter c_μ :

$$\mathbf{C}^{[1]} = (1 - c_\mu) \mathbf{C}^{[0]} + c_\mu \sum_{i=1}^{\mu} w_i \mathbf{y}_{i:\lambda}^{[1]} (\mathbf{y}_{i:\lambda}^{[1]})^T \quad (\text{rank-}\mu \text{ update}).$$

*if points are too close
we will shrink the
cov. matrix and get stuck.
We usually do 20 starts of the
algorithm
if we get stuck we restart
with bigger pp.*

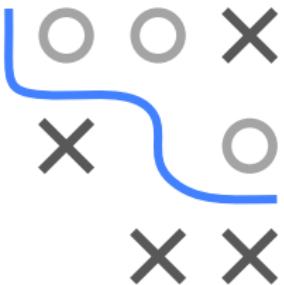
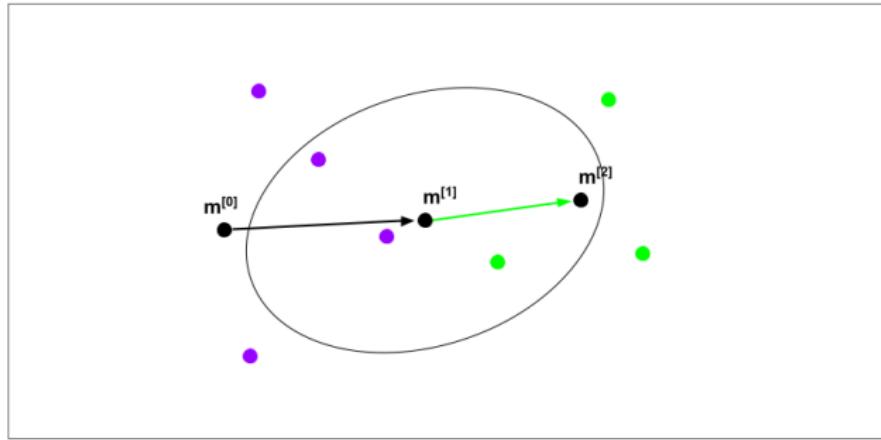
CMA-ES: BASIC METHOD - ITERATION 2

- ➊ Sample from distribution for new generation



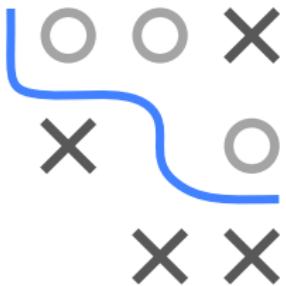
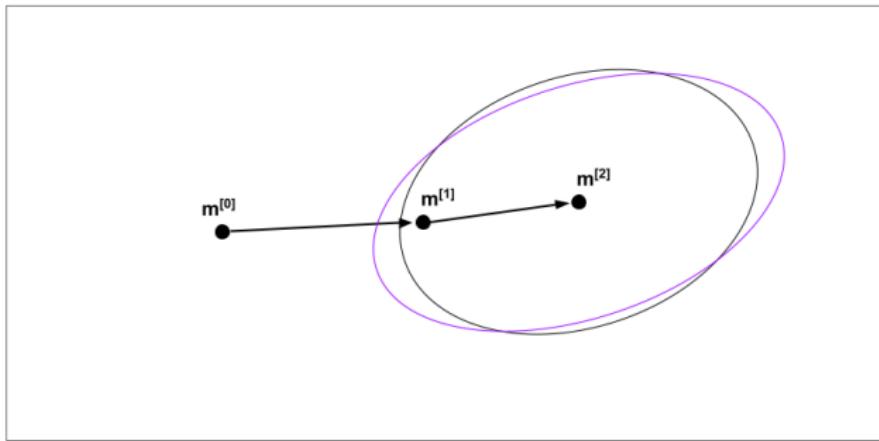
CMA-ES: BASIC METHOD - ITERATION 2

- ② Selection and recombination of $\mu < \lambda$ best-performing offspring
- ③ Update mean



CMA-ES: BASIC METHOD - ITERATION 2

④ Update covariance matrix

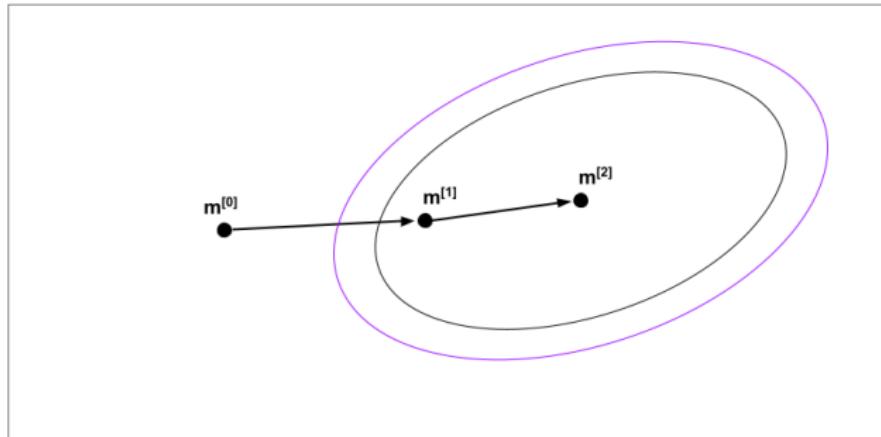


CMA-ES: BASIC METHOD - ITERATION 2

- ⑤ **Update step-size** exploiting correlation in history of steps.

steps point in similar direction \implies increase step-size

steps cancel out \implies decrease step-size



kind of like
momentum



UPDATING C: FULL UPDATE

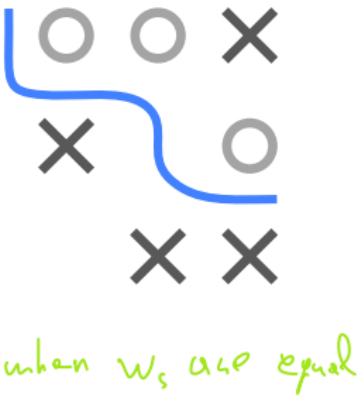
Full CMA update of \mathbf{C} combines rank- μ update with a rank-1 update using exponentially smoothed evolution path $\mathbf{p}_c \in \mathbb{R}^d$ of successive steps and learning rate c_1 :

$$\mathbf{p}_c^{[0]} = \mathbf{0}, \quad \mathbf{p}_c^{[t+1]} = (1 - c_1)\mathbf{p}_c^{[t]} + \sqrt{\frac{c_1(2 - c_1)}{\sum_{i=1}^{\mu} w_i^2}} \mathbf{y}_w \quad \text{fitness?}$$

Final update of \mathbf{C} is

$$\mathbf{C}^{[t+1]} = (1 - c_1 - c_\mu \sum_j w_j) \mathbf{C}^{[t]} + c_1 \underbrace{\mathbf{p}_c^{[t+1]} (\mathbf{p}_c^{[t+1]})^\top}_{\text{rank-1}} + c_\mu \underbrace{\sum_{i=1}^{\mu} w_i \mathbf{y}_{i:\lambda}^{[t+1]} (\mathbf{y}_{i:\lambda}^{[t+1]})^\top}_{\text{rank-}\mu}$$

- Correlation between generations used in rank-1 update
- Information from entire population is used in rank- μ update

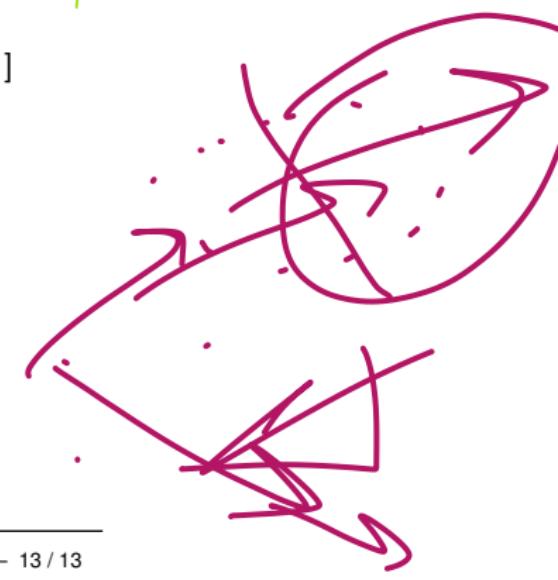
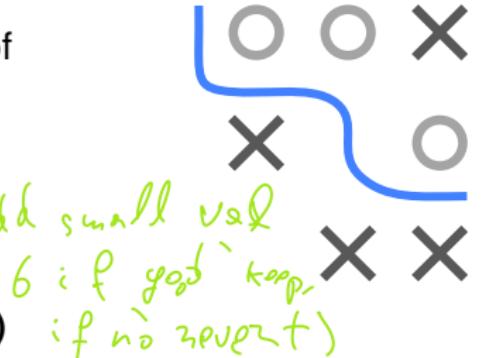
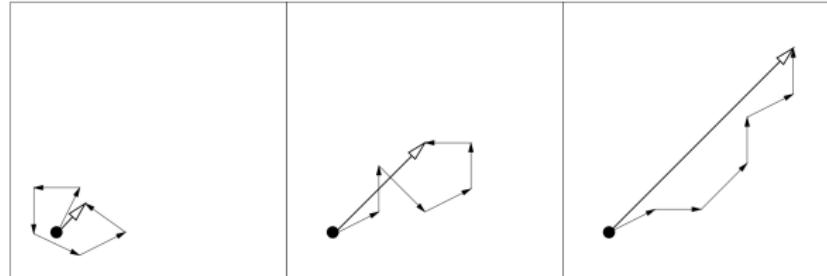


UPDATING σ : METHODS STEP-SIZE CONTROL

- **1/5-th success rule:** increases the step-size if more than 20 % of the new solutions are successful, decrease otherwise
- **σ -self-adaptation:** mutation is applied to the step-size and the better - according to the objective function value - is selected
- **Path length control via cumulative step-size adaptation (CSA)**

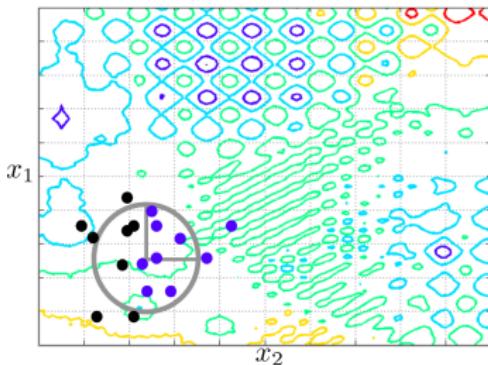
Intuition:

- Short cumulative step-size \triangleq steps cancel \rightarrow decrease $\sigma^{[t+1]}$
- Long cumulative step-size \triangleq corr. steps \rightarrow increase $\sigma^{[t+1]}$



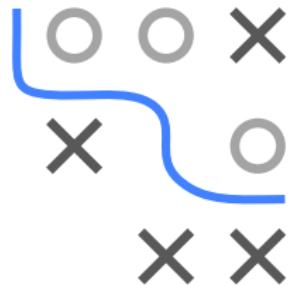
Optimization in Machine Learning

Evolutionary Algorithms CMA-ES Wrap Up



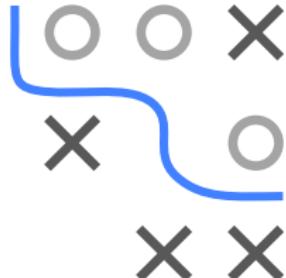
Learning goals

- Advantages & Limitations
- IPOPT-CMA-ES
- Benchmark



CMA-ES: WRAP UP

Algorithm CMA-ES



CMA-ES: WRAP UP - DEFAULT VALUES

Related to selection and recombination:

- λ : offspring number, population size $4 + \lfloor 3 \ln d \rfloor$
- μ : parent number, solutions involved in mean update $\lfloor \lambda/2 \rfloor$
- w_i : recombination weights (preliminary convex shape) $\ln \frac{\lambda+1}{2} - \ln i$, for $i = 1, \dots, \lambda$

Related to \mathbf{C} -update:

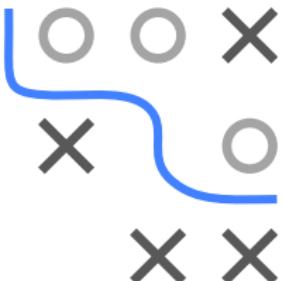
- $1 - c_C$: decay rate for evolution path, cumulation factor $1 - \frac{4+\mu_w/d}{d+4+2\mu_w/d}$
- c_1 : learning rate for rank-one update of \mathbf{C} $\frac{2}{(d+1.3)^2 + \mu_w}$
- c_μ : learning rate for rank- μ update of \mathbf{C} $\min\left(1 - c_1, 2 \cdot \frac{\mu_w - 2 + 1/\mu_w}{(d+2)^2 + \mu_w}\right)$

Related to σ -update:

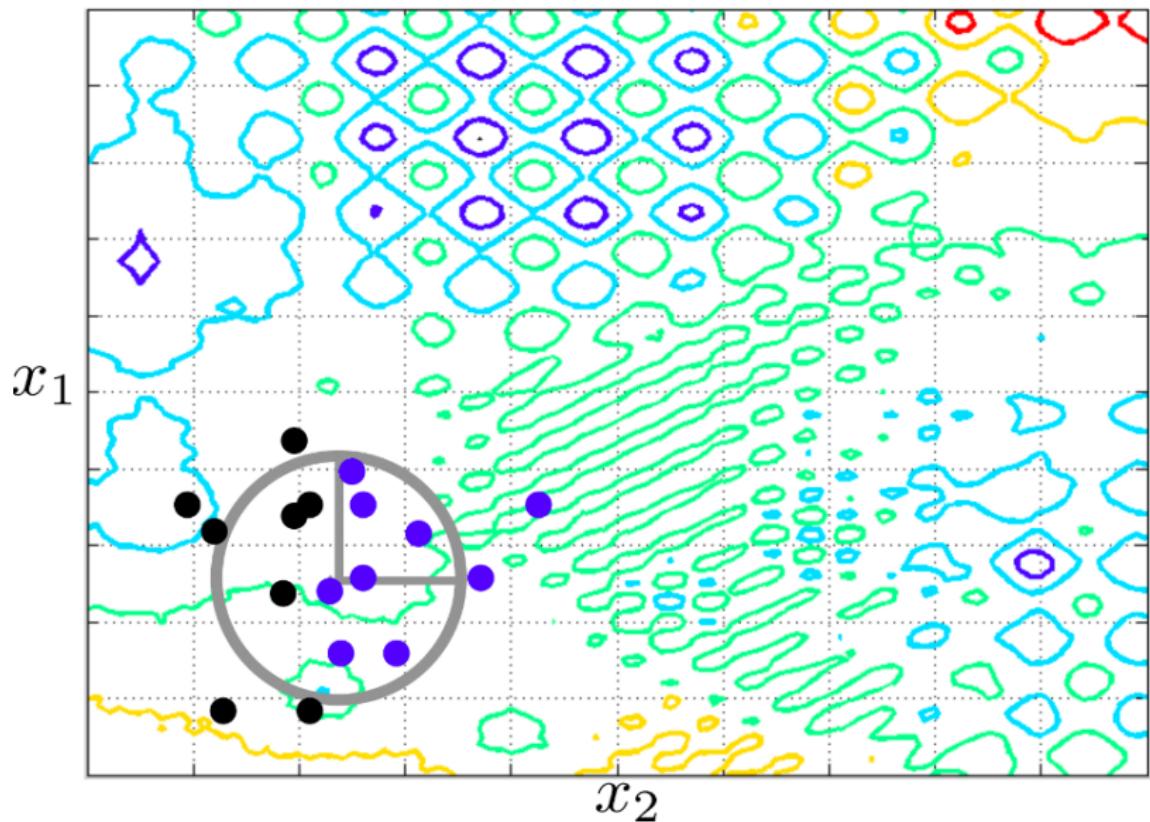
- $1 - c_\sigma$: decay rate for evolution path $1 - \frac{\mu_w + 2}{d + \mu_w + 5}$
- d_σ : damping for σ -change $1 + 2 \max\left(0, \sqrt{\frac{\mu_w - 1}{d + 1}} - 1\right) + c_\sigma$

with $\mu_w = \left(\frac{\|\mathbf{w}\|_1}{\|\mathbf{w}\|_2}\right) = \frac{(\sum_{i=1}^\mu |w_i|)^2}{\sum_{i=1}^\mu w_i^2} = \frac{1}{\sum_{i=1}^\mu w_i^2}$ and typical default parameter values.

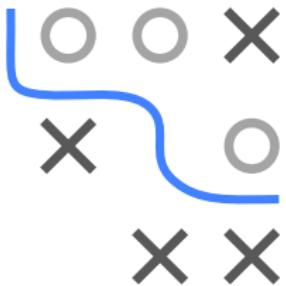
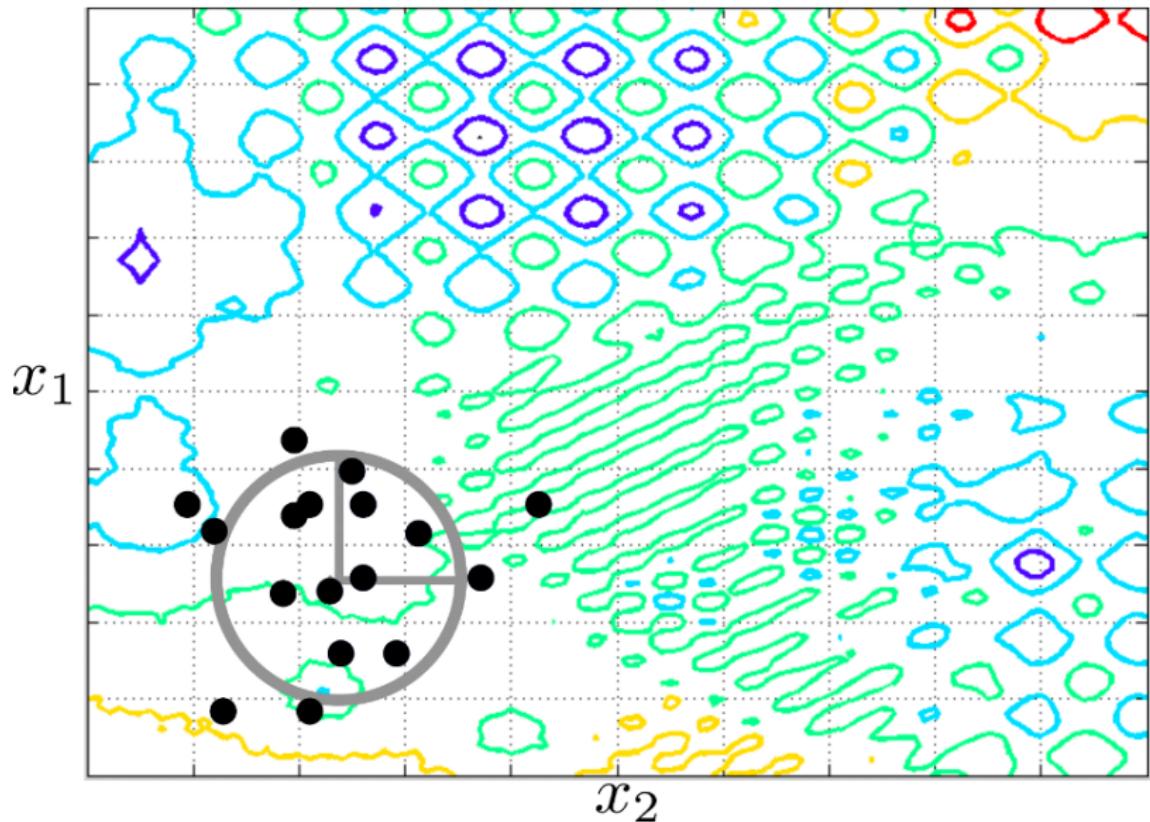
- μ_w can be extended to λ instead of μ weights, allowing negative weights for the remaining $\lambda - \mu$ points ("active covariance adaption").



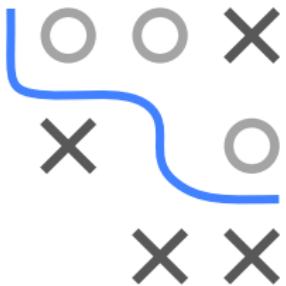
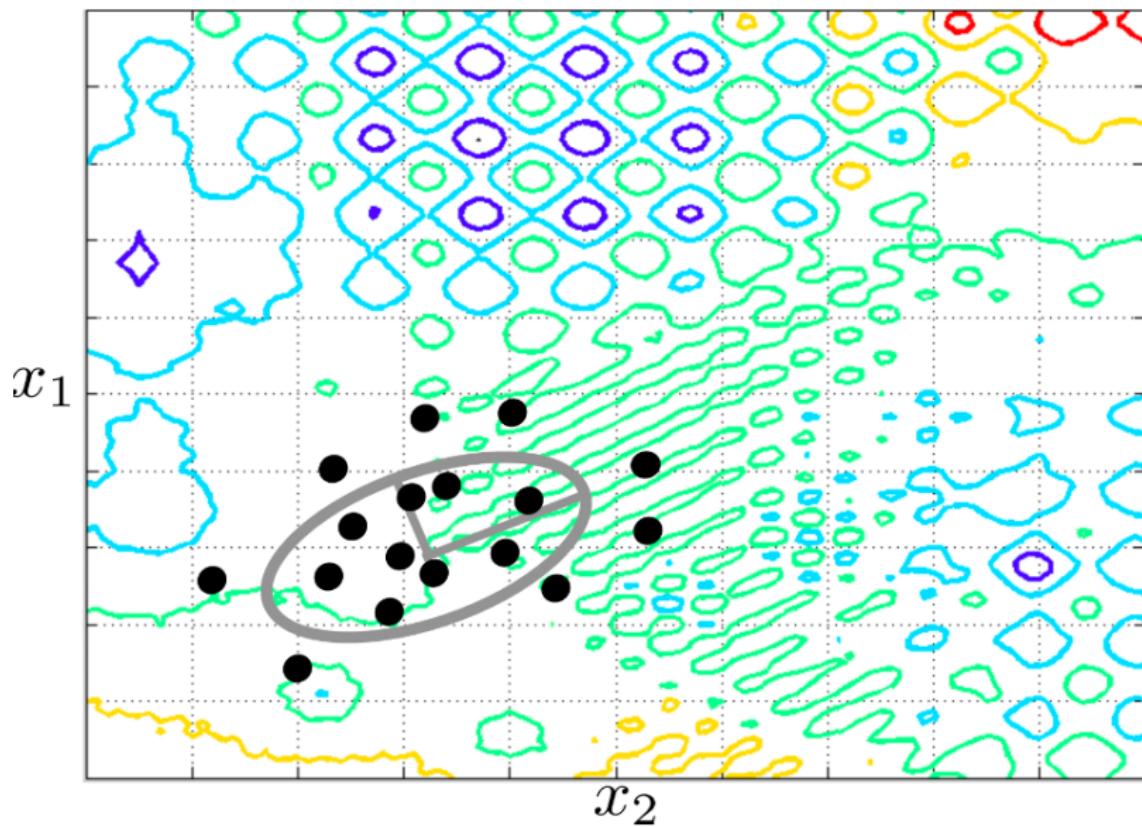
CMA-ES: WRAP UP



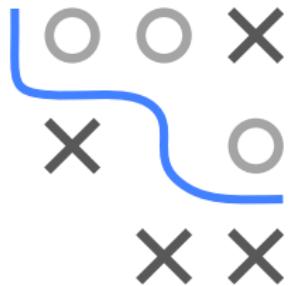
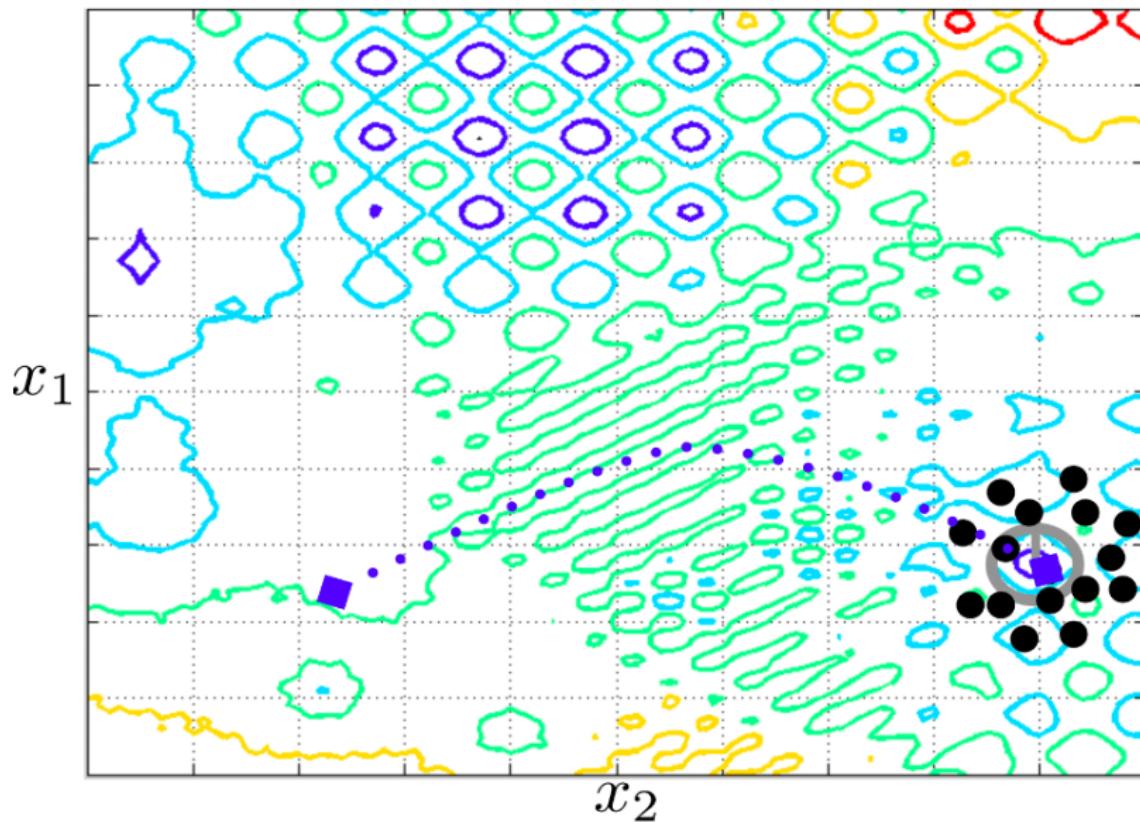
CMA-ES: WRAP UP



CMA-ES: WRAP UP



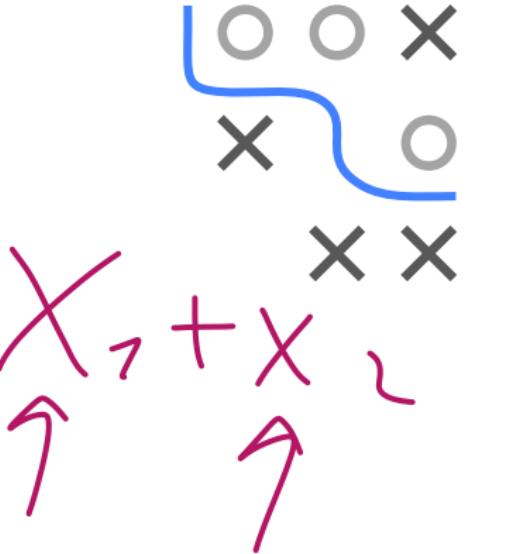
CMA-ES: WRAP UP



CMA-ES: WRAP UP - ADVANTAGES

CMA-ES can outperform other strategies in following cases:

- Non-separable problems (parameters of the objective function are dependent)
- Derivative of the objective function is not available
- High-dimensional problems (large d)
- Very large search space
- Useful in case “classical” search methods like quasi-Newton methods (BFGS) or conjugate gradient methods fail due to a non-convex or rugged search landscape (e.g. outliers, noise, local optima, sharp bends).



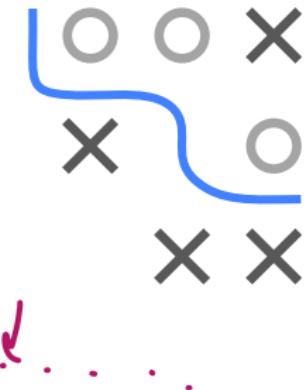
$$f(x_1) + f(x_2)$$

A hand-drawn diagram illustrating a non-separable classification problem. A blue line separates two classes of points: circles and crosses. Some points are misclassified. Below the plot, handwritten equations show the addition of variables x_1 and x_2 , with arrows pointing to each variable.

CMA-ES: WRAP UP - LIMITATIONS

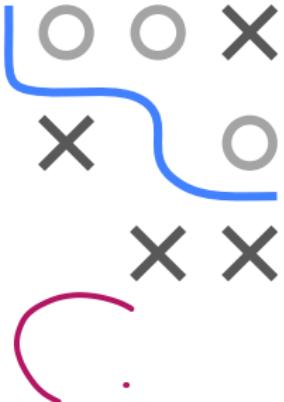
CMA-ES can be outperformed by other strategies in following cases:

- Partly separable problems (i.e. optimization of n -dimensional objective function can be divided into a series of d optimizations of every single parameter)
- Derivative of the objective function is easily available → Gradient Descend / Ascend
- Low dimensional problems (small d)
- Problems that can be solved by using a relative small number of function evaluations (e.g. $< 10d$ evaluations)



CMA-ES: IPOP

- Many special forms and extensions of the “basic” CMA-ES exist
- CMA-ES efficiently minimizes unimodal objective functions and is in particular superior on ill-conditioned, non-separable problems
- Default population size λ_{default} has been tuned for unimodal functions and however can get stuck in local optima on multi-modal functions, such that convergence to global optima is not guaranteed
- It could be shown that increasing the population size improves the performance of the CMA-ES on multi-modal functions
- **IPOP-CMA-ES** is a special form of restart-CMA-ES, where the *population size is increased for each restart (IPOP)* *Double it and give it to the next iteration*
- By increasing the population size the search characteristic becomes more global after each restart
- For the restart strategy CMA-ES is stopped whenever some stopping criterion is met, and an independent restart is launched with the population size increased by a factor of 2 (values between 1.5 and 5 are reasonable).



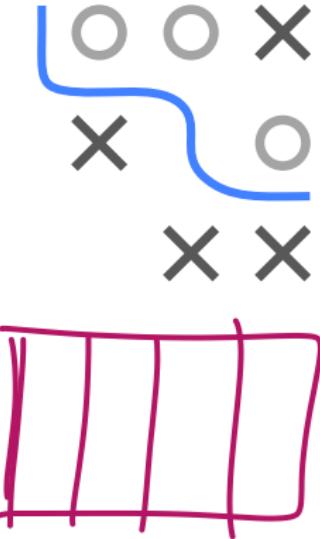
CMA-ES: WRAP UP - BENCHMARK EAS

Example: Black-box optimization of 25 benchmark functions under thoroughly defined experimental and recording conditions for the 2005 IEEE Congress on Evolutionary Computation: Session on Real-Parameter Optimization.

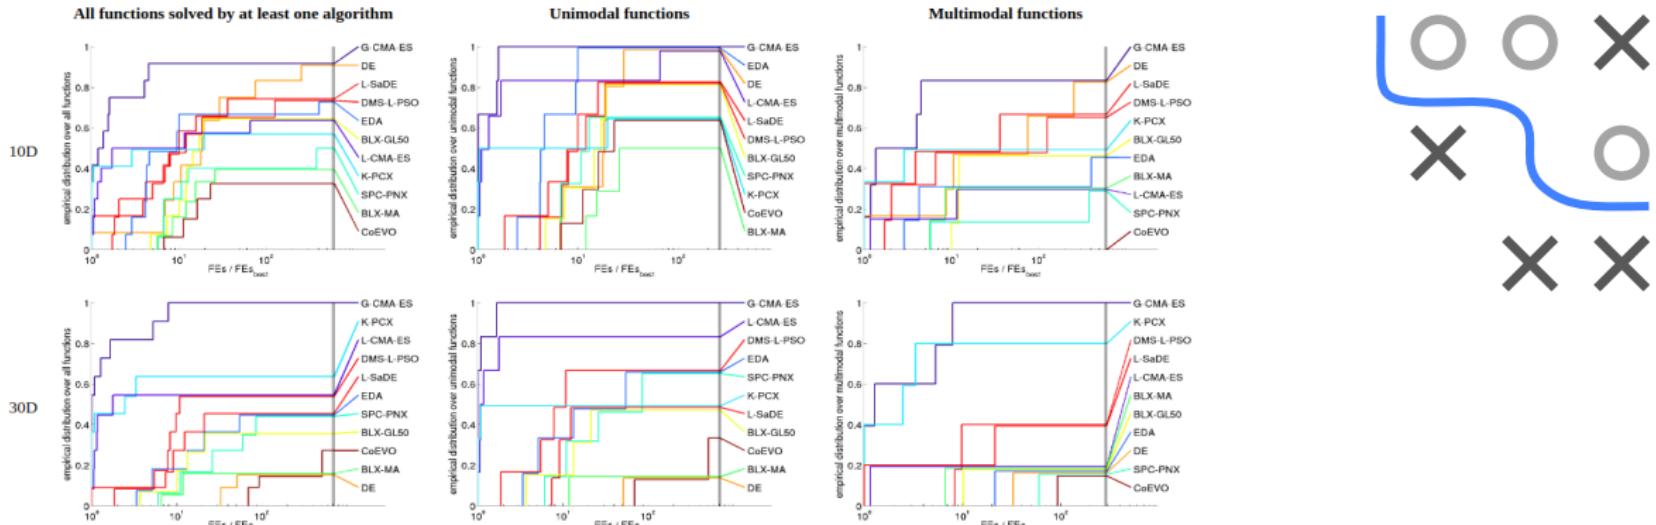
17 papers were submitted, 11 were accepted, thereunder hybrid methods.

Two of the Algorithms:

- L-CMA-ES (Auger and Hansen. 2005a): A CMA evolution strategy with small population size and small initial step-size to emphasize on local search characteristics. Independent restarts are conducted until the target function value is reached or the maximum number of function evaluations is exceeded.
- G-CMA-ES (Auger and Hansen. 2005b): A CMA evolution strategy restarted with increasing population size (IPOP). Independent restarts are conducted with increasing population size until the target function value is reached or the maximum number of function evaluations is exceeded. With the initial small population size the algorithm converges fast, with the succeeding larger population sizes the global search performance is emphasized in subsequent restarts.



CMA-ES: WRAP UP - BENCHMARK EAS



- Comparison of performance results from 11 algorithms for search space dimension 10 and 30 on different function subsets
- Expected number of function evaluations (FEs) to reach the target function value is normalized by the value of the best algorithm on the respective function FEs_{best}
- Calculation of the empirical cumulative distribution function of FEs / FEs_{best} for each algorithm over different sets of functions in 10 and 30D
- Small values for FEs / FEs_{best} and therefore large values of the graphs are preferable.

