

Lecture 3: Properties of Estimators

Bias · Variance · MSE · Consistency · Sufficiency · Cramér–Rao

We use estimators every day. Are they any good?

We already use estimators (Lecture 1, plug-in principle):

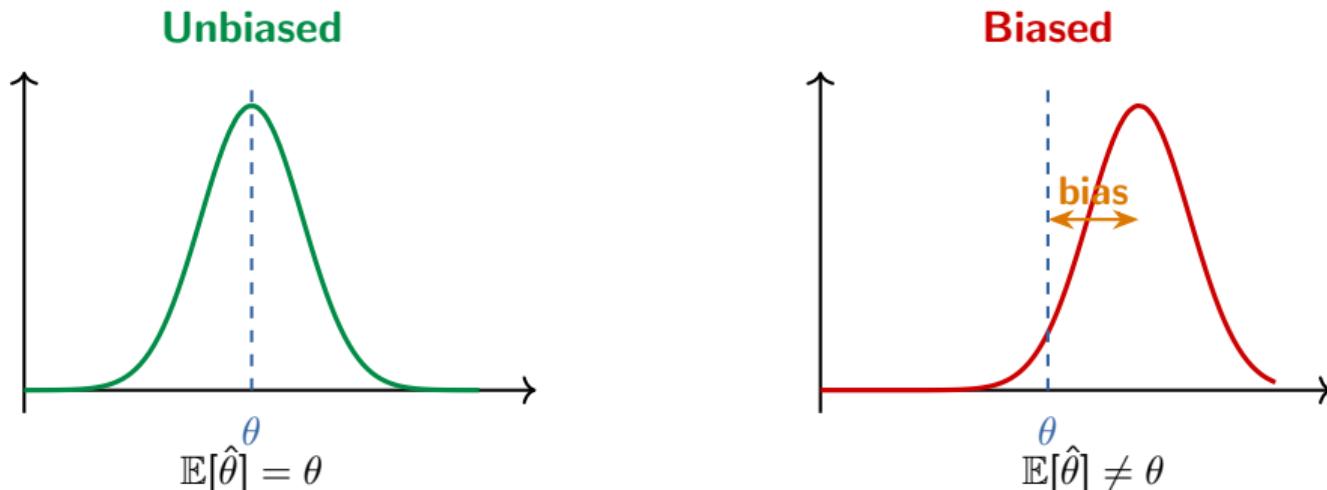
$$\bar{X} \text{ for } \mu, \quad S^2 \text{ for } \sigma^2, \quad \hat{p} = \frac{\text{count}}{n} \text{ for } p$$

But how do we **judge** an estimator?

Is it close to the truth? How much does
it jump around? Can we do better?

Bias: Is the Estimator Centered on the Truth?

$$\text{Bias}(\hat{\theta}) = \mathbb{E}[\hat{\theta}] - \theta$$



If $\text{Bias}(\hat{\theta}) = 0$ for all θ , the estimator is **unbiased**.

Worked Example: Is \bar{X} Unbiased for μ ?

Let X_1, \dots, X_n be i.i.d. with $\mathbb{E}[X_i] = \mu$. Is $\hat{\mu} = \bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ unbiased?

Step 1: Compute $\mathbb{E}[\hat{\mu}]$:

$$\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \cdot n\mu = \mu$$

Worked Example: Is \bar{X} Unbiased for μ ?

Let X_1, \dots, X_n be i.i.d. with $\mathbb{E}[X_i] = \mu$. Is $\hat{\mu} = \bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ unbiased?

Step 1: Compute $\mathbb{E}[\hat{\mu}]$:

$$\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \cdot n\mu = \mu$$

Step 2: Check bias:

$$\text{Bias}(\bar{X}) = \mathbb{E}[\bar{X}] - \mu = \mu - \mu = 0 \quad \checkmark \text{ Unbiased!}$$

Recipe for any estimator:

- (1) Compute $\mathbb{E}[\hat{\theta}]$
- (2) Subtract the true θ
- (3) If the result is 0, it's unbiased.

Worked Example: Why Dividing by n Is Biased

We want to estimate $\sigma^2 = \text{Var}(X_i)$. Natural guess: $\hat{\sigma}_n^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$.

Trick: rewrite each $(X_i - \bar{X})$ by adding and subtracting the true mean μ :

$$X_i - \bar{X} = \underbrace{(X_i - \mu)}_{\text{deviation from truth}} - \underbrace{(\bar{X} - \mu)}_{\text{estimation error}}$$

Squaring and summing gives the **key identity**:

$$\sum_{i=1}^n (X_i - \bar{X})^2 = \sum_{i=1}^n (X_i - \mu)^2 - n(\bar{X} - \mu)^2$$

Worked Example: Why Dividing by n Is Biased

We want to estimate $\sigma^2 = \text{Var}(X_i)$. Natural guess: $\hat{\sigma}_n^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$.

Trick: rewrite each $(X_i - \bar{X})$ by adding and subtracting the true mean μ :

$$X_i - \bar{X} = \underbrace{(X_i - \mu)}_{\text{deviation from truth}} - \underbrace{(\bar{X} - \mu)}_{\text{estimation error}}$$

Squaring and summing gives the **key identity**:

$$\sum_{i=1}^n (X_i - \bar{X})^2 = \sum_{i=1}^n (X_i - \mu)^2 - n(\bar{X} - \mu)^2$$

Take expectations (using $\mathbb{E}[(X_i - \mu)^2] = \sigma^2$ and $\text{Var}(\bar{X}) = \sigma^2/n$):

$$\mathbb{E}\left[\sum(X_i - \mu)^2\right] = n\sigma^2 \quad (n \text{ terms, each } \sigma^2)$$

$$\mathbb{E}[n(\bar{X} - \mu)^2] = n \cdot \text{Var}(\bar{X}) = n \cdot \frac{\sigma^2}{n} = \sigma^2 \quad (\text{one "lost" degree of freedom})$$

$$\Rightarrow \mathbb{E}\left[\sum(X_i - \bar{X})^2\right] = n\sigma^2 - \sigma^2 = (n - 1)\sigma^2$$

Bessel's Correction: The Fix

From the previous slide: $\mathbb{E}[\sum(X_i - \bar{X})^2] = (n-1)\sigma^2$, so:

$$\mathbb{E}[\hat{\sigma}_n^2] = \mathbb{E}\left[\frac{1}{n} \sum(X_i - \bar{X})^2\right] = \frac{(n-1)\sigma^2}{n} \neq \sigma^2 \quad \text{Biased!}$$

It **underestimates** by σ^2/n . Why? We used \bar{X} instead of μ , “using up” one degree of freedom.

Bessel's correction: Divide by $n-1$ instead of n :

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2 \quad \mathbb{E}[S^2] = \sigma^2 \quad \checkmark \text{ Unbiased!}$$

Intuition: We estimated μ from the same data, so the residuals $(X_i - \bar{X})$ are “too small” on average. Dividing by $n-1$ corrects for this.

Bias: Summary

Estimator	Bias	Unbiased?
$\bar{X} = \frac{1}{n} \sum X_i$ for μ	0	Yes
$\hat{\sigma}_n^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$ for σ^2	$-\frac{\sigma^2}{n}$	No
$S^2 = \frac{1}{n-1} \sum (X_i - \bar{X})^2$ for σ^2	0	Yes
$\hat{p} = \frac{\sum X_i}{n}$ for p (Bernoulli)	0	Yes

Dividing by n instead of $n-1$ **underestimates** the true variance.
Bessel's correction ($n-1$) fixes this. Recall Lecture 2!

Unbiasedness Alone Isn't Enough

Consider estimating $\mu = \mathbb{E}[X_i]$ from X_1, \dots, X_n .

Surprising fact: $\tilde{\mu} = X_1$ is also **unbiased!**

$$\mathbb{E}[X_1] = \mu \quad \Rightarrow \quad \text{Bias}(X_1) = 0 \quad \checkmark$$

Unbiasedness Alone Isn't Enough

Consider estimating $\mu = \mathbb{E}[X_i]$ from X_1, \dots, X_n .

Surprising fact: $\tilde{\mu} = X_1$ is also **unbiased**!

$$\mathbb{E}[X_1] = \mu \quad \Rightarrow \quad \text{Bias}(X_1) = 0 \quad \checkmark$$

But it's a terrible estimator — it ignores X_2, \dots, X_n entirely!

Three statisticians go deer hunting.

The first one shoots **1 meter to the left** of the deer.

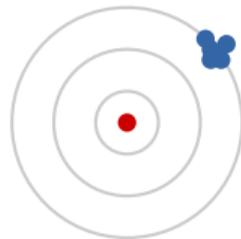
The second one shoots **1 meter to the right**.

The third one shouts: “We got it!”

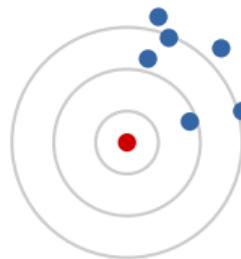
On average they hit the target — **unbiased!** But not very useful...

Lesson: Unbiasedness only says $\mathbb{E}[\hat{\theta}] = \theta$. It says nothing about how much $\hat{\theta}$ **varies**. We need more: **variance** and **MSE**.

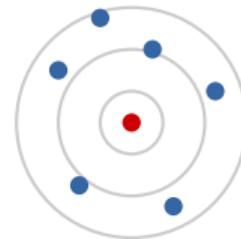
The Dartboard Analogy



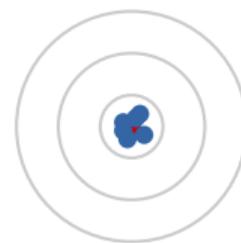
High bias, low var
Precise but inaccurate



High bias, high var
Worst of both worlds



Low bias, high var
Accurate but imprecise



Low bias, low var
The goal!

Bullseye = true θ . Blue dots = estimates from repeated samples.

Variance of an Estimator

The **variance** measures how much $\hat{\theta}$ wobbles across samples: $\text{Var}(\hat{\theta}) = \mathbb{E}[(\hat{\theta} - \mathbb{E}[\hat{\theta}])^2]$

Why is $\text{Var}(\bar{X}) = \sigma^2/n$ and not σ^2/n^2 ?

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \quad (\frac{1}{n} \text{ comes out as } \frac{1}{n^2})$$

$$= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \quad (\text{independent} \Rightarrow \text{variances add})$$

$$= \frac{1}{n^2} \cdot n\sigma^2 = \boxed{\frac{\sigma^2}{n}} \quad (n \text{ terms cancel one } n)$$

$$\text{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}} \quad (\text{standard error} = \sqrt{\text{Var}})$$

Mean Squared Error: The Total Error

Bias tells us about the **aim**, variance about the **spread**. Can we combine them?

$$\text{Mean Squared Error: } \text{MSE}(\hat{\theta}) = \mathbb{E}[(\hat{\theta} - \theta)^2]$$

The average squared distance from the estimate to the truth.

The trick: add and subtract $\mathbb{E}[\hat{\theta}]$ to decompose the error:

$$\hat{\theta} - \theta = \underbrace{(\hat{\theta} - \mathbb{E}[\hat{\theta}])}_{\text{random fluctuation}} + \underbrace{(\mathbb{E}[\hat{\theta}] - \theta)}_{\text{bias (a constant!)}}$$

This splits the total error into two pieces: the **random part** (how much $\hat{\theta}$ moves around its own mean) and the **systematic part** (how far that mean is from the truth).

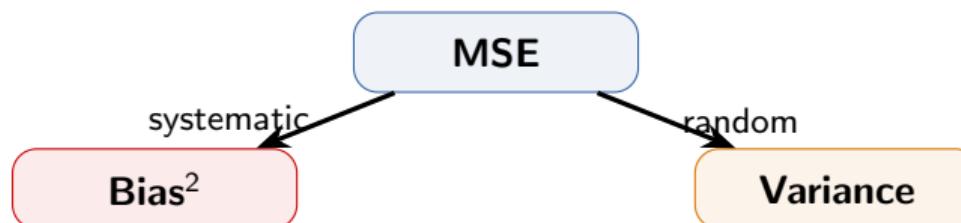
MSE = Bias² + Variance: The Proof

Now square $\hat{\theta} - \theta = (\hat{\theta} - \mathbb{E}[\hat{\theta}]) + (\mathbb{E}[\hat{\theta}] - \theta)$ and take expectations:

$$\mathbb{E}[(\hat{\theta} - \theta)^2] = \mathbb{E}\left[(\hat{\theta} - \mathbb{E}[\hat{\theta}])^2\right] + 2\underbrace{(\mathbb{E}[\hat{\theta}] - \theta)}_{\text{constant}} \cdot \underbrace{\mathbb{E}[\hat{\theta} - \mathbb{E}[\hat{\theta}]]}_{=0 \text{ (always!)}} + (\mathbb{E}[\hat{\theta}] - \theta)^2$$

The cross term vanishes because $\hat{\theta} - \mathbb{E}[\hat{\theta}]$ has mean zero **by definition**.

$$\boxed{\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta}) + \text{Bias}^2(\hat{\theta})}$$



Unbiased means $\text{MSE} = \text{Var}$, but a biased estimator can still win if its variance is low enough.

When Biased Beats Unbiased

Example: Estimating σ^2 from $X_1, \dots, X_n \sim N(\mu, \sigma^2)$.

Estimator	Bias	Variance	MSE
$S^2 = \frac{1}{n-1} \sum (X_i - \bar{X})^2$	0	$\frac{2\sigma^4}{n-1}$	$\frac{2\sigma^4}{n-1}$
$\hat{\sigma}_n^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$	$-\frac{\sigma^2}{n}$	$\frac{2(n-1)\sigma^4}{n^2}$	$\frac{(2n-1)\sigma^4}{n^2}$

Compare: $\frac{2n-1}{n^2}$ vs $\frac{2}{n-1}$ \Rightarrow $\hat{\sigma}_n^2$ has **lower MSE** for all $n \geq 2$!

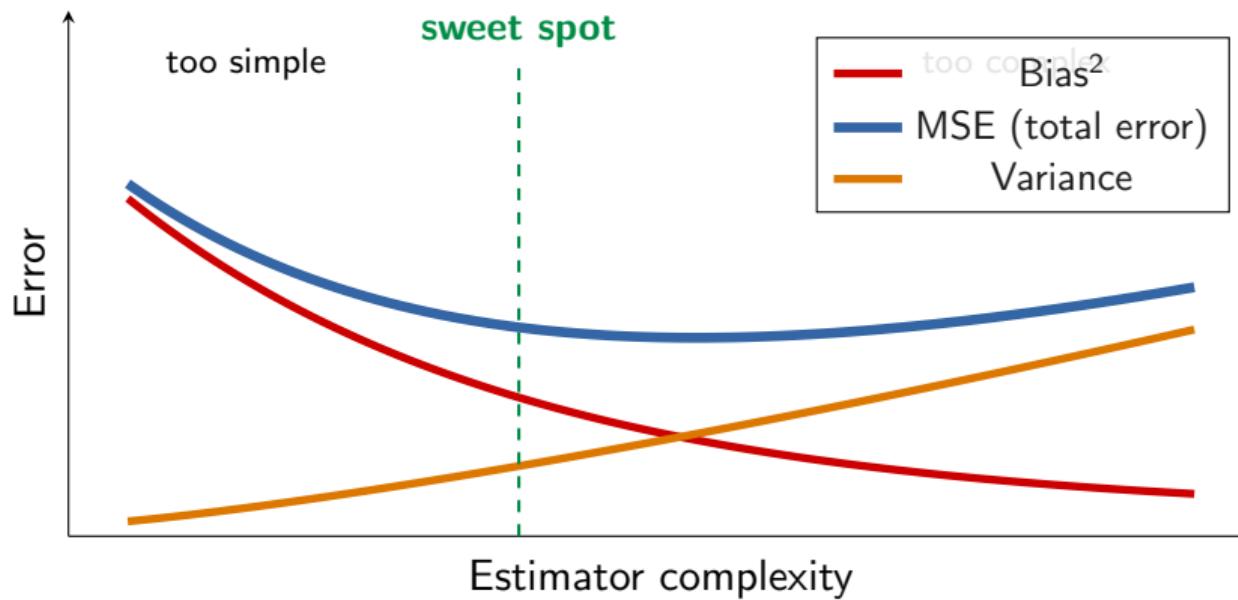
The biased estimator beats the unbiased S^2 because its variance reduction outweighs the small bias.

The Bias-Variance Tradeoff

You can't minimize bias and variance at the same time.

How do we find the **sweet spot**?

The Bias-Variance Tradeoff



Bias-Variance in Machine Learning

This tradeoff is **everywhere** in ML — it's the same principle in different disguises:

Setting	Too simple (high bias)	Too complex (high var)
Polynomial regression	Degree 1 (line)	Degree 20 (wiggly)
KNN	Large k (oversmoothed)	$k = 1$ (memorizes noise)
Decision tree	Shallow tree (underfits)	Deep tree (overfits)
Neural network	Too few neurons	Too many neurons
Regularization	Strong penalty (λ large)	No penalty ($\lambda = 0$)

Key insight: In all these cases, the total error (MSE, test loss) is minimized at an intermediate complexity. This is why we need **cross-validation**, **regularization**, and **held-out test sets** — to find the sweet spot empirically.

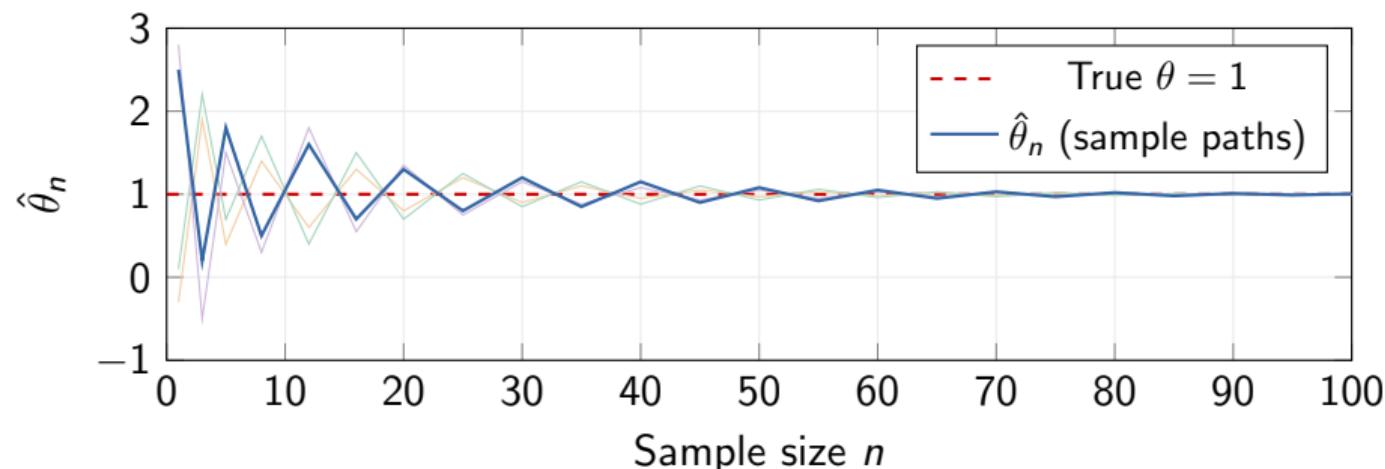
Consistency

Does our estimator converge to the truth
as we collect more and more data?

Consistency: Getting It Right Eventually

An estimator $\hat{\theta}_n$ is **consistent** if it converges to the truth as $n \rightarrow \infty$:

$$\hat{\theta}_n \xrightarrow{P} \theta \quad \text{i.e.,} \quad \Pr(|\hat{\theta}_n - \theta| > \varepsilon) \rightarrow 0 \text{ for all } \varepsilon > 0$$



Consistent vs Inconsistent: A Contrast

Consistent: $\hat{\mu} = \bar{X}_n$

- ▶ $\mathbb{E}[\bar{X}_n] = \mu$ (unbiased)
- ▶ $\text{Var}(\bar{X}_n) = \sigma^2/n \rightarrow 0$
- ▶ Uses **all** n observations
- ▶ More data \Rightarrow more precise

Not consistent: $\tilde{\mu} = X_1$

- ▶ $\mathbb{E}[X_1] = \mu$ (also unbiased!)
- ▶ $\text{Var}(X_1) = \sigma^2$ (constant!)
- ▶ Uses **only** the first observation
- ▶ Ignores all other data forever

Unbiased \neq consistent. X_1 is unbiased but NOT consistent.

Consistent \neq unbiased. $\hat{\sigma}_n^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$ is biased but IS consistent
(because its bias $\rightarrow 0$ and its variance $\rightarrow 0$).

Sufficient Conditions for Consistency

Chebyshev's inequality gives us a concrete tool:

$$\Pr(|\hat{\theta}_n - \theta| \geq \varepsilon) \leq \frac{\mathbb{E}[(\hat{\theta}_n - \theta)^2]}{\varepsilon^2} = \frac{\text{MSE}(\hat{\theta}_n)}{\varepsilon^2} = \frac{\text{Bias}^2 + \text{Var}}{\varepsilon^2}$$

$\text{Bias}(\hat{\theta}_n) \rightarrow 0$ as $n \rightarrow \infty$

$\text{Var}(\hat{\theta}_n) \rightarrow 0$ as $n \rightarrow \infty$



$\text{MSE} \rightarrow 0 \Rightarrow \Pr(|\hat{\theta}_n - \theta| \geq \varepsilon) \rightarrow 0 \Rightarrow \text{consistent!}$

Example: \bar{X}_n is consistent for μ : $\text{Bias} = 0$, $\text{Var} = \sigma^2/n \rightarrow 0$, so

$$\Pr(|\bar{X}_n - \mu| \geq \varepsilon) \leq \sigma^2/(n\varepsilon^2) \rightarrow 0.$$

This is precisely the **(Weak) Law of Large Numbers**: $\bar{X}_n \xrightarrow{P} \mu$.

Sufficiency

We have n data points. Do we really need **all** of them?

Can we **compress** without losing information?

Sufficiency: Can We Compress the Data?

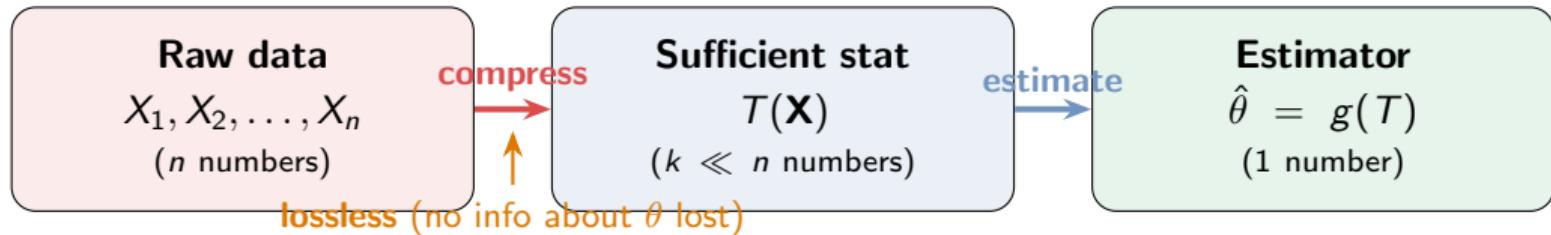
Example: $X_1, \dots, X_n \sim \text{Bern}(p)$. To estimate p :

- We only need $T = \sum X_i$ (total number of successes)
- The specific order (HHTHT vs THHTH) tells us nothing more about p

Definition: A statistic $T(\mathbf{X})$ is **sufficient** for θ if
the conditional distribution of $\mathbf{X} \mid T(\mathbf{X})$ does not depend on θ .

Intuition: Once you know T , the remaining randomness in the data is just noise —
it carries **no information** about θ . T is a “lossless summary.”

Sufficiency as Data Compression



Example

$$0, 1, 1, 0, 1, 1, 1, 0, 1, 0 \longrightarrow T = \sum X_i = 6 \longrightarrow \hat{p} = 6/10 = 0.6$$

Bernoulli

The order $(0, 1, 1, 0, 1, \dots)$ doesn't matter for estimating p — only the **total count** matters.

How to Check: Fisher–Neyman Factorization

Theorem: $T(\mathbf{X})$ is sufficient for θ if and only if:

$$f(\mathbf{x} \mid \theta) = g(T(\mathbf{x}), \theta) \cdot h(\mathbf{x})$$

where g depends on the data **only through** T , and h does not depend on θ .

Bernoulli worked example: $X_1, \dots, X_n \sim \text{Bern}(p)$, let $T = \sum X_i$.

$$f(\mathbf{x} \mid p) = \prod_{i=1}^n p^{x_i} (1-p)^{1-x_i} = \underbrace{p^{\sum x_i} (1-p)^{n-\sum x_i}}_{g(T, p)} \cdot \underbrace{1}_{h(\mathbf{x})}$$

Model	Sufficient statistic	Intuition
$\text{Bern}(p)$	$T = \sum X_i$	1 number for 1 parameter
$N(\mu, \sigma_0^2)$ (σ_0^2 known)	$T = \bar{X}$	1 number for 1 parameter
$N(\mu, \sigma^2)$ (both unknown)	$T = (\bar{X}, S^2)$	2 numbers for 2 parameters

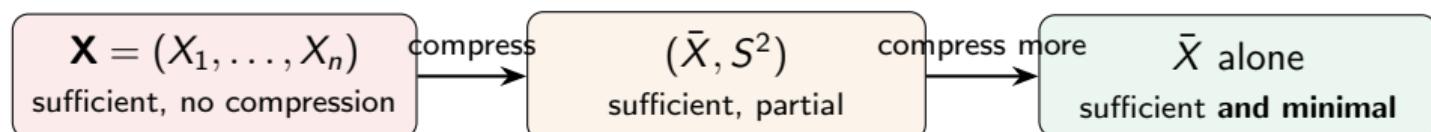
Minimal Sufficiency

The full data \mathbf{X} is always trivially sufficient. But can we compress **further**?

A sufficient statistic is **minimal** if it is a function of every other sufficient statistic.

It achieves the **maximum compression** without losing information about θ .

Example: For $X_1, \dots, X_n \sim N(\mu, \sigma_0^2)$ with σ_0^2 known:



Since only μ is unknown, S^2 carries no extra information — \bar{X} alone is enough.

The Rao–Blackwell Theorem

Why does sufficiency matter for estimation? Because it lets us **improve** any estimator:

Rao–Blackwell Theorem: Given *any* unbiased estimator $\tilde{\theta}$ and a sufficient statistic T , define $\hat{\theta} = \mathbb{E}[\tilde{\theta} | T]$. Then:

- (1) $\hat{\theta}$ is still **unbiased**: $\mathbb{E}[\hat{\theta}] = \mathbb{E}[\tilde{\theta}] = \theta$
- (2) $\hat{\theta}$ has **lower or equal variance**: $\text{Var}(\hat{\theta}) \leq \text{Var}(\tilde{\theta})$

Conditioning on a sufficient statistic **never hurts, often helps**.

The Rao–Blackwell Theorem

Why does sufficiency matter for estimation? Because it lets us **improve** any estimator:

Rao–Blackwell Theorem: Given *any* unbiased estimator $\tilde{\theta}$ and a sufficient statistic T , define $\hat{\theta} = \mathbb{E}[\tilde{\theta} | T]$. Then:

- (1) $\hat{\theta}$ is still **unbiased**: $\mathbb{E}[\hat{\theta}] = \mathbb{E}[\tilde{\theta}] = \theta$
- (2) $\hat{\theta}$ has **lower or equal variance**: $\text{Var}(\hat{\theta}) \leq \text{Var}(\tilde{\theta})$

Conditioning on a sufficient statistic **never hurts, often helps**.

Worked example: $X_1, \dots, X_n \sim \text{Bern}(p)$, sufficient stat $T = \sum X_i$.

$$\underbrace{\tilde{p} = X_1}_{\text{naive: unbiased, } \text{Var} = p(1-p)} \xrightarrow{\mathbb{E}[\cdot | T]} \underbrace{\hat{p} = \mathbb{E}[X_1 | T] = T/n = \bar{X}}_{\text{improved: unbiased, } \text{Var} = p(1-p)/n} \quad \times \text{n better!}$$

What Does $\mathbb{E}[\tilde{\theta} \mid T]$ Actually Mean?

Concrete example: $X_1, X_2, X_3 \sim \text{Bern}(p)$, $T = X_1 + X_2 + X_3$, $\tilde{p} = X_1$.

Suppose someone tells you $T = 2$ (two successes). Which data vectors give $T = 2$?

(X_1, X_2, X_3)	T	$\tilde{p} = X_1$	Equally likely?
$(1, 1, 0)$	2	1	Yes
$(1, 0, 1)$	2	1	Yes
$(0, 1, 1)$	2	0	Yes

What Does $\mathbb{E}[\tilde{\theta} \mid T]$ Actually Mean?

Concrete example: $X_1, X_2, X_3 \sim \text{Bern}(p)$, $T = X_1 + X_2 + X_3$, $\tilde{p} = X_1$.

Suppose someone tells you $T = 2$ (two successes). Which data vectors give $T = 2$?

(X_1, X_2, X_3)	T	$\tilde{p} = X_1$	Equally likely?
(1, 1, 0)	2	1	Yes
(1, 0, 1)	2	1	Yes
(0, 1, 1)	2	0	Yes

$$\mathbb{E}[X_1 \mid T = 2] = \frac{1 + 1 + 0}{3} = \frac{2}{3} = \frac{T}{n} \quad \checkmark$$

“Condition on T ” means: average $\tilde{\theta}$ over all data configurations that produce the same value of T . The noise (which specific X_i 's are 1 vs 0) gets averaged away. Only the useful part (T) survives.

Why Does Rao–Blackwell Work?

The key is the **law of total variance** (a.k.a. Eve's law):

$$\text{Var}(\tilde{\theta}) = \underbrace{\mathbb{E}[\text{Var}(\tilde{\theta} | T)]}_{\text{"useless" noise } \geq 0} + \text{Var}\left(\underbrace{\mathbb{E}[\tilde{\theta} | T]}_{\hat{\theta}}\right)$$

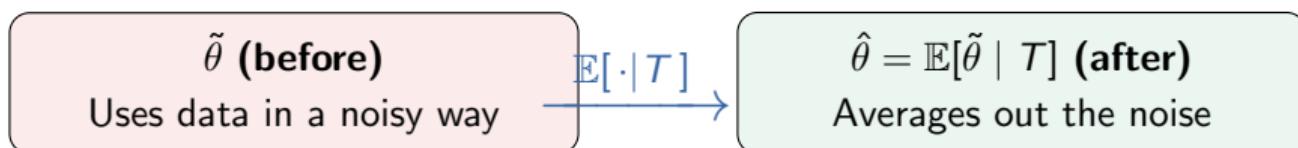
Why Does Rao–Blackwell Work?

The key is the **law of total variance** (a.k.a. Eve's law):

$$\text{Var}(\tilde{\theta}) = \underbrace{\mathbb{E}\left[\text{Var}(\tilde{\theta} | T)\right]}_{\text{"useless" noise } \geq 0} + \text{Var}\left(\underbrace{\mathbb{E}[\tilde{\theta} | T]}_{\hat{\theta}}\right)$$

Since the first term ≥ 0 , we immediately get:

$$\text{Var}(\tilde{\theta}) \geq \text{Var}(\hat{\theta})$$



Intuition: T captures all the useful information about θ . Conditioning on T removes the “useless” randomness (the part that doesn’t tell us about θ). What’s left is a cleaner estimator.

Finding Minimal Sufficient Statistics

Theorem (Likelihood Ratio Criterion): $T(\mathbf{X})$ is minimal sufficient iff for all \mathbf{x}, \mathbf{y} :

$$T(\mathbf{x}) = T(\mathbf{y}) \iff \frac{f(\mathbf{x} | \theta)}{f(\mathbf{y} | \theta)} \text{ does not depend on } \theta$$

Bernoulli example: $X_1, \dots, X_n \sim \text{Bern}(p)$.

$$\frac{f(\mathbf{x} | p)}{f(\mathbf{y} | p)} = \frac{p^{\sum x_i} (1-p)^{n-\sum x_i}}{p^{\sum y_i} (1-p)^{n-\sum y_i}} = \left(\frac{p}{1-p} \right)^{\sum x_i - \sum y_i}$$

Free of $p \iff \sum x_i = \sum y_i$. So $T = \sum X_i$ is **minimal sufficient** for p . ✓

Recipe: Write the likelihood ratio $f(\mathbf{x} | \theta)/f(\mathbf{y} | \theta)$.

Find which function of the data must match for the ratio to lose its θ -dependence.

That function is the minimal sufficient statistic.

The Exponential Family: A Unifying Framework

All our examples — Bernoulli, Normal, Poisson, Exponential — share one structure:

$$f(x | \theta) = h(x) \exp\left(\eta(\theta) T(x) - A(\theta)\right)$$

Distribution	Natural param $\eta(\theta)$	$T(x)$	Suff. stat (n obs)
Bern(p)	$\log \frac{p}{1-p}$	x	$\sum X_i$
$N(\mu, \sigma_0^2)$ (σ_0^2 known)	μ/σ_0^2	x	$\sum X_i$
Pois(λ)	$\log \lambda$	x	$\sum X_i$
Exp(λ)	$-\lambda$	x	$\sum X_i$

Pattern: For single-parameter families, $T(x) = x$. The sufficient statistic for n observations is always $\sum T(X_i)$ — straight from the factorization theorem!

Why Exponential Families Are Special

Nearly every nice property we've discussed is **automatic** in exponential families:

Sufficiency: $T(\mathbf{X}) = \sum T(X_i)$ is sufficient **and minimal**

Completeness: the natural sufficient statistic is **complete** (see below)

Regularity: all conditions for the Cramér–Rao bound (coming soon) are satisfied

Optimal estimators exist: we'll see this when we reach the CR bound

Completeness means: if $\mathbb{E}_\theta[g(T)] = 0$ for all θ , then $g(T) = 0$ a.s. \rightarrow **no non-trivial unbiased estimator of zero** based on T .

Lehmann–Scheffé: An unbiased estimator based on a **complete** sufficient statistic is the **unique best** unbiased estimator (UMVUE). For exp. families, $\sum T(X_i)$ is always complete \Rightarrow UMVUE exists!

Homework

1. Show that \bar{X} is unbiased for μ and compute its MSE.
2. Show that $\hat{\sigma}_n^2 = \frac{1}{n} \sum (X_i - \bar{X})^2$ is biased for σ^2 . Find the bias.
3. Suppose you shrink \bar{X} toward 0: $\hat{\mu}_c = c\bar{X}$ for $0 < c < 1$.
Find the bias, variance, and MSE as functions of c .
For what value of c is MSE minimized? Is the optimal estimator biased?
4. Use the factorization theorem to show that $T = \sum X_i$ is a sufficient statistic for λ when $X_1, \dots, X_n \sim \text{Poisson}(\lambda)$.

Questions?