Lecture 7: Sampling Distributions

Monte Carlo Simulation · CLT in Action · Standard Errors

Previously, on Lectures 5 & 6...

MLE: $\hat{\theta} = \arg\max \ell(\theta)$. Find the parameter that makes data most likely.

MAP: $\hat{\theta} = \arg\max[\ell(\theta) + \log P(\theta)]$. MLE with a prior = regularization.

Fisher info: $I(\theta) = -E[\ell''(\theta)]$. Curvature of the log-likelihood.

CR bound: $\mathrm{Var}(\hat{\theta}) \geq 1/(nI(\theta))$. A floor on how precise we can be.

Asymptotic normality: $\hat{\theta}_{\mathrm{MLE}} \dot{\sim} N(\theta_0,\ 1/(nI(\theta_0)))$ for large n.

Today: You computed $\hat{\theta} = 3.2$. But how **reliable** is this number?

If you repeated the experiment, would you get 3.2 again? 3.1? 5.7?

Your Estimator Is Random

Different sample $\Rightarrow$ different $\hat{\theta}$

The Thought Experiment

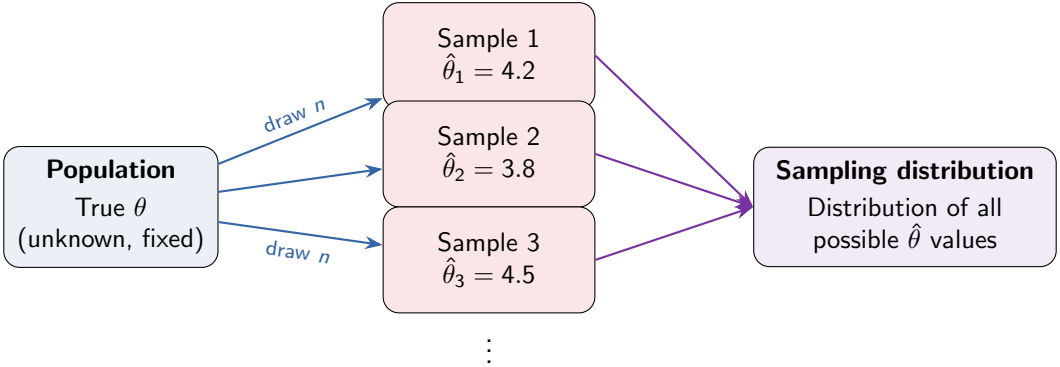

The Thought Experiment

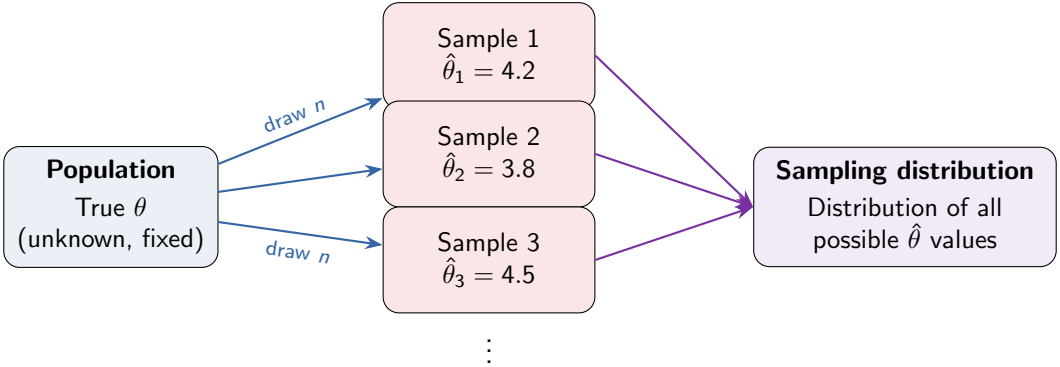

Key insight: $\hat{\theta}$ is a *function of random data*, so $\hat{\theta}$ itself is a random variable.

Its distribution is called the **sampling distribution**. Understanding it is the foundation of all inference.

But We Only Have One Sample!

In real life, you run the experiment **once**. You get **one** $\hat{\theta}$.

The problem:

We can't actually repeat the experiment 1,000 times to see the sampling distribution.

Today: solutions 1 and 2. Next lecture: solution 3 (the bootstrap).

But We Only Have One Sample!

In real life, you run the experiment **once**. You get **one** $\hat{\theta}$.

The problem:

We can't actually repeat the experiment 1,000 times to see the sampling distribution.

Three solutions:

1. **Theory** — derive it mathematically
2. **Simulation** — fake it on a computer
3. **Bootstrap** — resample the data itself

But We Only Have One Sample!

In real life, you run the experiment **once**. You get **one** $\hat{\theta}$.

The problem:

We can't actually repeat the experiment 1,000 times to see the sampling distribution.

Three solutions:

1. **Theory** — derive it mathematically
2. **Simulation** — fake it on a computer
3. **Bootstrap** — resample the data itself

Today: solutions 1 and 2. Next lecture: solution 3 (the bootstrap).

Seeing It: Monte Carlo Simulation

Let the computer repeat the experiment for you

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

Why "Monte Carlo"? Named after the famous casino. We use *randomness* to learn about randomness.

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

Step 1: Generate a random sample $X_1, \ldots, X_n \sim F$.

Why "Monte Carlo"? Named after the famous casino. We use *randomness* to learn about randomness.

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

Step 1: Generate a random sample $X_1, \ldots, X_n \sim F$.

Step 2: Compute the estimator: $\hat{\theta} = T(X_1, \ldots, X_n)$.

Why "Monte Carlo"? Named after the famous casino. We use *randomness* to learn about randomness.

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

Step 1: Generate a random sample $X_1, \ldots, X_n \sim F$.

Step 2: Compute the estimator: $\hat{\theta} = T(X_1, \ldots, X_n)$.

Step 3: Repeat Steps 1–2 a total of R times (e.g., $R = 10{,}000$).

Why "Monte Carlo"? Named after the famous casino. We use *randomness* to learn about randomness.

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

↓

Step 1: Generate a random sample $X_1, \ldots, X_n \sim F$.

↓

Step 2: Compute the estimator: $\hat{\theta} = T(X_1, \ldots, X_n)$.

↓

Step 3: Repeat Steps 1–2 a total of R times (e.g., $R = 10{,}000$).

↓

Step 4: Plot the histogram of $\hat{\theta}_1, \hat{\theta}_2, \ldots, \hat{\theta}_R$. This *is* the sampling distribution.

The Monte Carlo Recipe

Setup: Pick a true population distribution F with known θ.

↓

Step 1: Generate a random sample $X_1, \ldots, X_n \sim F$.

↓

Step 2: Compute the estimator: $\hat{\theta} = T(X_1, \ldots, X_n)$.

↓

Step 3: Repeat Steps 1–2 a total of R times (e.g., $R = 10{,}000$).

↓

Step 4: Plot the histogram of $\hat{\theta}_1, \hat{\theta}_2, \ldots, \hat{\theta}_R$. This *is* the sampling distribution.

Why "Monte Carlo"? Named after the famous casino. We use *randomness* to learn about randomness.

Example: Sample Mean from $N(5, 4)$

Population: $X_i \sim N(\mu = 5, \sigma^2 = 4)$, $n = 10$, estimator: $\hat{\mu} = \bar{X}$

Population $N(5, 4)$

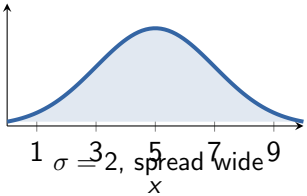

$\sigma = 2$, spread wide

simulate $R = 10{,}000$

Sampling dist. of $\bar{X}$

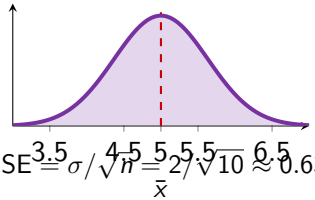

$\text{SE} = \sigma/\sqrt{n} = 2/\sqrt{10} \approx 0.63$

The sampling distribution is **much narrower** than the population — averaging reduces the spread by $\sqrt{n}$.

What Happens as n Grows?

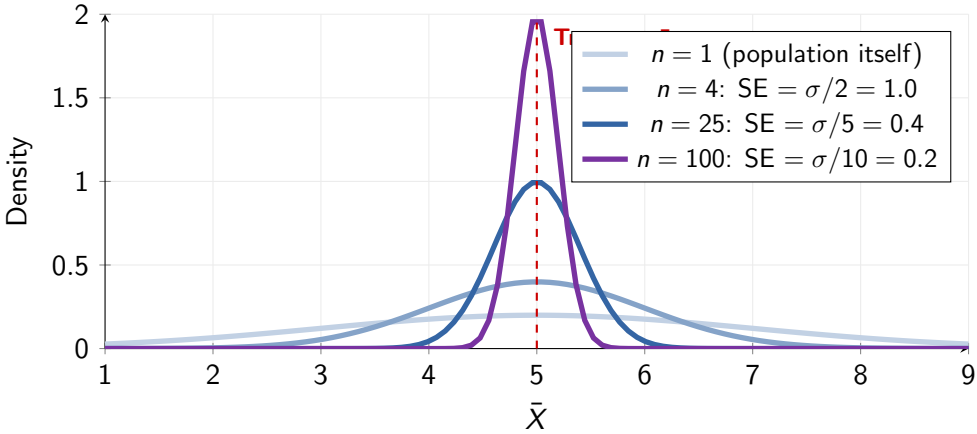

More data $\Rightarrow$ the sampling distribution **concentrates** around the true value. This is **consistency**.

It Works for Any Estimator

Not just $\bar{X}$! **Every** estimator has a sampling distribution.

Sample Mean $\bar{X}$

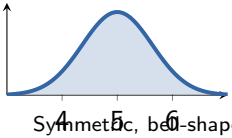

Symmetric, bell-shaped

Sample Median

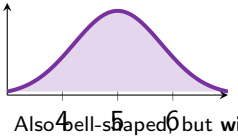

Also bell-shaped but **wider**

Sample Variance S^2

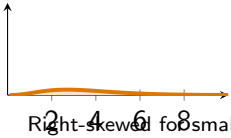

Right-skewed for small n

It Works for Any Estimator

Key observation

Observation: Different estimators have **different** sampling distributions.

The mean is narrower than the median (for Normal data) — it's a *more efficient* estimator.

Recall from Lecture 4: the Cramér–Rao bound tells us the *best possible* variance.

Coming up: We'll see that the **Central Limit Theorem** explains *why* sampling distributions tend to look Normal — and **Fisher information** determines how narrow they are.

Example: MLE for Exp(λ)

Population: $X_i \sim \text{Exp}(\lambda = 2)$, MLE: $\hat{\lambda} = 1/\bar{X}$

Population Exp(2)

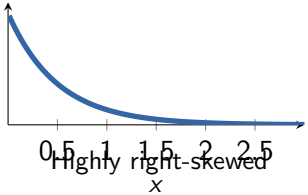

Highly right-skewed

$R = 10{,}000$, $n = 20$

Sampling dist. of $\hat{\lambda}_{\text{MLE}}$

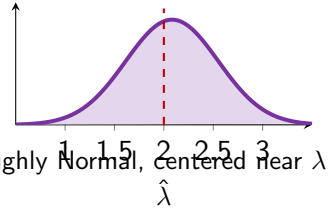

Roughly Normal, centered near $\lambda = 2$

The population is **exponential** (skewed), but the MLE is **approximately Normal**. Why? **CLT!**

The Central Limit Theorem

Why everything looks Normal

CLT: The Most Important Theorem in Statistics

Central Limit Theorem: If $X_1, \ldots, X_n$ are i.i.d. with mean μ and variance $\sigma^2 < \infty$, then

$$\frac{\bar{X}-\mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0,1) \quad \text{as } n \to \infty$$

Equivalently: $\bar{X} \dot{\sim} N\left(\mu, \frac{\sigma^2}{n}\right)$ for large n.

CLT: The Most Important Theorem in Statistics

Central Limit Theorem: If $X_1, \ldots, X_n$ are i.i.d. with mean μ and variance $\sigma^2 < \infty$, then

$$\frac{\bar{X}-\mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0,1) \quad \text{as } n \to \infty$$

Equivalently: $\bar{X} \dot{\sim} N\left(\mu, \frac{\sigma^2}{n}\right)$ for large n.

The magic: It doesn't matter what the population looks like — uniform, exponential, Bernoulli, bimodal, anything.
As long as the variance is finite, $\bar{X}$ becomes Normal.

CLT in Action: Uniform Population

$X_i \sim \text{Uniform}(0, 1)$ — completely flat, yet $\bar{X}$ becomes Normal

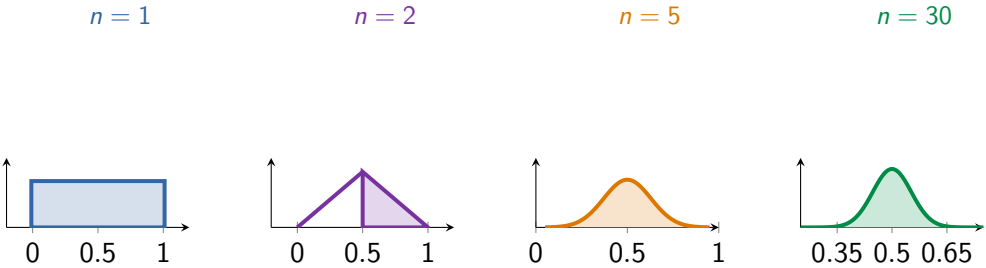

CLT in Action: Exponential Population

$X_i \sim \text{Exp}(1)$ — strongly right-skewed, yet $\bar{X}$ becomes Normal (needs larger n)

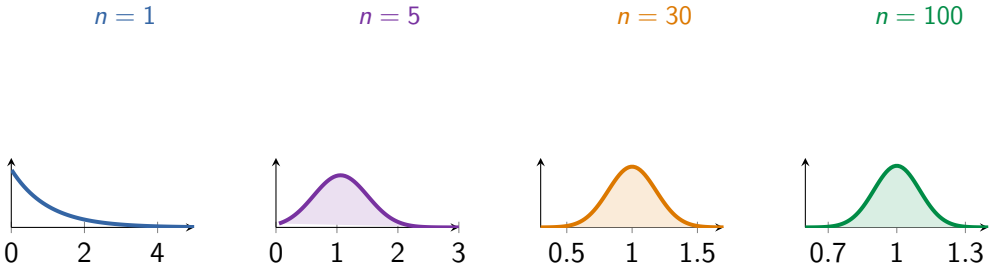

How Large Must n Be?

Symmetric
(Uniform, Normal)

$n \geq 5$ is often enough

Mildly skewed
(Exponential, Poisson)

$n \geq 25$–30 usually works

Heavily skewed
(Pareto, lognormal)

May need $n \geq 100$ or more

How Large Must n Be?

Symmetric
(Uniform, Normal)
$n \geq 5$ is often enough

Mildly skewed
(Exponential, Poisson)
$n \geq 25$–30 usually works

Heavily skewed
(Pareto, lognormal)
May need $n \geq 100$ or more

The common "$n \geq 30$" rule is a useful heuristic, not a law.
For symmetric data, even $n = 5$ is fine. For heavy tails, even $n = 100$ may not suffice.
When in doubt — **simulate!** That's what Monte Carlo is for.

CLT for MLEs: Even More Powerful

The CLT doesn't just apply to $\bar{X}$. For **any** MLE (under regularity conditions):

$$\hat{\theta}_{\text{MLE}} \quad \dot{\sim} \quad N\left(\theta_0, \ \frac{1}{n \cdot I(\theta_0)}\right)$$

The MLE is approximately Normal, centered at the truth,
with variance determined by the **Fisher information**.

CLT for MLEs: Even More Powerful

The CLT doesn't just apply to $\bar{X}$. For **any** MLE (under regularity conditions):

$$\hat{\theta}_{\text{MLE}} \quad \dot{\sim} \quad N\left(\theta_0, \ \frac{1}{n \cdot I(\theta_0)}\right)$$

The MLE is approximately Normal, centered at the truth,
with variance determined by the **Fisher information**.

Bernoulli(p)	**Poisson(λ)**	**Exp(λ)**
$\hat{p} \dot{\sim} N\left(p, \ \frac{p(1-p)}{n}\right)$	$\hat{\lambda} \dot{\sim} N\left(\lambda, \ \frac{\lambda}{n}\right)$	$\hat{\lambda} \dot{\sim} N\left(\lambda, \ \frac{\lambda^2}{n}\right)$

Each follows from $I(\theta)$: Bernoulli has $I(p) = 1/(p(1-p))$, Poisson has $I(\lambda) = 1/\lambda$, etc.

Standard Error

The most important number after the estimate itself

SD $\neq$ SE: The Most Confused Pair in Statistics

Standard Deviation (SD)

Spread of the **data**

$$\sigma = \sqrt{E[(X-\mu)^2]}$$

"How variable are individual observations?"

Does not shrink with n.
More data doesn't make individuals less variable.

Standard Error (SE)

Spread of the **estimator**

$$\mathrm{SE}(\bar{X}) = \sigma/\sqrt{n}$$

"How variable is $\hat{\theta}$ across repeated experiments?"

Shrinks with n.
More data $\Rightarrow$ more precise estimate.

SD $\neq$ SE: The Most Confused Pair in Statistics

Standard Deviation (SD)

Spread of the **data**

$$\sigma = \sqrt{E[(X-\mu)^2]}$$

"How variable are individual observations?"

Does not shrink with n.
More data doesn't make individuals less variable.

Standard Error (SE)

Spread of the **estimator**

$$\mathrm{SE}(\bar{X}) = \sigma/\sqrt{n}$$

"How variable is $\hat{\theta}$ across repeated experiments?"

Shrinks with n.
More data $\Rightarrow$ more precise estimate.

Example: Heights of adults: SD $\approx$ 7 cm (people vary).
Average of 100 heights: SE $= 7/\sqrt{100} = 0.7$ cm (the average is precise).

Standard Error of the Sample Mean

$$\mathrm{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

The **standard error** is the standard deviation of the sampling distribution.

It measures how much $\bar{X}$ would vary if we repeated the experiment.

Standard Error of the Sample Mean

$$\mathrm{SE}(\bar{X}) \quad = \quad \frac{\sigma}{\sqrt{n}}$$

The **standard error** is the standard deviation of the sampling distribution.

It measures how much $\bar{X}$ would vary if we repeated the experiment.

Derivation:

$$\mathrm{Var}(\bar{X}) = \mathrm{Var}\left(\frac{1}{n}\sum_{i=1}^{n} X_i\right) = \frac{1}{n^2}\sum_{i=1}^{n}\mathrm{Var}(X_i) = \frac{1}{n^2}\cdot n\sigma^2 = \frac{\sigma^2}{n}$$

Standard Error of the Sample Mean

$$\mathrm{SE}(\bar{X}) \quad = \quad \frac{\sigma}{\sqrt{n}}$$

The **standard error** is the standard deviation of the sampling distribution.

It measures how much $\bar{X}$ would vary if we repeated the experiment.

Derivation:

$$\mathrm{Var}(\bar{X}) = \mathrm{Var}\left(\frac{1}{n}\sum_{i=1}^{n} X_i\right) = \frac{1}{n^2}\sum_{i=1}^{n}\mathrm{Var}(X_i) = \frac{1}{n^2}\cdot n\sigma^2 = \frac{\sigma^2}{n}$$

$$\Rightarrow \quad \mathrm{SE}(\bar{X}) = \sqrt{\mathrm{Var}(\bar{X})} = \frac{\sigma}{\sqrt{n}}$$

Uses i.i.d. assumption: observations are independent, so variances **add**.

The $\sqrt{n}$ Law: Precision Is Expensive

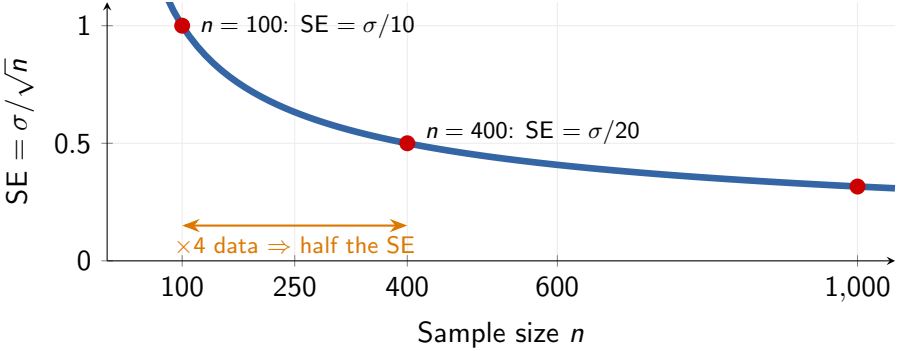

The $\sqrt{n}$ Law: Precision Is Expensive

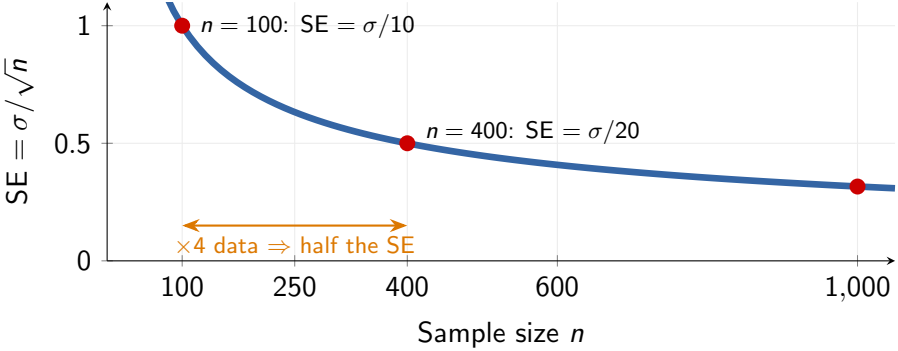

To **halve** the SE, you must **quadruple** the sample size.
Going from $n = 100$ to $n = 10{,}000$ only reduces SE by $10\times$.
This is why sample size planning matters!

SE for Any MLE: The General Formula

From Lecture 4 (Cramér–Rao) and Lecture 5 (MLE asymptotics):

$$\mathrm{SE}(\hat{\theta}_{\mathrm{MLE}}) \approx \frac{1}{\sqrt{n \cdot I(\theta)}}$$

The Fisher information $I(\theta)$ determines the precision of the MLE.

More informative data (larger I) $\Rightarrow$ smaller SE.

SE for Any MLE: The General Formula

From Lecture 4 (Cramér–Rao) and Lecture 5 (MLE asymptotics):

$$\mathrm{SE}(\hat{\theta}_{\mathrm{MLE}}) \approx \frac{1}{\sqrt{n \cdot I(\theta)}}$$

The Fisher information $I(\theta)$ determines the precision of the MLE.

More informative data (larger I) $\Rightarrow$ smaller SE.

Problem: This formula contains the *unknown* θ!

SE for Any MLE: The General Formula

From Lecture 4 (Cramér–Rao) and Lecture 5 (MLE asymptotics):

$$\mathrm{SE}(\hat{\theta}_{\mathrm{MLE}}) \approx \frac{1}{\sqrt{n \cdot I(\theta)}}$$

The Fisher information $I(\theta)$ determines the precision of the MLE.

More informative data (larger I) $\Rightarrow$ smaller SE.

Problem: This formula contains the *unknown* θ!

Solution: Plug-in. Replace θ with $\hat{\theta}$:

$$\widehat{\mathrm{SE}}(\hat{\theta}) = \frac{1}{\sqrt{n \cdot I(\hat{\theta})}}$$

Estimated standard error — this is what we report in practice.

SE for Any MLE: The General Formula

From Lecture 4 (Cramér–Rao) and Lecture 5 (MLE asymptotics):

$$\mathrm{SE}(\hat{\theta}_{\mathrm{MLE}}) \approx \frac{1}{\sqrt{n \cdot I(\theta)}}$$

The Fisher information $I(\theta)$ determines the precision of the MLE.

More informative data (larger I) $\Rightarrow$ smaller SE.

Problem: This formula contains the *unknown* θ!

Solution: Plug-in. Replace θ with $\hat{\theta}$:

$$\widehat{\mathrm{SE}}(\hat{\theta}) = \frac{1}{\sqrt{n \cdot I(\hat{\theta})}}$$

Estimated standard error — this is what we report in practice.

Known vs Estimated SE

If σ is known:

$$\mathrm{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

Rare in practice.
Use the z-distribution (Normal).

The t-distribution has heavier tails than the Normal — it accounts for the
extra uncertainty from estimating σ. As $n \to \infty$, $t_{n-1} \to N(0, 1)$.

Known vs Estimated SE

If σ is known:

$$\mathrm{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

Rare in practice.
Use the z-distribution (Normal).

If σ is unknown (typical):

$$\widehat{\mathrm{SE}}(\bar{X}) = \frac{S}{\sqrt{n}}$$

Replace σ with $S = \sqrt{\frac{1}{n-1}\sum(X_i - \bar{X})^2}$.
Use the t-distribution.

Known vs Estimated SE

If σ is known:

$$\mathrm{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

Rare in practice.
Use the z-distribution (Normal).

If σ is unknown (typical):

$$\widehat{\mathrm{SE}}(\bar{X}) = \frac{S}{\sqrt{n}}$$

Replace σ with $S = \sqrt{\frac{1}{n-1}\sum(X_i - \bar{X})^2}$.
Use the t-distribution.

The t-distribution has heavier tails than the Normal — it accounts for the
extra uncertainty from estimating σ. As $n \to \infty$, $t_{n-1} \to N(0,1)$.

Connecting the Dots

From data to uncertainty

Fisher Information $\leftrightarrow$ Sampling Distribution

High Fisher info

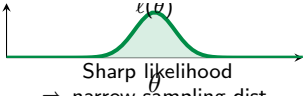

Sharp likelihood
$\Rightarrow$ narrow sampling dist.
$\Rightarrow$ small SE

Low Fisher info

Flat likelihood
$\Rightarrow$ wide sampling dist.
$\Rightarrow$ large SE

Fisher Information $\leftrightarrow$ Sampling Distribution

High Fisher info

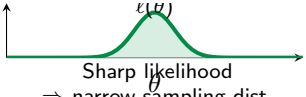

Sharp likelihood
$\Rightarrow$ narrow sampling dist.
$\Rightarrow$ small SE

Low Fisher info

Flat likelihood
$\Rightarrow$ wide sampling dist.
$\Rightarrow$ large SE

The chain: More curvature ($I(\theta)$ $\uparrow$) $\Rightarrow$ Smaller variance ($1/(nI)$ $\downarrow$) $\Rightarrow$ Smaller SE $\Rightarrow$ More precise $\hat{\theta}$

This is why Fisher information is called "information" — it literally measures how much the data tells us about θ.

The Full Picture

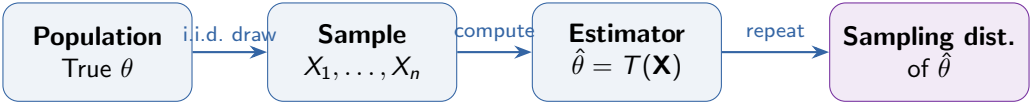

Two paths to quantify uncertainty: **analytical** (formulas, today) and **computational** (bootstrap, next lecture).

The Full Picture

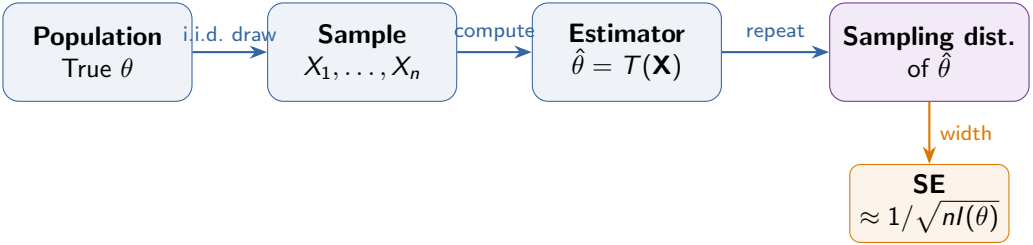

Two paths to quantify uncertainty: **analytical** (formulas, today) and **computational** (bootstrap, next lecture).

The Full Picture

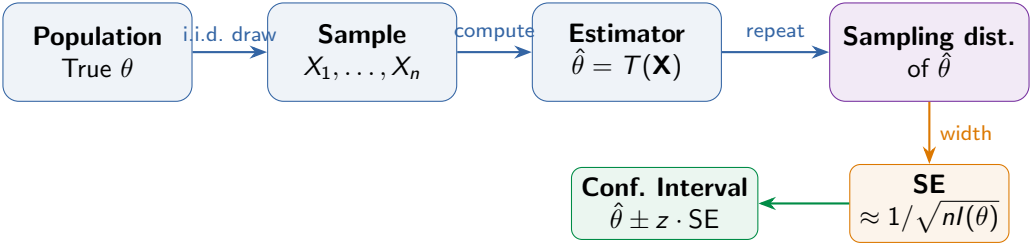

Two paths to quantify uncertainty: **analytical** (formulas, today) and **computational** (bootstrap, next lecture).

The Full Picture

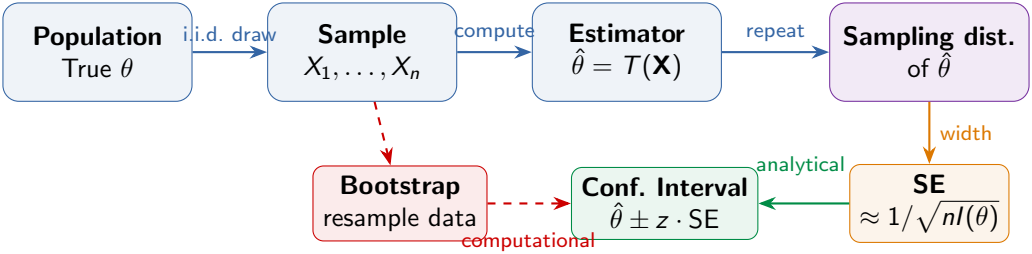

The Full Picture

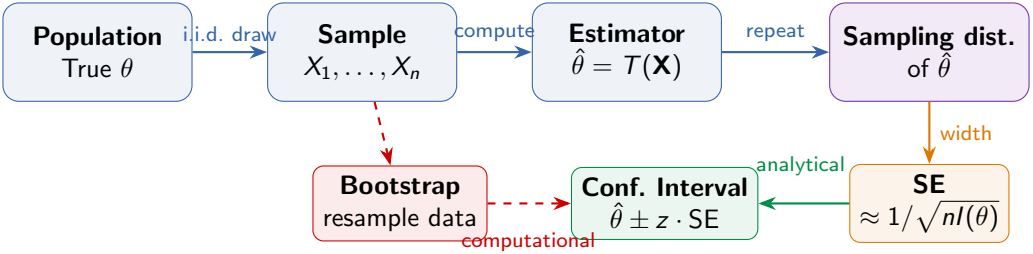

Two paths to quantify uncertainty: **analytical** (formulas, today) and **computational** (bootstrap, next lecture).

Simulation in Python

Verify everything with code

Monte Carlo: Verifying the Sampling Distribution

```python
import numpy as np
import matplotlib.pyplot as plt

# True parameters
mu, sigma, n = 5.0, 2.0, 25
R = 10_000   # number of repetitions

# Simulate R experiments
xbar = [np.mean(np.random.normal(mu, sigma, n))
        for _ in range(R)]

# Results
print(f"Mean of xbar: {np.mean(xbar):.3f}")  # ≈ 5.0
print(f"SD of xbar:  {np.std(xbar):.3f}")  # ≈ 0.4 = 2/√25

plt.hist(xbar, bins=50, density=True)
plt.title("Sampling distribution of X-bar")
```

Run this code. Change n, σ, the distribution. **See** the sampling distribution form.

Verifying the CLT

```
# CLT for Exponential: mean=1, var=1
for n in [1, 5, 30, 100]:
  xbar = [np.mean(np.random.exponential(1, n))
        for _ in range(10_000)]
  plt.subplot(2, 2, [1,5,30,100].index(n)+1)
  plt.hist(xbar, bins=50, density=True)
  plt.title(f"n={n}, SD={np.std(xbar):.3f}")

# Watch it go from skewed (n=1) to Normal (n=30+)
# Check: SD should be ≈ 1/sqrt(n)
```

Verifying the CLT

```
# CLT for Exponential: mean=1, var=1
for n in [1, 5, 30, 100]:
  xbar = [np.mean(np.random.exponential(1, n))
          for _ in range(10_000)]
  plt.subplot(2, 2, [1,5,30,100].index(n)+1)
  plt.hist(xbar, bins=50, density=True)
  plt.title(f"n={n}, SD={np.std(xbar):.3f}")

# Watch it go from skewed (n=1) to Normal (n=30+)
# Check: SD should be ≈ 1/sqrt(n)
```

Try it yourself:

1. Replace `exponential` with `uniform`, `binomial`, or `pareto`
2. For each: at what n does the histogram "look Normal"?
3. Overlay a $N(\mu, \sigma^2/n)$ curve — does it match?

Summary: The Life of an Estimator

$\hat{\theta}$ **is random:** Different samples give different estimates. The sampling distribution captures th

Monte Carlo: Simulate many experiments to *see* the sampling distribution. Always verify theory wi

CLT: $\bar{X} \dot{\sim} N(\mu, \sigma^2/n)$ regardless of the population shape. The MLE is also asymptotically Norm

SD $\neq$ SE: SD measures data spread. SE measures estimator precision. SE shrinks with n.

$\sqrt{n}$ **law:** Halving the SE requires quadrupling n. Precision has diminishing returns.

SE for MLEs: SE $\approx 1/\sqrt{nI(\theta)}$. Fisher information determines precision.

Plug-in SE: In practice, replace θ with $\hat{\theta}$. For $\bar{X}$: replace σ with S (use t-distribution).

Next: Two ways to use the sampling distribution: confidence intervals (formulas) and bootstrap (sim

Practical: Monte Carlo Simulations

1. **Sampling distribution of $\bar{X}$:**
   - Draw $R = 10{,}000$ samples of size $n = 25$ from $N(10, 9)$
   - Plot the histogram of $\bar{X}$ values. Overlay the $N(10, 9/25)$ density
   - Compute the simulated SE. Compare with $\sigma/\sqrt{n} = 0.6$
2. **CLT stress test:**
   - Repeat with Exp(1), Uniform$(0, 1)$, and Bernoulli(0.1)
   - For each: try $n = 5, 10, 30, 100$. When does it "look Normal"?
   - Hint: Bernoulli(0.1) needs larger n than Uniform. Why?
3. **Mean vs Median:**
   - Simulate both for $N(0, 1)$ with $n = 20$. Compare their SEs
   - Now try Cauchy data. Which estimator is more stable? Why?

Homework

1. Derive the SE of $\hat{p}$ for $X_1, \ldots, X_n \sim \text{Bernoulli}(p)$ in two ways:
   (a) directly from $\text{Var}(\bar{X}) = p(1-p)/n$, and
   (b) using the Fisher information $I(p) = 1/(p(1-p))$.

2. **Simulation:** Generate $R = 5{,}000$ samples of size $n = 50$ from $\text{Poisson}(\lambda = 3)$. For each, compute $\hat{\lambda} = \bar{X}$.
   (a) Plot the histogram and verify it looks $N(3, 3/50)$.
   (b) What fraction of $\hat{\lambda}$ values fall within $\pm 1.96 \cdot \text{SE}$ of $\lambda = 3$?

3. The SD of IQ scores is $\sigma = 15$. A researcher wants the SE of the sample mean to be at most 1 point.
   (a) What sample size is needed?
   (b) If she wants $\text{SE} \leq 0.5$, how much more data does she need?

4. **Challenge:** For $X_1, \ldots, X_n \sim \text{Uniform}(0, \theta)$, the MLE is $\hat{\theta} = X_{(n)}$. Show that $\text{SE}(\hat{\theta}) = \theta \cdot n/((n+1)\sqrt{n+2})$, not $1/\sqrt{nI(\theta)}$. Why doesn't the MLE asymptotic formula work here?

Recommended Visualizations & Resources

Interactive: Seeing Theory — Basic Probability (Brown University)

`seeing-theory.brown.edu/basic-probability` — drag sliders to change n and watch the sampling distribution form in real time.

Interactive: CLT Visualization (Online Stat Book)

`onlinestatbook.com/stat_sim/sampling_dist` — choose any population shape and see the sampling distribution for different n.

Video: 3Blue1Brown — But what is the Central Limit Theorem?

Beautiful visual explanation of *why* CLT works, using convolutions. Watch after the lecture.

Video: StatQuest — Standard Error

Clear explanation of the difference between SD and SE, with worked examples.

Reading: Wasserman, "All of Statistics" — Ch. 3 & 5

Convergence, CLT, and the delta method. Concise, mathematically clean, excellent reference.

Questions?

Next: Lecture 8 — Confidence intervals and the bootstrap