

From N-grams to LSTMs

N-grams · Word2Vec · RNNs · LSTMs · Motivation for Attention

The core problem

“The cat sat on the mat”



?

Representation

How do we turn words
into numbers that
capture meaning?

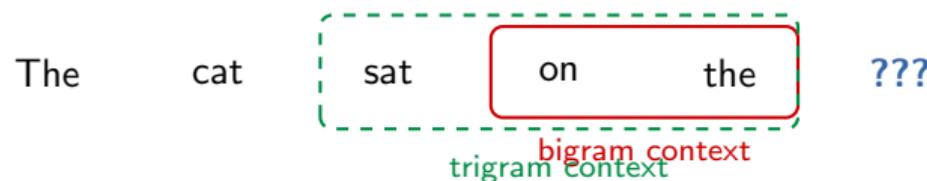
Sequence modeling

How do we capture con-
text and word order?



N-gram language models

$$P(w_t \mid w_{t-n+1}, \dots, w_{t-1}) = \frac{\text{count}(w_{t-n+1} \dots w_t)}{\text{count}(w_{t-n+1} \dots w_{t-1})}$$



Bigram:

$$P(\text{mat} \mid \text{the}) = \frac{\text{count}(\text{"the mat"})}{\text{count}(\text{"the"})}$$

Trigram:

$$P(\text{mat} \mid \text{on, the}) = \frac{\text{count}(\text{"on the mat"})}{\text{count}(\text{"on the"})}$$

Idea: estimate probabilities by counting how often word sequences appear in a large corpus

N-gram limitations

12 words apart — trigram window can't see "cat"

"The **cat** that sat on the mat near the door by the window **was** _____"

1. Fixed context window

A trigram only sees 2 previous words. Long-range dependencies are invisible.

2. Curse of dimensionality

Vocabulary $V = 50k \Rightarrow$ possible 5-grams: $50k^5 = 3 \times 10^{23}$. Most never observed.

3. Data sparsity

"armadillo aardvark" has count 0 in most corpora. Smoothing helps but doesn't solve it.

4. No generalization

Knowing "dog sat on" doesn't help predict after "cat sat on." Words are discrete symbols.

We need a way to represent words so that similar words share information.

One-hot encoding

Vocabulary

1. cat

2. dog

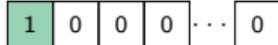
3. sat

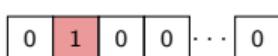
4. the

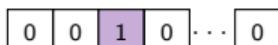
:

50k. zoo

One-hot vectors

cat = 

dog = 

sat = 

50,000 dimensions

No similarity

$$\text{cat} \cdot \text{dog} = 0$$

$$\text{cat} \cdot \text{banana} = 0$$

Every pair is equally different!

Huge & sparse

50k-dim vectors with exactly one non-zero entry

We need: dense, low-dimensional vectors
where similar words have similar representations

Word2Vec — the distributional hypothesis

“You shall know a word by the company it keeps”
— J. R. Firth, 1957

Words appearing in similar contexts have similar meanings:

“The movie was surprisingly **good** and **the audience** loved it”

“The movie was surprisingly **great** and **the audience** loved it”

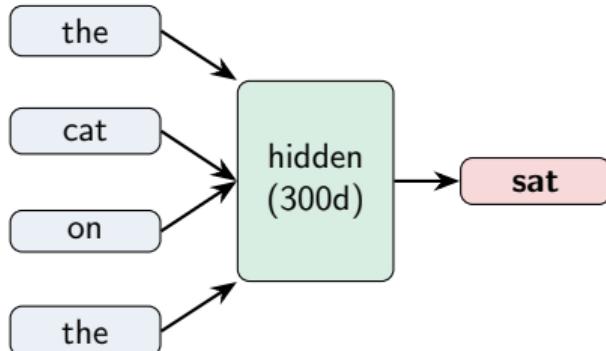
“The movie was surprisingly **excellent** and **the audience** loved it”

good, great, and excellent share context ⇒
their vectors should be close in embedding space

Word2Vec — CBOW & Skip-gram

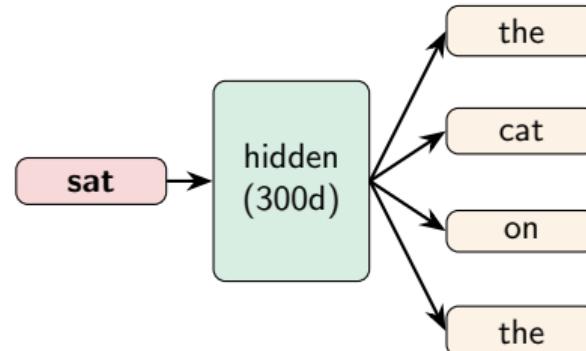
CBOW

Context → Center word



Skip-gram

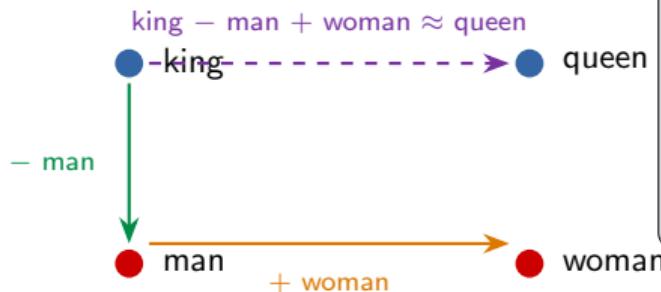
Center word → Context



Result: the hidden layer weights *are* the word embeddings — dense vectors (100–300 dims)

Word2Vec — emergent properties

Vector arithmetic in embedding space



More analogies:

Paris – France + Italy \approx Rome
bigger – big + small \approx smaller
walking – walk + swim \approx swimming

Critical limitation: one vector per word

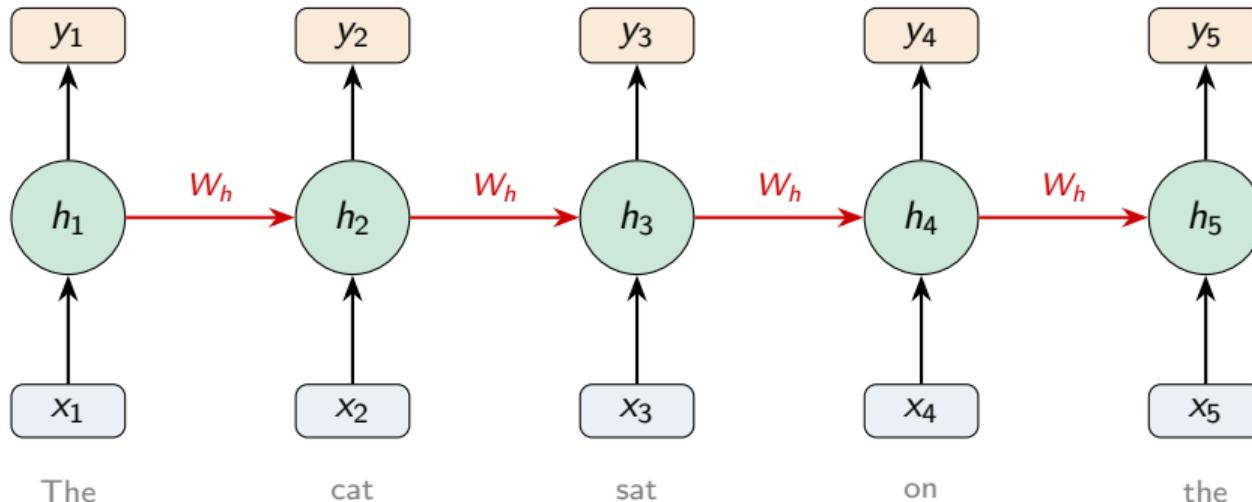
"I went to the river **bank**" vs. "I deposited money at the **bank**"

Both map to the *same* vector — no polysemy, no context-dependence

We need representations that change based on context. Enter: recurrent neural networks.

Recurrent Neural Networks (RNNs)

Unrolled RNN across time steps

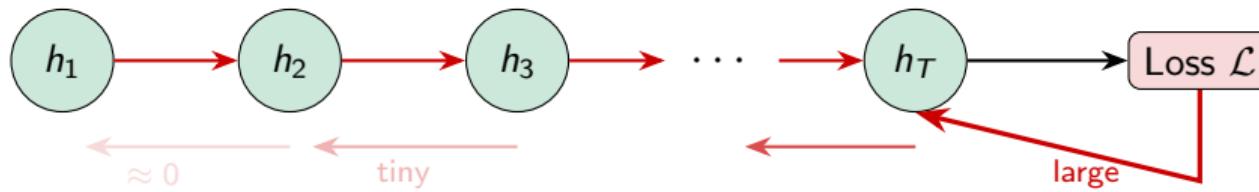


$$h_t = \tanh(W_h \cdot h_{t-1} + W_x \cdot x_t + b)$$

Key: hidden state h_t is a “memory” — no fixed window!

The vanishing gradient problem

Backpropagation through time



$$\frac{\partial \mathcal{L}}{\partial h_1} = \frac{\partial \mathcal{L}}{\partial h_T} \cdot \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}}$$

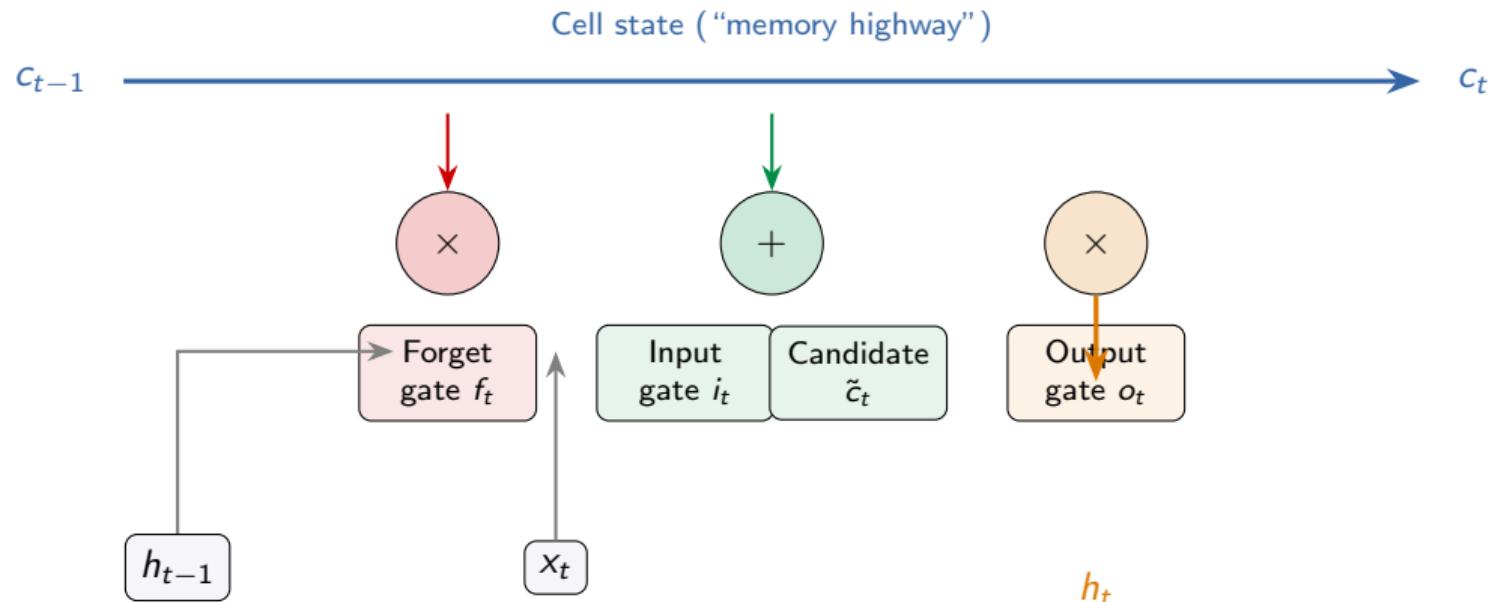
Product of many values $< 1 \Rightarrow$ gradient vanishes

“**The cat** that I saw yesterday in the garden near the old oak tree **was** cute”

The gradient from “**was**” can’t reach “**cat**” — the RNN *forgets* early tokens

Solution attempt: LSTM (1997) and GRU (2014)

LSTM — Long Short-Term Memory

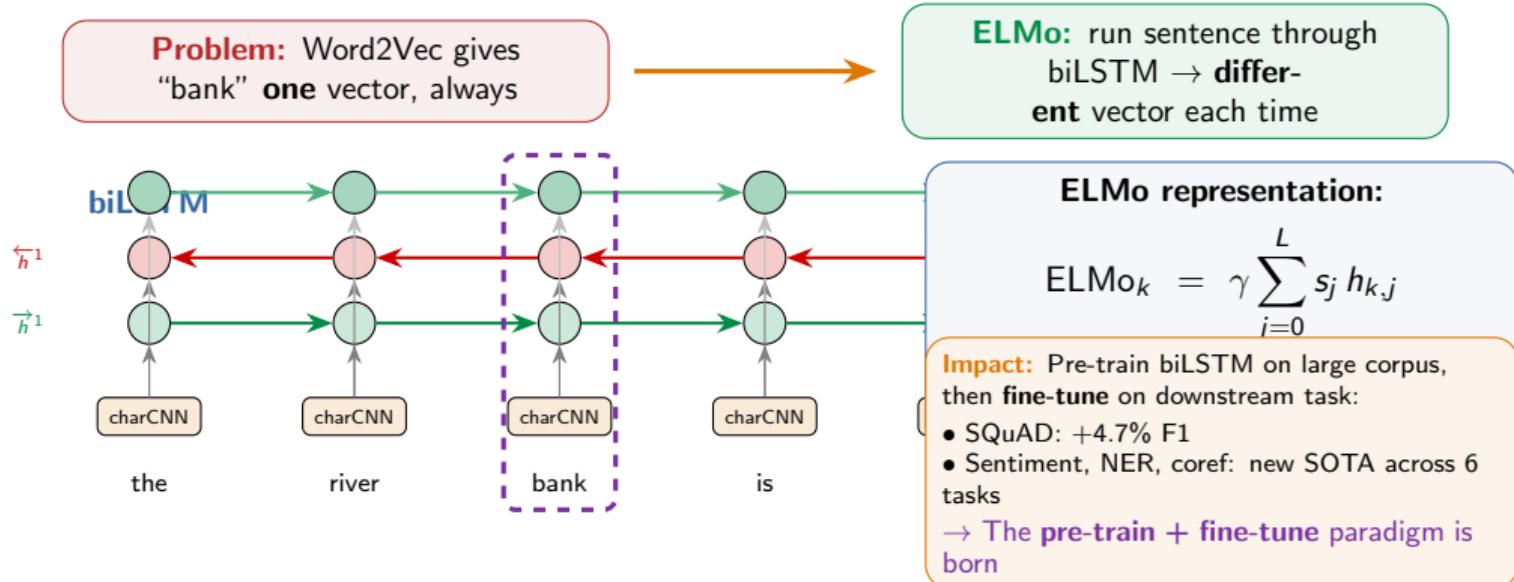


Cell state flows with minimal modification — gradient highway

Gates learn what to forget, store, and output at each step

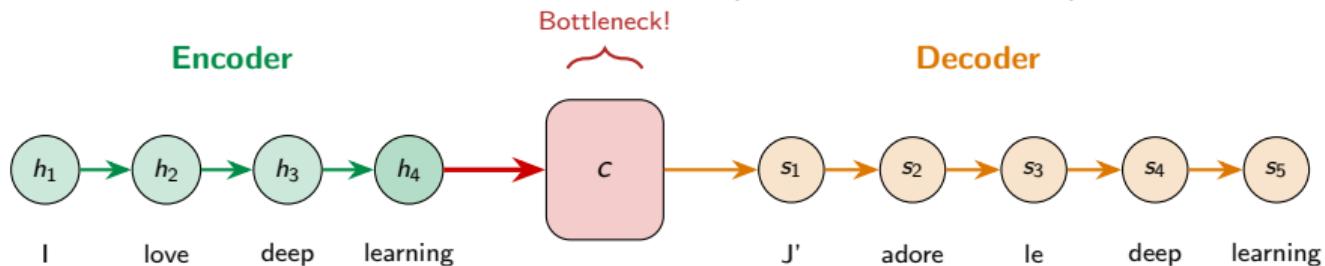
ELMo — context changes meaning

Peters et al. (2018): “Deep contextualized word representations”



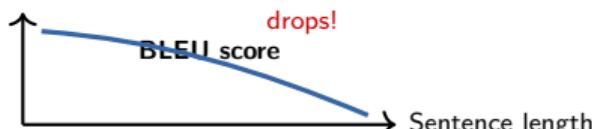
Sequence-to-sequence (Encoder–Decoder)

Machine translation with LSTMs (Sutskever et al., 2014)



The entire input sentence is compressed into a single fixed-size vector c !

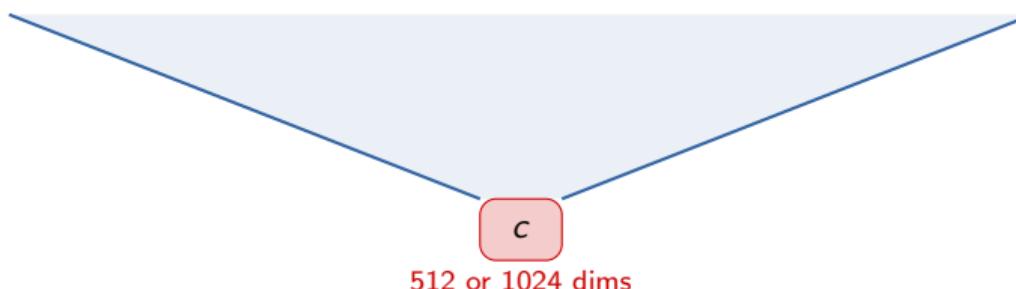
Works for short sentences, but performance degrades for long ones.



What if the decoder could look back at all encoder states, not just c ?
→ This is attention!

Why LSTMs are not enough — the bottleneck

"The quick brown fox jumped over the lazy dog that
was sleeping by the fireplace in the old cottage"



Information lost:

Early tokens get overwritten by later ones as h_t is updated sequentially

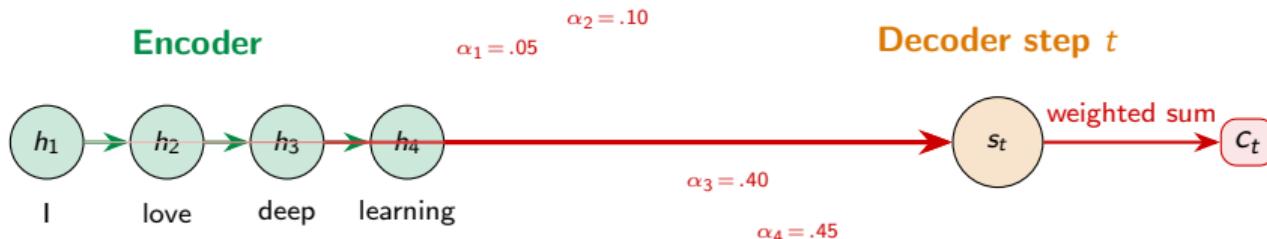
Fixed capacity:

Whether the input is 5 tokens or 500, it must fit in the same vector

Fundamental issue: you can't losslessly compress arbitrary-length sequences into a fixed-size vector

Attention — looking back at the source

Bahdanau et al. (2015): “Jointly Learning to Align and Translate”



$$\alpha_{ti} = \text{softmax}(e_{ti}) \quad \text{where} \quad e_{ti} = a(s_{t-1}, h_i) \quad c_t = \sum_i \alpha_{ti} h_i$$

No more bottleneck!

Each decoder step gets a **different** weighted view of all encoder states

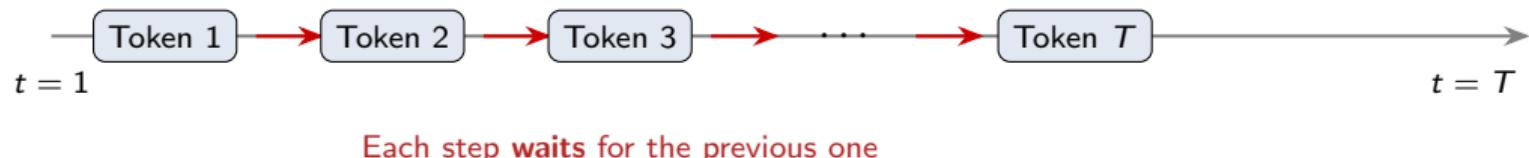
Alignment is learned

The model learns which source words to focus on for each output word

What if we applied this idea to every token attending to *all other tokens*? → **Self-attention**

Why LSTMs are not enough — sequential processing

LSTM: must process tokens one at a time



What if we could process all tokens in parallel?



LSTM: $O(T)$ sequential steps
Can't utilize GPU parallelism
Training is *slow*

Transformer: $O(1)$ parallel steps
Fully utilizes GPU cores
Training is *fast*

Why LSTMs are not enough — long-range dependencies

“The **trophy** didn’t fit in the **suitcase** because **it** was too _____”

If **it** = **trophy**:

“it was too **big**”

If **it** = **suitcase**:

“it was too **small**”

Path length from “it” to its referent:

LSTM: $O(n)$ steps

Information must pass through
every intermediate hidden state

Signal degrades over distance

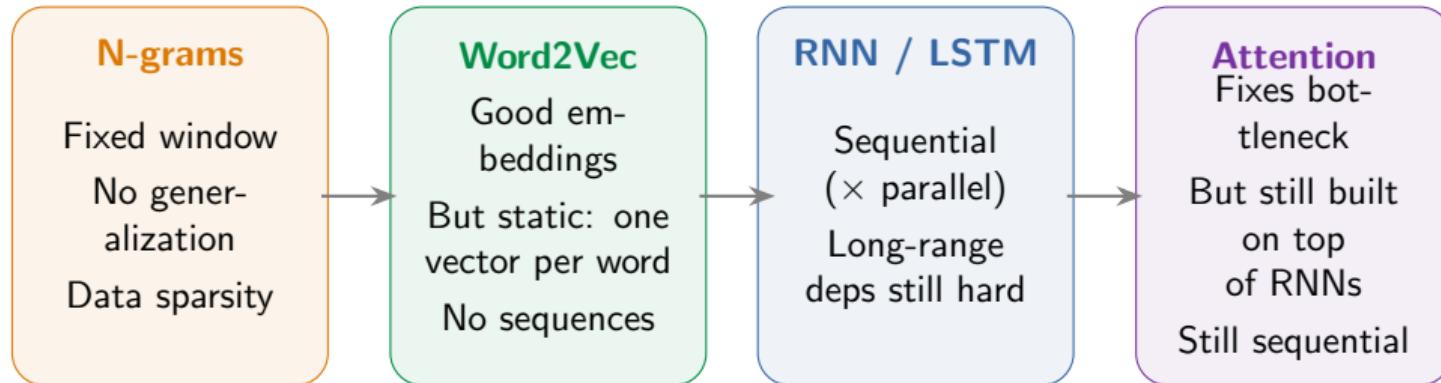
Attention: $O(1)$ steps

Direct connection be-
tween any two tokens

No degradation over distance

Attention: “let me directly look at any token I need” — regardless of distance

The stage is set



What if attention is **all you need**? Drop the RNN entirely.

Parallel processing + direct access to all positions + contextual representations

⇒ **Self-Attention** and the **Transformer** (Vaswani et al., 2017)

Word2Vec — training tricks

$$\text{Skip-gram objective: } \max_{\theta} \sum_{t=1}^T \sum_{\substack{-c \leq j \leq c \\ j \neq 0}} \log P(w_{t+j} \mid w_t)$$

Problem: full softmax

Normalizing over all V words per update — too slow!

Negative sampling

Sample k random “negatives”
Binary: real context vs. noise

$$\log \sigma(v'_{w_0} \cdot v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n} [\log \sigma(-v'_{w_i} \cdot v_{w_I})] \quad k = 5-20, \quad P_n \propto f(w)^{3/4}$$

Subsampling

Downsample frequent words (“the”, “a”, “is”)

Window size

Large → semantic similarity
Small → syntactic similarity

Dimensionality

Typical: 100–300 dims
More data → more dims

Word2Vec trained on Google News: 3M vocabulary, 300-dim vectors, 100B tokens

Beyond Word2Vec — GloVe & FastText

GloVe (Pennington et al., 2014)

Co-occurrence matrix X_{ij} + factorize:
 $w_i^T \tilde{w}_j + b_i + \tilde{b}_j = \log X_{ij}$

Global statistics + local context

FastText (Bojanowski et al., 2017)

Words as bags of char n-grams:
“where” → {wh, whe, her, ere, re}

Handles OOV words & morphology

GloVe

FastText

Word2Vec			
Training signal	Local context	Global co-occur.	Local + subword
OOV words	No	No	Yes
Morphology	No	No	Yes
Speed	Fast	Fast	Medium

All three share the same limitation: static embeddings — one vector per word regardless of context

Further reading

N-grams & Statistical NLP

- Jurafsky & Martin, *Speech and Language Processing*, Ch. 3 — N-gram Language Models

Word Embeddings

- Mikolov et al. (2013), “Efficient Estimation of Word Representations in Vector Space”
- Mikolov et al. (2013), “Distributed Representations of Words and Phrases” (neg. sampling)

RNNs, Attention & Sequence Models

- Hochreiter & Schmidhuber (1997), “Long Short-Term Memory”
- Sutskever et al. (2014), “Sequence to Sequence Learning with Neural Networks”
- Bahdanau et al. (2015), “Neural Machine Translation by Jointly Learning to Align and Translate”
- Luong et al. (2015), “Effective Approaches to Attention-based Neural Machine Translation”
- Peters et al. (2018), “Deep contextualized word representations” (ELMo)

Questions?

Next: The Transformer Architecture