

From N-grams to LSTMs

N-grams · Word2Vec · RNNs · LSTMs · Motivation for Attention

Outline

Statistical language models

Word embeddings

Recurrent models

Seq2seq & attention

Beyond Word2Vec & wrap-up

The core problem

“The cat sat on the mat”



Representation

How do we turn words into numbers that capture meaning?


Sequence modeling

How do we capture context and word order?



N-gram language models

$$P(w_t \mid w_{t-n+1}, \dots, w_{t-1}) = \frac{\text{count}(w_{t-n+1} \cdots w_t)}{\text{count}(w_{t-n+1} \cdots w_{t-1})}$$



Bigram:

$$P(\text{mat} \mid \text{the}) = \frac{\text{count}(\text{"the mat"})}{\text{count}(\text{"the"})}$$

Trigram:

$$P(\text{mat} \mid \text{on, the}) = \frac{\text{count}(\text{"on the mat"})}{\text{count}(\text{"on the"})}$$

Idea: estimate probabilities by counting how often word sequences appear in a large corpus

N-gram limitations

12 words apart — trigram window can't see "cat"

"The **cat** that sat on the mat near the door by the window **was** _____"

1. Fixed context window

A trigram only sees 2 previous words. Long-range dependencies are invisible.

2. Curse of dimensionality

Vocabulary $V = 50k \Rightarrow$ possible 5-grams: $50k^5 = 3 \times 10^{23}$. Most never observed.

3. Data sparsity

"armadillo aardvark" has count 0 in most corpora. Smoothing helps but doesn't solve it.

4. No generalization

Knowing "dog sat on" doesn't help predict after "cat sat on." Words are discrete symbols.

We need a way to represent words so that similar words share information.

One-hot encoding

Vocabulary

1. cat

2. dog

3. sat

4. the

⋮

50k. zoo



One-hot vectors

cat =

1	0	0	0	⋯	0
---	---	---	---	---	---

dog =

0	1	0	0	⋯	0
---	---	---	---	---	---

sat =

0	0	1	0	⋯	0
---	---	---	---	---	---


50,000 dimensions

No similarity

$$\text{cat} \cdot \text{dog} = 0$$

$$\text{cat} \cdot \text{banana} = 0$$

Every pair is equally different!

Huge & sparse

50k-dim vectors with exactly one non-zero entry

We need: dense, low-dimensional vectors
where similar words have similar representations

Word2Vec — the distributional hypothesis

“You shall know a word by the company it keeps”

— J. R. Firth, 1957

Words appearing in similar contexts have similar meanings:

“The movie was surprisingly **good** and the audience loved it”

“The movie was surprisingly **great** and the audience loved it”


“The movie was surprisingly **excellent** and the audience loved it”

good, **great**, and **excellent** share context $\Rightarrow$
their vectors should be close in embedding space

Word2Vec — CBOW & Skip-gram


CBOW

Context → Center word



Skip-gram


Center word → Context



Result: the hidden layer weights *are* the word embeddings — dense vectors (100–300 dims)

Word2Vec — emergent properties

Vector arithmetic in embedding space



More analogies:

Paris - France + Italy $\approx$ Rome
bigger - big + small $\approx$ smaller
walking - walk + swim $\approx$ swimming


Critical limitation: one vector per word

“I went to the river **bank**” vs. “I deposited money at the **bank**”
Both map to the *same* vector — no polysemy, no context-dependence

We need representations that change based on context. Enter: recurrent neural networks.

Recurrent Neural Networks (RNNs)

Unrolled RNN across time steps




$$h_t = \tanh(W_h \cdot h_{t-1} + W_x \cdot x_t + b)$$

Key: hidden state h_t is a “memory” — no fixed window!

The vanishing gradient problem

Backpropagation through time



$$\frac{\partial \mathcal{L}}{\partial h_1} = \frac{\partial \mathcal{L}}{\partial h_T} \cdot \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}}$$


Product of many values $< 1 \Rightarrow$ gradient vanishes

“**The cat** that I saw yesterday in the garden near the old oak tree **was** cute”

The gradient from “was” can’t reach “cat” — the RNN *forgets* early tokens

Solution attempt: LSTM (1997) and GRU (2014)

LSTM — Long Short-Term Memory



Cell state flows with minimal modification — gradient highway

Gates learn what to forget, store, and output at each step

LSTM — gate equations

Forget gate: $f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$

Input gate: $i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$ $\tilde{c}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$

Cell update: $c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$

Output gate: $o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$ $h_t = o_t \odot \tanh(c_t)$

Why it works: the cell state update is **additive** ($c_t = f_t \odot c_{t-1} + \dots$), not multiplicative like vanilla RNNs. Gradients flow through the $c_{t-1} \rightarrow c_t$ path with minimal decay — solving the vanishing gradient problem.


ELMo — context changes meaning

Peters et al. (2018): “Deep contextualized word representations”

Problem: Word2Vec gives “bank” **one** vector, always

ELMo: run sentence through biLSTM → **different** vector each time

biLSTM



ELMo representation:

$$\text{ELMo}_k = \gamma \sum_{j=0}^L s_j h_{k,j}$$


Impact: Pre-train biLSTM on large corpus, then **fine-tune** on downstream task:

- SQuAD: +4.7% F1
- Sentiment, NER, coref: new SOTA across 6 tasks

→ The **pre-train + fine-tune** paradigm is born


Sequence-to-sequence (Encoder-Decoder)

Machine translation with LSTMs (Sutskever et al., 2014)



The entire input sentence is compressed into a single fixed-size vector c !

Works for short sentences, but performance degrades for long ones.




What if the decoder could **look back** at all encoder states, not just c ?

→ This is attention!

Why LSTMs are not enough — the bottleneck

“The quick brown fox jumped over the lazy dog that was sleeping by the fireplace in the old cottage”



Information lost:

Early tokens get overwritten by later ones as h_t is updated sequentially


Fixed capacity:

Whether the input is 5 tokens or 500, it must fit in the same vector

Fundamental issue: you can't losslessly compress arbitrary-length sequences into a fixed-size vector

Attention — looking back at the source

Bahdanau et al. (2015): “Jointly Learning to Align and Translate”



$$e_{ti} = v_a^\top \tanh(W_a s_{t-1} + U_a h_i) \quad \alpha_{ti} = \text{softmax}(e_{ti}) \quad c_t = \sum_i \alpha_{ti} h_i$$

No more bottleneck!

Each decoder step gets a **different** weighted view of all encoder states

Alignment is learned

The model learns which source words to focus on for each output word

What if we applied this idea to every token attending to *all other* tokens? → **Self-attention**

Why LSTMs are not enough — sequential processing

LSTM: must process tokens one at a time



Each step **waits** for the previous one

What if we could process all tokens in parallel?



All processed **simultaneously** on GPU

LSTM: $O(T)$ sequential steps
Can't utilize GPU parallelism
Training is *slow*

Transformer: $O(1)$ parallel steps
Fully utilizes GPU cores
Training is *fast*

Why LSTMs are not enough — long-range dependencies

“The **trophy** didn't fit in the **suitcase** because **it** was too _____”

If **it** = **trophy**:
“it was too **big**”

If **it** = **suitcase**:
“it was too **small**”

Path length from “it” to its referent:

LSTM: $O(n)$ steps

Information must pass through
every intermediate hidden state

Signal degrades over distance


Attention: $O(1)$ steps

Direct connection be-
tween any two tokens

No degradation over distance

Attention: “let me directly look at any token I need” — regardless of distance

The stage is set



What if attention is **all you need**? Drop the RNN entirely.
Parallel processing + direct access to all positions + contextual representations
⇒ **Self-Attention** and the **Transformer** (Vaswani et al., 2017)

Word2Vec — training tricks

Skip-gram objective:
$$\max_{\theta} \sum_{t=1}^T \sum_{\substack{-c \leq j \leq c \\ j \neq 0}} \log P(w_{t+j} \mid w_t)$$

Problem: full softmax

Normalizing over all V words
per update — too slow!



Negative sampling

Sample k random “negatives”
Binary: real context vs. noise

$$\log \sigma(v'_{w_0} \cdot v_{w_l}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n} [\log \sigma(-v'_{w_i} \cdot v_{w_l})] \quad k=5-20, \quad P_n \propto f(w)^{3/4}$$

Subsampling

Downsample frequent words
("the", "a", "is")

Window size

Large → semantic similarity
Small → syntactic similarity

Dimensionality

Typical: 100–300 dims
More data → more dims

Word2Vec trained on Google News: 3M
vocabulary, 300-dim vectors, 100B tokens

Beyond Word2Vec — GloVe & FastText

GloVe (Pennington et al., 2014)

Co-occurrence matrix X_{ij} + factorize:

$$w_i^T \tilde{w}_j + b_i + \tilde{b}_j = \log X_{ij}$$

Global statistics + local context

FastText (Bojanowski et al., 2017)

Words as bags of char n -grams:

<where> $\rightarrow$ <wh, whe, her, ere, re>

Handles OOV words & morphology

	Word2Vec	GloVe	FastText
Training signal	Local context	Global co-occur.	Local + subword
OOV words	No	No	Yes
Morphology	No	No	Yes
Speed	Fast	Fast	Medium

All three share the same limitation: static embeddings — one vector per word regardless of context

Further reading

N-grams & Statistical NLP

- Jurafsky & Martin, *Speech and Language Processing*, Ch.3 — N-gram Language Models
- Chen & Goodman (1999), “An Empirical Study of Smoothing Techniques”

Word Embeddings

- Mikolov et al. (2013), “Efficient Estimation of Word Representations in Vector Space”
- Mikolov et al. (2013), “Distributed Representations of Words and Phrases” (neg. sampling)
- Pennington et al. (2014), “GloVe: Global Vectors for Word Representation”
- Bojanowski et al. (2017), “Enriching Word Vectors with Subword Information”

RNNs, Attention & Sequence Models

- Hochreiter & Schmidhuber (1997), “Long Short-Term Memory”
- Sutskever et al. (2014), “Sequence to Sequence Learning with Neural Networks”
- Bahdanau et al. (2015), “Neural Machine Translation by Jointly Learning to Align and Translate”
- Luong et al. (2015), “Effective Approaches to Attention-based Neural Machine Translation”
- Peters et al. (2018), “Deep contextualized word representations” (ELMo)

Questions?

Next: The Transformer Architecture