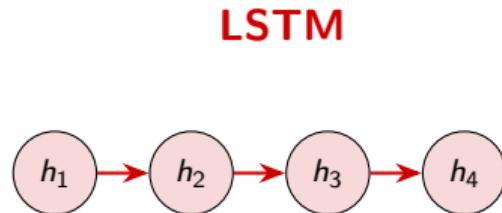


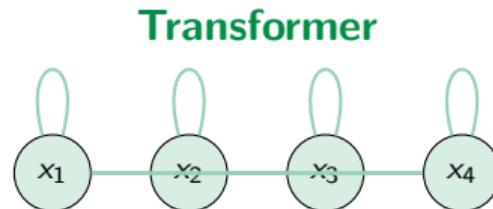
# The Transformer Architecture

Self-Attention · Multi-Head Attention · Positional Encoding · Encoder–Decoder

# The key idea



Sequential:  $O(n)$  path



Parallel:  $O(1)$  path

**"Attention Is All You Need"** (Vaswani et al., 2017)

**Parallel**

Process all tokens simultaneously

**Direct access**

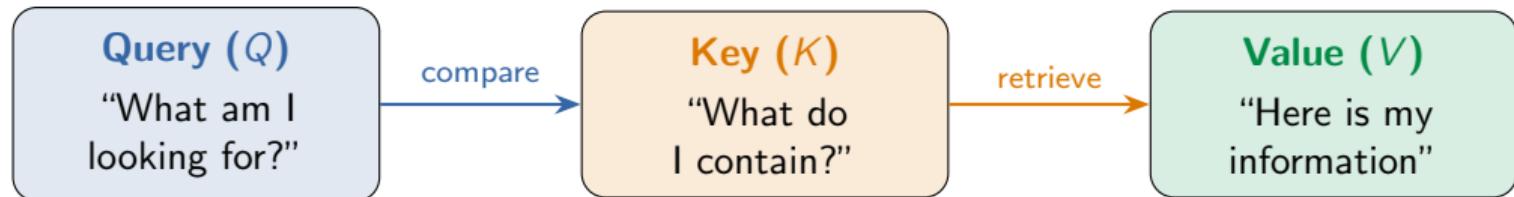
Any token can attend to any other

**Contextual**

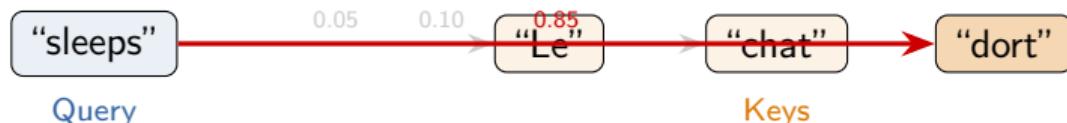
Representations change with context

# Attention — intuition

Think of it as a soft dictionary lookup

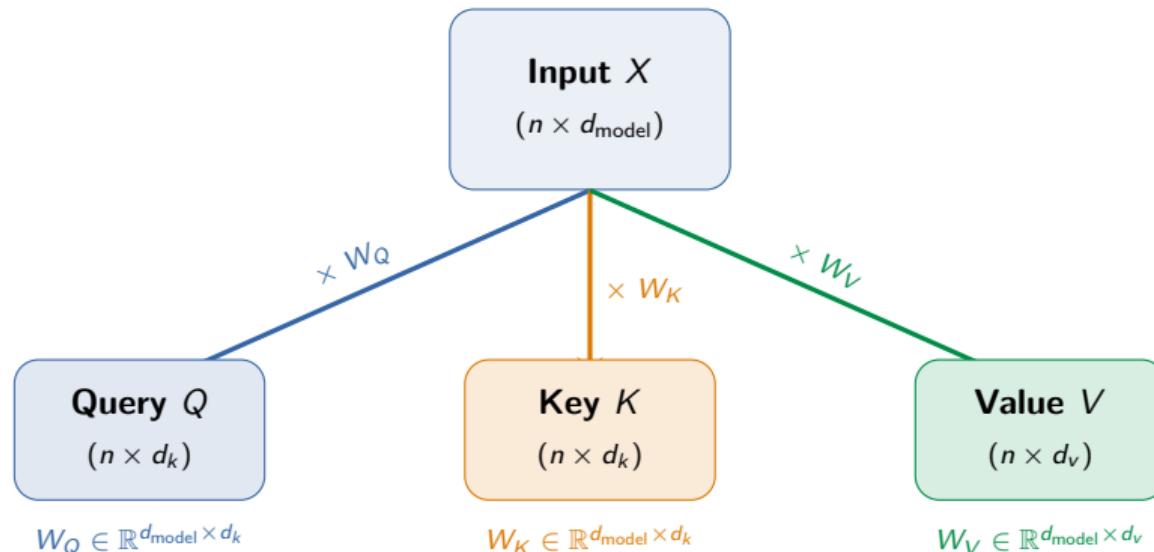


Example: translating “Le chat dort” → “The cat sleeps”



Output = weighted sum of values:  $0.05 \cdot v_{\text{Le}} + 0.10 \cdot v_{\text{chat}} + 0.85 \cdot v_{\text{dort}}$

# Q, K, V — where do they come from?



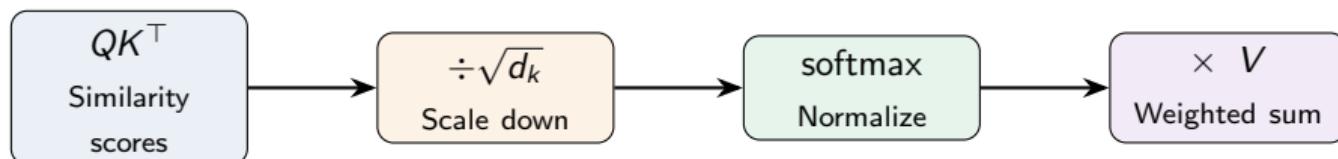
**Same input, three different views — the model learns**  
what to query, what to advertise, and what to send

**Self-attention:**  $Q, K, V$  all from same  $X$

**Cross-attention:**  $Q$  from decoder,  
 $K, V$  from encoder

## Scaled dot-product attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V$$



Why  $\sqrt{d_k}$ ?

Without scaling, dot products grow with  $d_k$ , pushing softmax into saturation (near 0 or 1) where gradients vanish

**Dimensions:**

$Q: (n \times d_k)$     $K: (m \times d_k)$   
 $V: (m \times d_v)$   
 $QK^\top: (n \times m)$    Output:  $(n \times d_v)$

Each output row is a **weighted average** of value vectors, where weights come from query-key similarity

# Attention — worked example

Tokens: The cat sat

	$QK^\top / \sqrt{d_k}$			Weights			
	The	cat	sat				
The	1.2	0.5	0.3	.49	.28	.23	
cat	0.4	2.1	0.8	<span>softmax</span>	.12	.63	.25
sat	0.2	0.9	1.5		.15	.30	.55

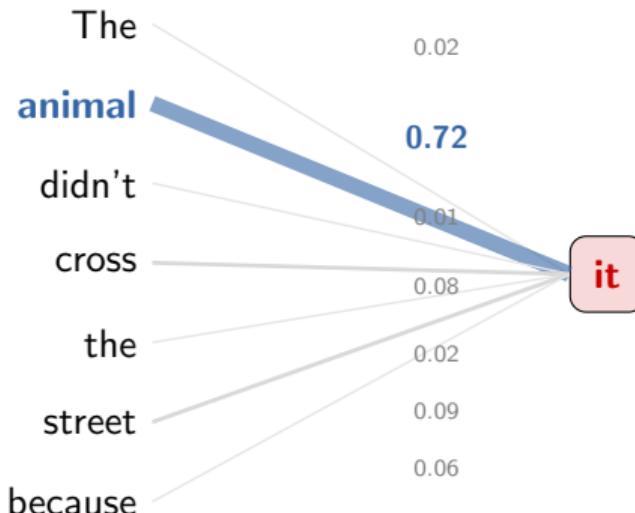
Row for “cat”: attends 63% to itself, 25% to “sat”, 12% to “The”

Output for “cat” =  
 $0.12 \cdot v_{\text{The}} + 0.63 \cdot v_{\text{cat}} + 0.25 \cdot v_{\text{sat}}$

Each output is a **context-aware** representation — unlike Word2Vec, the same word gets different embeddings in different contexts

# Attention learns meaningful relationships

"The **animal** didn't cross the street because **it** was too tired"



## Coreference resolved!

"it" attends most strongly to "animal" — the model learned that "it" refers to "animal"

## No explicit rule

This emerges from training — the model learns which tokens are relevant for each query

Different attention heads specialize in different relationships (see next slides)

# Self-attention vs. cross-attention

## Self-Attention

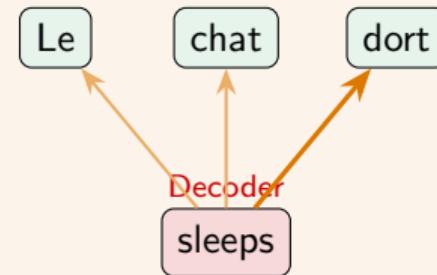


$Q, K, V$  all come from  
the **same** sequence

Used in: encoder, decoder (masked)

## Cross-Attention

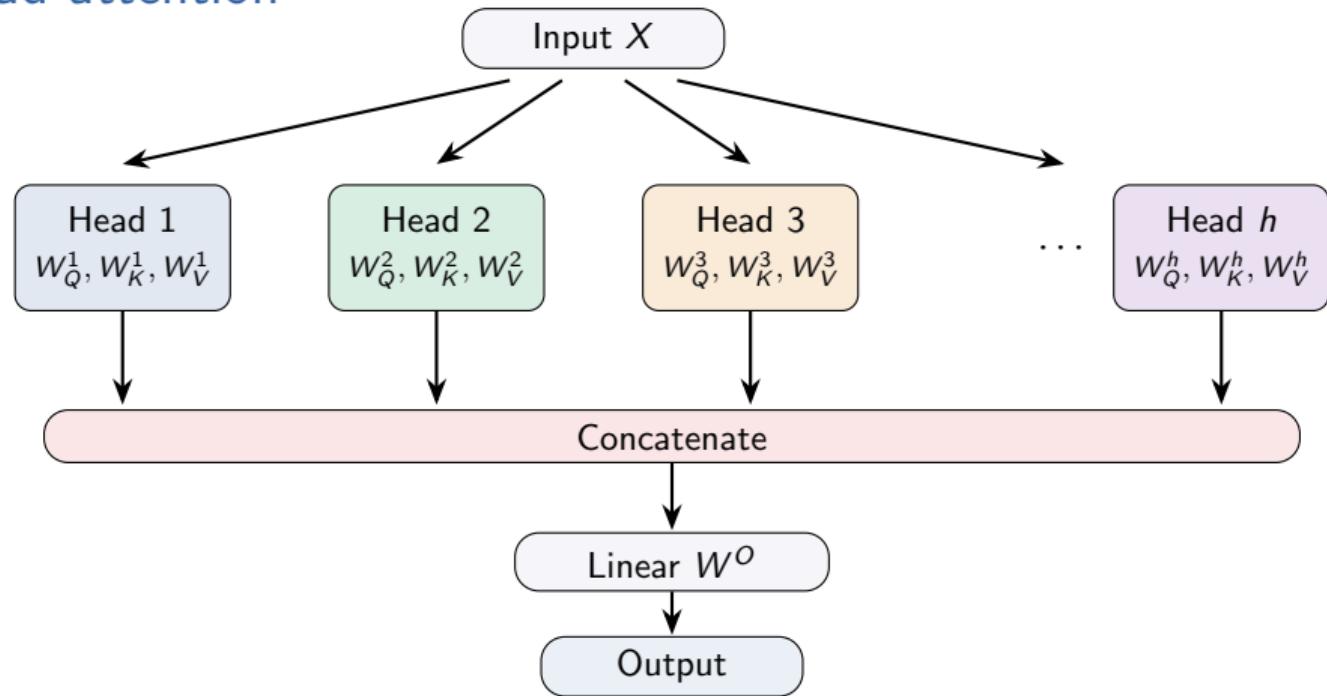
Encoder



$Q$  from decoder  
 $K, V$  from encoder

Used in: decoder (enc-dec models)

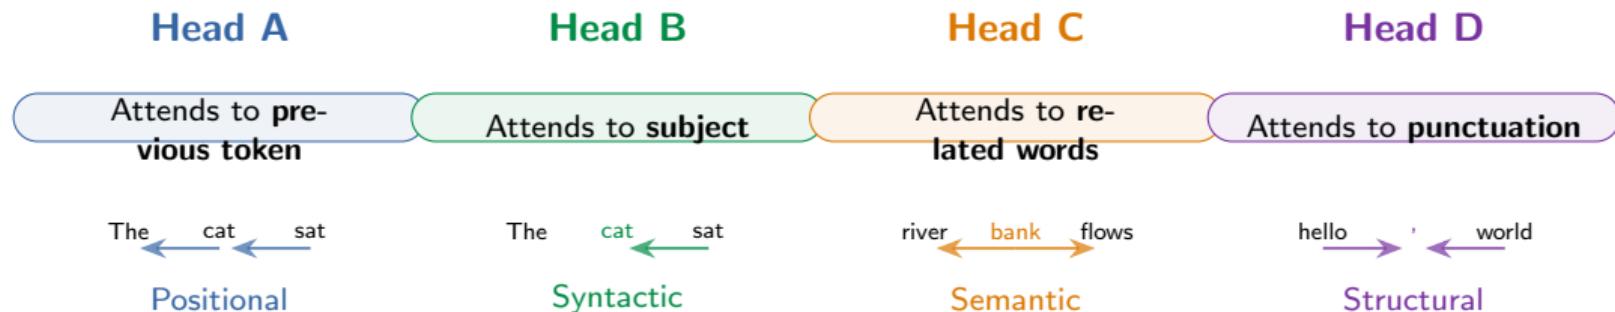
## Multi-head attention



$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$
$$\text{where } \text{head}_i = \text{Attention}(QW_Q^i, KW_K^i, VW_V^i)$$

# What different heads learn

Different heads specialize in different types of relationships:

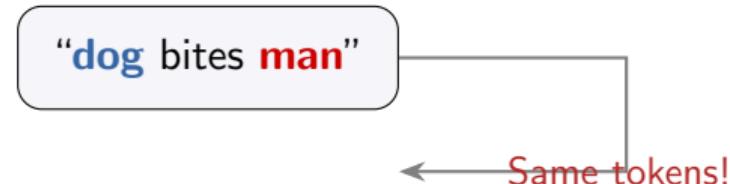


**Multiple heads** = multiple “perspectives” on the same input.

Typical:  $h = 8$  (BERT-base) or  $h = 12\text{--}96$  (larger models).  $d_k = d_{\text{model}}/h$

Total compute is the same as single-head attention with full  $d_{\text{model}}$ , since each head uses  $d_k = d_{\text{model}}/h$

## Positional encoding — why we need it



**Problem:** Self-attention computes dot products between token pairs.

It's **permutation-invariant** — swapping token order doesn't change the output!

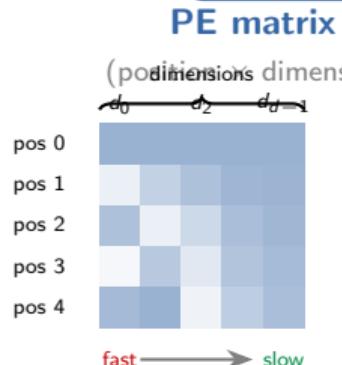
But word order matters: these two sentences mean very different things.

**Solution:** Add **positional information** to the input embeddings

$$\text{input}_i = \text{token\_embedding}_i + \text{position\_encoding}_i;$$

# Sinusoidal positional encoding

$$\begin{aligned} \text{PE}(\text{pos}, 2i) &= \sin\left(\frac{\text{pos}}{10000^{2i/d}}\right) \\ \text{PE}(\text{pos}, 2i + 1) &= \cos\left(\frac{\text{pos}}{10000^{2i/d}}\right) \end{aligned}$$



## Unique pattern

Each position gets a distinct encoding

## Relative positions

$\text{PE}_{\text{pos}+k}$  can be expressed as linear function of  $\text{PE}_{\text{pos}}$

## Multi-scale

Low dims = fine position, high dims = coarse

## No learned params

Fixed, deterministic, works for any length

$$\text{input}_i = \text{embedding}(x_i) + \text{PE}(i) \quad (\text{element-wise addition})$$

# Modern positional encodings

## Sinusoidal

Vaswani et al., 2017

Added to embeddings

Absolute position

Fixed (not learned)

Used by: original Transformer

## RoPE

Su et al., 2021

Rotates  $Q$  and  $K$  vectors

Relative position

No extra parameters

Used by: LLaMA, Mistral

## ALiBi

Press et al., 2022

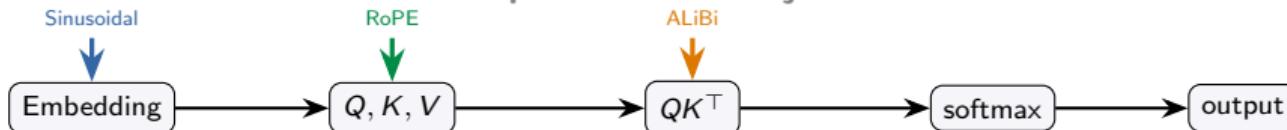
Bias in attention scores

Linear distance penalty

No extra parameters

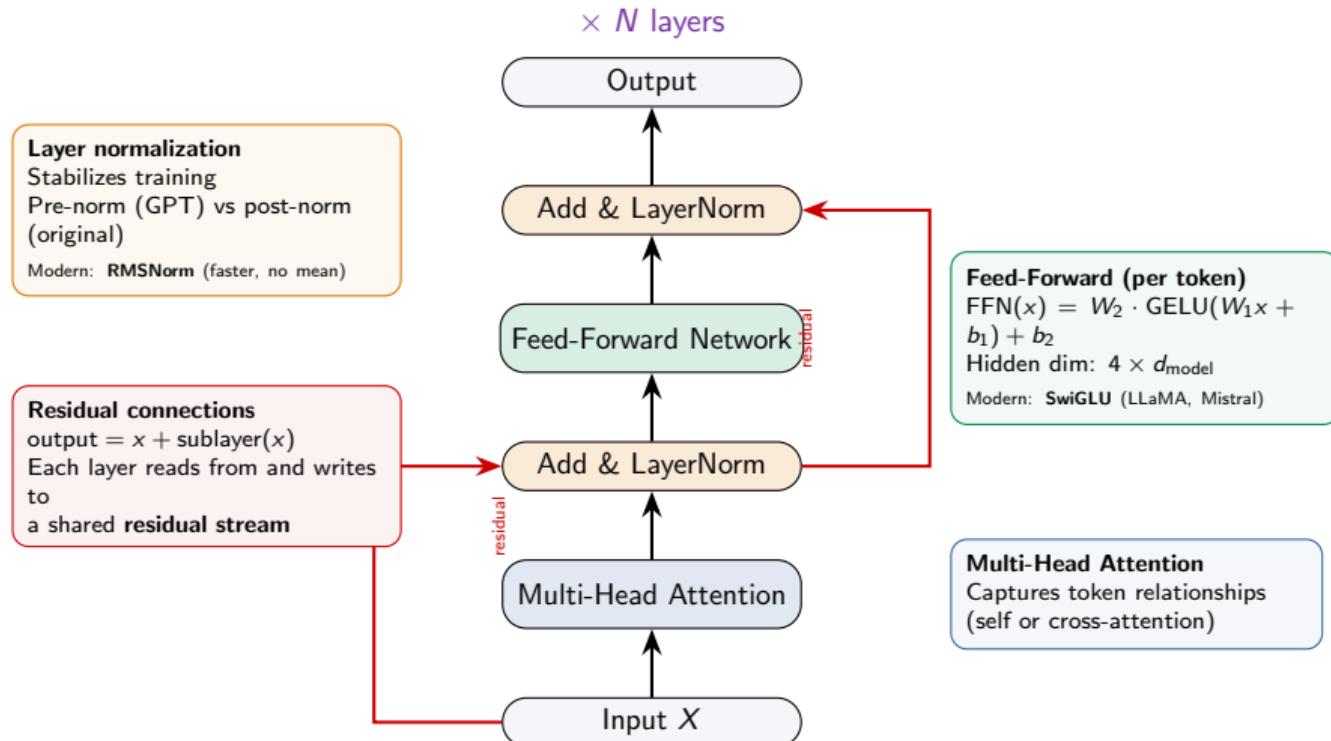
Used by: BLOOM, MPT

Where position info is injected:

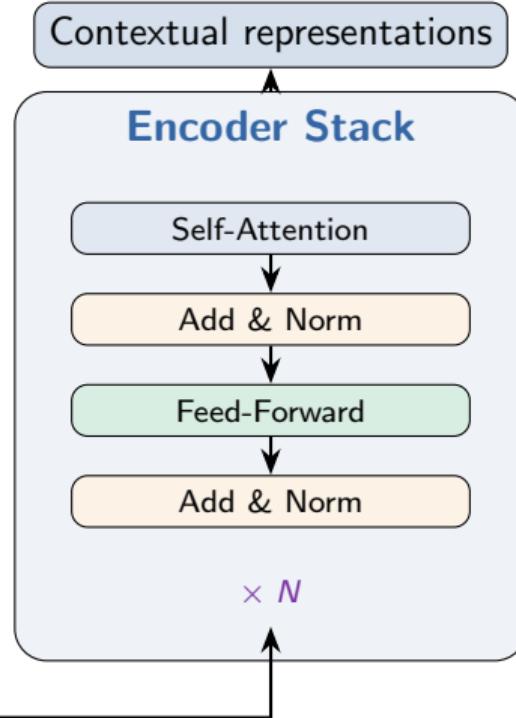
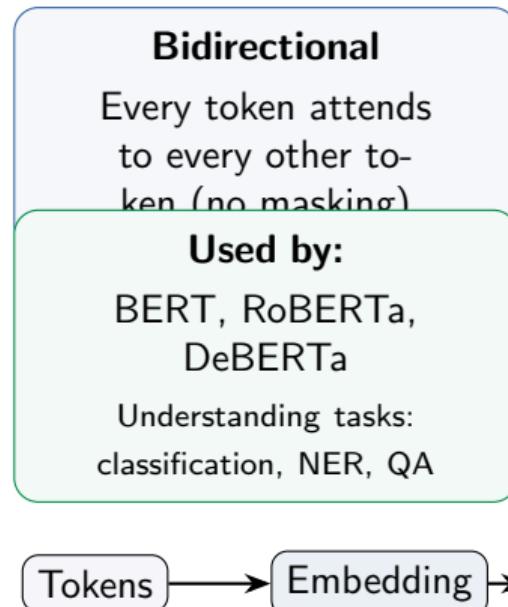


**RoPE** is the most popular today — it naturally encodes relative position and supports length extrapolation with techniques like YaRN

# The Transformer block

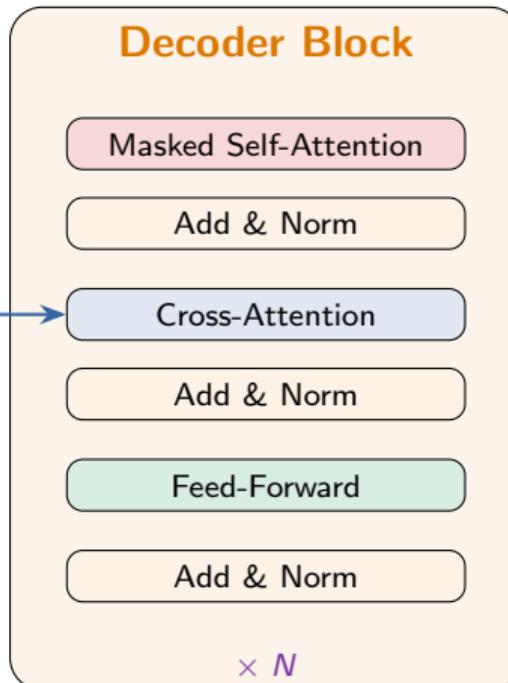


# The Transformer encoder



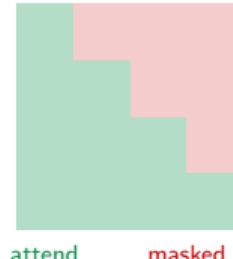
# The Transformer decoder

From encoder  
( $K, V$ )



**Decoder-only** models (GPT) skip cross-attention — only masked self-attention + FFN

## Causal Mask



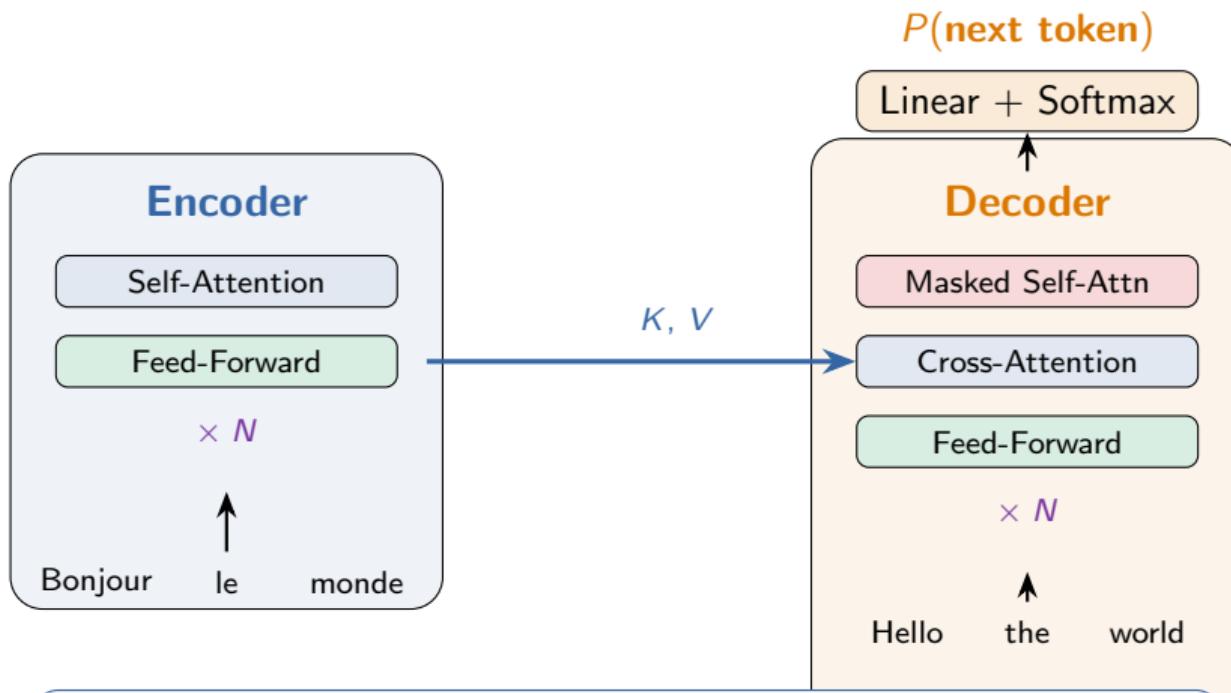
Token  $t$  can only attend to tokens  $< t$

## Used by:

GPT (decoder-only)  
T5 decoder  
Generation tasks

Dropout ( $p=0.1$ ) applied after attention weights and after each sublayer during training

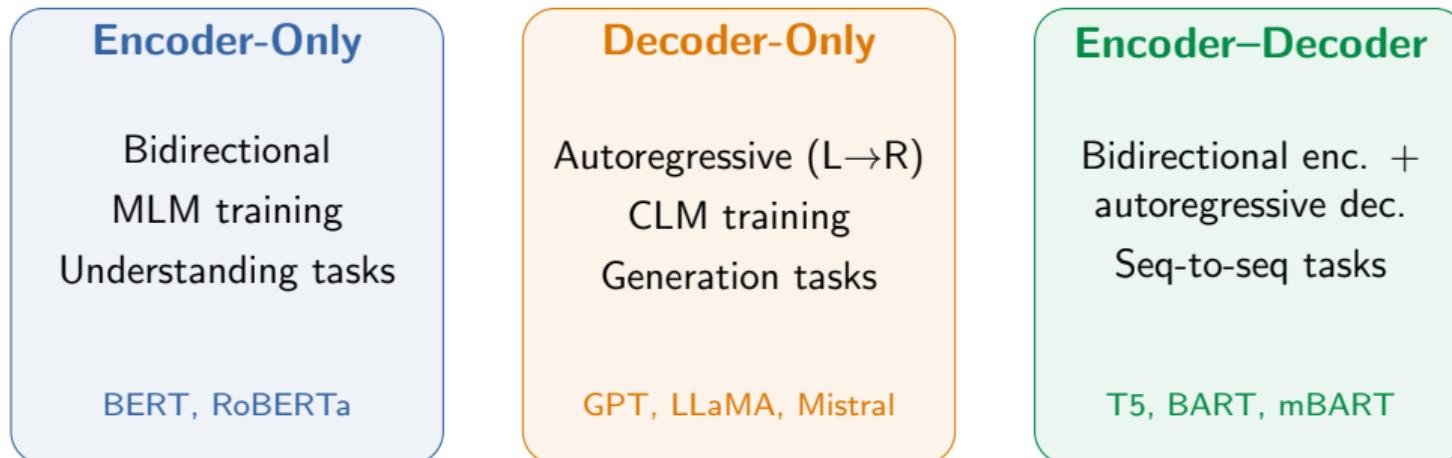
# Encoder–Decoder architecture



**Used for:** machine translation, summarization, question answering

**Models:** T5, BART, mBART, original Transformer

# Three Transformer architectures



	Enc-Only	Dec-Only	Enc-Dec
Direction	Bidirectional	Left-to-right	Both
Best for	Classification, NER	Text generation	Translation, summary
Today's trend	Less common	<b>Dominant</b>	Niche

The trend since GPT-3 (2020): decoder-only models at scale can do almost everything, including understanding tasks.

# How Transformers learn

## Causal LM (decoder)

"The cat sat on the \_\_\_\_\_"



Transformer



softmax over vocab

$$P(\text{"mat"}) = 0.42$$

$$\mathcal{L} = - \sum_{t=1}^T \log P(x_t \mid x_{<t})$$

Cross-entropy on **every** position

Teacher forcing: use true tokens  
as input, not model's predictions

## Masked LM (encoder)

"The [MASK] sat on the mat"



Transformer



softmax over vocab

$$P(\text{"cat"}) = 0.67$$

$$\mathcal{L} = - \sum_{i \in \mathcal{M}} \log P(x_i \mid x_{\setminus \mathcal{M}})$$

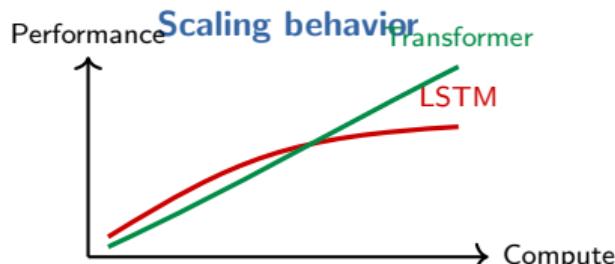
Loss only on **masked** positions

15% of tokens masked: 80% [MASK],  
10% random, 10% unchanged

Since GPT-3: **causal LM is the dominant paradigm** —  
one training objective, scales to any task via prompting

# Why Transformers won

	LSTM	Transformer
Parallelizable	No ( $h_t$ needs $h_{t-1}$ )	Yes (all tokens at once)
Long-range path	$O(n)$ steps	$O(1)$ (direct attention)
Context	Limited by hidden size	Full context window
Scaling	Diminishing returns	Log-linear improvement
Training speed	Slow (sequential)	Fast (GPU-friendly)



**The scaling insight:**  
Transformers reliably improve with  
more data, parameters, and compute  
This enabled the LLM revolution:  
GPT-3, PaLM, LLaMA, Claude, ...

# Computational complexity

## Self-Attention

$O(n^2 \cdot d)$  time

$O(n^2)$  memory

( $QK^\top$  is an  $n \times n$  matrix)

## Feed-Forward

$O(n \cdot d^2)$  time

$O(d^2)$  memory

(two  $d \times 4d$  weight matrices)

Which dominates?

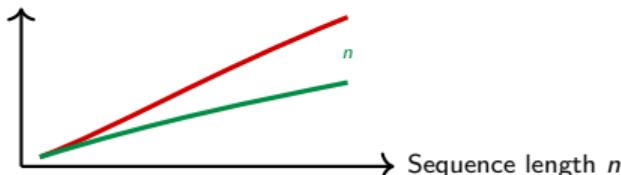
### Short context ( $n < d$ )

FFN dominates

E.g.,  $n=512, d=4096$

Attention:  $0.5M \cdot d$  vs FFN:  $512 \cdot 16M$

### The quadratic wall



### Long context ( $n > d$ )

**Attention dominates**

E.g.,  $n=128K, d=4096$

#### Doubling sequence length:

- Attention:  $4\times$  more compute
- KV cache memory:  $2\times$  more
- Attention matrix:  $4\times$  more memory

128K tokens  $\rightarrow$  16 **billion** entries per head per layer

This quadratic bottleneck motivates FlashAttention, sparse patterns, and state space models

## Transformer dimensions in practice

Model	Layers	$d_{\text{model}}$	Heads	$d_{\text{ff}}$	Params
BERT-base	12	768	12	3072	110M
BERT-large	24	1024	16	4096	340M
GPT-2	12	768	12	3072	117M
GPT-3	96	12288	96	49152	175B
LLaMA-2 7B	32	4096	32	11008	7B
LLaMA-2 70B	80	8192	64	28672	70B

$$d_k = \frac{d_{\text{model}}}{h}$$

Head dimension

$$d_{\text{ff}} \approx 4 \times d_{\text{model}}$$

FFN hidden size

$$\text{Params} \approx 12 \cdot N \cdot d^2$$

Rough estimate

# Further reading

## Attention & Transformers

- Vaswani et al. (2017), “Attention Is All You Need” — the original Transformer paper
- Bahdanau et al. (2015), “Neural Machine Translation by Jointly Learning to Align and Translate”
- Ian Alhammar, “The Illustrated Transformer”

## Positional Encodings

- Su et al. (2021), “RoFormer: Enhanced Transformer with Rotary Position Embedding” (RoPE)

## Architecture Variants & Surveys

- Devlin et al. (2019), “BERT: Pre-training of Deep Bidirectional Transformers” (encoder-only)
- Radford et al. (2018/2019), “Improving/Language Models are Unsupervised Multi-task Learners” (GPT/GPT-2)
- Lin et al. (2022), “A Survey of Transformers” — comprehensive taxonomy

# Questions?

Next: Tokenization