

Pre-training & Fine-tuning

CLM · MLM · NSP · SFT · RLHF · DPO · GRPO · LoRA

Outline

Pre-training objectives

Fine-tuning & instruction tuning


Alignment & RLHF

Preference optimization: DPO & GRPO

Parameter-efficient fine-tuning

Putting it together

The two-stage paradigm



Part I

Pre-training Objectives

How models learn from raw text

Causal Language Modeling (CLM)

$$\mathcal{L}_{\text{CLM}} = - \sum_{i=1}^T \log P(x_i \mid x_1, \dots, x_{i-1}; \theta)$$



Each token can only see tokens to its left

Advantages:

- Natural for text generation
- Every token contributes to loss
- Simple and scalable

Disadvantage:


- Unidirectional — can't use right context
- "The ____ sat on the mat" — CLM can't peek right to resolve this

Models: GPT, GPT-2, GPT-3, LLaMA, Mistral, Claude — *all modern LLMs*


Masked Language Modeling (MLM)

$$\mathcal{L}_{\text{MLM}} = - \sum_{i \in \mathcal{M}} \log P(x_i \mid x_{\setminus \mathcal{M}}; \theta)$$

$\mathcal{M}$ = set of masked positions (15% of tokens)



BERT's 80/10/10 rule for selected tokens:



Prevents pre-train/fine-tune mismatch:

[MASK] never appears in downstream tasks

Models: BERT, RoBERTa, DeBERTa | **Best for:** understanding tasks (NLI, NER, QA)

Next Sentence Prediction (NSP) & Sentence Order Prediction (SOP)

NSP (BERT)

Is sentence B the actual next sentence after A?

[CLS] A went to the store [SEP]
He bought some milk [SEP]

IsNext (50%) / NotNext (50%)

Too easy! Random sentences differ in *topic*. Model learns topic detection, not coherence. RoBERTa dropped NSP entirely.

SOP (ALBERT)

Are consecutive sentences in the correct order?

Positive: (A, B) natural order
Negative: (B, A) swapped order


Same topic — must understand coherence

Harder task: can't rely on topic difference. Forces understanding of logical ordering between sentences.

Lesson: auxiliary objectives must be *genuinely challenging* to be useful. NSP was mostly abandoned after 2019.

Contrastive learning

Learn by comparing: pull similar pairs together, push dissimilar apart



$$\mathcal{L}_{\text{InfoNCE}} = -\log \frac{\exp(\text{sim}(z_i, z_i^+)/\tau)}{\sum_{j=1}^N \exp(\text{sim}(z_i, z_j)/\tau)}$$

where z_i^+ = positive pair, others are negative

SimCSE

Sentence embeddings
(same input, different dropout)

CLIP

Image-text alignment
(match image to caption)

DPR

Dense passage retrieval
(question-passage matching)

Pre-training objectives compared

Objective	Direction	Loss on	Models	Best for
CLM	Left→Right	All tokens	GPT, LLaMA	Generation
MLM	Bidirectional	15% masked	BERT, RoBERTa	Understanding
NSP/SOP	Sentence-level	[CLS]	BERT, ALBERT	Sentence pairs
Contrastive	Embedding	Pairs	CLIP, SimCSE	Retrieval
RTD	Bidirectional	All tokens	ELECTRA	Efficient NLU


ELECTRA (bonus):

Small generator creates fake tokens; discriminator detects them.
All tokens provide signal.

Modern trend:

CLM dominates at scale.
GPT-3 showed autoregressive models can handle understanding tasks too via prompting.

The modern pre-training pipeline



Scale examples:

Model	Tokens	Params	Compute
GPT-3	300B	175B	3640 petaflop-days
LLaMA	1.4T	65B	2048 A100s × 21 days
LLaMA-2	2T	70B	1.7M GPU-hours


Pre-training is **expensive** (\$1–100M+) but done **once**.
Fine-tuning is cheap (\$10–1000) and done per task.

Part II


Fine-tuning & Alignment

From base model to useful assistant

From base model to useful model




Concrete example:



Surprising: a 1.3B InstructGPT was preferred by humans over the 175B base GPT-3!
Fine-tuning quality matters more than raw model size.

Supervised Fine-Tuning (SFT)



Instruction tuning

Train on diverse (instruction, response) pairs across many tasks

Instruction: Translate to French.
Input: "The cat sat on the mat."
Output: "Le chat était assis sur le tapis."

Instruction: Is this review positive?
Input: "I loved this movie!"
Output: "Yes, this is a positive review."

FLAN (Google)

60+ tasks $\Rightarrow$ strong zero-shot generalization to unseen tasks

Instruction: Summarize this article.
Input: [long article text]
Output: "Researchers found that..."


Instruction: Write a Python function that sorts a list.
Output: `def sort_list(lst):...`

InstructGPT (OpenAI)

1.3B params, instruction-tuned
> 175B base GPT-3

Instruction tuning is typically **Step 1** before alignment (RLHF/DPO)

The alignment problem



Goal: **Helpful** + **Harmless** + **Honest** (Anthropic's HHH)

The challenge: you can't write a loss function for "be helpful and safe."
Human preferences are nuanced, subjective, and context-dependent.

Solution: learn from human feedback — let humans judge and the model learn from their preferences

Constitutional AI — alignment from AI feedback

Bai et al. (2022), Anthropic: replace (most) human feedback with **AI feedback**. Define a “constitution” of principles; let the model *critique* and revise its own outputs.

Phase 1: Supervised (RLAIF)

1. Model generates response
2. Model *critiques* it using a principle (e.g., “Is this harmful?”)
3. Model *revises* its response
4. Train SFT on revised outputs

Phase 2: RL (RLAIF)

Train a reward model using **AI-generated** preference labels instead of human annotators.

Then run RLHF as usual.

Example constitutional principle: “Choose the response that is most helpful while being *harmless* and *honest*. Avoid responses that are toxic, racist, sexist, or encourage illegal activity.”

Advantages:

Scales cheaply (no human labelers)
Principles are explicit and auditable
Less harmful than pure RLHF

Limitations:

AI judge inherits model biases
Quality depends on constitution design
Still needs some human oversight

RL basics for language models


Reinforcement Learning $\longleftrightarrow$ Language Model

RL Concept	LLM Mapping	Notation
State s	Prompt + tokens so far	$(x, y_{<t})$
Action a	Next token	y_t
Policy $\pi(a s)$	LLM (token probs)	$\pi_\theta(y_t x, y_{<t})$
Trajectory τ	Full generated response	$y = (y_1, \dots, y_T)$
Reward R	Human preference score	$R_\lambda(x, y) \in \mathbb{R}$

Why RL? The reward (human preference) is **non-differentiable**
— can't backprop through “did the human like it?”

RL handles non-differentiable reward signals via policy gradient methods

RLHF — Step 1: Train a reward model



Bradley-Terry: $P(y_w \succ y_l) = \sigma(R_\phi(x, y_w) - R_\phi(x, y_l))$

Train to predict which response humans prefer

Pairwise preferences


Rankings are less noisy
than absolute scores

Scale: $\sim 50k$ preference pairs

RM can be smaller
than policy

(e.g., 6B RM for 175B policy)

RLHF — Step 2: PPO optimization



4 models in memory simultaneously:

Policy π_θ

Reference π_{ref}

Reward R_ϕ

Value V_ψ


Complex · Unstable (PPO hyperparameters) · Expensive (4 models) · Reward hacking risk

PPO — the clipped surrogate objective

$$\mathcal{L}^{\text{CLIP}} = \mathbb{E}_t \left[\min \left(\underbrace{r_t(\theta) \hat{A}_t}_{\text{standard PG}}, \underbrace{\text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t}_{\text{clipped}} \right) \right]$$

where $r_t(\theta) = \frac{\pi_\theta(y_t | x, y_{<t})}{\pi_{\text{old}}(y_t | x, y_{<t})}$ (probability ratio)

What does clip



If $\hat{A}_t > 0$ (good action):

Policy wants to increase r_t
but clipping caps at $1 + \epsilon$

Prevents overly aggressive updates

If $\hat{A}_t < 0$ (bad action):

Policy wants to decrease r_t
but clipping caps at $1 - \epsilon$

Prevents catastrophic forgetting

Typical: $\epsilon = 0.2$ · $\hat{A}_t$ computed via GAE (generalized advantage estimation)

Reward hacking & the KL penalty

Problem: the reward model is imperfect.
The policy can find **adversarial outputs**
that score high on R_ϕ but are gibberish or degenerate to humans.

Examples of reward hacking:

Repetition: "This is great! Great!
Great!
Really great! So great!" scores
high
on positivity rewards

Length gaming: longer responses
tend to score higher $\Rightarrow$ model
becomes verbose

Solution: KL divergence penalty

$$r_{\text{total}} = R_\phi(x, y) \\ - \beta \cdot D_{\text{KL}}[\pi_\theta \parallel \pi_{\text{ref}}]$$


β too small $\Rightarrow$ reward hacking
 β too large $\Rightarrow$ no learning

Typical β is 0.01-0.1

In practice, KL is computed **per token**:
$$D_{\text{KL}} = \sum_{t=1}^T \log \frac{\pi_\theta(y_t | x, y_{<t})}{\pi_{\text{ref}}(y_t | x, y_{<t})}$$

This keeps the policy close to the SFT model — it can improve but not drift too far
The KL penalty is what makes RLHF work in practice. Without it, training collapses within hours.

RLHF — the complete pipeline



DPO — Direct Preference Optimization

Key insight (Rafailov et al., 2023): the optimal RLHF policy has a **closed-form solution**.

Express the reward implicitly through the policy — no RL needed!

$$\mathcal{L}_{\text{DPO}} = -\mathbb{E} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_{\theta}(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

y_w = preferred, y_l = dispreferred, π_{ref} = frozen SFT model

RLHF

3 stages
4 models in memory
Unstable RL training
Complex hyperparameters

DPO

1 stage (supervised loss)
2 models (policy + reference)
Stable training
Just one hyperparameter (β)



DPO matches or exceeds PPO quality: 61% win rate vs. PPO's 57% on summarization

Used by: Zephyr, Tulu, many open-source models | Equivalent to RLHF

Step 1: RLHF objective $\max_{\pi_{\theta}} \mathbb{E}_{x,y} [R(x,y)] - \beta \cdot D_{\text{KL}}[\pi_{\theta} \parallel \pi_{\text{ref}}]$

Step 2: Closed-form optimal policy $\pi^*(y|x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y|x) \exp\left(\frac{R(x,y)}{\beta}\right)$

Step 3: Rearrange to express reward via policy $R(x,y) = \beta \log \frac{\pi^*(y|x)}{\pi_{\text{ref}}(y|x)} + \beta \log Z(x)$

Step 4: Substitute into Bradley-Terry preference model $P(y_w \succ y_l) = \sigma(R(x, y_w) - R(x, y_l))$

Result — DPO loss:


$$\mathcal{L}_{\text{DPO}} = -\mathbb{E} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_{\theta}(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

Intuition: increase $\pi_{\theta}(y_w|x)$ relative to π_{ref} , decrease $\pi_{\theta}(y_l|x)$ relative to π_{ref}

The implicit reward is $\hat{R}(x,y) = \beta \log \frac{\pi_{\theta}(y|x)}{\pi_{\text{ref}}(y|x)}$ — the policy is the reward model

GRPO — Group Relative Policy Optimization


DeepSeek's innovation: eliminate the value network from PPO
G samples



$$\mathcal{L}_{\text{GRPO}} = -\frac{1}{G} \sum_{i=1}^G \min\left(r_i(\theta) \hat{A}_i, \text{clip}(r_i(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_i\right) - \beta \cdot D_{\text{KL}}[\pi_{\theta} \parallel \pi_{\text{ref}}]$$

Same PPO clipping, but advantages $\hat{A}_i$ come from **group statistics**, not a value network

How group advantage works:



Rule-based rewards:

Correctness
Does the answer
match ground truth?

Format
Does it follow
the required format?

Code tests
Do unit tests
pass or fail?

**No human labelers
needed — rewards
are automatic!**

Alignment methods compared

Method	Models needed	Stability	Data	Best for
RLHF (PPO)	4 (policy, ref, RM, value)	Unstable	Preferences	General alignment
DPO	2 (policy, ref)	Stable	Preferences	Simple alignment
GRPO	2 (policy, ref)	Medium	Verifiable	Math/code reasoning
CAI (RLAIF)	2 (policy, ref) + AI judge	Stable	AI-generated	Scalable safety


Rules of thumb:

Need maximum control over alignment? $\Rightarrow$ **RLHF**


Have preference data, want simplicity? $\Rightarrow$ **DPO**

Training a reasoning model with verifiable answers? $\Rightarrow$ **GRPO**

The cost problem



LoRA — Low-Rank Adaptation



$$h = W_0 x + \frac{\alpha}{r} \cdot B \cdot A \cdot x$$

B initialized to **zero** $\Rightarrow$ starts
with pre-trained weights

GPT-3 175B with LoRA:

Trainable params:
4.7M (0.003%)
VRAM: 350 GB (vs. 1.2 TB)
Quality: matches full fine-tuning

At inference:

Merge $W = W_0 + BA$
 $\Rightarrow$ **zero** extra latency!
Hot-swap LoRA modules per task

QLoRA — Quantized LoRA

Dettmers et al. (NeurIPS 2023): store frozen weights W_0 in **4-bit** precision, train LoRA adapters in 16-bit. Fine-tune a **65B** model on a **single 48GB GPU**.

4-bit NormalFloat

Quantization type optimized for normally distributed weights.
Better than INT4.

Double quantization

Quantize the quantization constants themselves.
Saves ~ 0.37 bits per parameter extra.

Paged optimizers

Offload optimizer states to CPU RAM on OOM.
Uses NVIDIA unified memory.

Memory comparison (65B model fine-tuning):

Method	VRAM	GPUs needed
Full fine-tuning (16-bit)	> 780 GB	$\sim 10 \times$ A100
LoRA (16-bit base)	~ 130 GB	$2 \times$ A100

Key result: QLoRA matches 16-bit LoRA quality on all benchmarks.
Guanaco (QLoRA-tuned LLaMA 65B) reached **99.3%** of ChatGPT quality on Vicuna benchmark.

Dettmers et al. (2023), "QLoRA: Efficient Finetuning of Quantized Language Models" — NeurIPS 2023

PEFT methods compared

Method	Approach	Params	Inference	Quality
LoRA	Low-rank update to W	$\sim 0.01\%$	No overhead	Excellent
QLoRA	LoRA + 4-bit base weights	$\sim 0.01\%$	No overhead	Excellent
Adapters	New layers between blocks	$\sim 1\text{--}5\%$	+20–30%	Excellent
Prefix tuning	Learnable prefix to K, V	$\sim 0.1\%$	Slight	Good
Prompt tuning	Soft prompt embeddings	$\sim 0.01\%$	Slight	Good
Full FT	Update all weights	100%	Baseline	Best

Adapter module:

$$x \rightarrow W_{\text{down}} \rightarrow \text{ReLU} \rightarrow W_{\text{up}} + x$$

Bottleneck inserted between layers

Cannot merge $\Rightarrow$ permanent overhead

Why LoRA dominates:


Zero inference overhead (merge)

Fewer params than adapters

Modular (hot-swap per task)

Simple to implement

The full picture



Real-world pipelines:

ChatGPT: Pre-train (GPT-3.5/4) → SFT (demonstrations) → RLHF (PPO + human preferences)

Llama 2 Chat: Pre-train (2T tokens) → SFT (27.5k examples) → RLHF (iterative 5 rounds)

DeepSeek-R1: Pre-train → Cold-start SFT → GRPO (rule-based rewards) → SFT on distilled data → GRPO again

Zephyr: Pre-train (Mistral 7B) → SFT (UltraChat) → DPO (UltraFeedback preferences)

Practical guide

Which method should I use?

Base model + limited data?

⇒ **LoRA + SFT**

Efficient, fits on one GPU

Math / code reasoning?

⇒ **GRPO**

Rule-based rewards, no RM needed

Text generation?

⇒ **CLM pre-training** (GPT-style)

Then SFT + alignment

Need alignment with values?

⇒ **DPO** (simple) or
RLHF (max control)

Requires preference data

**Understanding
tasks (NER, QA)?**

⇒ **MLM pre-training** (BERT-style)

Then fine-tune on task data

Don't want to train at all?

⇒ **Prompting** (next lecture!)

Zero-shot, few-shot, CoT

Further reading

Pre-training

- Devlin et al. (2019), “BERT” — masked language modeling

RLHF & Alignment

- Ouyang et al. (2022), “Training language models to follow instructions with human feedback” (InstructGPT)
- Rafailov et al. (2023), “Direct Preference Optimization” (DPO)
- Shao et al. (2024), “DeepSeekMath: Pushing the Limits of Mathematical Reasoning” (GRPO)

Parameter-Efficient Fine-Tuning

- Hu et al. (2022), “LoRA: Low-Rank Adaptation of Large Language Models”
- Dettmers et al. (2023), “QLoRA: Efficient Finetuning of Quantized Language Models”
- Hounsby et al. (2019), “Parameter-Efficient Transfer Learning for NLP” (Adapters)

Questions?

Next: Prompting — Zero-shot, Few-shot, Chain-of-Thought