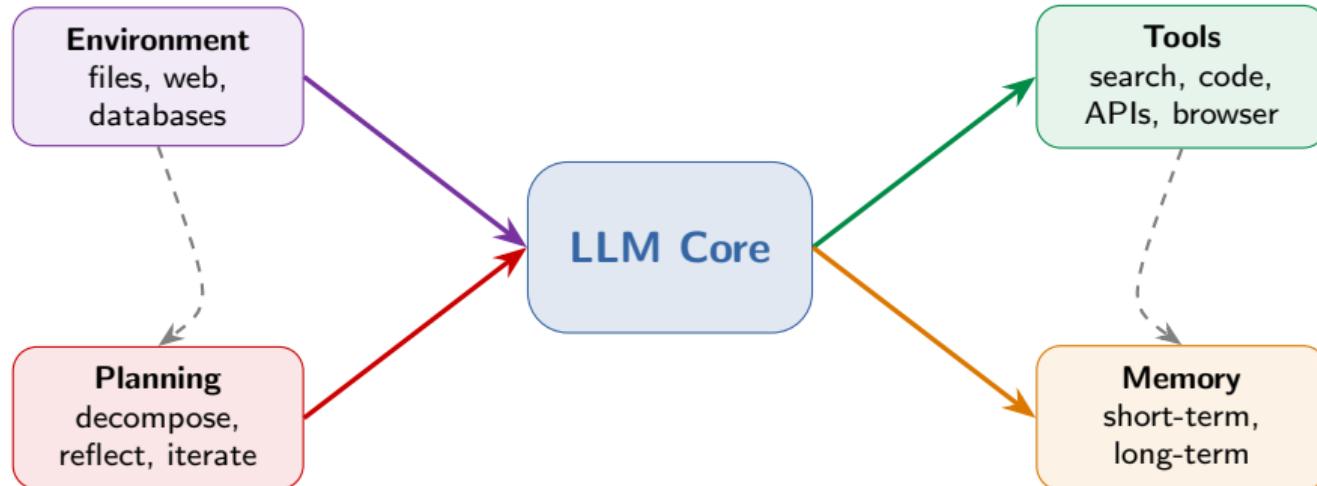


AI Agents & Tool Use

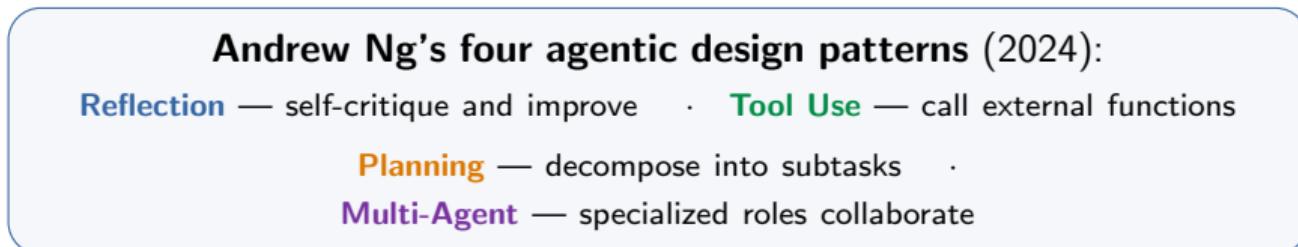
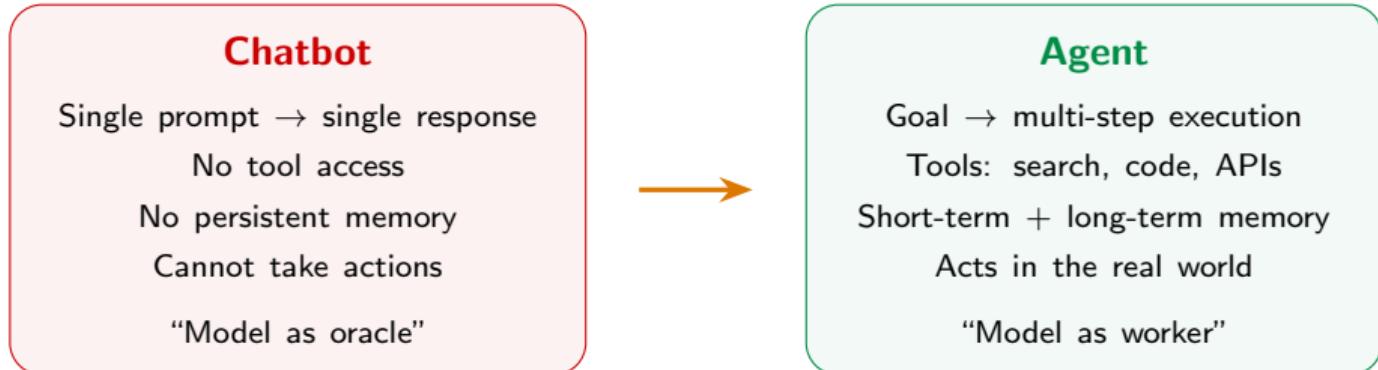
ReAct · Function Calling · Multi-Agent · MCP

What is an AI agent?

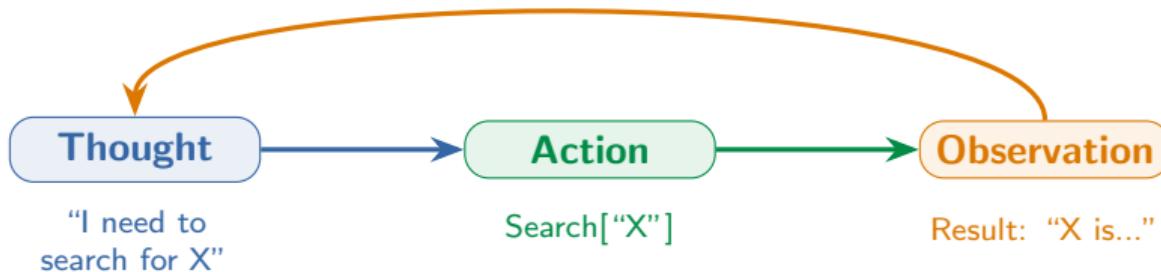


An **agent** is an LLM that can autonomously **plan**, **use tools**, and **act** in a loop until a goal is achieved — not just generate a single response.

Chatbot vs. agent



ReAct: Reasoning + Acting



Thought: I need to find when the Eiffel Tower was built.

Action: Search["Eiffel Tower construction date"]

Observation: Construction began in 1887 and was completed in 1889.

Thought: Now I have the answer.

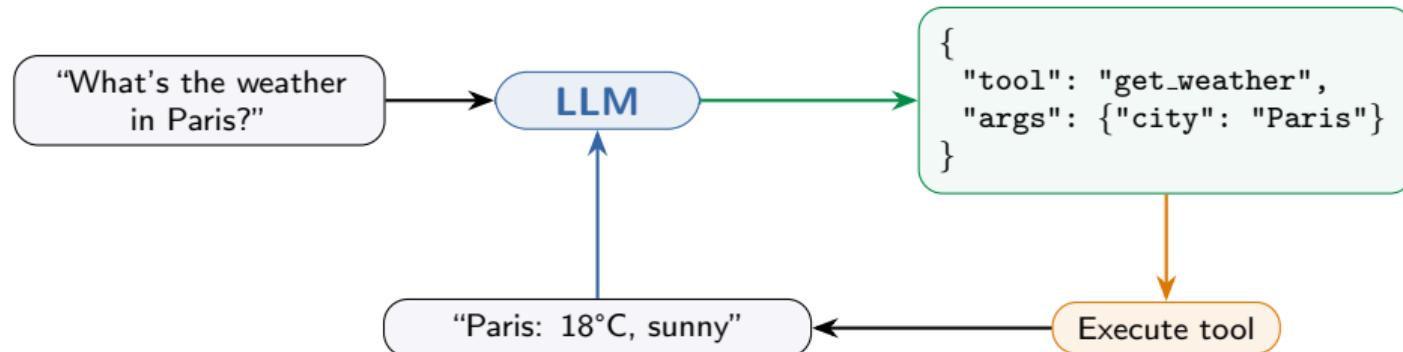
Action: Finish["The Eiffel Tower was built 1887–1889."]

Yao et al., ICLR 2023 (Oral) · ALFWorld: +34% over RL baselines

Reasoning helps plan and handle exceptions; actions provide grounded information.

ReAct is now the **default loop** in virtually every agent framework.

Function calling



OpenAI (June 2023): `function_call` API, parallel calls, structured outputs

Anthropic: typed tool schemas, parallel tool use

Google: strict JSON schema enforcement via Vertex AI

Models are **fine-tuned** on (instruction, tool-call, result) triples to learn *when* to call, *which* tool to select, and *how* to format arguments. Structured JSON output is now standard.

Types of tools

Search / Retrieval

Web search APIs
Vector DB (RAG)
Wikipedia, knowledge bases

Code Execution

Python interpreter
Sandboxed runners
Jupyter notebooks

External APIs

Weather, databases
Calendars, e-commerce
CRM, payments

File System

Read / write / edit files
Directory traversal
Document parsing

Browser / Web

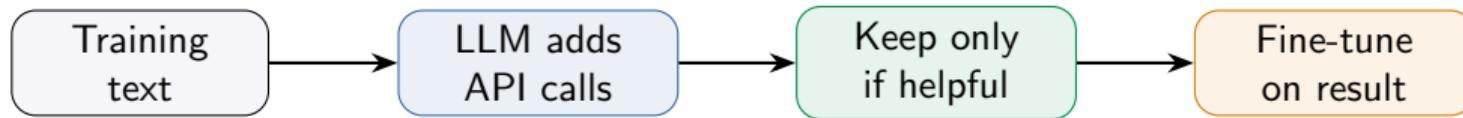
Headless browser
Screenshot analysis
Form filling, navigation

Math / Reasoning

Calculator
Symbolic math
Wolfram Alpha

LLMs as Tool Makers (Cai et al., 2023): a powerful LLM *creates* reusable tools (Python functions), then a cheaper LLM *uses* them. GPT-4 maker + GPT-3.5 user \approx GPT-4 quality at lower cost.

Toolformer: self-supervised tool learning



Original: "The population of Toronto is 2,794,356."

Annotated: "The population of Toronto is [QA("population of Toronto") → 2,794,356]."

Kept because the API call *reduced* perplexity of the next tokens (i.e., helped prediction)

Schick et al., NeurIPS 2023 (Meta) · **6.7B**
model with tools > GPT-3 175B without

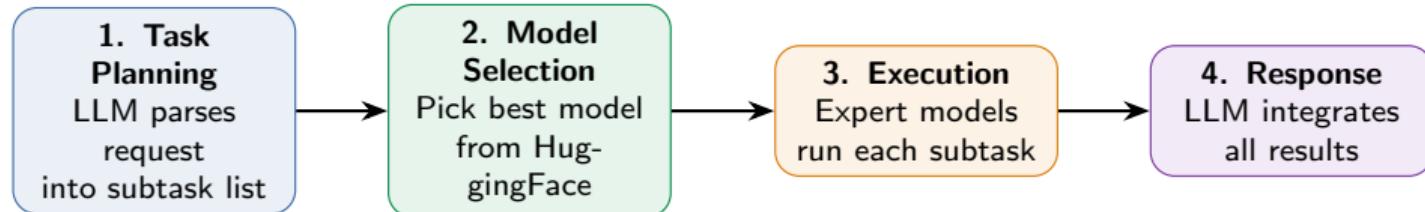
Tools: calculator, QA, search, translation, calendar

Calculator used in 97.0% of math examples (doubling performance)

Key insight: the model teaches *itself* when and how to use tools, using perplexity as the signal.
No human annotation of tool calls needed — fully self-supervised.

Planning & decomposition

HuggingGPT (Shen et al., NeurIPS 2023)



User: "Describe this image and read any text in it."

Plan: 1. Image captioning (BLIP-2) → 2. OCR (TrOCR) → 3. Combine results

Plan-and-Execute

Create full plan upfront,
then execute step by step.
Can re-plan if a step fails.

Inner Monologue

Agent maintains a scratchpad
of reasoning, tracks progress,
and self-corrects along the way.

Memory systems

Short-term

Current conversation
in the context window

Limited by context length
(4K–200K tokens)

Lost when session ends

Long-term

Persistent storage across
sessions via vector DBs

Pinecone,
Chroma, Weaviate
Retrieved by relevance
(RAG-style)

Episodic

Memories of past
interactions
and outcomes

“Last time I tried X,
it failed because Y”

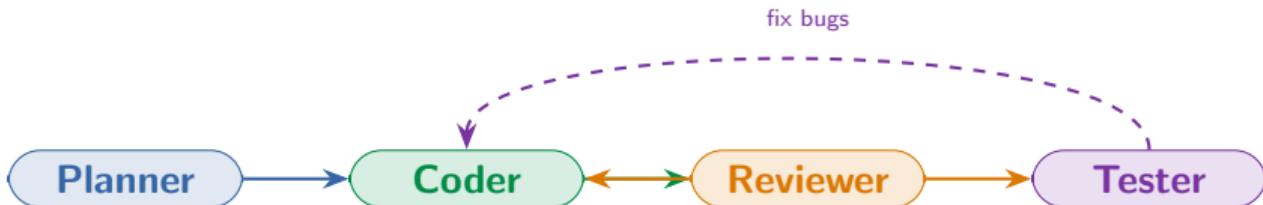
Enables learning
from experience

MemGPT (Packer et al., 2023): OS-inspired virtual memory for LLMs

Main context (“RAM”) = limited prompt buffer
↔ External storage (“disk”) = unbounded archive

Agent manages data movement between tiers us-
ing function calls: `search_memory()`, `save_memory()`

Multi-agent systems



AutoGen (Microsoft, 2023)

Customizable conversable agents, flexible topologies

Merged into MS Agent Framework (2025)

CrewAI (2024)

Role-based orchestration
100K+ agent runs/day
60% of Fortune 500

ChatDev (Qian et al., 2023)

Virtual software company:
CEO, CTO, programmer,
tester, reviewer

Multi-agent debate (Du et al., 2023): multiple LLMs debate and critique each other's answers
⇒ improves factual accuracy and reasoning. Diversity of perspectives reduces hallucination.

Agent frameworks

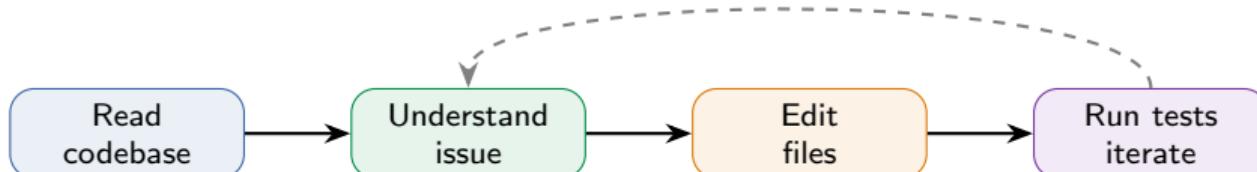
Framework	Focus	Key Abstraction
LangChain / LangGraph	General-purpose	Graph-based state machines
LlamaIndex	Data & RAG	300+ data connectors
Semantic Kernel (MS)	Enterprise .NET/Python	Plugins, planners
Claude Agent SDK	Building on Claude	Agent loops, tool use
OpenAI Assistants API	Managed infrastructure	Threads, runs, tools
CrewAI	Multi-agent enterprise	Crews, roles, tasks

Common abstractions: Tools (callable functions) · Agents (LLM + tools + instructions)
Memory (short/long-term state) · Chains/Graphs
(Trend: LangChain pivoted from chains to Lang-Graph for complex agent workflows (cycles, state).) · Planners (decomposition)

LlamaIndex dominates RAG-heavy apps.

CrewAI dominates multi-agent enterprise.

Code agents



Devin (Cognition, Mar 2024)

“First AI software engineer”

SWE-Bench: 13.9% → GA Dec 2024

ARR: \$1M → \$73M in 9 months

SWE-Agent (Princeton, Apr 2024)

Open-source, custom ACI

Agent-Computer Interface

designed for software engineering

OpenHands (2024)

Open-source platform

SOTA: 72% SWE-Bench Verified

(Claude Sonnet 4.5 + ext. thinking)

Claude Code (Anthropic, 2025)

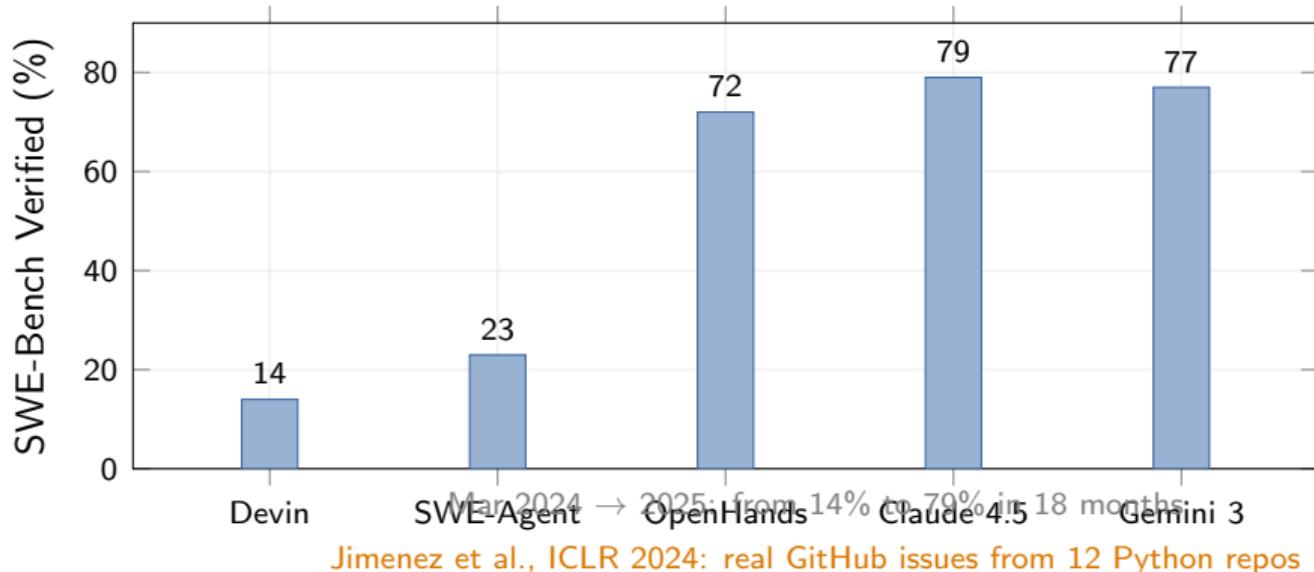
CLI-based coding agent

Read, edit, run, iterate

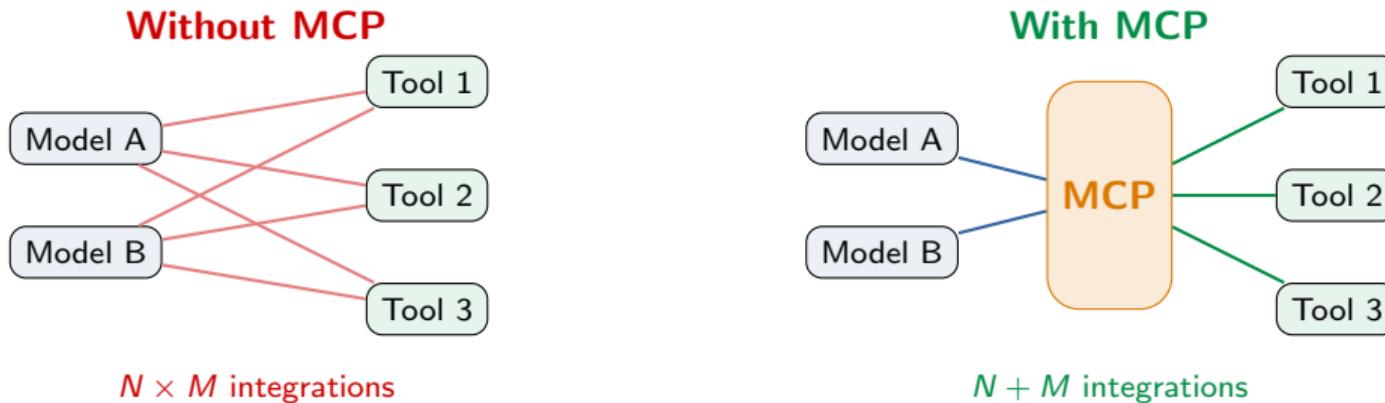
Built on Claude Agent SDK

GitHub Copilot Coding Agent (May 2025):
async agent for features, bugs, tests, refactoring.

SWE-Bench: the code agent benchmark



MCP: Model Context Protocol



$N \times M$ integrations

With MCP

MCP

$N + M$ integrations

Anthropic, November 2024 · Like
LSP for AI: standardizes tool integration

JSON-RPC 2.0 over stdio/HTTP · Primitives: **Tools, Resources, Prompts**

OpenAI adopted (Mar 2025) · Google adopted
(Apr 2025) · 97M+ monthly SDK downloads

Challenges

Reliability

95% per-step accuracy
⇒ 60% for 10 steps
One bad step derails the entire workflow

Current: ~80%

Need: ~99%

Hallucinated calls

Non-existent tools
Wrong parameter values
Incorrect types
Hard to catch semantically
(syntax may be valid)

Cost & latency

Complex task = 10–50+ API calls
Cost: 10–100× a single-shot response
End-to-end: minutes to hours

Safety: autonomous agents act in the real world — editing files, sending emails, making purchases, browsing the web. Unintended consequences are a genuine risk.

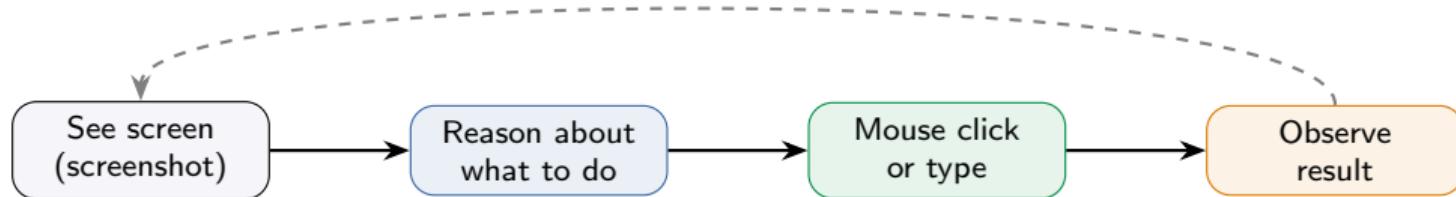
Evaluation is hard: multi-step workflows are non-deterministic, have multiple valid paths, and success depends on the process, not just the final answer.

Agent benchmarks

Benchmark	What it measures	Scale	Key number
SWE-Bench	Real GitHub issues	2.3K tasks	Best: 79%
AgentBench	8 environments (OS, DB, ...)	29 LLMs	Large open/closed gap
WebArena	Web browsing tasks	812 tasks	4 realistic domains
GAIA	General AI assistant	466 Q	Human: 92%, Best: 65%
ToolBench	API tool use	16K APIs	49 categories

GAIA gap: humans score 92%, best agents score 65%.
Unlike text benchmarks (which are near-saturated), agent benchmarks reveal
a **large gap** between human and AI capability on real-world tasks.

Computer use agents



Claude Computer Use
(Anthropic, Oct 2024)

OSWorld:
14.9% → 72.5%
(~5x in 16 months)

Operator / CUA
(OpenAI, Jan 2025)

GPT-4o vision + RL
Raw pixels +
virtual mouse

Project Mariner
(Google, Dec 2024)

GUI control via
multimodal models

The paradigm shift: from “model as oracle” to “model as worker.”

Give a goal, the agent plans + acts + observes + iterates until done.

Browsers and operating systems are becoming **agent-native** platforms.

Practical guide

Simple Q&A?

→ Standard LLM
(no agent needed)

Need external data?

→ RAG or single
tool call

Multi-step workflow?

→ ReAct agent with
tools + memory

Complex software task?

→ Code agent
(SWE-Agent,
Claude Code,
OpenHands)

Need specialization?

→ Multi-agent system
(CrewAI, AutoGen)

Standardize tools?

→ MCP servers
(build once, use
everywhere)

Rule of thumb: use the simplest approach that works.

Single LLM call → tool call → ReAct loop →
multi-agent. Add complexity only when needed.

Questions?

All DL4NLP topics complete!