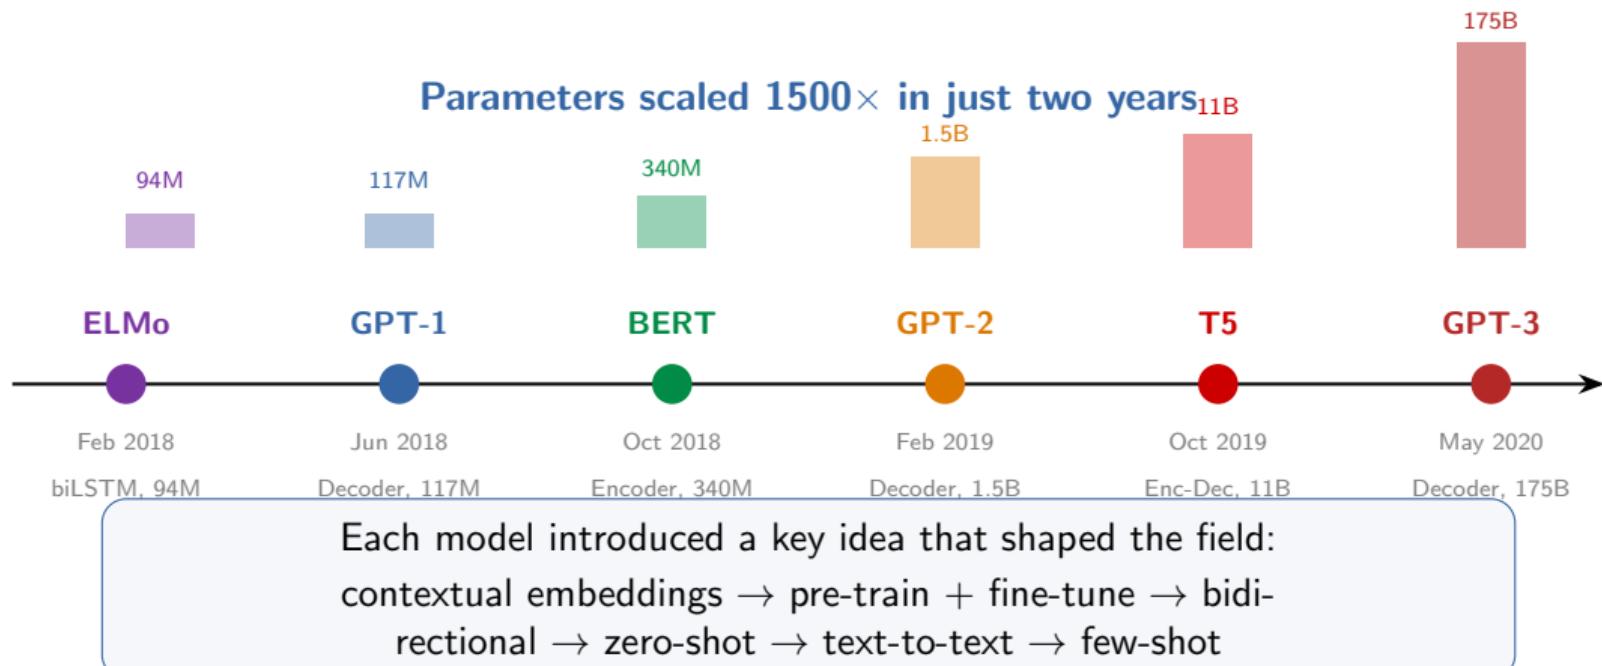


## Early Notable Models

ELMo · GPT · BERT · GPT-2 · T5 · GPT-3

# The timeline: 2018–2020



# ELMo — contextual embeddings (2018)

**Before ELMo:** Word2Vec / GloVe give each word a **single fixed vector**  
“bank” has the same embedding in “river bank” and “bank account”

## Architecture:

Character CNN →  
2-layer biLSTM (4096 units)

Forward: predict next token  
Backward: predict prev token

Final: learned weighted  
sum of all layers

## Innovation:

Same word → **different vectors**  
depending on context!

Lower layers: syntax  
Higher layers: semantics

Trained on 1B words

“I went to the **bank** to deposit money” → finance vector

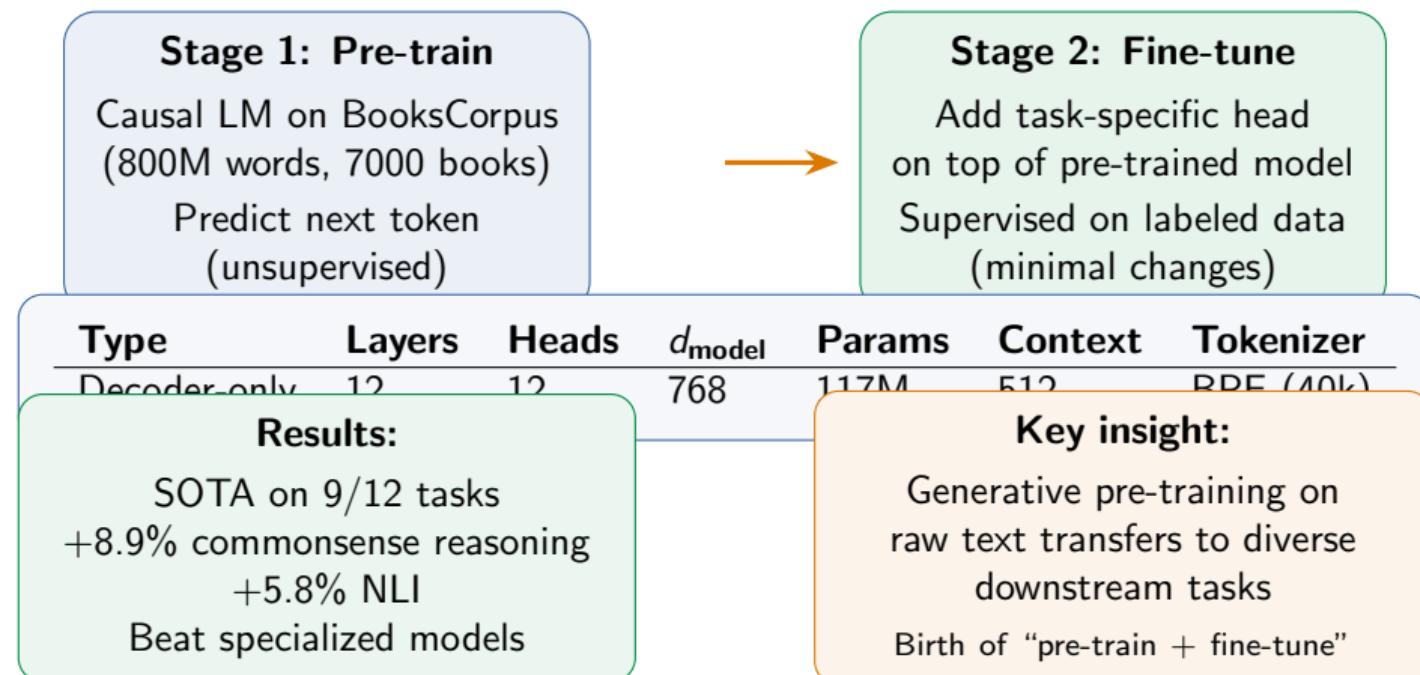
“I sat on the river **bank**” → geography vector

The *same* word gets *different* representations — this was revolutionary

**Impact:** proved that pre-trained contextual representations transfer to downstream tasks.  
Directly inspired GPT and BERT. But ELMo used LSTMs — transformers would do it better.

# GPT-1 — the pre-train + fine-tune paradigm (2018)

“Improving Language Understanding by Generative Pre-Training” (OpenAI)



# BERT — bidirectional is better (2018)

**GPT's limitation:** decoder-only = unidirectional (left-to-right only).  
“The \_\_\_\_\_ sat on the mat” — GPT can't use right context to fill the blank!

**BERT's solution: Masked Language Modeling (bidirectional)**

**GPT (unidirectional):**

The [?] sat on the mat  
← only sees “The”



**BERT (bidirectional):**

The [MASK] sat on the mat  
← sees “The” AND  
“sat on the mat”

	<b>Layers</b>	<b>Hidden</b>	<b>Heads</b>	<b>Params</b>	<b>Vocab</b>
<b>BERT-Base</b>	12	768	12	110M	30,522 (WordPiece)
<b>BERT-Large</b>	24	1024	16	340M	30,522 (WordPiece)

Trained on BooksCorpus (800M words) + English Wikipedia (2.5B words) = 3.3B words total

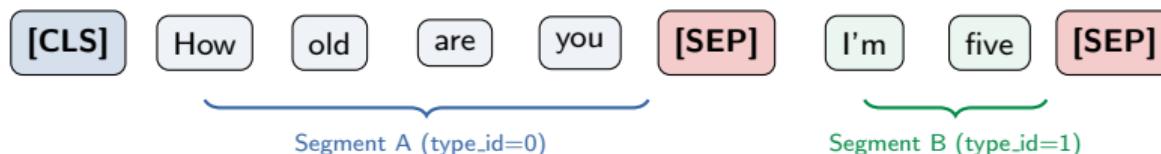
**Encoder-only:** no autoregressive generation. Designed for **understanding** tasks.

# BERT — input format and special tokens

Single sentence:



Sentence pair:



Input = Token Embedding + Segment Embedding + Position Embedding

## [CLS]

Classification token.  
Its final hidden state is  
the sequence  
representation  
for classification tasks.

## [SEP]

Separator token.  
Marks boundaries between  
sentences. Used with  
segment embeddings.

## [MASK]

Masking token.  
Replaces 15% of tokens  
during MLM pre-training.  
Never seen at  
fine-tune time.

Tokenizer: WordPiece (30,522 vocab) · Subword: “playing” → “play” + “##ing” · Max length: 512

# BERT — pre-training objectives

## Masked LM (MLM)

Randomly mask

15% of tokens

Model predicts original token  
using bidirectional context

Of the 15% selected:

80% → [MASK]

10% → random token

10% → unchanged

"The cat [MASK] on the mat"  
→ predict "sat"

## Next Sentence Prediction

Given sentence A and B,  
predict: is B the actual next  
sentence after A?

50% real pairs (IsNext)

50% random pairs (NotNext)

Binary classification on [CLS]

[CLS] A [SEP] B [SEP]  
→ IsNext / NotNext

**Later found:** NSP was too easy (topic detection, not coherence).

RoBERTa (2019) dropped NSP entirely and got **better** results.

ALBERT replaced it with Sentence Order Prediction (SOP).

Training: 4 Cloud TPUs (16 chips), 4 days · Estimated cost: ~\$500 (2018 prices!)

# BERT — why everyone lost their minds

BERT set SOTA on 11 NLP tasks simultaneously (October 2018)

## GLUE benchmark:

Previous SOTA: 72.8

BERT-Large: **80.5**

+7.7 points absolute improvement

## MultiNLI:

Previous SOTA: 82.1

BERT: **86.7**

## SQuAD v1.1 (QA):

Previous SOTA: 91.7 F1

BERT-Large: **93.2 F1**

First to surpass human-level (91.2)

## CoNLL NER:

Previous SOTA: 92.2 F1

BERT: **92.8 F1**

**The BERT effect:** a single pre-trained model, fine-tuned with just a linear head, beat years of task-specific engineering on virtually every NLU benchmark.

BERT became the default starting point for NLP from late 2018 through 2020

Spawned: RoBERTa, ALBERT, DistilBERT, DeBERTa, SpanBERT, SciBERT, BioBERT, ...

# Three architectures, three strengths

## Encoder-only

Bidirectional  
attention

Best for:  
understanding

NLI, NER, QA,  
classification

BERT, RoBERTa

## Decoder-only

Causal (L→R)  
attention

Best for:  
generation

Text, code,  
chat, reasoning

GPT, LLaMA, Mistral

## Encoder-Decoder

Bidirectional enc +  
causal dec

Best for:  
seq-to-seq

Translation,  
summarization

T5, BART, mBART

**2024 landscape:** decoder-only dominates  
(GPT-4, Claude, LLaMA, Mistral, Gemini).

With enough scale + prompting, decoder-only handles understanding tasks too.

BERT-style encoders still used for: embeddings, retrieval, NER, classification at small scale

# GPT-2 — scaling up, zero-shot emerges (2019)

“Language Models are Unsupervised Multitask Learners” (OpenAI)

Variant	Params	Layers	Hidden
Small	124M	12	768
Medium	355M	24	1024
Large	762M	36	1280
	<b>1.5B</b>	<b>48</b>	<b>1600</b>

## vs. GPT-1:

13× more parameters  
Pre-LayerNorm (more stable)  
Context: 1024 tokens (vs. 512)  
WebText: 40 GB

**Zero-shot breakthrough:**  
SOTA on 7/8 LM benchmarks  
with **no fine-tuning at all**  
Coherent multi-paragraph text

**“Too dangerous to release”:** OpenAI withheld the full model for 9 months.

Staged release: Small (Feb) → Medium  
(May) → Large (Aug) → XL (Nov 2019)

First major debate about responsible AI release. Mostly a PR event in hindsight.

**Key lesson:** scale alone produces qualitatively new capabilities. Model “still underfits WebText.”

# T5 — everything is text-to-text (2019)

**Unified framework:** cast every NLP task as “text in, text out”

Same model · same loss · same hyperparameters · different task prefixes

## Translation:

“translate English to German: Hello”  
→ “Hallo”

## Summarization:

“summarize: [long article]”  
→ “[short summary]”

## Classification:

“sst2 sentence: The movie was great”  
→ “positive”

## Question answering:

“question: What is the capital? context: ...”  
→ “Paris”

Variant	Params	Enc/Dec layers	Trained on
Small	60M	6 / 6	C4 (750 GB)
Base	220M	12 / 12	C4
Large	770M	24 / 24	C4

**Encoder-decoder** with span corruption objective · Introduced **C4 dataset** (750 GB cleaned web text)  
Massive systematic study: compared architectures, objectives, data sizes, transfer approaches

# T5 — the definitive transfer learning study

T5 paper systematically compared every design choice:

## Architectures

Encoder-decoder  
Decoder-only  
Prefix LM

Winner: enc-dec

## Objectives

Autoregressive (CLM)  
BERT-style (MLM)  
Span corruption  
Prefix LM

Winner: span corruption

## Data quality

Unfiltered web  
Filtered web (C4)  
Different domains

Winner: clean + large

## Transfer

Fine-tune all params  
Adapters  
Multi-task  
Gradual unfreezing

Winner: fine-tune all

## Key findings:

1. Encoder-decoder > decoder-only for most tasks (at equal compute)
2. Span corruption (denoising) > standard MLM > autoregressive
3. Scaling model size + data size + training steps all help, roughly equally
4. Data quality matters more than quantity — C4's cleaning was crucial

# GPT-3 — the leap to 175B (2020)

“Language Models are Few-Shot Learners” (OpenAI, NeurIPS 2020)

Layers	Heads	$d_{\text{model}}$	Head dim	Context	Params
96	96	12,288	128	2048	<b>175 billion</b>

**100× larger than GPT-2, 1500× larger than GPT-1**

## Training data:

300B tokens total  
Common Crawl (60%)  
WebText2 (22%)  
Books (16%) + Wikipedia (3%)

## Training cost:

355 GPU-years  
~\$4.6 million  
10,000 V100 GPUs for weeks

**8 model sizes** spanning 3 orders of magnitude:

Ada (125M) → Babbage (1.3B) → Curie (6.7B) → **Davinci (175B)**

**June 2020:** OpenAI launched the GPT-3 API — first major “LLM as a service”

Text in, text out. No training needed. This began the era of API-based AI.

# GPT-3 — in-context learning

No gradient updates. No fine-tuning. Just text in the prompt.

## Zero-shot

Task description only.  
No examples.

“Translate English to French: cheese →”

Relies on pre-training

## One-shot

1 demonstration example.

“sea otter → loutre  
de mer  
cheese →”

Learn format from 1 example

## Few-shot

10–100 examples  
in the prompt.

“sea otter → loutre  
peppermint → menthe  
cheese →”

Near fine-tuned quality


**Key finding:** few-shot performance **scales much faster with model size** than zero-shot. Larger models are dramatically better at learning from context.

CoQA (QA): 85.0 F1 few-shot · Competitive with fine-tuned SOTA on many benchmarks

Translation, arithmetic, code generation, common sense — all from prompting alone

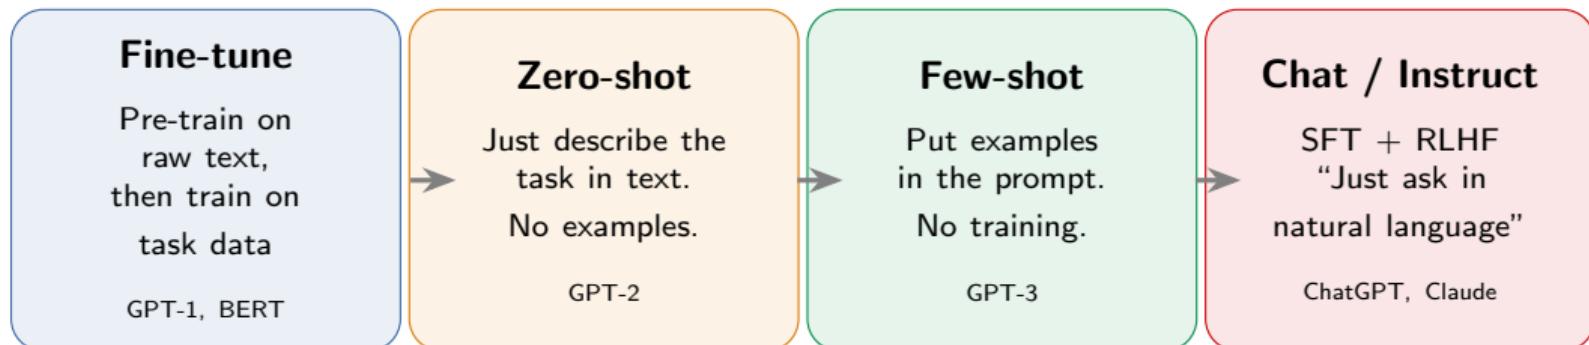
# GPT-3 — few-shot scales with model size

Accuracy vs. model size on various tasks



# The evolution: from fine-tuning to prompting

How we use language models evolved as they scaled



**Trend:** task-specific data requirements **decreased** as model scale **increased**

10k labeled examples → 100 examples → 10 examples → 0 examples (just instructions)

**T5's text-to-text anticipated this:** if every task is text-in/text-out, the interface naturally becomes a prompt. Scale made it work without training.

# Data matters: the data scaling story

Model	Dataset	Size	Key characteristic
ELMo	1B Word Benchmark	1B words	Curated news text
GPT-1	BooksCorpus	800M words	7,000 unpublished books
BERT	Books + Wikipedia	3.3B words	Two curated sources
GPT-2	WebText	40 GB	Reddit-filtered web (3+ karma)
T5	C4	750 GB	Cleaned Common Crawl
GPT-3	Mixed (5 sources)	300B tokens	Weighted sampling strategy
LLaMA	Mixed (8 sources)	1.4T tokens	Publicly available data only
LLaMA 2	Unknown	2T tokens	Undisclosed mix

## Quality > quantity:

T5 showed filtered C4 beats unfiltered Common Crawl.

GPT-3 upsampled high-quality sources (Wikipedia 3×).

## Diversity matters:

GPT-3 mixed: web, books, Wikipedia, code.  
Broad coverage → broader capabilities.

From 1B words (2018) to 2T tokens (2023) — a  $2000\times$  increase in training data

## Summary: what each model taught us

Model	Year	Architecture	Params	Key contribution
ELMo	2018	biLSTM	94M	Contextual embeddings (same word → different vectors)
GPT-1	2018	Decoder	117M	Pre-train + fine-tune paradigm
BERT	2018	Encoder	340M	Bidirectional (MLM), SOTA on 11 tasks
GPT-2	2019	Decoder	1.5B	Zero-shot emerges from scale
T5	2019	Enc-Dec	11B	Text-to-text unification, systematic study
GPT-3	2020	Decoder	175B	Few-shot in-context learning, API launch

### The themes that emerged:

1. Scale unlocks new capabilities (GPT-1 → 2 → 3)
2. Pre-training on raw text transfers to any task (GPT-1, BERT)
3. Architecture choice matters less than scale (T5 study)
4. Data quality is as important as quantity (T5, GPT-3)
5. Prompting can replace fine-tuning at sufficient scale (GPT-3)

# Further reading

## GPT Series

- Radford et al. (2018), "Improving Language Understanding by Generative Pre-Training" (GPT-1)
- Radford et al. (2019), "Language Models are Unsupervised Multitask Learners" (GPT-2)

~~• Brown et al. (2020). "Language Models are Few-Shot Learners" (GPT-3)~~

## BERT & Variants

- Devlin et al. (2019), "BERT: Pre-training of Deep Bidirectional Transformers"
- Liu et al. (2019), "RoBERTa: A Robustly Optimized BERT Pretraining Approach"

## T5 & Encoder-Decoder

- Raffel et al. (2020), "Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer" (T5)
- Lewis et al. (2020), "BART: Denoising Sequence-to-Sequence Pre-training"

# Questions?

Next: Prompting — Zero-shot, Few-shot, Chain-of-Thought