

## From N-grams to LSTMs

N-grams · Word2Vec · RNNs · LSTMs · Motivation for Attention

# The core problem

“The cat sat on the mat”



?

## Representation

How do we turn words  
into numbers that  
capture meaning?

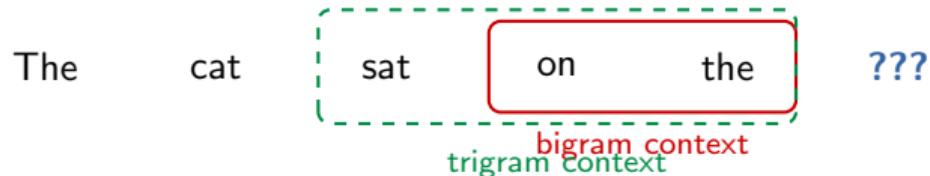
## Sequence modeling

How do we capture con-  
text and word order?



# N-gram language models

$$P(w_t \mid w_{t-n+1}, \dots, w_{t-1}) = \frac{\text{count}(w_{t-n+1} \dots w_t)}{\text{count}(w_{t-n+1} \dots w_{t-1})}$$



## Bigram:

$$P(\text{mat} \mid \text{the}) = \frac{\text{count}(\text{"the mat"})}{\text{count}(\text{"the"})}$$

## Trigram:

$$P(\text{mat} \mid \text{on, the}) = \frac{\text{count}(\text{"on the mat"})}{\text{count}(\text{"on the"})}$$

**Idea:** estimate probabilities by counting how often word sequences appear in a large corpus

# N-gram limitations

12 words apart — trigram window can't see "cat"

"The **cat** that sat on the mat near the door by the window **was** \_\_\_\_\_"

## 1. Fixed context window

A trigram only sees 2 previous words. Long-range dependencies are invisible.

## 2. Curse of dimensionality

Vocabulary  $V = 50k \Rightarrow$  possible 5-grams:  $50k^5 = 3 \times 10^{23}$ . Most never observed.

## 3. Data sparsity

"armadillo aardvark" has count 0 in most corpora. Smoothing helps but doesn't solve it.

## 4. No generalization

Knowing "dog sat on" doesn't help predict after "cat sat on." Words are discrete symbols.

We need a way to represent words so that similar words share information.

# One-hot encoding

## Vocabulary

1. cat

2. dog

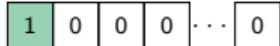
3. sat

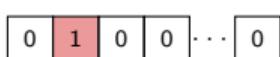
4. the

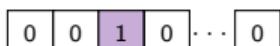
:

50k. zoo

## One-hot vectors

cat = 

dog = 

sat = 

50,000 dimensions

### No similarity

$\text{cat} \cdot \text{dog} = 0$

$\text{cat} \cdot \text{banana} = 0$

Every pair is equally different!

### Huge & sparse

50k-dim vectors with exactly one non-zero entry

We need: dense, low-dimensional vectors  
where similar words have similar representations

# Word2Vec — the distributional hypothesis

“You shall know a word by the company it keeps”  
— J. R. Firth, 1957

Words appearing in similar contexts have similar meanings:

“The movie was surprisingly **good** and **the audience** loved it”

“The movie was surprisingly **great** and **the audience** loved it”

“The movie was surprisingly **excellent** and **the audience** loved it”

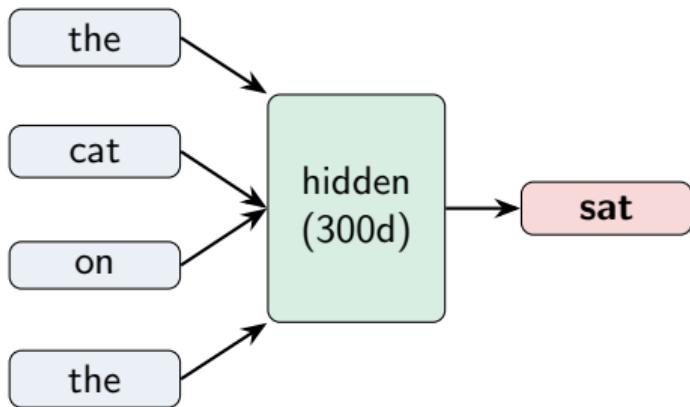


**good**, **great**, and **excellent** share context ⇒  
their vectors should be close in embedding space

# Word2Vec — CBOW & Skip-gram

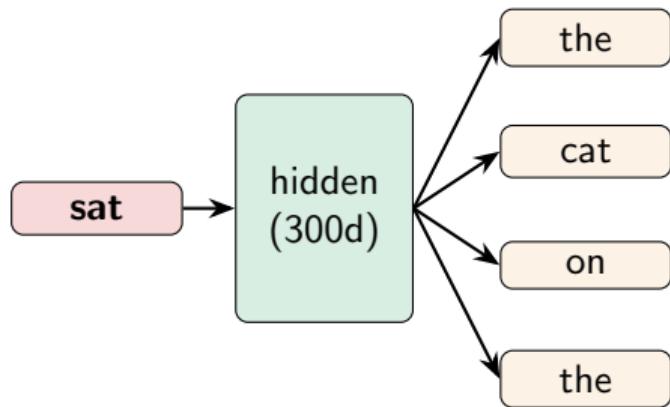
## CBOW

Context → Center word



## Skip-gram

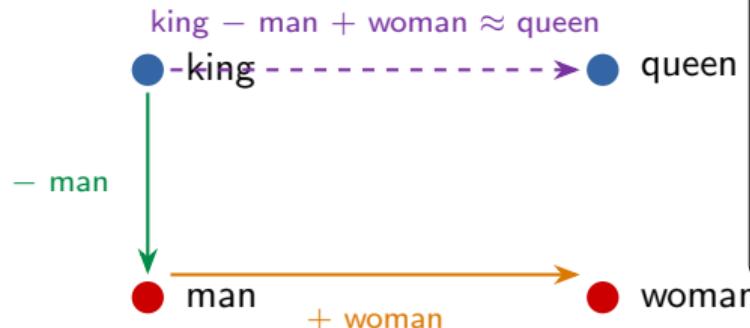
Center word → Context



**Result:** the hidden layer weights *are* the word embeddings — dense vectors (100–300 dims)

# Word2Vec — emergent properties

## Vector arithmetic in embedding space



### More analogies:

Paris – France + Italy  $\approx$  Rome  
bigger – big + small  $\approx$  smaller  
walking – walk + swim  $\approx$  swimming

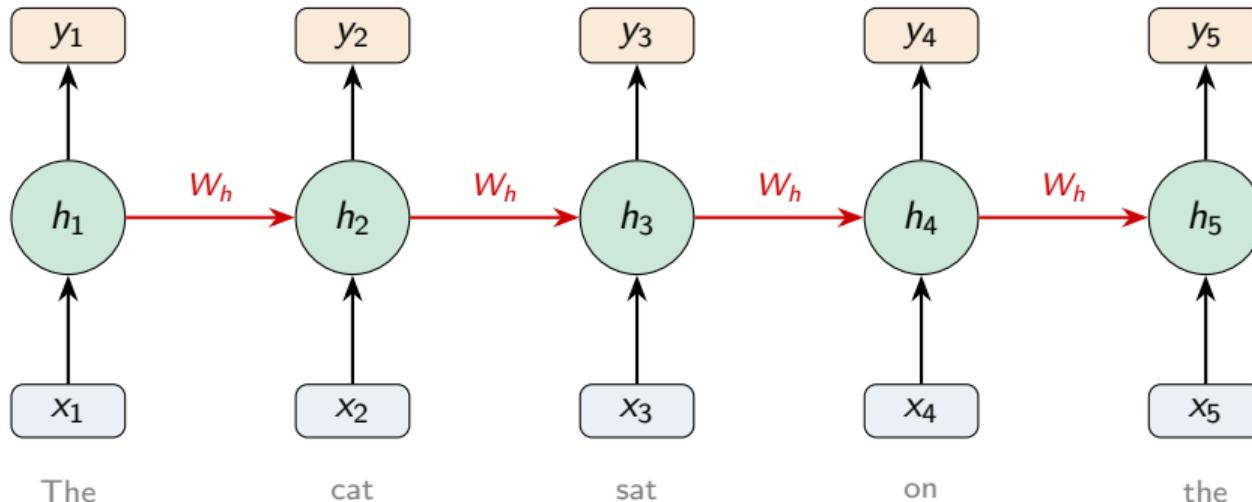
### Critical limitation: one vector per word

"I went to the river **bank**" vs. "I deposited money at the **bank**"  
Both map to the same vector — no polysemy, no context-dependence

We need representations that change based on context. Enter: recurrent neural networks.

# Recurrent Neural Networks (RNNs)

Unrolled RNN across time steps

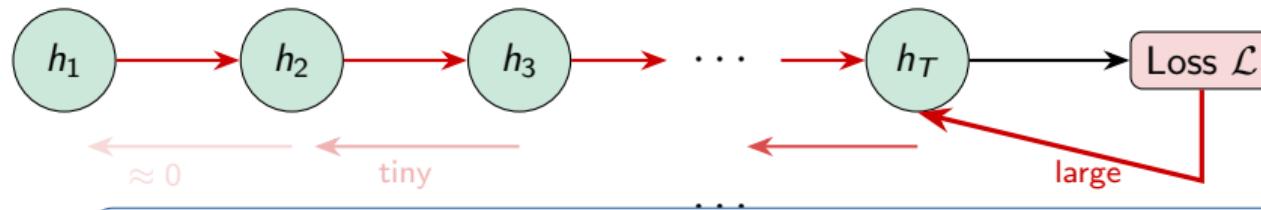


$$h_t = \tanh(W_h \cdot h_{t-1} + W_x \cdot x_t + b)$$

**Key:** hidden state  $h_t$  is a “memory” — no fixed window!

# The vanishing gradient problem

## Backpropagation through time



$$\frac{\partial \mathcal{L}}{\partial h_1} = \frac{\partial \mathcal{L}}{\partial h_T} \cdot \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}}$$

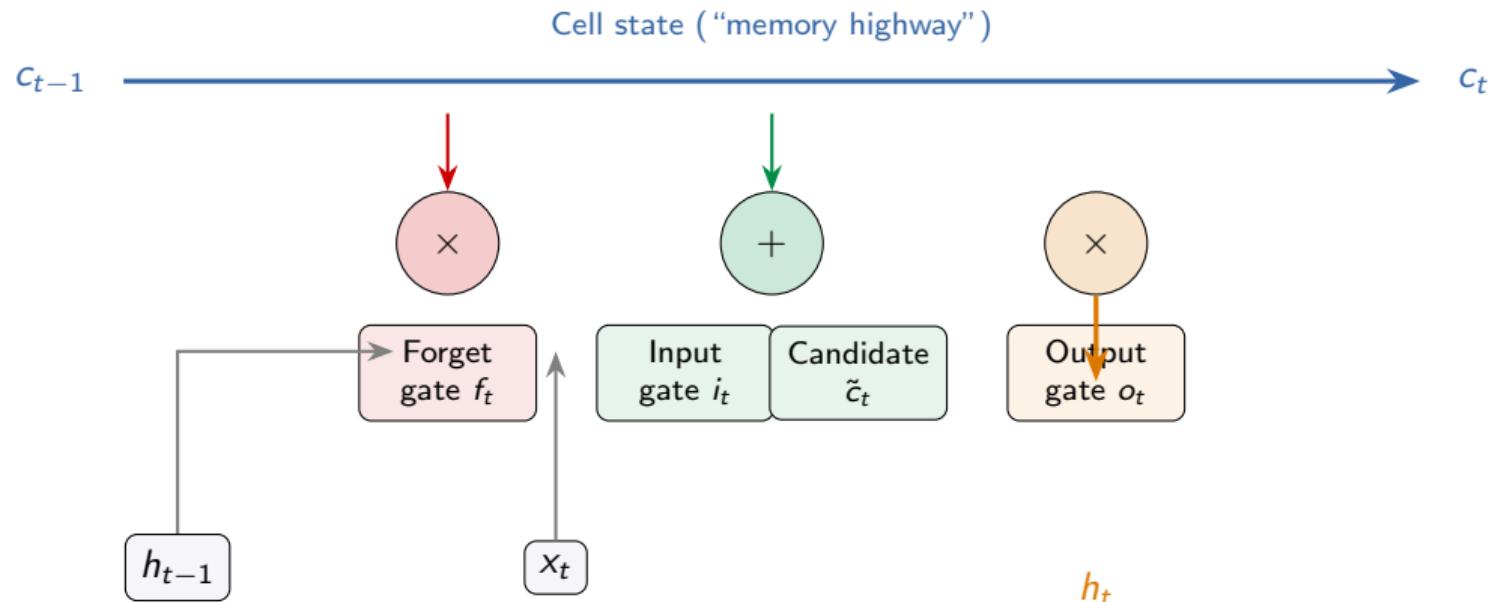
Product of many values  $< 1 \Rightarrow$  gradient vanishes

“**The cat** that I saw yesterday in the garden near the old oak tree **was** cute”

The gradient from “**was**” can’t reach “**cat**” — the RNN *forgets* early tokens

Solution attempt: LSTM (1997) and GRU (2014)

# LSTM — Long Short-Term Memory

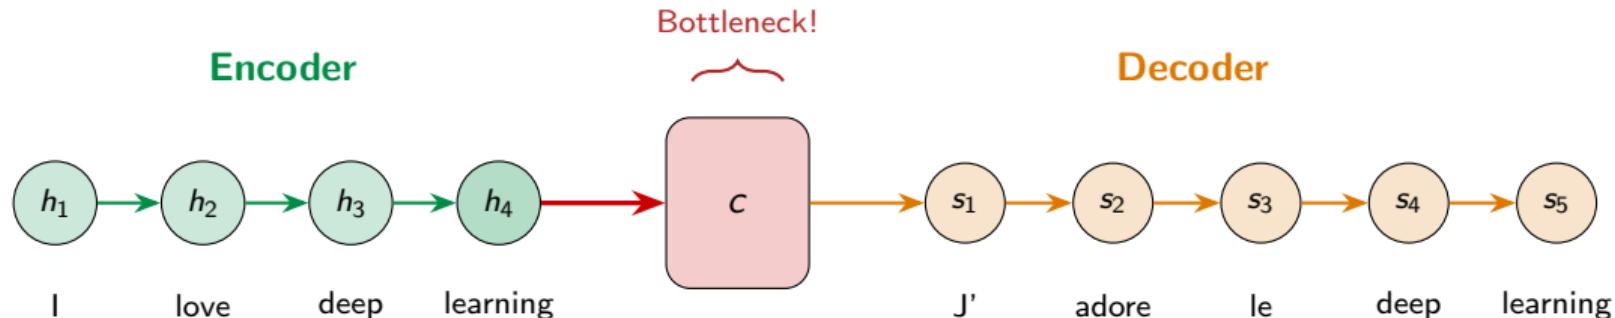


**Cell state** flows with minimal modification — gradient highway

**Gates** learn what to forget, store, and output at each step

# Sequence-to-sequence (Encoder–Decoder)

Machine translation with LSTMs (Sutskever et al., 2014)



The entire input sentence is compressed into a single fixed-size vector  $c$ !

Works for short sentences, but performance degrades for long ones.

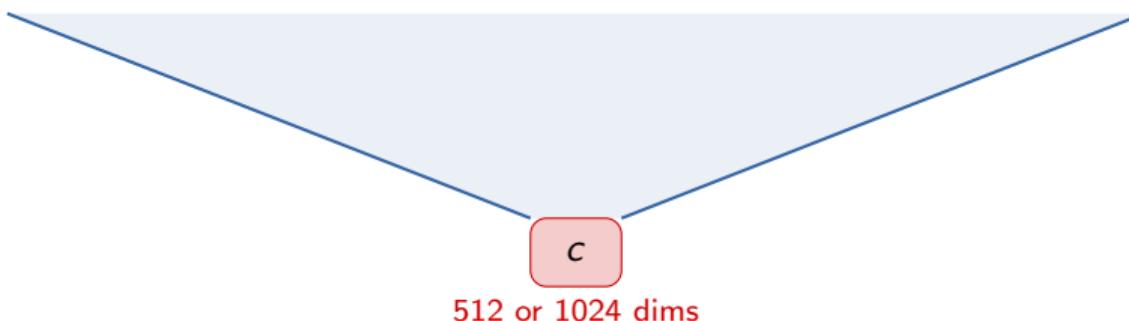


What if the decoder could look back at all encoder states, not just  $c$ ?

→ This is attention!

## Why LSTMs are not enough — the bottleneck

“The quick brown fox jumped over the lazy dog that  
was sleeping by the fireplace in the old cottage”



### Information lost:

Early tokens get overwritten by later ones as  $h_t$  is updated sequentially

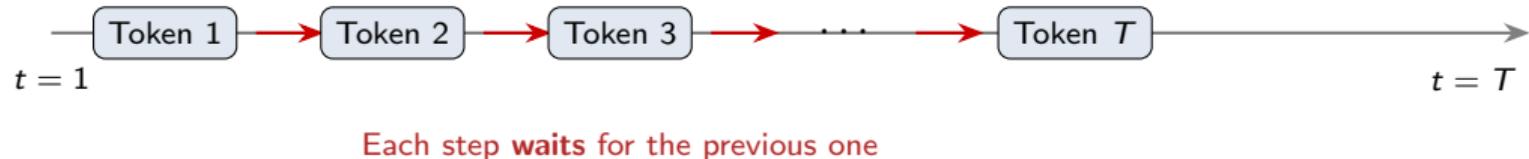
### Fixed capacity:

Whether the input is 5 tokens or 500, it must fit in the same vector

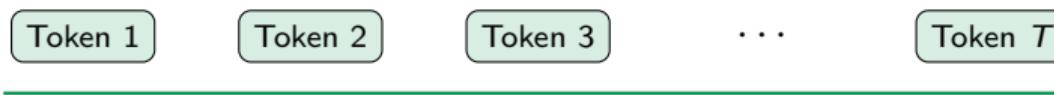
**Fundamental issue:** you can't losslessly compress arbitrary-length sequences into a fixed-size vector

# Why LSTMs are not enough — sequential processing

LSTM: must process tokens one at a time



What if we could process all tokens in parallel?



**LSTM:**  $O(T)$  sequential steps

Can't utilize GPU parallelism

Training is *slow*

**Transformer:**  $O(1)$  parallel steps

Fully utilizes GPU cores

Training is *fast*

# Why LSTMs are not enough — long-range dependencies

“The **trophy** didn't fit in the **suitcase** because **it** was too \_\_\_\_\_”

If **it** = **trophy**:

“it was too **big**”

If **it** = **suitcase**:

“it was too **small**”

Path length from “it” to its referent:

**LSTM:  $O(n)$  steps**

Information must pass through  
every intermediate hidden state

Signal degrades over distance

**Attention:  $O(1)$  steps**

Direct connection be-  
tween any two tokens

No degradation over distance

**Attention:** “let me directly look at any token I need” — regardless of distance

# The stage is set

## N-grams

- Fixed window
- No generalization
- Data sparsity
- Can't scale

## Word2Vec

- Good embeddings
- But static (one vector per word)
- No sequence modeling

## RNN / LSTM

- Sequential ( $\times$  parallel)
- Bottleneck vector
- Long-range deps still hard



We need: **parallel processing, direct access to all positions, contextual representations**

⇒ **Self-Attention** and the **Transformer** (Vaswani et al., 2017)

# Questions?

Next: The Transformer Architecture