

# Scaling Laws

Kaplan · Chinchilla · Emergent Abilities · Inference-Time Scaling

# What are scaling laws?

**Discovery:** language model performance follows **smooth power-law relationships** as you scale up three key variables — predictably, over many orders of magnitude.

## Model Size $N$

Number of parameters  
(non-embedding)

117M → 175B → 1T+

## Data Size $D$

Number of training  
tokens

1B → 300B → 15T

## Compute $C$

Total FLOPs spent  
on training

$10^{18} \rightarrow 10^{24}$  FLOPs

**Power law:**  $L = \left(\frac{K}{X}\right)^\alpha$  where  $L = \text{loss}$ ,  $X \in \{N, D, C\}$

On a log-log plot, this is a straight line:  $\log L = -\alpha \log X + \text{const}$

**Why this matters:** you can **predict** the performance of a 100B model by extrapolating from smaller experiments. This saves millions of dollars.

Kaplan et al. (2020), "Scaling Laws for Neural Language Models" — OpenAI

# Kaplan et al. (2020): the power-law relationships

## Model size:

$$L(N) = \left(\frac{N_c}{N}\right)^{0.076}$$

$$N_c = 8.8 \times 10^{13}$$

10× params →  
19% less loss

## Data size:

$$L(D) = \left(\frac{D_c}{D}\right)^{0.095}$$

$$D_c = 5.4 \times 10^{13}$$

10× data → 24% less loss

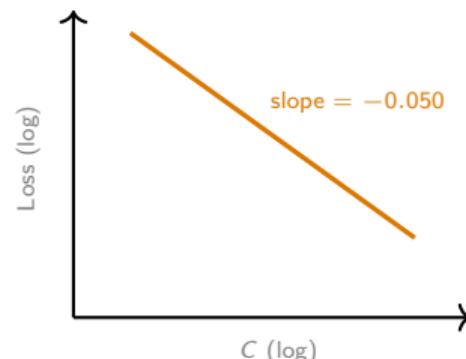
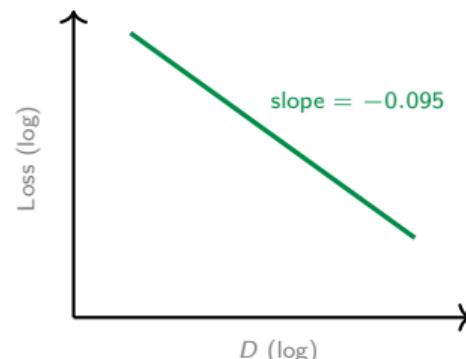
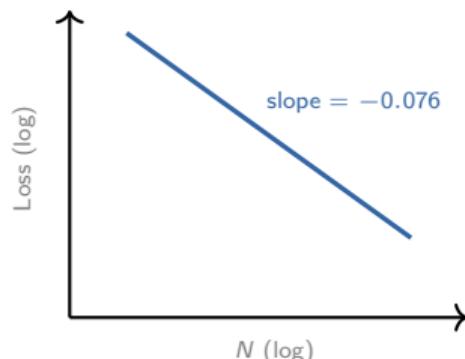
## Compute:

$$L(C) = \left(\frac{C_c}{C}\right)^{0.050}$$

$$C_c = 3.1 \times 10^8 \text{ PF-days}$$

10× compute →  
12% less loss

On a log-log plot, each relationship is a straight line:



# Kaplan: how to spend your compute budget

Given a fixed compute budget  $C$ , how should you split it between  $N$  and  $D$ ?

Since  $C \approx 6ND$ , making  $N$  bigger means less data (fewer passes), and vice versa.

$$N_{\text{opt}} \propto C^{0.73} \quad D_{\text{opt}} \propto C^{0.27}$$

“Train big models on relatively little data”

**10× more compute:**

Increase model size by  $\sim 5.4\times$

Increase data by only  $\sim 1.9\times$

$\approx 1.7$  tokens per parameter

**Key insight:**

Larger models are more **sample-efficient** — they reach lower loss with fewer tokens

So: prioritize model size!

**GPT-3 followed Kaplan:** 175B parameters trained on only 300B tokens  
 $\approx 1.7$  tokens per parameter. **This turned out to be sub-optimal.**

Problem: Kaplan counted only non-embedding parameters, biasing results at smaller scales

The compute approximation:  $C \approx 6ND$

$$C \approx 6 \cdot N \cdot D$$

Total FLOPs  $\approx 6 \times \text{parameters} \times \text{training tokens}$

Where does the 6 come from?

$\times 2$

Each parameter involves  
a multiply + accumulate  
 $= 2$  FLOPs per param  
per token (forward pass)

$\times 3$

Backward pass costs  
 $\approx 2 \times$  the forward pass  
Forward:  $1 \times$ ,  
Backward:  $2 \times$   
Total:  $3 \times$  forward

$= 6$

$2 \text{ (FLOPs/param/token)} \times 3 \text{ (fwd + bwd)}$   
 $= 6$  FLOPs per parameter  
per token

**Example — GPT-3:**  $N = 175B$ ,  $D = 300B$  tokens

$$C \approx 6 \times 175 \times 10^9 \times 300 \times 10^9 = 3.15 \times 10^{23} \text{ FLOPs} \approx 3,640 \text{ PetaFLOP-days}$$

**The fundamental trade-off:** for fixed  $C$ , increasing  $N$  means decreasing  $D$  (and vice versa).

Scaling laws tell us the **optimal split** that minimizes loss.

# Chinchilla (2022): Kaplan was wrong

**Hoffmann et al. (DeepMind):** trained  
400+ models from 70M to 16B params.

Found that Kaplan's recommendation to prioritize model size was **sub-optimal**.

**Three independent approaches, same conclusion:**

## Approach 1

Fix model sizes,  
vary number of tokens.  
Find minimum loss

for each compute level

## Approach 2

IsoFLOP profiles:  
fix compute budget,  
vary model size.

Find optimal  $N$  per  $C$

## Approach 3

Fit parametric model:  
 $L(N, D) = E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$   
Minimize analytically.

$$N_{\text{opt}} \propto C^{0.50} \quad D_{\text{opt}} \propto C^{0.50}$$

**“Scale model size and data equally”**

$$D_{\text{opt}} \approx 20 \times N$$
$$\approx 20 \text{ tokens per parameter}$$

Hoffmann et al. (2022), “Training Compute-Optimal Language Models” — NeurIPS 2022

# Chinchilla vs. Gopher: the empirical proof

	Same compute budget	Gopher	Gopherically different	Chinchilla
Parameters	280B		70B	
Training tokens	300B		1.4T	
Tokens / parameter		~1.07		~20
Training compute	$5.76 \times 10^{23}$ FLOPs		$5.76 \times 10^{23}$ FLOPs	
MMLU accuracy		60.0%		67.5%

Parameters:

Gopher: 280B

Chinchilla: 70B

Training data:

300B

Chinchilla: 1.4T

A  $4\times$  smaller model trained on  $4\times$  more data, using the same compute, uniformly outperformed Gopher, GPT-3 (175B), Jurassic-1 (178B), and MT-NLG (530B)

Inference cost is also  $4\times$  cheaper ( $4\times$  fewer parameters to serve)

# Kaplan vs. Chinchilla: the key differences

	Kaplan (2020)	Chinchilla (2022)
$N_{\text{opt}}$ scaling	$C^{0.73}$	$C^{0.50}$
$D_{\text{opt}}$ scaling	$C^{0.27}$	$C^{0.50}$
Tokens per parameter	$\sim 1.7$	$\sim 20$
Philosophy	"Train big, stop early"	"Scale $N$ and $D$ equally"
Parameter counting	Non-embedding only	All parameters
Experimental scale	Smaller models	Up to 16B (validated at 70B)

## Parameter counting

Kaplan excluded embedding params, biasing results at smaller scales.

## Why do they disagree?

### Experimental scale

Kaplan used smaller models.  
Chinchilla trained 400+ models up to 16B

### LR schedule

Different hyperparameter tuning approaches led to different optima

**Consensus today:** Chinchilla scaling ( $D \approx 20N$ ) is the accepted baseline.

But in practice, labs now **over-train** beyond even Chinchilla-optimal (see later).

# The Chinchilla loss model

$$L(N, D) = E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

$$E = 1.69$$

Irreducible loss  
(entropy of natural)

$$\frac{A}{N^\alpha}$$

$$A = 406.4$$
$$\alpha = 0.34$$

$$\frac{B}{D^\beta}$$

$$B = 410.7$$
$$\beta = 0.28$$

**Optimize:**

Given  $C = 6ND$ ,  
minimize  $L$

**Interpretation:** loss has three components — an irreducible floor  $E$ ,  
a penalty for too few parameters ( $A/N^\alpha$ ),  
and a penalty for too little data ( $B/D^\beta$ ).

The optimal model balances both penalties equally.

**Result:** minimizing  $L$  subject to  $C = 6ND$  gives

$$N_{\text{opt}} = G \cdot \left(\frac{C}{6}\right)^a, \quad D_{\text{opt}} = G^{-1} \cdot \left(\frac{C}{6}\right)^b \quad \text{where } a \approx b \approx 0.5$$

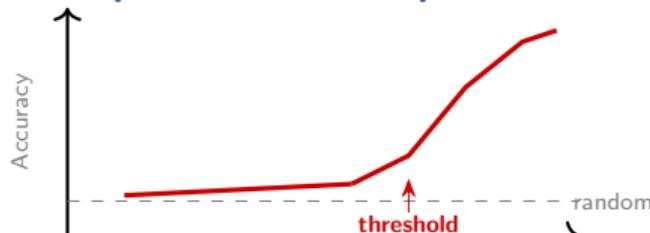
The near-equal exponents ( $a \approx b$ ) mean  $N$  and  $D$  should scale at the same rate

# Emergent abilities (Wei et al., 2022)

An ability is emergent if it is absent in smaller models but appears in larger ones

— it cannot be predicted by extrapolating from smaller-scale performance.

## The phase transition pattern:



Ability	Threshold
Multi-digit arithmetic	~100B params
Word unscrambling	~100B params
Chain-of-thought	~100B params
Instruction following	~100B+
College exams (MMLU)	~100B+

Wei et al. documented **137 emergent abilities** across BIG-Bench and other benchmarks.

Near-random performance until a critical scale, then a sharp jump.

**But is this real?** Are these genuine phase transitions in capability, or an artifact of how we measure performance?

Wei et al. (2022), "Emergent Abilities of Large Language Models" — TMLR

# Are emergent abilities a mirage?

**Schaeffer et al. (2023):** emergent abilities are an artifact of **metric choice**, not fundamental changes in model behavior. NeurIPS 2023 Oral.

## Exact-match accuracy

Discontinuous metric:  
either 100% correct or 0%

$"2 + 3 = 5"$  → correct  
 $"2 + 3 = 6"$  → wrong

Shows sudden emergence

→  
same model!

## Token-level log-likelihood

Continuous metric:  
measures partial knowledge

Probability of correct token  
increases *smoothly* with scale

Shows gradual improvement

**Key finding:** 92%+ of BIG-Bench “emergent” abilities use nonlinear metrics.  
Switch to continuous metrics → smooth scaling. No phase transition.

**Current resolution:** capability improves **smoothly** (continuous metrics),  
but practical task success can appear **sudden** (discontinuous metrics).  
Both perspectives are “correct” — it depends on what you measure.

Schaeffer et al. (2023), “Are Emergent Abilities of Large Language Models a Mirage?” — NeurIPS 2023

# Scaling laws beyond language

**Power-law scaling is not language-specific.** The same functional form  $L = (K/X)^\alpha$  holds across domains with remarkably similar exponents.

## Language

GPT, LLaMA  
Text generation  
 $\alpha_N \approx 0.08$

## Vision

ViT models  
Image classification  
Zhai et al. 2022

## Video

Video generation  
Autoregressive  
Henighan 2020

## Multimodal

Image + text  
CLIP, etc.  
Shukor 2025

## A new scaling axis: inference-time compute

**Snell et al. (2024):** instead of making models bigger, spend more compute at **inference time** — search over reasoning paths, verify with reward models.

A smaller model + test-time compute can outperform a **14 $\times$  larger** model.

**OpenAI o1 (2024):** AIME math competition 15.6%  $\rightarrow$  **83.3%** by scaling inference compute. o1 “thinks longer” on harder problems (chain-of-thought at test time).

Three axes of scaling: pre-training compute (Kaplan/Chinchilla), data quality, inference compute

# The over-training trend

**Chinchilla optimizes training efficiency.**

But in deployment, you pay per inference.

A ~~smaller~~ model trained ~~longer~~ takes ~~more~~ tokens gives ~~better~~ inference quality Chinchilla (~~2021~~) serving cost.

	Params	Total tokens	Time	Training cost	Inference cost
GPT-3 (2020)	175B	300B	1.7	12× under-trained	
Gopher (2021)	280B	300B	1.1	19× under-trained	
Chinchilla (2022)	70B	1.4T	20	1× optimal	
LLaMA-1 7B (2023)	7B	1.0T	143	7× over-trained	
LLaMA-1 65B (2023)	65B	1.4T	22	~1× near-optimal	
LLaMA-2 70B (2023)	70B	2.0T	29	1.4× over-trained	
LLaMA-3 8B (2024)	8B	15T	1,875	94× extreme over-trained	

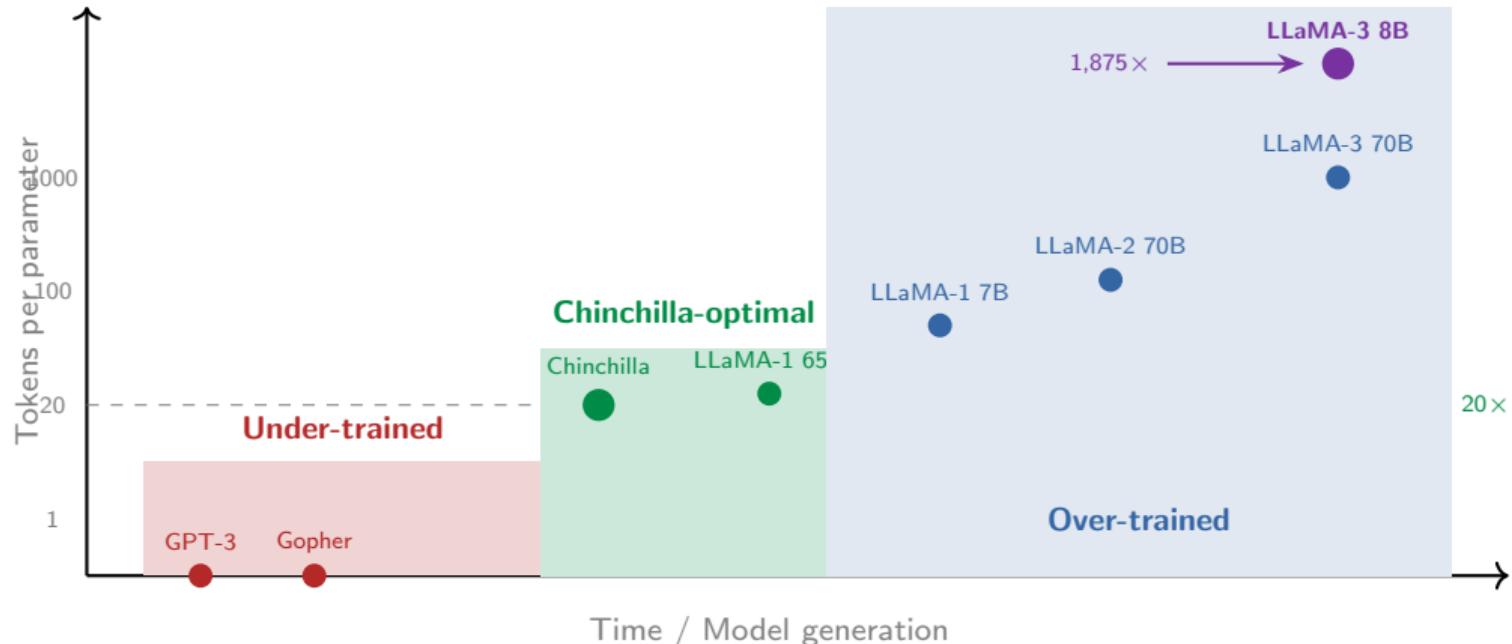
**Key insight:** LLaMA-1 7B (trained on 1T tokens) was competitive with **GPT-3 175B**  
~~on many benchmarks, 25× smaller → 25× cheaper to serve!~~

**The principle:** training is a one-time cost.

Inference is paid **every request, forever.**

Over-training = spend more on training to save on inference at scale.  
~~Loss keeps improving well past Chinchilla-optimal, just with diminishing returns~~

# The three eras of compute allocation



# The scaling playbook

- 1. Performance scales as a power law** with model size, data, and compute.

$L(N) \propto N^{-0.076}$ ,  $L(D) \propto D^{-0.095}$ ,  $L(C) \propto C^{-0.050}$  (Kaplan 2020)

- 2. Scale  $N$  and  $D$  equally** for compute-optimal training ( $D \approx 20N$ ).

Chinchilla 70B beat Gopher 280B with the same compute budget (Hoffmann 2022)

- 3. In practice, over-train** smaller models for cheaper inference.

LLaMA-3 8B: 15T tokens (1,875 tok/param). Training cost is one-time; inference is forever.

- 4. New capabilities emerge at scale** (whether real or metric-dependent).

137 emergent abilities documented (Wei 2022). Debate ongoing (Schaeffer 2023).

- 5. Inference-time compute is the new frontier.**

Smaller model + more thinking at test time can beat 14 $\times$  larger model (Snell 2024).

# Questions?

Next: Hallucination & Grounding