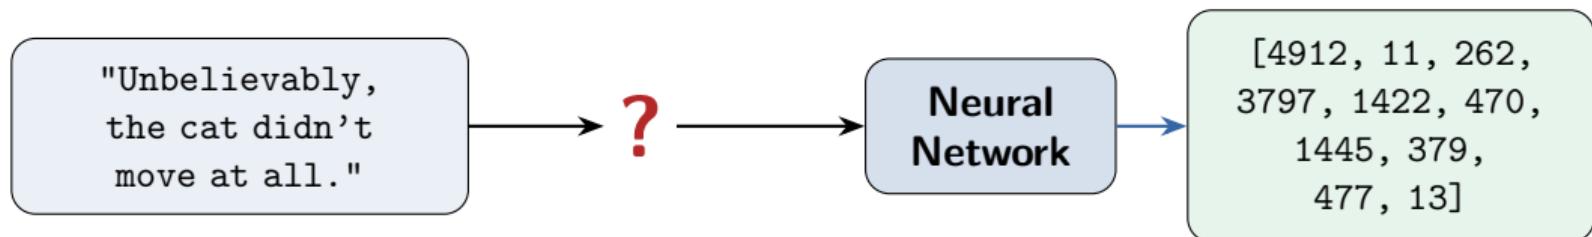


Tokenization

Word · Character · Subword · BPE · WordPiece · SentencePiece

Models need numbers, not text



Tokenization is the first step in any NLP pipeline:
split raw text into discrete units (tokens) and map each to an integer ID.
The choice of tokenizer affects **everything** downstream.

Three levels of granularity

Word-level

```
["Unbelievably",  
 "the", "cat",  
 "didn't", "move"]
```

Vocab size: ~100k+

OOV problem

Subword

```
["Un", "believ",  
 "ably", "", "the",  
 "cat", "didn", "'t"]
```

Vocab size: ~30k–50k

The sweet spot

Character

```
["U", "n", "b", "e",  
 "l", "i", "e", "v",  
 "a", "b", "l", "y", ...]
```

Vocab size: ~256

Very long sequences



Used by all modern LLMs

Word-level tokenization and the OOV problem

Vocabulary (fixed at training):

the, cat, sat, on, mat,
dog, run, happy, sad, ...

Size: 50,000–200,000 words

At test time:

"The **cryptocurrency** market
plummeted after the
CEO's tweet."



Words not in vocab → [UNK] [UNK]
market [UNK] after the [UNK] tweet.

Problems:

- New words, names, typos → [UNK]
- Huge vocabulary → huge embedding matrix
- Morphology lost: "run", "runs", "running" are unrelated tokens

Character-level: no unknowns, but...

Input: "The cat sat on the mat."

The | cat | sat | on | the | mat | .

7 tokens (word-level)

Drawbacks:

- Sequences are $\sim 4\text{--}5\times$ longer
- Self-attention is $O(n^2)$, so cost grows quadratically
- Each character carries little semantic meaning
- Harder to learn long-range dependencies

T|h|e| |c|a|t| |s|a|t| |o|n| |t|h|e| |m|a|t|.

23 tokens (character-level)

Advantages:

- Tiny vocabulary (~ 256)
- Zero unknown tokens
- Works for any language
- Handles typos, code, URLs

Too fine-grained on its own — but the idea of starting from characters inspires **subword** methods.

Subword tokenization: the key insight

Common words stay whole: the, cat, and

Rare words are split into known pieces: un + believ + ably

“playing” → play + ing

“unhappiness” → un + happi + ness

“ChatGPT” → Chat + G + PT

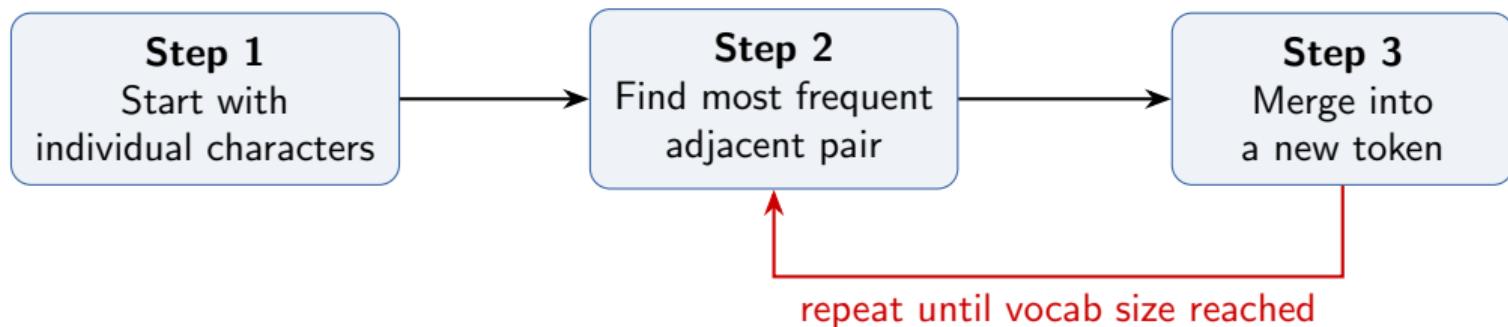
“brrrr” → br + rr + r

Vocab size ~30k–50k
able sequence lengths

- No [UNK] tokens
- Reason-
- Morphology is partially captured

Byte-Pair Encoding (BPE): the idea

Training: learn merge rules from a corpus. **Inference:** apply merge rules to new text.



Originally a **data compression** algorithm (Gage, 1994).
Adopted for NLP by Sennrich et al. (2016). Used
by **GPT**, **GPT-2**, **RoBERTa**, **BART**, **LLaMA**.

BPE: worked example

Corpus: hug (10), pug (5), pun (12), bun (4), hugs (5)

Initial vocab: h, u, g, p, n, b, s

Splits: hug (10) pug (5) pun (12) bun (4)
hugs (5)

Merge 1: most frequent pair = (u,g) freq =

$$10+5+5 = 20$$

Splits: **h ug** (10) **p ug** (5) **p un** (12) **b un** (4)
h ug s (5)

Merge 2: most frequent pair = (u,n) freq =

$$12+4 = 16$$

Splits: **h ug** (10) **p ug** (5) **p un** (12) **b un** (4)
h ug s (5)

Merge 3: most frequent pair = (h,ug) freq =

$$10+5 = 15$$

Splits: **hug** (10) **p ug** (5) **p un** (12) **b un** (4)
hug s (5)

Merge rules (ordered):

1. u + g → ug
2. u + n → un
3. h + ug → hug

⋮

Vocab after 3 merges:

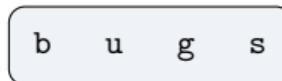
h, u, g, p, n, b, s,
ug, un, hug

BPE: tokenizing new text

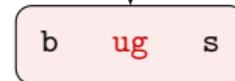
Given the learned merge rules, tokenize a new word:

Tokenize “bugs”:

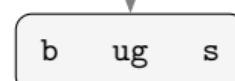
Start:



Rule 1 (u+g):

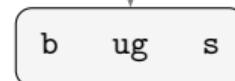


Rule 2 (u+n):



no match

Rule 3 (h+ug):



no match

Key point:

Apply merge rules in the *same order* they were learned.

Words not seen during training are still tokenized — just split into smaller known pieces.

Result: ["b", "ug", "s"] → IDs [5, 7, 6]

Byte-level BPE (GPT-2, GPT-3, LLaMA)

Standard BPE

Base vocab = Unicode characters

Vocab size = ~30k–50k

Unknown chars → [UNK]

Problem: Chinese, emoji, etc.
can hit unknown characters

upgrade
→

Byte-level BPE

Base vocab = **256 byte values**

Vocab size = 256 + merges
(GPT-2: 50,257 total)

No [UNK] ever

Any byte sequence is representable

Every text is ultimately a sequence of bytes (UTF-8 encoding).

By starting from **bytes** instead of characters, BPE can tokenize
any input: English, Chinese, Arabic, code, emoji,
binary data — all with the same vocabulary.

WordPiece (BERT, DistilBERT)

BPE

Merge criterion:

most frequent pair

Greedy count-based

Used by: GPT, LLaMA, etc.

WordPiece

Merge criterion:

pair that **maximizes likelihood** of the training corpus

Used by: BERT, DistilBERT

$$\text{score}(a, b) = \frac{\text{freq}(ab)}{\text{freq}(a) \times \text{freq}(b)}$$

Notation: WordPiece marks *continuation* subwords with ##

“unbelievably” → ["un", "##believ", "##ably"]

BPE instead marks *word-initial* subwords (e.g., GPT-2 uses G = space prefix)

Unigram LM and SentencePiece

Unigram LM

Start with a *large* vocab
Iteratively **remove** tokens
that hurt likelihood the least
Top-down (vs BPE's bottom-up)
Used by: T5, ALBERT, XLNet
Kudo, 2018

SentencePiece

Not an algorithm — a **library**
Treats input as raw byte stream
(no pre-tokenization needed)
Supports both **BPE** and **Uni-gram**
Language-agnostic: no need for
space-based word splitting
Kudo & Richardson, 2018

BPE = bottom-up (merge frequent pairs) vs

Unigram = top-down (prune unlikely tokens)

Both converge to similar subword vocabularies in practice.

Why tokenization matters: LLM quirks

Bad at arithmetic

"12345" → ["123", "45"]

The model never sees the individual digits together!

Poor non-English efficiency

English "hello" = 1 token

Korean "annyeong" = 3–5 tokens

Same meaning, 3–5× the cost!

Sensitive to formatting

"Hello World" and

"Hello World" produce different token sequences.

Can't count letters

"How many r's in strawberry?"

"straw" + "berry" — the model can't see individual letters.

Many apparent "reasoning failures" of LLMs are actually **tokenization artifacts**.

The model literally cannot see what you think it sees.

Visualizing: the same sentence, different tokenizers

Input: “The cat sat on the unbelievably soft mat”

GPT-2 (E The : _cat _sat _on _the _un believ ably _soft _mat **10 tok**

BERT (Wor the ec cat sat on the un ##bel ##ie ##va ##bly soft mat **12 tok**

Character: The_cat_sat_on_the_unbelievably_soft_mat **39 tok**

Same input, very different representations. “unbelievably”
= **3 tokens** (GPT-2), **5 tokens** (BERT), **12 tokens** (char).

Special tokens

[CLS]
Classification token (BERT)

[SEP]
Separator between segments

[PAD]
Padding to equal length

[UNK]
Unknown token (fallback)

⟨BOS⟩
Beginning of sequence

⟨EOS⟩
End of sequence

[MASK]
Masked position (BERT MLM)

Special tokens are **not** in the original text — they're added by the tokenizer to give the model structural signals: where sequences begin/end, what to predict, etc.

Comparison of tokenization methods

Method	Direction	Criterion	Vocab size	Used by
BPE	Bottom-up	Frequency	30k–50k	GPT, LLaMA
WordPiece	Bottom-up	Likelihood	30k	BERT
Unigram	Top-down	Likelihood	30k–50k	T5, XLNet
Byte BPE	Bottom-up	Frequency	50k–100k	GPT-2/3/4
Character	—	—	256	ByT5

In practice, the differences between BPE, WordPiece, and Unigram are **small**. What matters most: vocab size, training corpus, and whether byte-level is used.

Byte-level BPE is the current default for new large language models.

Practical: tokenizers in action

```
# GPT-4 tokenizer
import tiktoken
enc = tiktoken.encoding_for_model(
    "gpt-4")
tokens = enc.encode(
    "Hello world!")
# [9906, 1917, 0]
enc.decode(tokens)
# "Hello world!"
```

```
# BERT tokenizer
from transformers import
    AutoTokenizer
tok = AutoTokenizer.from_pretrained(
    "bert-base-uncased")
tok.tokenize(
    "unbelievably")
# ["un", "##bel", "##ie",
#  "##va", "##bly"]
```

Try it yourself: <https://platform.openai.com/tokenizer>

Paste any text and see how GPT tokenizes it. Pay attention to:
numbers, non-English text, code, and whitespace.

Further reading

Subword Tokenization

- Sennrich et al. (2016), “Neural Machine Translation of Rare Words with Subword Units” (BPE)
- Kudo & Richardson (2018), “SentencePiece: A simple and language independent subword tokenizer”
- Kudo (2018), “Subword Regularization: Improving Neural Network Translation Mod-

WordPiece & Byte-Level

- Schuster & Nakajima (2012), “Japanese and Korean Voice Search” (WordPiece)
- Wang et al. (2020), “Neural Machine Translation with Byte-Level Subwords”

Practical Resources

- HuggingFace Tokenizers library documentation —
huggingface.co/docs/tokenizers
- OpenAI Tokenizer tool — platform.openai.com/tokenizer

Questions?

Next: Evaluation — Perplexity, BLEU, ROUGE