

From hand-rolled boosting to the libraries that win

In [19] you built gradient boosting yourself: fit each tree to the *negative gradient* of the loss, shrink, repeat. The algorithm is finished.

What is left is **engineering** - the libraries that dominate tabular ML:

- ▶ **XGBoost** - a regularized objective + a fast, careful implementation;
- ▶ **LightGBM** - leaf-wise growth, smart sampling, blazing speed;
- ▶ **CatBoost** - categorical features done right.

Same gradient descent in function space - made **fast**, **regularized**, and **production-ready**.

Outline

XGBoost: the regularized objective and the structure score

Engineering: why it is fast

The modern trio: LightGBM and CatBoost

Using them well

Stacking: combining different models

XGBoost

A regularized objective, and a closed-form score for how good a tree is.

XGBoost: extreme gradient boosting

A scalable tree-boosting system (Chen & Guestrin, 2016):

- ▶ a **regularized objective** - penalties baked into the loss;
- ▶ **approximate / histogram** split finding for large data;
- ▶ row + column **subsampling**; **sparsity-aware** splits;
- ▶ parallel split search and **GPU** support.

Often the top performer on tabular benchmarks - **if properly tuned**. Many knobs, and the defaults are rarely optimal.

The plan, in plain words

Before any algebra, here is the whole idea:

1. **Score** how good a tree is by how much it would cut the loss - call it the *structure score*.
2. To get that score, **approximate the loss as a parabola** at the current prediction (a Newton step: use the slope *and* the curvature).
3. **Solve** that parabola for the best value in each leaf.
4. **Grow greedily**: keep a split only if it raises the score by more than the cost of adding a leaf.

The next slides just make each step precise. Hold on to this picture: **score a tree, then split to improve the score.**

A regularized objective

At step m we add one tree b with leaf values $c_1, \dots, c_T$. XGBoost minimizes the loss *plus* complexity penalties:

$$\text{obj} = \sum_{i=1}^n L(y_i, f^{[m-1]}(x_i) + b(x_i)) + \underbrace{\gamma T}_{\# \text{ leaves}} + \underbrace{\frac{1}{2}\lambda \|c\|_2^2}_{L_2} + \underbrace{\alpha \|c\|_1}_{L_1}.$$

- ▶ γ penalizes the **number of leaves** (tree size);
- ▶ λ (L_2) and α (L_1) shrink the **leaf values**.

Contrast [19]: plain gradient boosting put *no* penalty inside the objective. Here regularization is part of what every tree optimizes. (Software: `gamma`, `lambda`, `alpha`.)

Second-order Taylor of the loss

Expand the loss around the current prediction $f^{[m-1]}$, to **second order**:

$$L(y, f^{[m-1]} + b) \approx L(y, f^{[m-1]}) + g b + \frac{1}{2} h b^2,$$

with gradient and Hessian

$$g = \frac{\partial L}{\partial f^{[m-1]}}, \quad h = \frac{\partial^2 L}{\partial (f^{[m-1]})^2}.$$

g is exactly the **negative pseudo-residual** from [19]. XGBoost's one real addition is the **curvature** h . *Intuition*: picture a parabola hugging the loss - slope g points downhill, bend h says how big a step to trust (a Newton step).

Group the objective by leaf

Drop the constant $L(y, f^{[m-1]})$ and use that every point in leaf t shares one value c_t :

$$\text{obj} \approx \sum_{t=1}^T \left[G_t c_t + \frac{1}{2} (H_t + \lambda) c_t^2 \right] + \gamma T,$$

$$G_t = \sum_{i \in \text{leaf } t} g_i, \quad H_t = \sum_{i \in \text{leaf } t} h_i.$$

Intuition: every point in a leaf gets the **same** correction c_t , so a leaf is just **one number to choose** - and the objective is now a simple quadratic in it.

Optimal leaf value = the “structure score”

Minimizing each quadratic $G_t c_t + \frac{1}{2}(H_t + \lambda)c_t^2$ gives the best leaf value

$$c_t^* = -\frac{G_t}{H_t + \lambda}.$$

Plugging back, the best objective for a *fixed tree shape* is the **structure score**

$$\text{obj}^* = -\frac{1}{2} \sum_{t=1}^T \frac{G_t^2}{H_t + \lambda} + \gamma T.$$

Intuition: a leaf's best step = total error to fix (G_t) per unit of confidence ($H_t + \lambda$); the score adds up each leaf's payoff - **lower is a better tree**.

Numbers: $G_t = -3$, $H_t = 4$, $\lambda = 1 \Rightarrow c_t^* = -(-3)/(4+1) = 0.6$.


The split-gain criterion

How good is a split? By how much it **improves the structure score**:

$$\text{gain} = \frac{1}{2} \left[\frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda} \right] - \gamma.$$

- *Intuition*: do the two children together score better than the parent, by more than the cost γ of a new leaf?
- Take the split only if $\text{gain} > 0$; γ is the **minimum gain to keep it**.
- This - not Gini/variance - is what XGBoost maximizes at each node.

Numbers: parent $(G=-3, H=4)$, $\lambda=1$, split into $L(-3, 2)$, $R(0, 2)$: $\text{gain} = \frac{1}{2} \left[\frac{9}{3} + \frac{0}{3} - \frac{9}{5} \right] - \gamma = 0.6 - \gamma$ (keep only if $\gamma < 0.6$); then $c_L^*=1.0$, $c_R^*=0$.



$$G = G_L + G_R, \quad H = H_L + H_R$$

Why it is fast

The same gradient boosting, engineered for speed and scale.

Tree growing and pruning

- ▶ Grow each tree **level by level**, up to `max_depth`.
- ▶ A split is kept only if its **gain** (previous slide) is positive.
- ▶ After growing, **prune** back any split with $\text{gain} < 0$ - the γ threshold.


Depth-limited growth + post-pruning by gain keeps trees shallow and regularized.

Approximate (histogram) split finding

Testing every possible split is slow. Instead, bucket each feature into $\sim \ell$ **percentile bins** and evaluate only the **bin edges**.

- **Global**: propose the bins once per tree.
- **Local**: re-propose at every node (more accurate, more work).

“Histogram-based gradient boosting” - XGBoost `tree_method="hist"`; the default in LightGBM and sklearn `HistGradientBoosting`.



Candidates sit at percentile edges (denser where data is dense). Source: Chen & Guestrin (2016).

Sparsity-aware splits

Real data is full of **missing values** and **zeros** (one-hot columns).

XGBoost gives every split a learned **default direction**: rows missing (or zero) on the split feature all go that way.

- ▶ Handles missing values **natively** - no imputation needed.
- ▶ Skips the absent entries during split search, so sparse data is also **faster**.

Subsampling, parallelism, GPU

- ▶ **Row subsampling** (`subsample`) - stochastic gradient boosting.
- ▶ **Column subsampling** (`colsample_bytree/bylevel/bynode`) - like a random forest's `mtry`; faster and more diverse trees.
- ▶ Boosting is **sequential**, but the **split search inside one tree** parallelizes over sorted feature blocks - and moves well to **GPU**.

Subsampling is both a speed-up and a regularizer.

DART: dropout for trees

Borrow **dropout** from deep learning. When adding a new tree, compute the update against a **random subset** of the existing trees (drop the rest), then rescale so the ensemble does not overshoot.

- ▶ $p_{\text{drop}} = 0$: ordinary gradient boosting.
- ▶ $p_{\text{drop}} = 1$: trees trained independently, equally weighted $\approx$ a random forest.

p_{drop} is a smooth dial from **boosting** to **random forest** (XGBoost booster="dart").

The modern trio

XGBoost is the baseline. LightGBM (the favorite) and CatBoost push it further.

LightGBM: my favorite

For everyday tabular work, this is what I reach for first:

- ▶ **very fast** and **low memory** (histogram + leaf-wise);
- ▶ excellent **out-of-the-box defaults**;
- ▶ **native categorical** support;
- ▶ scales smoothly to **large** data.

Instructor's pick

My default first model on tabular data - the next frames are why.

Ke et al. (2017), *LightGBM*.

Me and LightGBM ❤️



<https://youtu.be/UqmqC10ZRwg>



Leaf-wise (best-first) growth

XGBoost grows **level-wise by default** (a whole depth level at a time). LightGBM instead splits the **single leaf with the largest gain**, anywhere in the tree.

- ▶ Unbalanced, deeper trees; **converges in fewer rounds**.
- ▶ Can **overfit** - so the main complexity knob is `num_leaves` (keep it $< 2^{\text{max_depth}}$), plus `min_data_in_leaf`.

Intuition: spend each split where it helps most, instead of splitting every leaf just because it is at the same depth.

Source: LightGBM docs.



GOSS: gradient-based one-side sampling

Random subsampling wastes the **informative** rows. GOSS samples smarter, by gradient:

- ▶ keep the top $a \cdot n$ rows with the **largest** **|gradient|** (still-hard cases);
- ▶ randomly sample $b \cdot n$ of the rest;
- ▶ **up-weight** those sampled rows by $\frac{1-a}{b}$ so the gain estimate stays unbiased.

Intuition: rows the model already gets right (tiny gradient) teach it little - spend the budget on the hard rows, but reweight the easy sample so the data still looks unbiased. Defaults $a = 0.2$, $b = 0.1$: each split sees $\sim 30\%$ of the rows.

Ke et al. (2017).

EFB: exclusive feature bundling

In sparse, high-dimensional data many features are **mutually exclusive** - never nonzero at the same time (think one-hot columns).

Bundle such features into one feature and build a **single histogram** for the bundle - nearly lossless.

- ▶ Histogram building drops from $\mathcal{O}(np)$ to $\mathcal{O}(n \cdot \text{\#bundles})$.
- ▶ Optimal bundling is NP-hard; a **greedy** grouping works well in practice.

Intuition: one-hot columns are mostly zeros and rarely fire together - pack them into one column and you lose almost nothing while doing far less work.

Ke et al. (2017).

CatBoost: the categorical specialist

Built around **categorical** features:

- ▶ **Ordered target statistics** - encode a category by its target mean computed on *past* rows only, so there is **no target leakage**;
- ▶ **Ordered boosting** - a permutation trick that removes the **prediction-shift** bias of ordinary boosting;
- ▶ **Symmetric (oblivious) trees** - the same split across a whole level: fast to evaluate, naturally regularized.

Intuition: encoding a category with the *whole* dataset's average target peeks at the answer (leakage); using only earlier rows is an honest running average - that is what "ordered" means. Prokhorenkova et al. (2018).

Which one, when?

	XGBoost	LightGBM	CatBoost
Tree growth	level-wise	leaf-wise	symmetric
Row sampling (option)	random	random or GOSS	random
Native categoricals	no*	yes	yes (best)
Speed on large n	good	fastest	good
Reach for it when	solid baseline	default for tabular	many categoricals

* XGBoost can also grow leaf-wise and gradient-subsample, and has experimental categorical support. The table shows *defaults* - except row sampling, which is *off* by default in all three (cells name the available strategy). **All three are the same gradient boosting** - the differences are engineering and categorical handling.

Using them well

Tuning order, domain constraints, and reading the model.

Hyperparameters that matter


Group	XGBoost	LightGBM
Learning	eta, n_estimators (+ early stop)	learning_rate, n_estimators
Complexity	max_depth, min_child_weight	num_leaves, min_data_in_leaf
Regularize	gamma, lambda, alpha	min_gain_to_split, lambda_l1/l2
Randomness	subsample, colsample_by*	bagging_fraction, feature_fraction

The knob differs: XGBoost sets tree size with max_depth; LightGBM with num_leaves (default 31). Defaults: XGB eta 0.3 / depth 6; LGBM lr 0.1.

The tuning order that matters

1. Set `n_estimators` *large* and use **early stopping**.
2. Trade **learning rate** against the number of trees.
3. **Tree complexity**: `max_depth` (XGB) / `num_leaves` (LGBM).
4. **Regularization**: `gamma`, `lambda`, `alpha`.
5. **Subsampling**: rows + columns.

Search with random or Bayesian optimization (see the hyperparameter-tuning deck).



Early stopping takes the round before validation stalls

- step 1.

Inject domain knowledge: constraints

Sometimes you *know* something the data should obey:

- ▶ **Monotonic constraints** (`monotone_constraints`): force a feature's effect to be non-decreasing (or non-increasing) - e.g. "higher dose $\Rightarrow$ no lower response".
- ▶ **Interaction constraints**: restrict which features may appear together on a path.


Constraints add **trust** and act as **free regularization** - use them when the direction or structure is known a priori.

Are they black boxes? Feature importance

Boosted models are **not** opaque:

- ▶ built-in importances: **gain** (loss reduction), cover, frequency;
- ▶ but **split-count importance is biased** toward high-cardinality features - prefer **permutation importance**.

For per-prediction attribution there is **SHAP** (fast *TreeSHAP* for trees) - we will cover it in its own session.



In code: LightGBM and XGBoost, side by side

LightGBM (the favorite)

```
from lightgbm import LGBMClassifier
from lightgbm import early_stopping
m = LGBMClassifier(
    n_estimators=2000,
    learning_rate=0.05,
    num_leaves=31)
m.fit(X_tr, y_tr,
      eval_set=[(X_val, y_val)],
      callbacks=[early_stopping(50)])
```

XGBoost (near-identical API)

```
from xgboost import XGBClassifier
m = XGBClassifier(
    n_estimators=2000,
    learning_rate=0.05,
    max_depth=6,
    early_stopping_rounds=50)
m.fit(X_tr, y_tr,
      eval_set=[(X_val, y_val)])
```

No install? `sklearn.ensemble.HistGradientBoostingClassifier` mirrors both.

Beyond binary classification

The same g, h machinery powers many tasks - just change the **objective**:

- **Regression** (reg:squarederror), **multiclass** (softmax), **ranking** (rank:* / LambdaMART), Poisson, quantile - even **custom** losses (supply your own g and h).

Imbalanced classes? Reweight the rare class with `scale_pos_weight` (XGBoost) or `is_unbalance / class_weight` (LightGBM), and judge with **AUC / PR**, not accuracy.

When boosting struggles

Powerful, but not magic:

- ▶ **No extrapolation** - a tree predicts a constant outside the training range; it cannot trend beyond what it has seen.
- ▶ **Needs tuning** - strong defaults exist, but top results need real HP search (a random forest is more forgiving).
- ▶ **Sequential** - trees are built one after another; less parallel than bagging.
- ▶ **Not for raw unstructured data** - images, audio, long text → neural nets; very sparse high-dimensional signals can favor linear models.

Rule of thumb: reach for boosting on **tabular** data; reach elsewhere when the data is not tabular.

Stacking


The third ensemble idea: let a meta-learner blend different model types - and close the chapter.

Stacking: a meta-learner over base models

Bagging and boosting combine *one* kind of model. **Stacking** combines **different** ones:

- ▶ train several diverse **base learners** (level 0) - e.g. LightGBM, a random forest, a linear model;
- ▶ feed their predictions to a **meta-learner** (level 1) that learns how best to *combine* them.

Intuition: different models make different mistakes; the meta-learner learns whom to trust, and when.



Stacking done right: out-of-fold predictions

The trap: if the meta-learner trains on base predictions of the *same* rows the base models were fit on, it sees **leaked**, over-optimistic inputs.

The fix: feed it **out-of-fold** (cross-validated) predictions - each row's base prediction comes from a model that did *not* train on it. StackingClassifier does this for you via cv.

```
from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import LogisticRegression

base = [("lgbm", LGBMClassifier()), ("rf", RandomForestClassifier())]
clf = StackingClassifier(estimators=base,
                        final_estimator=LogisticRegression(), cv=5)
clf.fit(X_tr, y_tr)    # cv=5 -> honest out-of-fold base predictions
```

Simpler cousins: voting and blending

- ▶ **Voting** - no meta-learner, just combine base predictions directly: **hard** (majority class) or **soft** (average predicted probabilities, usually better).
`VotingClassifier`.
- ▶ **Blending** - like stacking, but the meta-learner trains on a single **holdout** set instead of full cross-validation. Simpler and faster; uses a bit less data.

Stacking usually wins but costs the most (every base model + the meta-learner, with CV). Start with one strong model; stack only to chase the last bit of accuracy.

The ensemble landscape

	Bagging (random forest)	Boosting (GB / XGB / LGBM)	Stacking
Base models	same type, parallel	same type, sequential	different types
Combine by	averaging	weighted additive sum	a meta-learner
Mainly cuts	variance	bias	both
Trains	in parallel	sequentially	bases parallel, then meta
Use when	low-tuning baseline	top tabular accuracy	squeeze diverse models

Three ways to turn many so-so models into one strong one.

Recap: the trees and ensembles chapter

- ▶ **One tree** ([17]) - interpretable, but unstable.
- ▶ **Bagging** → **random forests** ([18]) - average many parallel trees, cut variance.
- ▶ **Boosting** ([19]) - add trees sequentially to fix errors, cut bias; gradient descent in function space.
- ▶ **The libraries** ([20]) - XGBoost (regularized objective + structure score), **LightGBM** ❤️ (leaf-wise + GOSS + EFB), CatBoost (categoricals); tune in order, add constraints, read importance.
- ▶ **Stacking** - a meta-learner over diverse models, to squeeze out the last gains.

Trees + ensembles dominate tabular data. Next: models for data where trees struggle - images, text, sequences.