


Before training: how far back does it 'hear'?



— vanilla RNN


— LSTM, forget bias 0 (PyTorch default)

— LSTM, bias 1 (Keras default)

— LSTM, bias 3 (gate open)

After training: recall the first token after T steps

accuracy, % (bar = mean, dots = 2 seeds)



--- chance (25%)