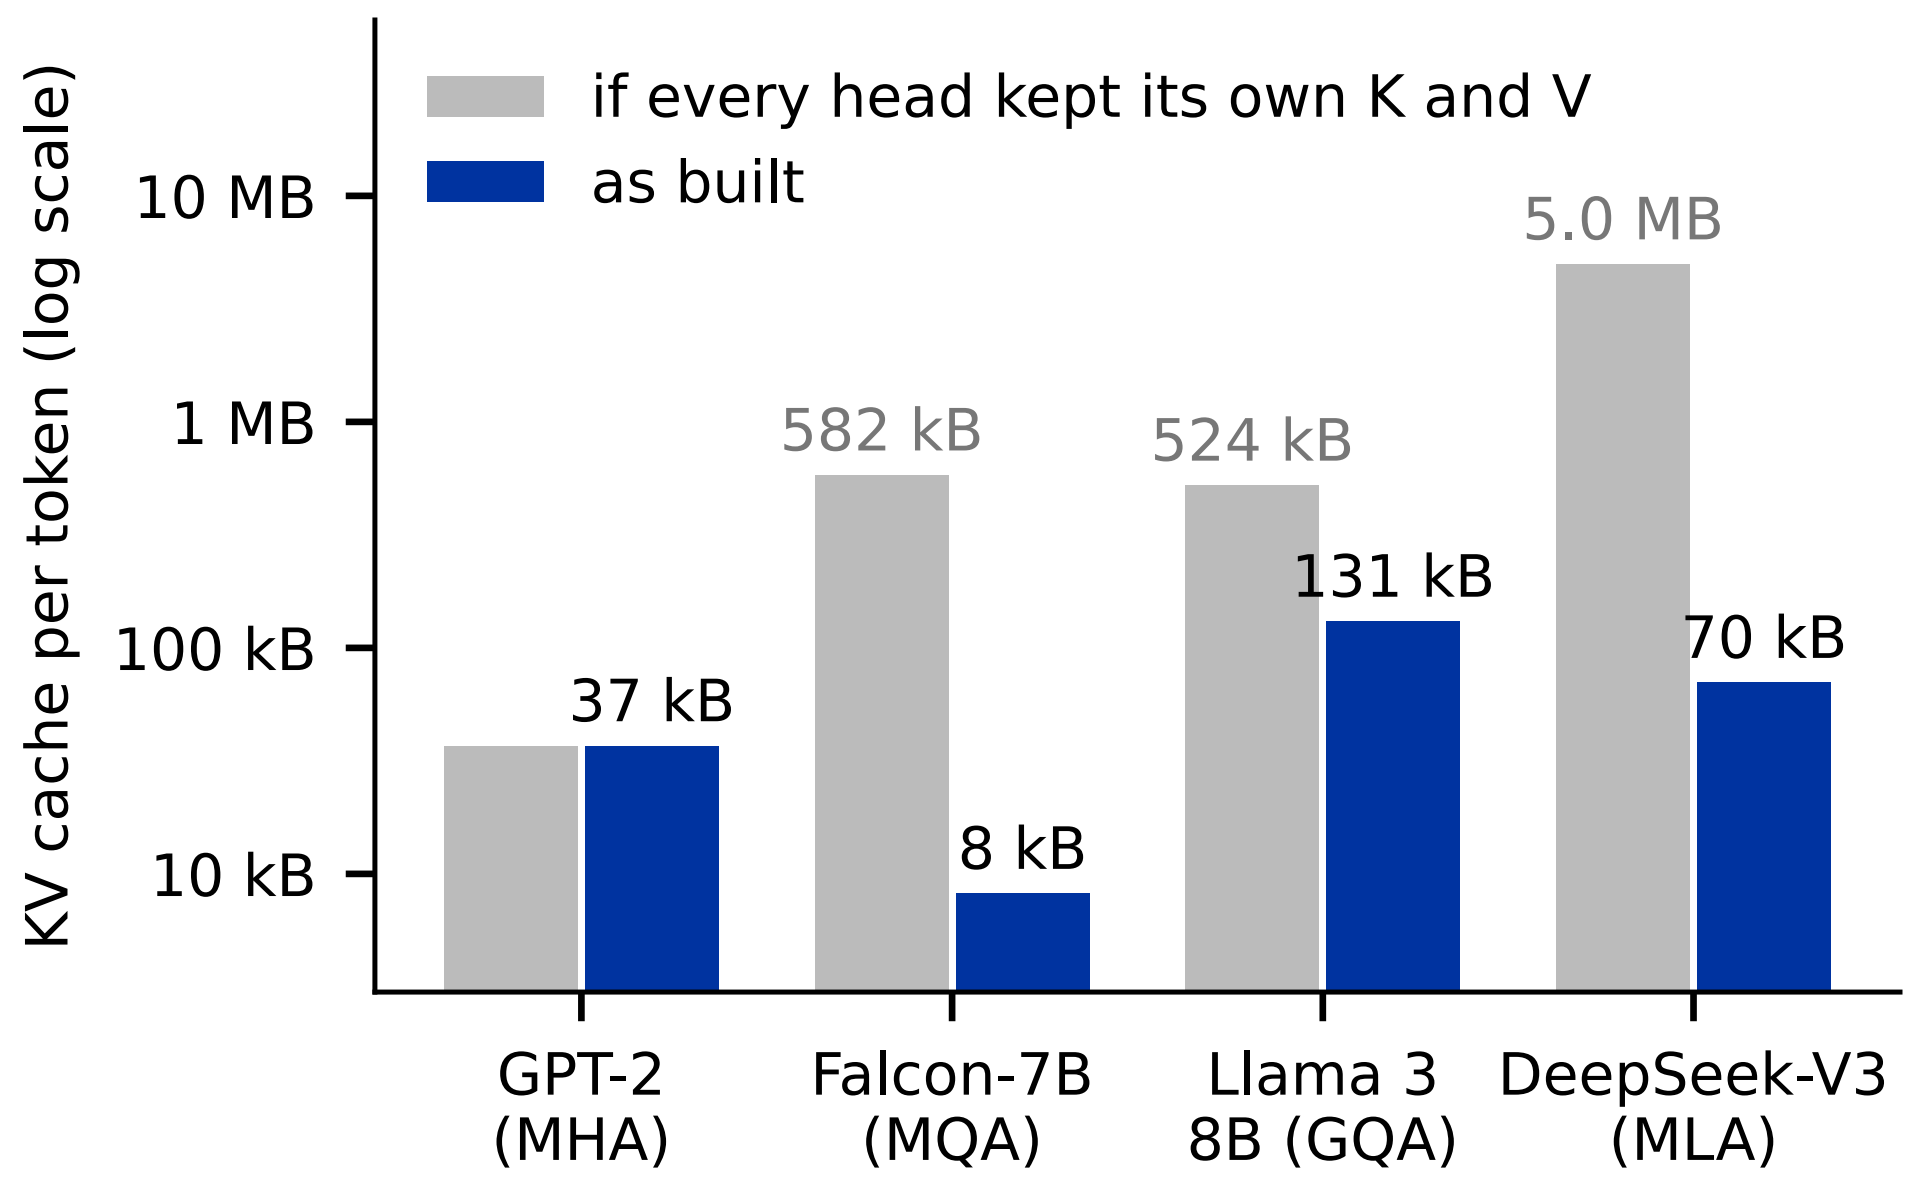
if every head kept its own K and V
as built
KV cache per token (log scale)
10 MB
1 MB
100 kB
10 kB
37 kB
582 kB
8 kB
524 kB
131 kB
5.0 MB
70 kB
GPT-2
(MHA)
Falcon-7B
(MQA)
Llama 3
8B (GQA)
DeepSeek-V3
(MLA)