

make it

use it

run and adapt it



Pretrain

next token on
~15T tokens



Scale

loss falls as
a power law



Post-train

base model
to assistant



Prompt

steer it,
no training



Evaluate

benchmarks
and judges



Quantize

8B model:
16 → 5 GB



LoRA

train < 1%
of weights