

Not every rung exists on every GPU

T4
(Turing, Colab free)

A100
(Ampere)

H100
(Hopper)

TF32

no

yes

yes

BF16 autocast

no

yes

yes

FP16 + gradient scaler

yes

yes

yes

torch.compile

yes

yes

yes

FlashAttention kernel

fallback

yes

yes

power-of-two shapes

yes

yes

yes

TF32 and BF16 arrived with Ampere; on Turing the FlashAttention kernel is unavailable and PyTorch uses its memory-efficient backend