

Lecture 21: The Road to Attention

The fix that worked, and the wall it hit

Parts of this chapter adapt LMU I2DL (slds-lmu), CC BY 4.0.

The disease, recalled

L20 ended on a number: influence fades as V^k , and $0.8^{29} \approx 0.0015$. Multiplying the same matrix into itself, 29 times, starved the gradient before it ever reached the first token.


Predict: what arithmetic operation would *not* shrink a signal no matter how many times you repeated it?

The disease, recalled

L20 ended on a number: influence fades as V^k , and $0.8^{29} \approx 0.0015$. Multiplying the same matrix into itself, 29 times, starved the gradient before it ever reached the first token.

Predict: what arithmetic operation would *not* shrink a signal no matter how many times you repeated it?

Addition. Repeated multiplication decays (or explodes); repeated addition just accumulates. Hold that thought - it is the first half of this lecture, and you have already seen it once (ResNet's skip connection, L17).



Multiplication: the disease. **Addition:** today's medicine.

Outline

LSTM: the additive fix

Tokens: how text becomes numbers

Language modeling

Seq2Seq and the wall

Epilogue: recurrence in 2026

A memory highway


1997's answer to a disease diagnosed, in full, in L20.

The cell state

The LSTM adds a second track running alongside the hidden state z : the **cell state c** . At every step it is only lightly modified - a little is removed, a little is added.


Sorting-line metaphor (ch1-6's world): the belt itself carries context past every station; each worker only adds or removes a little as the box goes by. Nobody rebuilds the box from scratch at every station.

Every horizontal line in the diagram is a full vector; circles are element-wise operations; boxes are learned layers.



Anatomy of one LSTM step

Same toy setup as L20's anatomy frame: a 3-word vocabulary, now a **2-d cell state** alongside the 2-d hidden state. Three gates and one candidate, each a function of $\mathbf{x}^{[t]}$ and $\mathbf{z}^{[t-1]}$; each gate is a sigmoid - a vector in $(0, 1)^2$ used as an **element-wise mask** ($\odot$).




FORGET: drop the old subject on a new one (LMU's gender-pronoun case) INPUT: write down a strong opinion word OUTPUT: what the readout needs right now

Predict: forget gate $[0.9, 0.1]$ meets cell state $[2.0, -1.0]$ - what survives?

Anatomy of one LSTM step

Same toy setup as L20's anatomy frame: a 3-word vocabulary, now a **2-d cell state** alongside the 2-d hidden state. Three gates and one candidate, each a function of $\mathbf{x}^{[t]}$ and $\mathbf{z}^{[t-1]}$; each gate is a sigmoid - a vector in $(0, 1)^2$ used as an **element-wise mask** ($\odot$).



FORGET: drop the old subject on a new one (LMU's gender-pronoun case) INPUT: write down a strong opinion word OUTPUT: what the readout needs right now

Predict: forget gate $[0.9, 0.1]$ meets cell state $[2.0, -1.0]$ - what survives?

$[0.9, 0.1] \odot [2.0, -1.0] = [1.8, -0.1]$: memory 1 survives almost untouched, memory 2 is nearly wiped. A gate is L14's sigmoid, rebranded as a soft switch.

For the record: the four equations

Reference, not exam material - shown once.

Forget gate: $\mathbf{f}^{[t]} = \sigma(\mathbf{W}_f^\top \mathbf{x}^{[t]} + \mathbf{V}_f^\top \mathbf{z}^{[t-1]} + \mathbf{b}_f)$


Input gate: $\mathbf{i}^{[t]} = \sigma(\mathbf{W}_i^\top \mathbf{x}^{[t]} + \mathbf{V}_i^\top \mathbf{z}^{[t-1]} + \mathbf{b}_i)$
 $\tilde{\mathbf{c}}^{[t]} = \tanh(\mathbf{W}^\top \mathbf{x}^{[t]} + \mathbf{V}^\top \mathbf{z}^{[t-1]} + \mathbf{b})$

Cell update: $\mathbf{c}^{[t]} = \mathbf{f}^{[t]} \odot \mathbf{c}^{[t-1]} + \mathbf{i}^{[t]} \odot \tilde{\mathbf{c}}^{[t]}$

Output gate: $\mathbf{o}^{[t]} = \sigma(\mathbf{W}_o^\top \mathbf{x}^{[t]} + \mathbf{V}_o^\top \mathbf{z}^{[t-1]} + \mathbf{b}_o)$ $\mathbf{z}^{[t]} = \mathbf{o}^{[t]} \odot \tanh(\mathbf{c}^{[t]})$

The punchline: $\mathbf{c}^{[t]} = \mathbf{f}^{[t]} \odot \mathbf{c}^{[t-1]} + \dots$ is **additive** - gradients flow through the cell state with minimal decay. ResNet solved *depth* with addition in 2015; LSTM solved *time* with addition in 1997 - same medicine, 18 years earlier.

Does it work? Proof + price



Same toy recall task as L20's empirical proof, same measurement method (RNNCell/LSTMCell unroll, per-step `retain_grad()`).

Honestly reported: at this budget the measured gap is not yet dramatic - untrained gates start close to "half open," and a hard recall task needs real training before the additive path pays off. The *architectural* argument (addition, not repeated multiplication) does not depend on this toy net converging.

Price: 4 weight sets instead of 1. Our $d_x=1, d_h=16$ vanilla cell: $16+256+16 = 288$ params. LSTM: $4 \times 288 = \mathbf{1,152}$.

GRU (Cho et al., 2014): a streamlined 3-gate cousin, similar results, a common default when 1,152 is too many.

Before any of this works...

...text must become numbers. The kit attention will need, built here.

The tokenization trilemma

The review sentence, split three ways:

"The pomegranates arrived fresh and sweet." split three ways

characters (41 tokens, vocab $\approx$ 100)

T	h	e	_	p	o	m	e	g	r	a	n	a	t	e	s	_	a	r	r	i	v	e	d	_	f	r	e	s	h	_	a	n	d	_	s	w	e	e	t	.
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

words (7 tokens, vocab explodes; "pomegranatey" = OOV)

The	pomegranates	arrived	fresh	and	sweet	.
-----	--------------	---------	-------	-----	-------	---

subwords (real BPE tokenizer) (10 tokens, vocab $\approx$ 100k)

The	_p	ome	gran	ates	_arrived	_fresh	_and	_sweet	.
-----	----	-----	------	------	----------	--------	------	--------	---

Characters: tiny vocabulary, but sequences get long and a letter carries no meaning on its own.
Words: meaningful units, but the vocabulary explodes - and "pomegranatey" was never in it (out-of-vocabulary, OOV). **Subwords:** the compromise every modern model uses - frequent chunks become tokens, rare words split into pieces.

Subwords on a real tokenizer

The same sentence through an actual production tokenizer (c1100k_base, the GPT-3.5/4 family), token boundaries and IDs visible:


10 tokens for a 6-word sentence -- vocabulary built from frequent character pairs merged together (BPE, one breath; full mechanics: ch8)



The vocabulary is chosen by merging frequent character pairs (byte-pair encoding, BPE) - one breath here; the full merge mechanics belong to the GenAI chapter. Try it live: **tiktok-enizer.vercel.app** - type Armenian and watch it fragment (next frame).

The Armenian tax

The chapter's Armenian line vs. its English gloss, through the *same* tokenizer - counts measured, not invented:




English-centric vocabularies fragment other scripts. Why it matters, twice over: longer sequences mean more steps, which means more fading (L20's disease again); and later, literally more money per API call - every token is billed.

One-hot and its poverty

Each token gets an ID; each ID, a one-hot vector. Callback: [03]'s one-hot encoding; L20's toy net used a 3-word version of exactly this.

In a 10,000-dimensional one-hot space, **every pair of words is equally far apart**. "good" is as far from "great" as it is from "plastic" - the geometry carries zero meaning.




Embeddings

Cashing in L20's planted beat, **by name**: the anatomy frame showed one-hot $\times \mathbf{W}$ selecting a row of $\mathbf{W}$ - that row *is* the word's own vector. An **embedding table** is just that $\mathbf{W}$, grown to vocabulary size and learned.

Dense, learned vectors: similar words end up close: the famous king-man+woman $\approx$ queen arithmetic.

Trained once on huge corpora (word2vec, GloVe) and reused - a warm start, i.e. L18's transfer learning, for text. Limitation in one line: "bank" gets ONE vector regardless of context - the GenAI chapter scratches that itch. Explore: projector.tensorflow.org.

Illustrative schematic of a word-embedding space (2-D projection)
real embeddings: 100-1000 dimensions, learned from huge text corpora



One task to rule them all

Predict the next token. That is the whole objective.

The objective

$$P(y_1, \dots, y_T) = \prod_{i=1}^T P(y_i \mid y_1, \dots, y_{i-1})$$

Turns raw text into supervised data **for free**: every prefix is an input, every next token a label. No labeling budget, no annotators - just more text.

Said slowly, once: this exact objective, scaled up, is what GPT trains on.

"The pomegranates arrived fresh"



The → pomegranates

The pomegranates → arrived

The pomegranates arrived → fresh

every prefix → next-token pair, free supervision

The mechanics, click by click

Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 1 of 6

feeding **q**u a l i h r one character at a time

ILLUSTRATIVE -- not measured from a trained model

q u a l i h r

(distribution appears once the first character is fed)

no character fed yet -- the net has only its initial state

The mechanics, click by click


Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 2 of 6

feeding **q**ulh

one character at a time

ILLUSTRATIVE -- not measured from a trained model




The mechanics, click by click

Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 3 of 6

feeding "cu" one character at a time

ILLUSTRATIVE -- not measured from a trained model




The mechanics, click by click

Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 4 of 6

feeding **q**ulh

ILLUSTRATIVE -- not measured from a trained model




The mechanics, click by click

Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 5 of 6

feeding "c" one character at a time

ILLUSTRATIVE -- not measured from a trained model




The mechanics, click by click

Feed a word's characters one at a time; at every step the net outputs a probability distribution over the *next* character (course cheese joke, five letters, shown below). **(Click to advance.)**

step 6 of 6

feeding Պաշկը one character at a time

ILLUSTRATIVE -- not measured from a trained model



No repeated character here, so no "double letter" surprise to spring - one remark instead: the word's third letter also appears in several other words of the chapter's Armenian line. The model's bet after that letter depends on what came before it, not on the letter alone.

Demo: a char-LSTM learns to write Armenian

What a character-level LSTM trained on Armenian text typically produces, one checkpoint per click (illustrative progression, not a measured run): gibberish → letter statistics → word-like fragments → almost-literary → a seeded continuation. **(Click to advance.)**

checkpoint 1 of 5

ILLUSTRATIVE -- typical progression, not a measured run

a char-LSTM learns to write Armenian (typical progression)

step 0 (untrained)



թգխղրջ անիաուրծ խքացնեռ ընթղմա ցվատեղծ որրիմբ

pure noise -- random characters from the alphabet

Demo: a char-LSTM learns to write Armenian

What a character-level LSTM trained on Armenian text typically produces, one checkpoint per click (illustrative progression, not a measured run): gibberish → letter statistics → word-like fragments → almost-literary → a seeded continuation. **(Click to advance.)**

checkpoint 2 of 5

ILLUSTRATIVE -- typical progression, not a measured run

a char-LSTM learns to write Armenian (typical progression)

step ~100



անի որնա տեմա անու Լանիս որա Լամի անե

letter statistics emerge -- vowel-consonant rhythm, no real words yet

Demo: a char-LSTM learns to write Armenian

What a character-level LSTM trained on Armenian text typically produces, one checkpoint per click (illustrative progression, not a measured run): gibberish → letter statistics → word-like fragments → almost-literary → a seeded continuation. **(Click to advance.)**

checkpoint 3 of 5

ILLUSTRATIVE -- typical progression, not a measured run

a char-LSTM learns to write Armenian (typical progression)

step ~1,000



նա ամեն լավ խոսք չէր գտնում մեջքով վարդի աչքով

word-like fragments -- some real short words, invented ones in between

Demo: a char-LSTM learns to write Armenian

What a character-level LSTM trained on Armenian text typically produces, one checkpoint per click (illustrative progression, not a measured run): gibberish → letter statistics → word-like fragments → almost-literary → a seeded continuation. **(Click to advance.)**

checkpoint 4 of 5

ILLUSTRATIVE -- typical progression, not a measured run

a char-LSTM learns to write Armenian (typical progression)

step ~5,000



Ես տեսնում եմ այս ամենը լուռ ու մտածում իմ մասին

almost-literary -- real words, loose grammar, a plausible-sounding line

Demo: a char-LSTM learns to write Armenian

What a character-level LSTM trained on Armenian text typically produces, one checkpoint per click (illustrative progression, not a measured run): gibberish → letter statistics → word-like fragments → almost-literary → a seeded continuation. **(Click to advance.)**

checkpoint 5 of 5

ILLUSTRATIVE -- typical progression, not a measured run

a char-LSTM learns to write Armenian (typical progression)


final (seeded)



Ես այս ամենիկնչ նայում եմ լուռ ու մտքումս խոսում եմ ինքս ինձ հետ

seeded with the chapter's opening words -- the model continues the line

The loss curve, and a link



Loss falls as the checkpoints improve (schematic curve) - the same next-token objective from two frames ago, watched from the outside.

This demo is the Armenian edition of a famous post: Andrej Karpathy, "**The Unreasonable Effectiveness of Recurrent Neural Networks**" (karpathy.github.io/2015/05/21/rnn-effectiveness/) - char-RNNs writing Shakespeare, Linux source, and Wikipedia markup, all from the same one objective.

The most visual proof in this course that next-token prediction alone extracts real structure from raw text.


Two networks, one vector between them

Translation needs a different-length output. Here is the classic fix.

Encoder-decoder

Translation needs different input/output lengths. Fix: an **encoder** LSTM reads the source and hands ONE **context vector** C to a **decoder** LSTM that writes the target.

Running example: "I watch all of this in silence, and the connoisseurs speak inside me." goes in - and the decoder writes out the chapter's Armenian line, the sentence you just watched cost 104 tokens.




The bottleneck

Predict: the 42-word complaint review (L20's shoehorn-hacks frame - it lost its verdict to truncation already) squeezed into a single 300-dim context vector. What goes wrong?

The bottleneck

Predict: the 42-word complaint review (L20's shoehorn-hacks frame - it lost its verdict to truncation already) squeezed into a single 300-dim context vector. What goes wrong?

42 words: "...expecting the same deep red arils and rich tart sweetness... arrived bruised dry and honestly disappointing"




Lost: early words get overwritten as the state updates sequentially - "expecting", "childhood markets" fade first

Fixed capacity: 5 words or 500, the same-size C must hold them all

And the LSTM does NOT save us: the cell state is still fixed-size vector. The fix for *fading* memory did not fix *compressed* memory.


The cliffhanger

What if the decoder could look back at *all* the encoder states - and learn *where* to look?



That question has a name. It was answered between 2014 and 2017, and it ends this course's pre-history of sequence models. Next chapter: attention and transformers.

Recurrence, retirement, and a comeback



2017's transformer retired RNNs from mainstream NLP - but recurrence came back as state-space models (Mamba) and, fittingly, as **Hochreiter's own** xLSTM. Where RNNs still earn their keep: streaming/edge inference and small-data settings, where a transformer's full attention is overkill.

Recap

- ▶ Length varies, memory fades (L20) → LSTM fixes fading with an **additive** cell state - same medicine as ResNet, 18 years earlier - at $4\times$ the parameter cost.
- ▶ Text becomes numbers through **tokenization**: char / word / subword, and the Armenian line costs several times more tokens than its English gloss.
- ▶ One-hot is geometry-blind; **embeddings** are one-hot's learned, dense upgrade - the anatomy frame's "row of **W**," grown up.
- ▶ One objective - **predict the next token** - powers a char-LSTM writing Armenian and, at a different scale, GPT.
- ▶ **Seq2seq** squeezes an entire source sentence into one vector - and even the LSTM's fix for fading memory cannot fix a memory that is *compressed*.

Next: the decoder wants to look back at everything. Letting it do that - attention - is where the GenAI chapter begins.