

Plain PPO regresses on public NLP tasks; mixing in pretraining gradients (PPO-ptx) pays back the tax

public NLP benchmark score (schematic)

