

The deadly triad

Any two are safe. All three together, and there is no convergence guarantee at all.

bootstrapping

targets built from
your own estimates
(TD, Q-learning)

function approximation

one weight vector
shared across states
(a neural net)

danger

values can
diverge

off-policy

learning about a policy
you are not following
(the max, replay)

