



One MCTS iteration. Run it thousands of times, then play the most-visited move from the root.

1. Selection




walk down using UCT,
balancing wins against
how little you have tried

2. Expansion




add one child for a
move not yet explored

3. Simulation



play to the end
(random, or with a
learned policy)

4. Backpropagation



push the result back up
through every node
on the path