


Same problem, three answers: how many tokens does the language model see, and who decides?

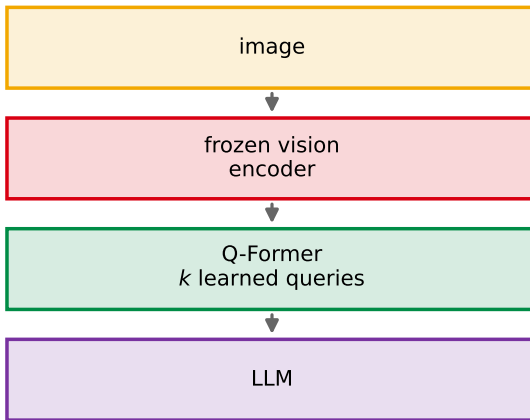
A - projector



576 tokens in,
576 tokens out

LLaVA

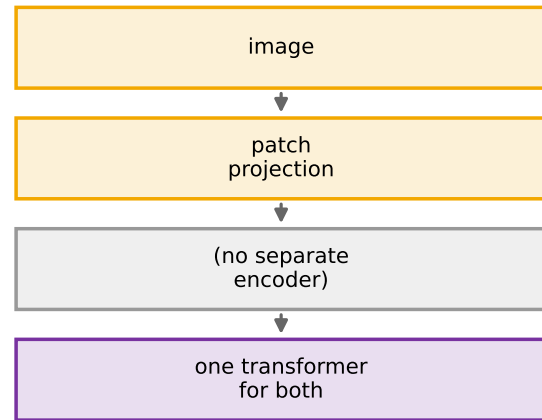
B - resampler



576 tokens in,
 $k = 32$ or 64 out

BLIP-2, Flamingo

C - early fusion



trained on both
from scratch

Fuyu, Chameleon