



All three are trying to reconcile ONE thing: a language model emits discrete symbols, an image is continuous.

1. fully discrete Chameleon



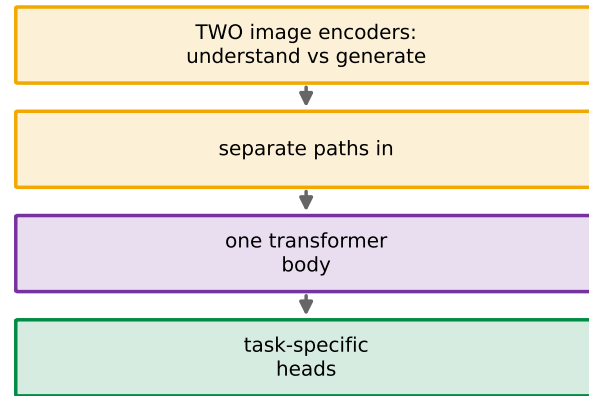
Clean. Pays the quantization cost.

2. hybrid Transfusion



No quantization loss. Two objectives to balance.

3. decoupled Janus



Understanding and drawing want different features.