


1. patchify
 24×24 , 4×4 patches



2. one linear projection



36 flattened patches

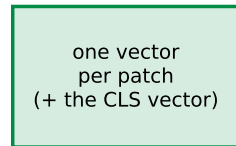
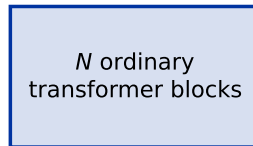
3. add position embeddings



now each patch knows where it was



4. prepend a classification token



Nothing on this row is vision-specific. Step 4 onward is ch9, unchanged.