


Each stage is its own transformer, conditioned on everything before it



*One model doing all three spends its capacity in the wrong place: gorgeous audio that says nothing.*