


Within each 80 ms frame, the TEXT token is predicted FIRST, and the audio is conditioned on it



*So the language model's fluency reaches the speech - and shifting this stream turns the same model into a recogniser or a synthesiser.*