



## The obvious design: distill semantics into level 1



one chain: levels 2-8 must model whatever level 1 left behind

Measured: phonetic discriminability up, audio quality DOWN. Not shipped.

## What Mimi ships: a split



in parallel, so acoustic detail need not live in the semantic quantizer's residual

They tried the obvious thing, measured it, and restructured (Défossez et al., 2024)