



The same question as the vision chapter: how many tokens does the language model see, and who decides?

A - projector
Qwen2-Audio



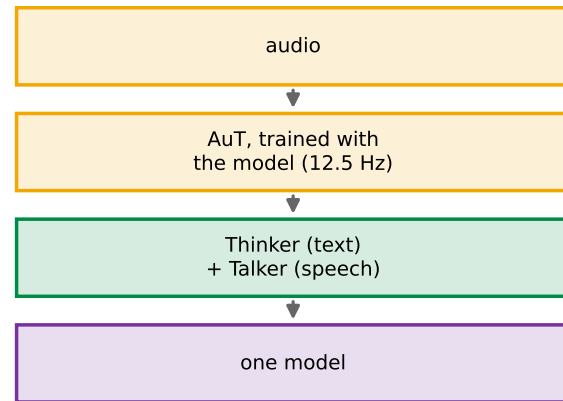
750 tokens per 30 s
(scales with length)

B - resampler
SALMONN



88 tokens per 30 s
(fixed, whatever is in it)

C - native
Qwen3-Omni



375 tokens per 30 s
and it can answer aloud