


Joint-embedding (CLIP, DINO)

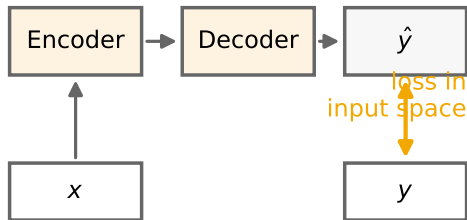
answers: do these match?

make them
agree



Generative (autoencoder, MAE, next token)

pays for every unpredictable pixel



JEPA (I-JEPA, V-JEPA)

predicts a target it invented

loss in
representation space

