


Bi-encoder

they never meet - so chunks precompute



Cross-encoder

accurate, but one pass per pair

