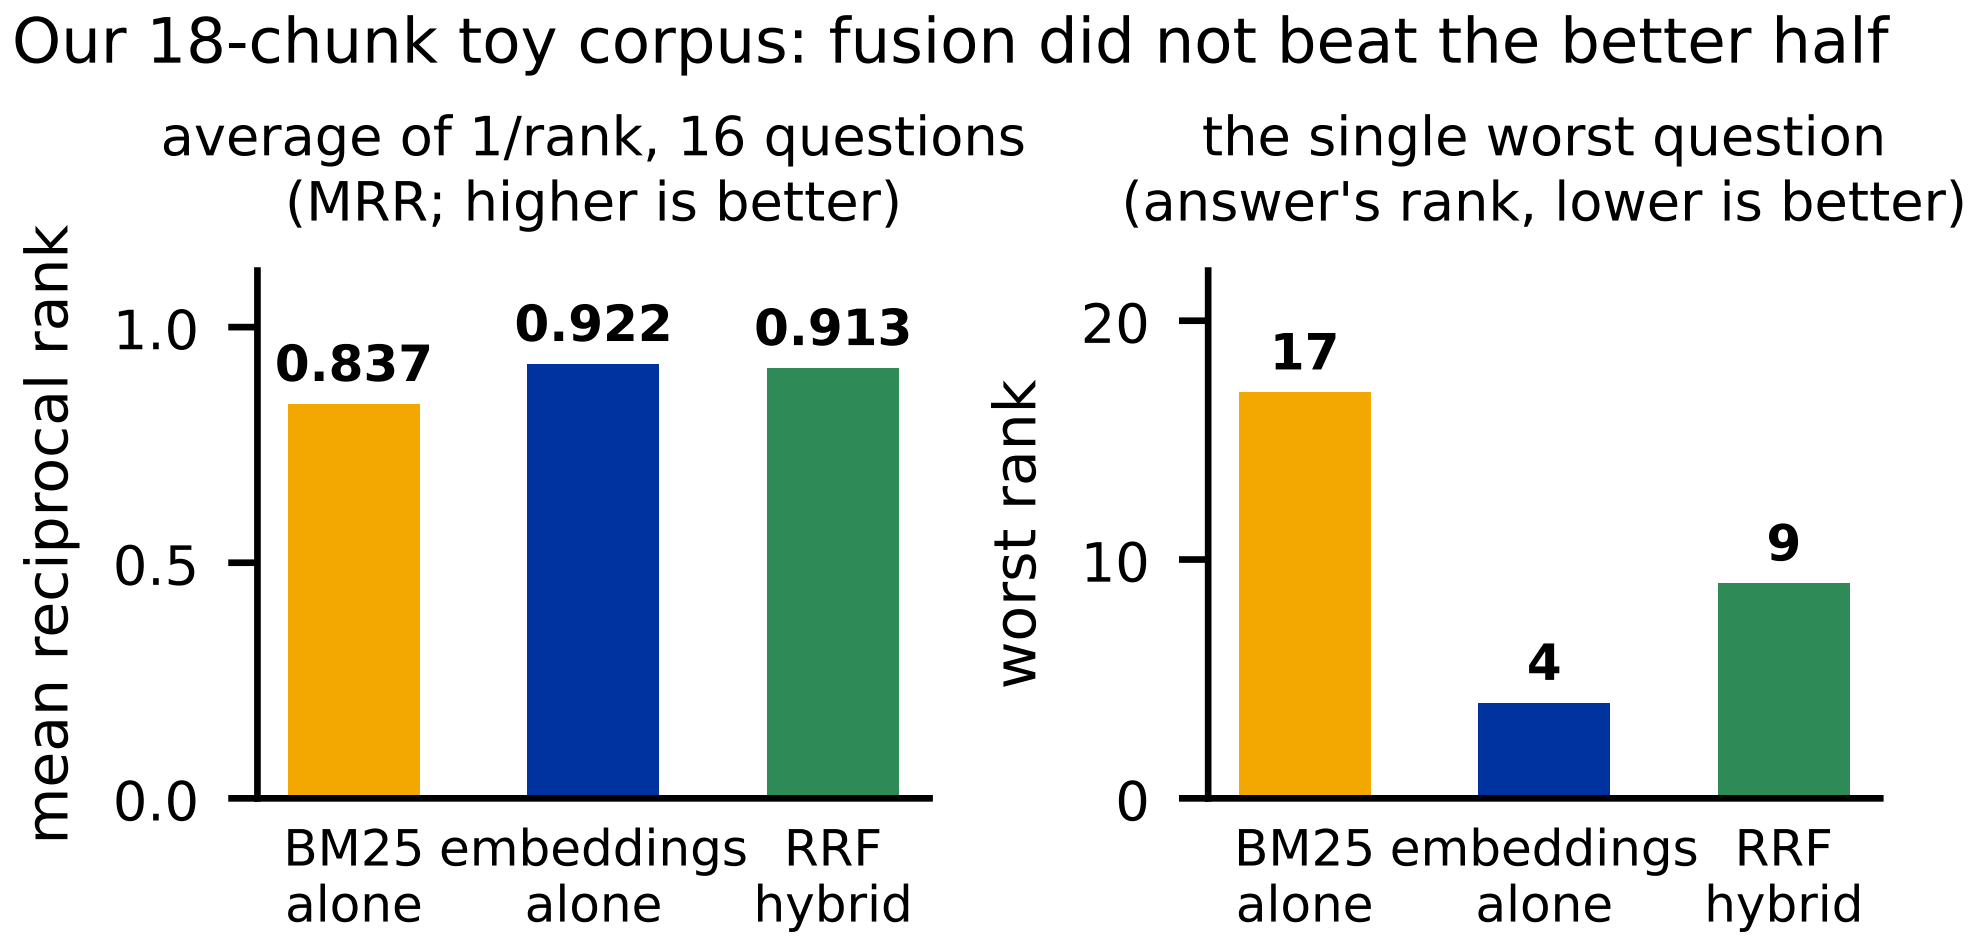
Our 18-chunk toy corpus: fusion did not beat the better half
average of 1/rank, 16 questions
(MRR; higher is better)
mean reciprocal rank
1.0
0.5
0.0
0.837
0.922
0.913
BM25 alone
embeddings alone
RRF hybrid
the single worst question
(answer's rank, lower is better)
worst rank
20
10
0
17
4
9
BM25 alone
embeddings alone
RRF hybrid