

**chunk the  
documents**



**retrieve  
top-k**



**assemble  
the prompt**



**generate  
the answer**



**what goes wrong**

the answer got  
cut in half

the right chunk  
is not in the top-k

two retrieved chunks  
contradict each other

the answer ignores  
the context

**who notices**

**nothing downstream  
can recover it**

**the retrieval metrics  
see this**

**no standard metric  
sees this**

**faithfulness  
sees this**