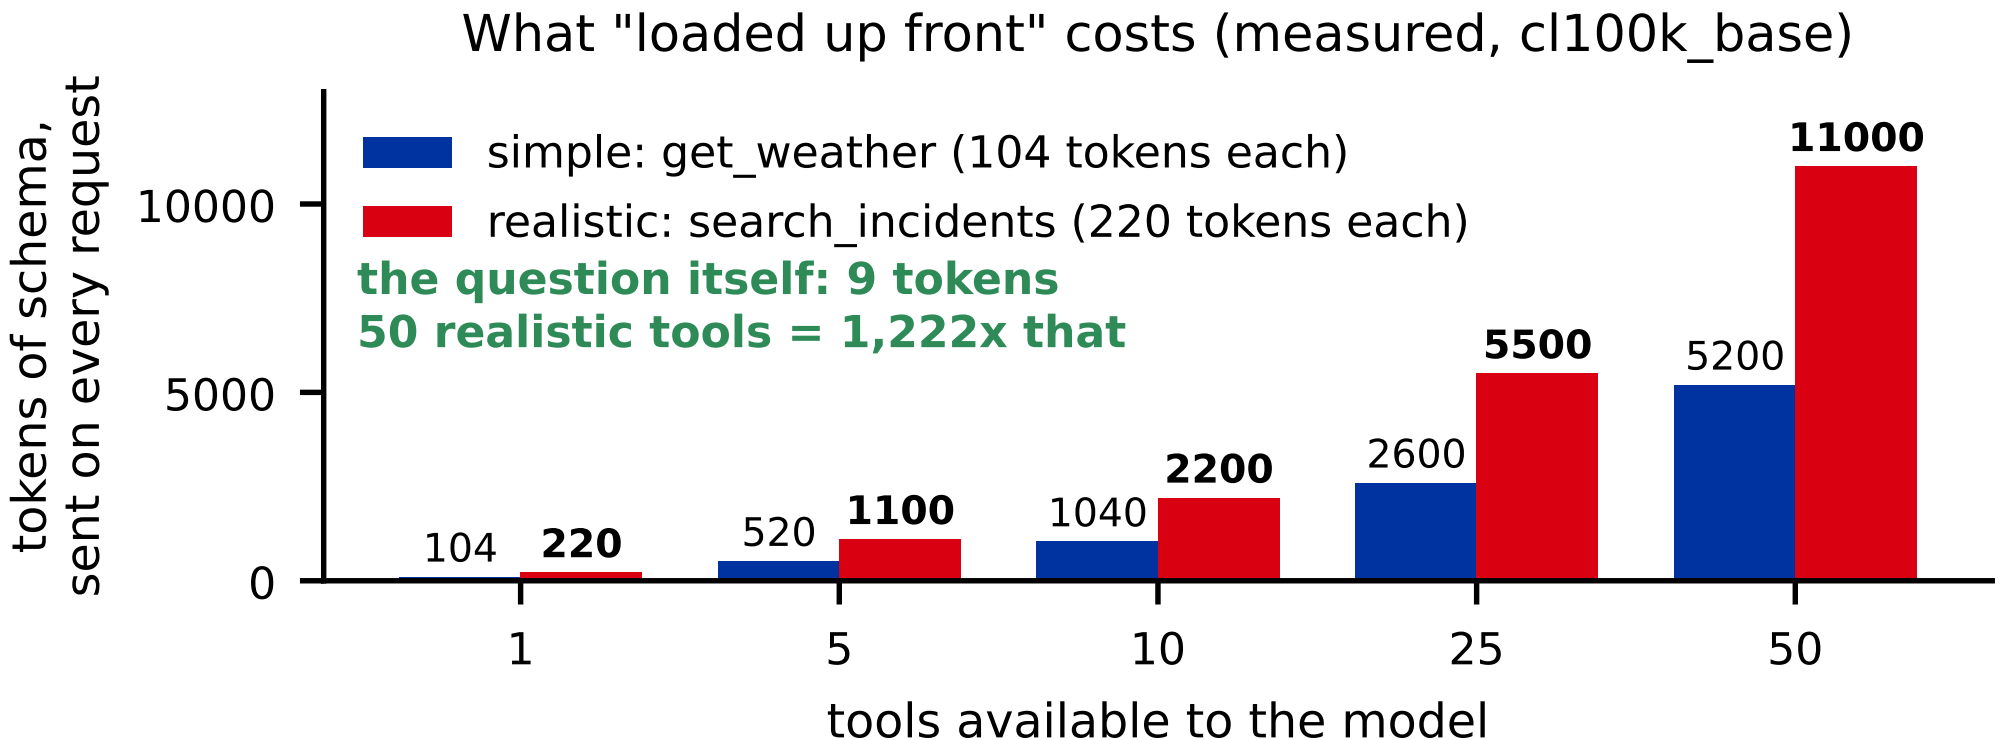
What "loaded up front" costs (measured, cl100k_base)
tokens of schema, sent on every request
simple: get_weather (104 tokens each)
realistic: search_incidents (220 tokens each)
the question itself: 9 tokens
50 realistic tools = 1,222x that
0
5000
10000
104
220
520
1100
1040
2200
2600
5500
5200
11000
1
5
10
25
50
tools available to the model