

	asks the model questions	reads its gradients	reads its weights and activations
the whole model	permutation importance PDP / ALE SHAP summary	<i>(rare)</i>	circuits probes sparse autoencoder features linear-model coefficients*
one prediction	LIME SHAP (one row) counterfactuals	saliency maps integrated gradients Grad-CAM	logit lens activation patching attribution graphs
	chapter 5	chapter 6 + the vision lecture	this chapter