


Trade one busy vector for a very wide, very empty one




trained on two things at once:

rebuild the input + keep almost every latent at zero

*Nothing supervises what the latents mean. Sparsity alone does the work.*

A real SAE, trained here: 8 features hidden in 4 dimensions



Every planted feature recovered (worst 0.84). Mean 3.6 of 32 latents active per input.