


Trade one busy vector for a very wide, very empty one




residual stream 24,576 latents - almost all exactly zero
768 numbers, all busy rebuilt vector

trained on two things at once:

rebuild the input + keep almost every latent at zero

Nothing supervises what the latents mean. Sparsity alone does the work.

A real SAE, trained here: 8 features hidden in 4 dimensions



Every planted feature recovered (worst 0.84). Mean 3.6 of 32 latents active per input.