

After effective RLHF training, an LLM should not give information about how to cheat on an exam.

TRUE

0%

After effective RLHF training, an LLM should be able to follow constraints like "give an answer in 15 words".

FALSE

0%

After effective RLHF training, an LLM will in general precisely state what it assumes when giving an answer.

FALSE DEPENDS

0%

After effective RLHF training, an LLM will educate the user if their question is based on an incorrect assumption.

FALSE DEPENDS

0%

The improvement of RLHF over SFT is about as large as the improvement of SFT over the raw model.

TRUE

0%

After effective RLHF training, an LLM will understand python code better.

FALSE

0%

After effective RLHF training, an LLM will hallucinate less compared to the SFT model.

DEPENDS

0%

The reward model takes a prompt as input and returns a number that tells us how well written the prompt is.

FALSE

0%

In RLHF, we train the LLM on input-output pairs.

TRUE

0%

Supervised instruction finetuning is expensive because labelers demonstrate desired output behavior.

TRUE

0%

A ranked list of 5 outputs gives rise to 10 output pairs for training the reward model.

TRUE

0%

The training objective for the PPO model has three terms: the reward term, one KL divergence term that ensures that we do not diverge too much from the pretraining objective and one log likelihood term for fitting the instruction-tuning training set.

FALSE

0%

The SFT model, the reward model and the PPO model are all initialized with a "raw" language model that is trained with the next word prediction objective on the pretraining corpus.

TRUE

0%

In RLHF training, an LLM learns to give answers in several different languages (not just English) because these languages occur in the training set.

FALSE