


RLHF: reward model + RL loop

2 models trained + RL loop + in-loop sampling



DPO: one classification loss

1 loss - no reward model - no RL - no sampling

