

Online inference (batch 1): adapters add latency, LoRA adds none

