

The pre-training data mix (15.6T tokens, curated & de-duplicated)

