


RLHF: supervised demos $\rightarrow$ a learned reward $\rightarrow$ RL against it



π_{SFT} frozen as the KL anchor for step 3