


Expert capacity = (tokens / experts) × capacity factor: the drop-vs-waste trade-off

Capacity factor 1.0 (2 slots)

overflow token dropped
→ skips layer via residual



Expert 1

2 routed

Expert 2


3 routed

Expert 3

1 routed

Capacity factor 1.5 (3 slots)

fits all tokens, but
empty slots = wasted compute



Expert 1

2 routed

Expert 2

3 routed

Expert 3

1 routed