

FlashAttention: up to ~3x end-to-end, exact result (7.6x on the attention op alone)

