

MLA (inherited from DeepSeek-V2): project K,V into a shared low-rank latent -- schematic

MLA: cache one small latent, rebuild per head

